

Reciprocal Mersenne Subseries

Support criteria, finite arithmetic and conditional orbit problems

WILL COOK*

18 September 2026

ABSTRACT

We prove a weighted summability criterion for the irrationality of reciprocal Mersenne subseries. For an integer base $b \geq 2$ and a finite nonempty prime set $\mathcal{D}$, let $h(a) = \prod_{p \in \mathcal{D}} p^{v_p(a)}$ be the part of a supported on these primes. Every infinite set $A \subseteq \mathbb{N}_{>0}$ satisfying

$$\sum_{a \in A} \frac{h(a)}{a(b^{h(a)} - 1)} < \infty$$

has irrational $\sum_{a \in A} (b^a - 1)^{-1}$. The base-two condition gives the conclusion at every integer base for every infinite subset of A , and it permits an explicit example with divergent reciprocal sum. The condition nevertheless excludes full support and the set of all odd exponents. The proof averages over multiples of a modulus divisible by a prescribed finite part of the support. A second average over dyadic observation lengths controls the errors from incomplete periods. Erdős already stated the weaker reciprocal-summable theorem without the coprimality restriction; its complete ordinary proof is also printed here. A finite estimate uniform in the moving modulus lets the weighted and positive-cover arguments share one observation window. This proves irrationality at every integer base for every infinite subset of the union of two sets satisfying the respective criteria.

We also give finite-denominator and achievement-set calculations, and formulate explicit conditional tests for the targets $1/2$ and $1/21$. These tests do not decide either target or the universal question. The classical full-support, periodic-support and strict-tail results used for comparison are attributed in the text. Existing work settles prime support at base 2 and squarefree support at every power-of-two base; those cases are not claimed as new.

How to read this document

Section 1 proves the support criteria and compares them with the literature. Sections 5–8 develop the finite arithmetic used in the rational-membership tests. Sections 9–12 distinguish the hypotheses of conditional implications from the finite evidence for them. Section 13 proves the arithmetic-sampling counterexample cited in the short paper. The global notation below is supplemented by definitions where each specialised recurrence is used.

* *Authorship and AI use.* Will Cook built and directed the research infrastructure and reviewed the claims when he could. AI agents did most of the research and drafting. Cook did not independently verify every claim.

A sufficient condition need not be necessary; an equivalent condition does not prove either side. Counterexamples concern the specified estimate, and a finite computation establishes only its stated range.

The short source tags distinguish the kind of evidence being cited. [Lean] identifies a cited formal-source declaration at the specified revision and follows the status recorded in the supplied material. The four support endpoints cited in Part I were replayed with their dependency closure under Lean 4.29.1 at public commit 065e09523286; their axiom audit reports only `propext`, `Classical.choice`, and `Quot.sound`. Other supplied files retain their individually stated status. [Cert] identifies a reported exact finite computation. [Math] identifies an ordinary argument, without asserting that no formal counterpart exists. [Cited] refers to the literature, and [Open] marks an assumption or conclusion not proved here. A statement may carry more than one tag. Parameter ranges are stated separately: a result at specified values or below a bound is not a result at arbitrarily large values.

The source links are for checking precise definitions and proof terms. They are not additional hypotheses and do not replace an explanation of the mathematics. No independent verification of every other proof is claimed. Neither the universal problem nor either rational-target question is resolved in this paper.

Notation used throughout

For $A \subseteq \mathbb{N}_{>0}$ and an integer $b \geq 2$, put

$$X_A(b) = \sum_{a \in A} \frac{1}{b^a - 1}, \quad c_A(n) = \#\{a \in A : a \mid n\} \quad (n \geq 1).$$

At base two write $w_n = (2^n - 1)^{-1}$, $R_n = \sum_{j > n} w_j$ and $\mathcal{A} = \{X_A(2) : A \subseteq \mathbb{N}_{>0}\}$. For $x \geq 0$ the real greedy remainder is defined by $r_0(x) = x$ and

$$r_n(x) = r_{n-1}(x) - \begin{cases} w_n, & w_n \leq r_{n-1}(x), \\ 0, & w_n > r_{n-1}(x). \end{cases}$$

The greedy support G_x consists of the selected indices; $G = G_{1/2}$. For the half target index 1 is skipped. Thus a recurrence begun at index 2 with remainder 1/2 is the same process, with its index shifted.

The finite quotient arguments use

$$q(M, d) = \left\lfloor \frac{2^M}{2^d - 1} \right\rfloor, \quad Q(D, M) = \sum_{d \in D} q(M, d).$$

For $M \geq k$ their nonnegative binary remainder and signed next-step expression are

$$S(D, k, M) = \max\{2^{M-k} - Q(D, M) - 1, 0\}, \\ H(D, k, n) = 2S(D, k, n-1) + 1 - c_D(n) \quad (n > k).$$

The maximum matters: subtraction in the corresponding natural-number source definition is truncated at zero. Statements that remove the maximum establish nonnegativity first.

The integer carry for a support omitting 1 is

$$\text{ihc}(A, N) = 2^N - \sum_{j=2}^{N+1} 2^{N+1-j} c_A(j), \quad C_A(N) = \text{ihc}(A, N) - 1.$$

Some finite-depth statements use $K_A(m) = \text{ihc}(A, m - 1)$. The recurring bound $B(m) = 2\lfloor\sqrt{m}\rfloor + 4$ is integral; bounds written with $\sqrt{m}$ rather than its floor are real-valued bounds. Quantities specific to one argument are defined in that argument. Some formal sources allow a support to be a subset of $\mathbb{N}$, including zero. In all reciprocal sums in this paper only positive exponents contribute: for such a support, $X_A(b)$ means $X_{A \cap \mathbb{N}_{>0}}(b)$. The expression $1/(b^0 - 1)$ is never evaluated.

Contents

I. Support theorems and their proofs	7
1 Support criteria and their proofs	7
1.1 Reciprocal-summable supports	7
1.2 A weighted condition on the support	8
1.3 A finite estimate on common observation scales	22
1.4 Positive divisor majorants	23
1.5 Combining the two support criteria	25
1.6 The universal problem and the half-value question	26
1.7 Known support theorems	27
1.8 Geometry and rational-target membership	30
II. Limitations and extensions of recorded methods	33
2 Limitations of the recorded methods	33
2.1 What the obstructions do and do not show	33
2.2 Uniqueness of the specified greedy words	34
2.3 Equivalent formulations of half-membership	35
2.4 The size of the required error bounds	36
2.5 What finite certificates decide	38
2.6 Why measure and category do not decide a rational target	39
2.7 Hypotheses of the CRT block-certificate argument	39
2.8 The distinct estimates still required	40
3 Unproved inputs for further arguments	40
3.1 A lower bound at reset ranks	41
3.2 Increasing the depths of exact finite sums	42
3.3 The known prime-support theorem	43
3.4 Equivalent membership conditions	44
3.5 Why zero windows do not split the support class	44
3.6 Comparison of the remaining inputs	44
4 Reading the detailed record	45
III. Detailed statement and arithmetic catalogue	46
5 The series, the greedy rule and finite identities	46
5.1 The object	46
5.2 What is proved	46
5.2.1 Full-support irrationality, for every base	46
5.2.2 Structured infinite supports	47

5.2.3	Topology and measure of the achievement set	48
5.3	The greedy rule for one half	48
5.4	Integer quotients and their remainder identity	49
5.5	A real-valued form of the quotient identity	50
5.6	Dynamics	51
5.7	Reported finite checks	52
6	Detailed results and their hypotheses	53
6.1	Conditional membership tests	53
6.2	Exact identities and reductions	61
6.3	Consequence theorems	73
6.4	Invariants	82
6.5	Obstructions and countermodels	93
6.6	Further finite and conditional results	97
6.6.1	The integer update in a middle interval	97
6.6.2	Finite approximations without compatibility	98
6.6.3	Removing one case from a sufficient tail condition	100
6.6.4	An eventual nonnegative margin	100
6.6.5	Rewriting the signed carry recurrence	101
6.6.6	Further restrictions on represented values	101
6.6.7	Conditional counterexamples at rational targets	101
6.6.8	Multiplicity of the stated two–three representations	102
6.6.9	Integer-quotient tests for one twenty-first	102
IV.	Scales, conditional interfaces and unresolved inputs	103
7	Finite ranges and unbounded parameters	103
8	Approximations to the greedy orbit	107
9	Which hypotheses remain unproved	111
9.1	The three #257 obligations, stated exactly	111
9.2	Reading a row	112
9.3	Reset estimates and their missing hypotheses	112
9.4	Finite quotient sums and unbounded depths	117
9.5	Universal support criteria and their limits	121
9.6	Three combined conditional arguments	125
9.7	Four limitations of the stated implications	126
10	Hypothesis-specific obstructions	127
10.1	Four tested approaches to the reset estimate	128
10.2	Excluding particular final dyadic intervals	130
10.3	Related counterexamples and restrictions	132
10.4	The scalar-localisation height obstruction	134
10.5	The final-skip band formula, and what it does not show	135

11 Unproved inputs for the recorded implications	136
11.1 Reset bounds and finite weighted divisor sums	137
11.2 Progress to larger endpoint depths	142
11.3 Prime support and the limits of one formal argument	144
11.4 Equivalent carry conditions at perfect-square depths	144
11.5 Sparse and dense supports under the certificate criterion	146
11.6 What the experiments and obstructions do not imply	147
 V. Outstanding obligations and provenance	 150
12 What is open, stated exactly	150
12.1 Notation fixed for this section	150
12.2 The universal statement	150
12.3 Whether one half is represented	151
12.3.1 Conditions equivalent to half-membership	151
12.4 Exact finite sums at unbounded depths	152
12.4.1 A capacity condition for exact sums at unbounded depths	153
12.5 Reset bounds and runs of right transitions	154
12.6 The remaining $-2, -1$ middle-cell exclusion	155
12.7 Whether one twenty-first is represented	155
12.8 Implication table	156
12.9 Possible inputs and their scope	157
12.10 Summary	157
 13 Logarithmic cost under arithmetic sampling	 158
13.1 The two finite costs	158
13.2 The arithmetic lower bound for every fractional cover	161
13.3 What survives without arithmetic restriction	162

I. Support theorems and their proofs

1 Support criteria and their proofs

1.1 Reciprocal-summable supports

For every infinite set A of positive integers with $\sum_{a \in A} 1/a < \infty$, the series

$$X_A(b) = \sum_{a \in A} \frac{1}{b^a - 1}$$

is irrational at every integer base $b \geq 2$. The proof below is ordinary mathematics and is the proof of the short paper's Theorem 1.2. Erdős's [1968 paper](#), p. 222 proves the pairwise-coprime case at every integer base and states the removal of coprimality without printing the details. The proof below gives the stated extension by averaging; no comparison with the omitted proof is possible. The linked pairwise-coprime theorem retains that extra hypothesis. The same paper says that $\sum 1/n_i < \infty$ could be replaced by a weaker but more complicated condition (p. 222), and suggests that for pairwise coprime supports Brun's method could probably replace it by $\sum_{n_i < x} 1/n_i = o(\log \log x)$ (p. 226). Theorem 1.1 below proves irrationality under one explicit weaker condition without coprimality; no identification with his suggestions is asserted.

The mechanism is to make positive displacements arbitrarily small. For $N > 0$, write $N = q_a a + r_a$, $0 \leq r_a < a$. The identity

$$\frac{b^N - 1}{b^a - 1} = b^{r_a} \sum_{j=0}^{q_a-1} b^{ja} + \frac{b^{r_a} - 1}{b^a - 1}$$

shows that

$$\Delta_{b,A}(N) := \sum_{a \in A} \frac{b^{N \bmod a} - 1}{b^a - 1} = (b^N - 1)X_A(b) - J_N, \quad J_N \in \mathbb{Z}.$$

The integer sum is finite, since $q_a = 0$ for $a > N$. Every summand of $\Delta_{b,A}(N)$ is nonnegative and an exponent $a > N$ makes it positive. Thus $X_A(b) = p/q$ would imply $\Delta_{b,A}(N) \geq 1/q$ for every $N > 0$.

Put $w_{b,d}(N) = b^{N \bmod d} / (b^d - 1)$. For a fixed positive modulus Q and $g = (Q, d)$, the orbit of $Qm \bmod d$ has length d/g . Its geometric sum gives

$$\lim_{X \rightarrow \infty} \frac{1}{X} \sum_{m=1}^X w_{b,d}(Qm) = M_{Q,b}(d) := \frac{g}{d(b^g - 1)} \leq \frac{1}{d}.$$

The important interchange is with the infinite support. Expanding the atom as $\sum_{r \geq 1} b^{-r} \mathbf{1}_{d|Qm+r}$ and counting positive multiples of d gives the bound

$$\frac{1}{X} \sum_{m=1}^X w_{b,d}(Qm) \leq \sum_{r \geq 1} b^{-r} \frac{Q+r}{d} \leq \frac{Q+2}{d}.$$

Here Q is fixed. Dominated convergence therefore identifies the limiting mean of $\sum_{d \in A} w_{b,d}(Qm)$ with $\sum_{d \in A} M_{Q,b}(d)$.

Now choose $Q_t = \text{lcm}(1, \dots, t)$. For each fixed d , eventually $d \mid Q_t$, so $M_{Q_t,b}(d) = (b^d - 1)^{-1}$. The bound $M_{Q_t,b}(d) \leq 1/d$ permits a second dominated-convergence argument:

$$\sum_{d \in A} M_{Q_t,b}(d) \rightarrow X_A(b).$$

Consequently the limiting mean of $\Delta_{b,A}(Q_t m)$ tends to zero. Given $\varepsilon > 0$, first choose t and then a sufficiently long finite average with mean below ε ; one term of that average is below ε . Its strict positivity contradicts the rational lattice. This proves the assertion. The exact all-base statement, with hypotheses $b \geq 2$, infinitude of A , and summability of $a \mapsto \mathbf{1}_A(a)/a$, is kernel-checked as `irrational_erdosSupportSeries_of_summable_reciprocal`. [Lean]

The order of the two limits matters: the observation length increases with Q_t fixed before t increases. No independence of different divisibility events is assumed. The mechanism also has a definite limit. At base two and full support, the contribution of exponents $a > N$ alone gives $\Delta_{2,\mathbb{N}_{>0}}(N) > 1 - 2^{-N} \geq 1/2$. Full-support irrationality therefore uses a different argument; its checked form appears below.

1.2 A weighted condition on the support

The reciprocal-summable case is a useful model, but powers of a fixed prime can make an exponent cheaper than its reciprocal suggests. For example, with $a = 2^k m$ and m odd, the weight below becomes $1/[m(b^{2^k} - 1)]$. The exponential denominator permits a large collection of odd factors at large k .

Theorem 1.1 (A weighted summability criterion). *Let $b \geq 2$ be an integer, let $\mathcal{D}$ be a finite nonempty set of primes, and let $A \subseteq \mathbb{N}_{>0}$ be infinite. Write*

$$h(a) = \prod_{p \in \mathcal{D}} p^{v_p(a)},$$

where $v_p(a)$ is the exponent of p in a ; thus $h(a)$ is the $\mathcal{D}$ -part of a . If

$$W_{b,\mathcal{D}}(A) := \sum_{a \in A} \frac{h(a)}{a(b^{h(a)} - 1)} < \infty, \tag{1}$$

then

$$X_A(b) = \sum_{a \in A} \frac{1}{b^a - 1}$$

is irrational. More precisely, for

$$\Delta_{b,A}(m) = \sum_{a \in A} \frac{b^{m \bmod a} - 1}{b^a - 1},$$

for every $\varepsilon > 0$ and every N there is an $m \geq N$ with $0 < \Delta_{b,A}(m) < \varepsilon$.

[Math](uniform)divisibility-weighted

Lean: `divisibilityWeightedClaim, weighted_displacement_cofinal_close_return`.

Examples and limits of the hypothesis. Squares, powers of two and factorials satisfy reciprocal summability and therefore satisfy the base-two weighted condition. Full support does not: for any fixed finite prime set P , its subset coprime to $\prod_{p \in P} p$ has weight $1/[a(b-1)]$ and divergent reciprocal sum. The set of all odd exponents also fails, by restricting to integers coprime to $2 \prod_{p \in P} p$. Inclusion–exclusion in the harmonic sum proves both divergences. This is a limitation of the sufficient condition, not an argument for rationality; full support is already known to be irrational.

Writing $a = hm$ with h its $\mathcal{D}$ -part gives the exact decomposition

$$W_{b,\mathcal{D}}(A) = \sum_{h \in \mathcal{H}_{\mathcal{D}}} \frac{1}{b^h - 1} \sum_{\substack{m \geq 1, hm \in A \\ (m, \prod_{p \in \mathcal{D}} p) = 1}} \frac{1}{m},$$

where $\mathcal{H}_{\mathcal{D}}$ consists of 1 and the positive integers with no prime factor outside $\mathcal{D}$. Each inner reciprocal sum must be finite; the displayed outer sum imposes the additional summability across different $\mathcal{D}$ -parts. For example, at base two with $\mathcal{D} = \{2\}$, choose in each layer $2^k, k \geq 1$, a finite set of odd factors whose reciprocal sum is at least 2^{2^k} . Every inner sum is finite, but the outer weighted sum diverges. Such finite sets exist by divergence of the harmonic sum over odd integers. Prime support also fails this criterion for every finite $\mathcal{D}$: its contribution from primes outside $\mathcal{D}$ is a constant multiple of their divergent reciprocal sum. These failures concern the sufficient hypothesis, not rationality of the resulting values. The following construction shows, nonetheless, that this condition is strictly weaker than reciprocal summability. At base two take

$$A_{\star} = \{2^k m : k \geq 1, m \text{ odd}, m \leq 2^{2^k}\}.$$

Write $H_N = \sum_{m=1}^N 1/m$ and $N_k = 2^{2^k}$. The reciprocal mass of the k th layer is

$$\rho_k = 2^{-k} \sum_{\substack{m \leq N_k \\ m \text{ odd}}} \frac{1}{m} = 2^{-k} (H_{N_k} - \frac{1}{2} H_{N_k/2}) \rightarrow \frac{\log 2}{2}.$$

Thus $\sum_{a \in A_{\star}} 1/a$ diverges. For $\mathcal{D} = \{2\}$, the weighted contribution of the same layer is

$$\frac{2^k}{2^{2^k} - 1} \rho_k.$$

Since $H_N \leq 1 + \log N$, these terms are at most $(1 + 2^k \log 2) 2^{1-2^k}$, which is summable. The theorem therefore proves irrationality at every integer base for every infinite subset of this set, although its reciprocal sum diverges.

Proof. We first record the rational obstruction and three finite estimates used to defeat it. If $m = qa + r, 0 \leq r < a$, then

$$\frac{b^m - 1}{b^a - 1} - \frac{b^r - 1}{b^a - 1} = \sum_{j=1}^q b^{m-ja}.$$

Consequently

$$\Delta_{b,A}(m) = (b^m - 1)X_A(b) - J_{b,A}(m), \quad J_{b,A}(m) = \sum_{\substack{a \in A \\ a \leq m}} \sum_{j=1}^{\lfloor m/a \rfloor} b^{m-ja} \in \mathbb{Z}. \quad (2)$$

The integer sum is finite. For $m \geq 1$, moreover, $\Delta_{b,A}(m) > 0$: since A is infinite, some $a \in A$ exceeds m , and its summand is positive. At $m = 0$ the displacement is zero. Thus, if $X_A(b) = u/v$ with $v > 0$, equation (2) gives $v\Delta_{b,A}(m) \in \mathbb{Z}$, whence

$$\Delta_{b,A}(m) \geq \frac{1}{v} \quad (m \geq 1). \quad (3)$$

It is therefore enough to construct arbitrarily small positive displacements.

Put

$$\phi_a(m) = \frac{b^{m \bmod a}}{b^a - 1}, \quad d_a(m) = \phi_a(m) - \frac{1}{b^a - 1}.$$

Then $d_a(m) \geq 0$ and $\Delta_{b,A}(m) = \sum_{a \in A} d_a(m)$. For $Q, a, T \geq 1$, put $g = (Q, a)$. The residues $tQ \bmod a$ have period a/g and visit $0, g, \dots, a - g$ once in a period, so

$$\sum_{t=1}^{a/g} \phi_a(tQ) = \frac{1 + b^g + \dots + b^{a-g}}{b^a - 1} = \frac{1}{b^g - 1}.$$

Complete periods and at most one incomplete period therefore give

$$\frac{1}{T} \sum_{t=1}^T d_a(tQ) \leq \frac{g}{a(b^g - 1)} + \frac{1}{T(b^g - 1)}. \quad (4)$$

We shall also use that $n/(b^n - 1)$ decreases for positive integers n ; indeed

$$n(b^{n+1} - 1) - (n+1)(b^n - 1) = b^n(n(b-1) - 1) + 1 > 0.$$

The part of the support beyond the observation range has a uniform bound. If $Y = QT$ and $a > Y \geq tQ$, then $tQ \bmod a = tQ$ and $d_a(tQ) \leq 22^{tQ-a}$. Hence

$$\sum_{a > Y} d_a(tQ) \leq 22^{tQ-Y}, \quad \sum_{t=1}^T 2^{tQ-Y} = \sum_{j=0}^{T-1} 2^{-jQ} \leq 2,$$

and therefore

$$\frac{1}{T} \sum_{t=1}^T \sum_{\substack{a \in A \\ a > QT}} d_a(tQ) \leq \frac{4}{T}. \quad (5)$$

Finally, let $\alpha_a \geq 0$ with $\sum_a \alpha_a/a < \infty$. For $Q, M \geq 1$, finite rearrangement and a geometric sum give

$$\sum_{j=M}^{2M-1} \frac{1}{2^j} \sum_{a \leq Q2^j} \alpha_a \leq 2Q \sum_{a \geq 1} \frac{\alpha_a}{a}. \quad (6)$$

Indeed, for each fixed a the admissible indices satisfy $2^j \geq a/Q$ and

$$\sum_{\substack{M \leq j < 2M \\ a \leq Q2^j}} 2^{-j} \leq \frac{2Q}{a}.$$

Only finitely many terms occur on the left, so no infinite interchange is hidden here.

We now choose the modulus and the observation lengths. Fix $\varepsilon > 0$ and a lower bound N . Select a finite $F \subseteq A$ such that

$$\sum_{a \in A \setminus F} \frac{h(a)}{a(b^{h(a)} - 1)} < \varepsilon, \quad (7)$$

and put $L = \text{lcm}(F)$, with $L = 1$ if F is empty. Write $p_* = \max \mathcal{P}$ and $r = |\mathcal{P}|$. For a large integer $H \geq 2p_*$, set

$$Q_0(H) = \prod_{p \in \mathcal{P}} p^{\lfloor \log_p H \rfloor}, \quad Q = LQ_0(H), \quad G = \left\lfloor \frac{H}{p_*} \right\rfloor.$$

Then $Q \leq LH^r$, $G \geq 1$, and every $a \in F$ divides Q , so its $d_a(tQ)$ contribution vanishes.

Consider first $a \in A \setminus F$ with $a \leq QT$ and $h(a) \leq H$. Every prime-power factor of $h(a)$ occurs in $Q_0(H)$, hence $h(a) \mid Q_0(H)$ and $(Q, a) \geq h(a)$. Applying (4) and the monotonicity above, the total contribution of these a is at most

$$\varepsilon + \frac{1}{T} \sum_{\substack{a \in A, a \leq QT \\ h(a) \leq H}} \frac{1}{b^{h(a)} - 1}. \quad (8)$$

The first term is bounded using (7).

For $h(a) > H$, one has $(Q, a) \geq G$. If every $\mathcal{P}$ -prime-power component of $h(a)$ is at most H , then $h(a) \mid Q_0(H)$ and the gcd is larger than H . Otherwise some $p^{v_p(a)} > H$, and the gcd contains $p^{\lfloor \log_p H \rfloor} > H/p \geq H/p_*$. Thus (4), enlarged from the support to all $a \leq QT$, bounds the large- $h(a)$ part by

$$\frac{G(1 + \log(QT)) + Q}{b^G - 1}; \quad (9)$$

we used $\sum_{a \leq QT} a^{-1} \leq 1 + \log(QT)$ and $\#\{a \leq QT\}/T = Q$.

Let

$$M(Q, T) = \frac{1}{T} \sum_{t=1}^T \Delta_{b,A}(tQ).$$

Combining (5), (8), and (9) gives

$$M(Q, T) \leq \varepsilon + \frac{1}{T} \sum_{\substack{a \in A, a \leq QT \\ h(a) \leq H}} \frac{1}{b^{h(a)} - 1} + \frac{G(1 + \log(QT)) + Q}{b^G - 1} + \frac{4}{T}. \quad (10)$$

After the second average, the competing errors will be of orders Q/M and GM/b^G . With L fixed, Q grows polynomially in H and G grows linearly, so an intermediate exponential scale makes both errors tend to zero. Choose

$$M = \lfloor b^{G/2} \rfloor$$

and average (10) over $T = 2^j$, $M \leq j < 2M$. Apply (6) with $\alpha_a = \mathbf{1}_A(a)/(b^{h(a)} - 1)$. Its hypothesis holds because

$$\sum_a \frac{\alpha_a}{a} = \sum_{a \in A} \frac{1}{a(b^{h(a)} - 1)} \leq W_{b,\mathcal{D}}(A).$$

We obtain the explicit finite diagonal estimate

$$\frac{1}{M} \sum_{j=M}^{2M-1} M(Q, 2^j) \leq \varepsilon + \frac{2QW_{b,\mathcal{D}}(A)}{M} + \frac{G(1 + \log Q + 2M \log 2) + Q}{b^G - 1} + 4 \cdot 2^{-M}. \quad (11)$$

As $H \rightarrow \infty$,

$$Q \leq LH^r, \quad G = H/p_* + O(1), \quad M \asymp b^{G/2}.$$

Consequently $Q/M \rightarrow 0$, $GM/b^G \rightarrow 0$, $(G \log Q + Q)/b^G \rightarrow 0$, and $2^{-M} \rightarrow 0$. For sufficiently large H , we may also arrange $Q \geq N$ and make the right side of (11) smaller than 2ε . The left side is a finite average of positive numbers $\Delta_{b,A}(tQ)$. One of them therefore satisfies

$$0 < \Delta_{b,A}(tQ) < 2\varepsilon, \quad tQ \geq Q \geq N.$$

The construction is valid for every positive ε ; applying it with $\varepsilon/2$ proves the stated cofinal bound with $\Delta_{b,A}(m) < \varepsilon$. This contradicts (3) if $X_A(b)$ is rational. $\square$

Since $b^{h(a)} - 1 \geq 2^{h(a)} - 1$ for every $b \geq 2$, one has $W_{b,\mathcal{D}}(A) \leq W_{2,\mathcal{D}}(A)$. Thus the base-two weighted hypothesis gives the conclusion simultaneously at every integer base. The hypothesis is also hereditary under passage to infinite subsets.

Averaging antecedents. Both averaging proofs in this section end by choosing a term no larger than a finite mean. [Duverney and Tachiya](#) choose their index in the same way in their refinement of the Chowla–Erdős method: they average a local coefficient mass over an arithmetic progression built by the Chinese remainder theorem, and a term attaining the minimum supplies the bounds that control the tail beyond it (Section 2, (2.3)–(2.9), pp. 5–6 of the author preprint). [Kaneko, Suzuki and Tachiya](#) measure the average decay of scaled tails by $R_c(q, x, z)$, their (1.7) on p. 3 of arXiv v1. Under a counting condition on the support and a bound on R_c , most scaled tails in a range are small, and arbitrarily small nonzero scaled tails give irrationality (Lemmas 1 and 2, pp. 6–8). The divisor-incidence coefficients $\#\{a \in A : a \mid n\}$ are positive on every multiple of $\min A$. For any splitting into two sequences the two coefficient supports therefore cover a set of positive lower density, which is incompatible with the support-counting conditions [5, Theorem 1(iii), p. 3, and Theorem 3(iv), p. 5] requiring both counts to be $o(x_n/z_n)$

with $z_n \geq 1$. Their remote-tail quantity $R_c(q, x, z)$ sums only offsets $j \geq z$, whereas the full scaled tail

$$\sum_{r \geq 1} \frac{\#\{a \in A : a \mid m + r\}}{b^r} = \sum_{a \in A} \phi_a(m) = X_A(b) + \Delta_{b,A}(m)$$

is at least $X_A(b)$, which gives no lower bound for R_c . The proofs above average the displacement instead, along multiples of a modulus that freezes a finite part of A ; the weighted proof adds the dyadic average (6). [Cited]

Evidence boundary. The exact interface of Theorem 1.1, including its fixed-base conclusion and hereditary all-base clause, is kernel-checked as `divisibilityWeightedClaim`. Its public dependency closure was replayed under Lean 4.29.1. This does not cover all infinite supports, supply a quantitative irrationality measure, or assert theorem priority.

Formalisation coverage and remaining dependencies

Every theorem, lemma, proposition and corollary of this record has a Lean statement of the same assertion, with the same hypotheses, checked by the Lean kernel using only the axioms `propext`, `Classical.choice` and `Quot.sound`, in the development at revision 181078b6b009. The snapshot link in the paragraph above identifies an earlier published checkpoint; the coverage statement here is about the development revision just named. Statements the record takes from the literature are attributed where they appear, and the [Cited] tag marks a clause that invokes an external result. Theorem 6.20 is attributed in this way to the periodic-weight irrationality theorem of Luca and Tachiya [3, Theorem A, p. 139]. No formal proof of their theorem is given here; Theorem 6.20 and the support statements following from it each carry a separate formal proof of their own assertion.

Concordance of statements and Lean declarations. Each result of this record that has a kernel-checked Lean statement of the same assertion is listed below with the declarations that jointly state it. Each name links to its declaration at revision 181078b6b009. Where the Lean statement is stronger than the printed one and implies it by an immediate specialisation, the entry says so. An entry marked *compared* was also checked independently: a restatement of the same declarations against Mathlib alone, together with its proof, was verified by Comparator (`leanprover/comparator`) in a clean continuous integration environment, in the run named by its number. Comparator trusts the restated statement, so the correspondence between the printed result and that statement is the one this concordance records.

Theorem 1.1 (the Lean statement is stronger): `divisibilityWeightedClaim`, `weighted_displacement_cofinal_close_return`. *Compared*, run 35544127144.

Theorem 1.2 (the Lean statement is stronger): `strengthenedPositiveCoverClaim`. *Compared*, run 35544127144.

Theorem 1.3 (the Lean statement is stronger): `mixedSupportClaim`, `arbitraryWeightMixedSupport_allBase_hereditary`. *Compared*, runs 35544127144, 35674034595.

Theorem 1.7: `paper_achievement_geometry`. *Compared*, run 35643815458.

Theorem 1.8 (the Lean statement is stronger): `volume_supportedMersenneAchievementSet_dichotomy`,

volume_supportedMersenneAchievementSet_eq_zero_of_compl_infinite,
 perfect_supportedMersenneAchievementSet, supportedMersenneDigitValue_injective.
Compared, run 35643815458.

Theorem 1.10: paper_greedy_survival. *Compared, run 35674034595.*

Proposition 2.2 (the Lean statement is stronger): half_agrees_greedy,
 eq_halfGreedyPrefixSupport_of_critical_crossing,
 remainder_lt_gap_iff_eq_integerGreedyBits.

Lemma 2.4 (the Lean statement is stronger):

integerHalfCarry_eq_scaled_residual_add_tail, binaryCoeffTail_nonneg,
 binaryCoeffTail_supportCoeff_le_two_sqrt_add_four. *Compared, runs 35624228171, 35643815458, 35674034595.*

Proposition 2.5: half_of_cofinal_absolute_carry, greedy_half_of_cofinal_upper_carry.
Compared, run 35674034595.

Lemma 2.6: square_depth_witness. *Compared, run 35674034595.*

Proposition 2.7: six_membership_conditions. *Compared, run 35674034595.*

Proposition 2.8 (the Lean statement is stronger): affineBinaryOrbit_mod_twoPow_eq,
 balancedPulse_weighted_pair, balancedPulse_endpoint_fanout,
 balancedPulse_label_card_lower_bound, balancedPulse_no_autonomous_decoder,
 fixedPrecisionTropicalNoGo.

Proposition 2.9: skipSum_den_dvd_prod, weighted_denominator_budget. *Compared, run 35643815458.*

Theorem 2.11: paper_one_sidedness, computablePred_halfFatalCertificateCode,
 half_not_mem_iff_exists_halfFatalCertificateCode,
 half_mem_iff_forall_not_halfFatalCertificateCode,
 finite_depth_survival_not_sufficient, existsFatalHalfGap_iff_exists_certificate,
 existsFatalHalfGap_of_certificate, certificate_of_existsFatalHalfGap,
 scaledMersenneWeight_cast, certifiedWordValue_cast, certifiedTailBound_cast,
 mersenneTail_eq_sum_add.

Theorem 5.4: irrational_erdosSum_full_support. *Compared, run 35624228171.*

Theorem (Purely periodic support): irrational_erdosSupportSeries_periodic. *Compared, run 35624228171.*

Theorem (Eventually periodic support):

irrational_erdosSupportSeries_eventuallyPeriodic. *Compared, run 35624228171.*

Theorem (Residue-class support): irrational_residueClass_positive_support. *Compared, run 35624228171.*

Theorem (Odd support): irrational_erdosSupportSeries_odd. *Compared, run 35624228171.*

Theorem 5.9: paper_achievement_geometry, mersenne_topology_quantitative. *Compared, run 35643815458.*

Theorem 5.10: mem_mersenneAchievementSet_iff_greedy_survival,
 greedy_survives_of_mem_mersenneAchievementSet,
 mem_mersenneAchievementSet_of_greedy_survival. *Compared, run 35674034595.*

Theorem 5.14: paper_master_identity_floors. *Compared, run 35643815458.*

Theorem 5.15: paper_real_quotient_core, paper_real_quotient_margins,
 row_constant_eq_tail, mersenne_constant_decimal. *Compared, run 35643815458.*

Theorem 5.16: paper_dynamics, paper_rowLower_existsUnique,
 paper_rowUpper_existsUnique, paper_greedySupport_isRowLower,
 paper_upperSupport_isRowUpper, paper_isRowLower_unique, paper_isRowUpper_unique,
 paper_greedySupport_greedy_rule, paper_greedySupport_mem, paper_greedy_step,
 paper_consecutive_not_both_divisible, paper_rowPulse_eq_indicators,
 paper_rowQuotient_eq_weightSum, paper_rowWeight_eq_floor, paper_rowTarget_eq.

- Theorem 6.1: `fatal_absorbing`. *Compared, run 35674034595.*
- Theorem 6.2: `paper_seam_limit_unconditional`,
`tendsto_seamGreedyFiniteValue_greedyHalfTargetValue`,
`tendsto_seamGreedyNormalizedRemainder`, `eventually_seamSupport_agrees`,
`seamScaledRem_eq_tailGreedyRemainder`, `mem_seamGreedySupport_iff_scaled`,
`abs_supportValue_sub_le_mersenneTail`.
- Theorem 6.3: `half_mem_mersenneAchievementSet_of_skipped_twoChannelCap`,
`half_mem_mersenneAchievementSet_of_skipped_dyadicCap`. *Compared, run 35674034595.*
- Theorem 6.4: `half_mem_mersenneAchievementSet_of_secondChannelSeparation`,
`half_mem_mersenneAchievementSet_of_secondChannelSeparationRat_from_seven`.
Compared, run 35674034595.
- Theorem 6.5: `straddle_all_depths_iff_mem`, `straddle_limiting_support_inputs`. *Compared, run 35643815458.*
- Theorem 6.6: `half_mem_mersenneAchievementSet_of_largestSkipLateStepSocket`,
`largestSkipLateAt_fourteen`, `seamGreedy_terminal_false_iff_upperOrMiddle`.
- Theorem 6.7: `half_mem_mersenneAchievementSet_of_middleProducerCardEscape`,
`half_mem_mersenneAchievementSet_of_middleProducerRowEscape`.
- Theorem 6.8 (the Lean statement is stronger):
`middleProducer_allRight_forces_carry_lt_tail`,
`middleProducer_allRight_forces_rational_skip`,
`producerCarry_insert_seamBelowSupport_eq_middleCoordinate`,
`binaryCoeffTail_nonneg`, `binaryCoeffTail_supportCoeff_coe_finset_le_card`.
- Theorem 6.9 (the Lean statement is stronger): `twoSided`.
- Theorem 6.10 (the Lean statement is stronger):
`half_mem_mersenneAchievementSet_of_upperResetDyadicBandEscape`,
`seamUpperResetDyadicBandEscape_through_thirty`.
- Theorem 6.11: `mobiusCenteredHalfCarry_nonneg_of_supportSeries_lt_half`,
`integerHalfCarry_eq_scaled_residual_add_tail`. *Compared, run 35674034595.*
- Theorem 6.12: `greedy_mobiusCenteredHalfCarry_nonneg`,
`greedy_half_infinite_of_mobiusCenteredHalfCarry_sqrtBound`,
`infinite_support_half_of_mobiusCenteredHalfCarry_sqrtBound`. *Compared, run 35674034595.*
- Theorem 6.13: `half_mem_mersenneAchievementSet_of_skippedFullShellNonnegative`,
`skippedSeamAlignmentZero_iff_skippedFullShellNonnegative`,
`half_mem_mersenneAchievementSet_of_skippedSeamEscape`.
- Theorem 6.14: `irrational_erdosSum_full_support`. *Compared, run 35624228171.*
- Theorem 6.15 (the Lean statement is stronger):
`irrational_erdosSupportSeries_pairwise_coprime`. *Compared, run 35624228171.*
- Theorem 6.16: `irrational_coeff_series_of_weighted_coeff_block_certificates`.
Compared, run 35624228171.
- Theorem 6.17: `irrational_erdosSum_of_lcm_gap`. *Compared, run 35624228171.*
- Theorem 6.18 (the Lean statement is stronger): `irrational_erdosSum_factorial_support`,
`irrational_erdosSum_two_pow_support`. *Compared, run 35624228171.*
- Theorem 6.19 (the Lean statement is stronger):
`erdosSupportSeries_multiples_eq_pow_base_full_support`,
`irrational_erdosSupportSeries_multiples`. *Compared, run 35624228171.*
- Theorem 6.20 (the Lean statement is stronger): `irrational_erdosSupportSeries_periodic`.
Compared, run 35624228171.
- Theorem 6.21 (the Lean statement is stronger):
`irrational_erdosSupportSeries_eventuallyPeriodic`. *Compared, run 35624228171.*

Theorem 6.22 (the Lean statement is stronger):

`irrational_erdosSupportSeries_residueClass, irrational_erdosSupportSeries_odd.`
Compared, run 35624228171.

Theorem 6.23 (the Lean statement is stronger):

`irrational_or_bpow_mul_eq_intCast_intWeightedErdosSeries_periodic, irrational_intWeightedErdosSeries_periodic_of_coeff_nonneg_of_frequently_ne_zero, irrational_intWeightedErdosSeries_periodic_of_coeff_nonpos_of_frequently_ne_zero.`
Compared, run 35624228171.

Theorem 6.24: `paper_mersenne_channel_survival,`

`paper_mersenne_channel_survival_of_coprime_scale,`
`paper_channel_factor_gcd_eq_one, paper_channel_factors_pairwise_coprime,`
`paperB_eq_divInt_paperA, paperB_eq_baseMobiusShadow, paperA_eq_mobiusNumerator,`
`channelProduct_coprime_mobiusNumerator_of_one_le,`
`upperHalfChannel_survivorProduct_dvd_den_of_one_le.` *Compared, run 35643815458.*

Theorem 6.25: `lcmHeight_scaledMobiusShadow_den_lower_bound,`

`upperHalfMersenneProduct_lower_bound, lcmHeight_scaledMobiusShadow_den_exact.`
Compared, run 35643815458.

Theorem 6.28 (the Lean statement is stronger): `tsum_moebius_div_two_pow_sub_one_eq_half.`

Corollary 6.29 (the Lean statement is stronger):

`tsum_negativeMobius_eq_half_add_positiveMobiusTail, half_lt_tsum_negativeMobius.`
Compared, run 35643815458.

Theorem 6.31 (the Lean statement is stronger):

`half_mem_mersenneAchievementSet_iff_greedySkippedSupport_infinite,`
`irrational_erdosBorweinMersenneConstant.` *Compared, runs 35643815458, 35674034595.*

Theorem 6.32 (the Lean statement is stronger):

`half_mem_mersenneAchievementSet_iff_not_seamGreedyEventuallyRight,`
`half_mem_mersenneAchievementSet_iff_unboundedTerminalFalse,`
`half_mem_mersenneAchievementSet_iff_unboundedUpperOrMiddle,`
`half_mem_mersenneAchievementSet_iff_cofinalTerminalFalse,`
`half_mem_mersenneAchievementSet_iff_exists_unboundedSkippedRanksAlong,`
`unboundedTerminalFalse_iff_greedyMersenneSkippedSupport_infinite,`
`half_mem_mersenneAchievementSet_iff_no_lastHalfGreedySkip.`

Lemma 6.33 (the Lean statement is stronger):

`prefix_add_mersenneTail_lt_half_of_eventually_right,`
`half_lt_upper_competitor_of_eventually_right.`

Lemma 6.34 (the Lean statement is stronger): `mersenneTail_eq_weight_add,`

`halfTwoChannelCap_lt_mersenneTail, mersenneTail_le_two_mul_weight,`
`two_mul_mersenneWeight_succ_lt, mersenneTail_lt_weight.` *Compared, run 35643815458.*

Theorem 6.35 (the Lean statement is stronger):

`mem_mersenneAchievementSet_iff_greedy_survival.` *Compared, run 35674034595.*

Lemma 6.36 (the Lean statement is stronger): `isStraddlePrefix_step_trichotomy,`

`half_step_forced.`

Lemma 6.37: `fatal_gap_excludes_every_representation, fatal_gap_endpoint_bounds,`

`depth_prefix_interval_disjoint, fatal_gap_within_prefix_interval.` *Compared, run 35643815458.*

Lemma 6.38 (the Lean statement is stronger):

`positiveMersenneSupportValue_coe_finset_ne_half,`
`half_ne_coe_finset_add_mersenneTail.` *Compared, run 35643815458.*

Lemma 6.39 (the Lean statement is stronger): `half_agrees_greedy.`

Theorem 6.40 (the Lean statement is stronger): `isLastHalfGreedySkip_iff_skip_and_fatal`, `half_mem_iff_every_actual_skip_survives`. *Compared*, run 35674034595.

Lemma 6.41 (the Lean statement is stronger): `seamGreedy_terminal_false_iff_upperOrMiddle`.

Lemma 6.42: `largest_false_rank_algebra`, `paper_largest_false_rank_algebra`.

Theorem 6.43: `exists_criticalDyadicBandIndex`, `dyadicBandEscape_iff_exists_critical`, `seamUpperResetCriticalBandEscape_iff`.

Theorem 6.44: `finalMiddleCell_neg_three_not_last`, `mobiusCenteredHalfCarry_add_two`.

Lemma 6.45: `halfGreedy_skipped_endpoint_trichotomy`.

Lemma 6.46: `paper_reverse_carry_word`, `reverse_carry_word_common_bound_sharp`.
Compared, run 35643815458.

Lemma 6.47 (the Lean statement is stronger): `linearDescender_eq_smul_eval`, `relationInvariantLinearChannels_det_eq_zero`. *Compared*, runs 35624228171, 35643815458.

Theorem 6.48: `paper_two_thirds_band`. *Compared*, run 35643815458.

Theorem 6.49: `paper_sharp_fatal_gap`. *Compared*, run 35643815458.

Lemma 6.50 (the Lean statement is stronger): `mersenneGap_pos`, `summable_mersenneGap_succ`, `mersenneGap_tail_le`, `tendsto_mersenneGap_tail_zero`. *Compared*, run 35643815458.

Lemma 6.51 (the Lean statement is stronger): `supportCoeff_insert_eq_add_indicator`, `supportCoeff_extend_true_eq_false_add_one_at_double`. *Compared*, runs 35624228171, 35643815458.

Theorem 6.52 (the Lean statement is stronger): `binaryCoeffSeries_rational_iff_exists_temperedBinaryOrbit`, `temperedBinaryOrbit_eq_scaledTail`. *Compared*, run 35643815458.

Lemma 6.53 (the Lean statement is stronger): `irrational_erdosSupportSeries_of_tail`, `irrational_erdosSupportSeries_tail_of_irrational`. *Compared*, run 35624228171.

Lemma 6.54 (the Lean statement is stronger): `divInt_le_nextDyadic_iff_excess_nonpos`, `greedyHalf_mem_nextMersenneDyadicSliver_iff_excess`. *Compared*, runs 35624228171, 35674034595.

Lemma 6.55 (the Lean statement is stronger): `dyadicResidual_denominator_sandwich`.
Compared, run 35624228171.

Lemma 6.56 (the Lean statement is stronger): `divisor_dvd_divInt_den`, `survivingDivisor_dvd_scaled_divInt_den`. *Compared*, run 35643815458.

Lemma 6.57 (the Lean statement is stronger): `primePowerLayer_comm`, `mixedPrimePowerLayerTwo_supportCoeffInt`. *Compared*, runs 35624228171, 35643815458.

Proposition 6.59: `paper_achievement_set_topology`. *Compared*, run 35643815458.

Theorem 6.60 (the Lean statement is stronger): `half_mem_mersenneAchievementSet_iff_no_existsFatalHalfGap`, `half_mem_mersenneAchievementSet_or_exists_fatal_gap`. *Compared*, run 35643815458.

Theorem 6.62 (the Lean statement is stronger): `prefixChoice_maximal`, `nextRemainder_trichotomy`.

Theorem 6.67: `CofinalExactLocalMersenneHalfRows`, `abs_exactLocalMersenneRowValue_sub_half_le`, `half_mem_mersenneAchievementSet_of_cofinalExactLocalRows`, `finite_boolSupport_ne_half`.

Theorem 6.68: `CofinalPositiveHalfGreedySkips`, `exactLocalMersenneHalfRow_of_positiveHalfGreedySkip`, `cofinalExactLocalMersenneHalfRows_of_positiveHalfGreedySkips`, `half_mem_mersenneAchievementSet_of_positiveHalfGreedySkips`, `cofinalPositiveHalfGreedySkips_iff_half_mem`.

- Theorem 6.69: `paper_compatible_bit_stable`, `paper_compatible_finite_row_conditions`, `paper_compatible_first_condition_gives_nonneg`, `paper_compatible_rows_agree_with_limit`.
- Theorem 6.70: `SkippedCoreCriticalQuotientSupply`, `cofinalExactLocalMersenneHalfRows_of_criticalQuotientSupply`, `half_mem_mersenneAchievementSet_of_criticalQuotientSupply`, `finite_boolSupport_ne_half`.
- Theorem 6.71: `ProtectedExactLocalMersenneRow`, `exists_laterProtectedExactLocalMersenneRow`, `cofinalExactLocalMersenneHalfRows_of_criticalQuotientSupply`.
- Theorem 6.72: `HalfGreedySkippedCriticalQuotientSupply`, `skippedCoreCriticalQuotientSupply_iff_halfGreedySkipped`.
- Theorem 6.73: `halfGreedy_precriticalSuffix_lt_of_next_skip`, `halfGreedySkippedCriticalQuotientSupply_of_precriticalSuffix`, `halfGreedySkippedPrecriticalSuffixSupply_iff_preTake`. *Compared, run 35674034595.*
- Theorem 6.74: `halfGreedy_precriticalSuffix_lt_of_future_skip_after_takenBlock`, `precriticalCrossingTax_of_futureThreshold`, `sub_two_le_two_pow_sub_four`. *Compared, runs 35643815458, 35674034595.*
- Theorem 6.75: `exists_exactRowStrictUpperFill_of_skippedCoreSharpCapacity`. *Compared, run 35643815458.*
- Theorem 6.76: `exactLocalMersenneHalfRow_two_mul_sub_two_of_skippedCoreSharpCapacity`. *Compared, run 35643815458.*
- Proposition 6.77: `localMersenneQuotient_eq_two_pow_sub_of_half_lt`, `exists_boolean_word_of_lt_two_pow`. *Compared, run 35643815458.*
- Theorem 6.78: `paper_exact_row_from_skipped_prefix`. *Compared, run 35643815458.*
- Theorem 6.79: `paper_critical_dyadic_band_index_unique`, `paper_critical_dyadic_boundary_is_smallest`, `paper_critical_dyadic_band_index_eq_top_of_le_two`, `paper_dyadic_band_escape_iff_single_test`. *Compared, run 35643815458.*
- Proposition 6.80: `seamUpperBranch_remainder_add_resetCharge_eq`, `seamUpperResetCharge_le`.
- Theorem 6.82: `HalfGreedySkippedSeamAlignmentZero`, `HalfGreedySkippedFullShellNonnegative`, `skipped_fullShell_neg_iff_alignment_and_seamRemainder_pos`, `greedyHalfFrozenMargin_fullShell_eq_neg_seamRemainder_of_alignment`, `skippedSeamAlignmentZero_iff_skippedFullShellNonnegative`.
- Theorem 6.83: `paper_seam_escape_forces_remainder_band`, `paper_seam_escape_implies_full_shell_nonnegative`, `paper_seam_escape_implies_half_membership`.
- Theorem 6.84: `greedy_mobiusCenteredHalfCarry_nonneg`, `greedy_half_infinite_of_mobiusCenteredHalfCarry_upperBound`. *Compared, run 35674034595.*
- Theorem 6.85 (the Lean statement is stronger): `paper_terminal_strip_forces_half_membership`, `paper_integerHalfCarry_eq_two_pow_sub_localPrefixQuotient`, `paper_exact_row_integerHalfCarry_eq_one`, `paper_terminal_strip_witness_six`, `paper_both_cofinal_statements_iff_half_membership`. *Compared, runs 35643815458, 35674034595.*

- Theorem 6.86: `GreedyHalfCarryCofinalStripReturn`,
`greedy_half_infinite_of_cofinalStripReturn`,
`halfGreedy_precriticalSuffix_lt_iff_futureSkipCoverage`.
- Theorem 6.87 (the Lean statement is stronger):
`paper_halfGreedyPrefixSupport_eq_greedy_inter_Icc`,
`paper_eventual_nonnegative_margin_equivalence`,
`paper_frozen_margin_normalised_value`, `paper_frozen_margin_normalised_monotone`,
`paper_eta_eq_coeffTail_sub_carry`, `paper_frozen_margin_normalised_tendsto`,
`paper_frozen_margin_limit_pos_iff`, `paper_greedyHalfRemainder_ne_dyadicCap`,
`paper_coeffTail_le_index_add_two`, `paper_effective_horizon_test`,
`paper_eta_hasRationalValue`. *Compared*, runs 35643815458, 35674034595.
- Theorem 6.88: `half_mem_mersenneAchievementSet_iff_greedySkippedSupport_infinite`,
`mem_mersenneAchievementSet_of_greedySkippedSupport_infinite`. *Compared*, run 35674034595.
- Theorem 6.89: `paper_next_floor_quotient`, `paper_next_floor_quotient_no_fixed_point`,
`paper_next_quotient_sum`, `paper_floor_quotient_geometric_sum`. *Compared*, run 35643815458.
- Theorem 6.90: `localPrefixQuotient_succ`. *Compared*, run 35643815458.
- Theorem 6.91: `paper_signed_endpoint_recurrence`, `paper_endpoint_term_counts_divisors`,
`paper_signed_endpoint_defect_succ`, `paper_repair_integer_eq_endpoint_defect`.
Compared, runs 35624228171, 35643815458.
- Theorem 6.93: `localBinarySuffix_two_mul_sub_two_lt_criticalCapacity_iff`. *Compared*,
run 35643815458.
- Theorem 6.94: `paper_unconditional_bound_one_extra_bit`,
`paper_sharper_additive_estimate`, `paper_capacity_band_exclusion`. *Compared*, run 35643815458.
- Theorem 6.95: `localBinarySuffix_two_mul_sub_one_lt_upperWindow_of_exact_below`,
`exists_exactRowStrictUpperExtension_two_mul_sub_one_of_exact_below`. *Compared*,
run 35643815458.
- Proposition 6.96: `two_mem_of_exact_localMersenneQuotient`. *Compared*, run 35643815458.
- Proposition 6.98: `abs_localMersennePrefixValue_sub_half_le`. *Compared*, run 35643815458.
- Proposition 6.99: `paper_finite_sum_inv_odd_den_odd`, `paper_finite_sum_inv_odd_ne_half`,
`paper_finiteErdosSum_den_odd`, `paper_finite_support_series_ne_half`,
`paper_half_representing_support_is_infinite`. *Compared*, runs 35624228171, 35643815458.
- Theorem 6.100: `remainder_lt_gap_iff_eq_integerGreedyBits`. *Compared*, run 35674034595.
- Proposition 6.101: `localMersenneWeightsFrom_gapDominates`,
`localMersenneWeights_gapDominates_even`, `localMersenneWeights_gapDominates_odd`.
Compared, run 35682858662.
- Proposition 6.102: `localMersenneHalfTarget_lower_word_eq_greedy_and_remainder_eq`.
Compared, run 35682858662.
- Theorem 6.103: `localMersenneGeometricQuotient`, `localMersenneQuotient_eq_geometric`,
`localBinarySuffix_two_mul_sub_two_lt_criticalCapacity_iff_geometric`,
`localBinarySuffix_two_mul_sub_two_lt_criticalCapacity_iff_geometricCore`.
- Theorem 6.104 (the Lean statement is stronger): `binaryCoeffTail_le`,
`binaryCoeffTail_div_pow_tendsto_zero`. *Compared*, run 35643815458.
- Theorem 6.105 (the Lean statement is stronger): `binaryCoeffSeries`, `IsTemperedBinaryOrbit`,
`not_irrational_binaryCoeffSeries_iff_exists_temperedBinaryOrbit`.

Theorem 6.106 (the Lean statement is stronger): `moebius_mul_supportCoeffAF`, `mobius_supportCoeff_eq_one_iff`, `mobius_supportCoeff_boolean`. Compared, run 35624228171.

Theorem 6.107 (the Lean statement is stronger): `erdosSupportSeries_two_eq_binaryCoeffSeries`, `supportCoeff_le_card_divisors`, `supportCoeff_le_self`, `erdosSupportSeries_rational_iff_exists_temperedCarry`. Compared, run 35624228171.

Theorem 6.108 (the Lean statement is stronger): `one_lt_reciprocalMass_of_dyadic_support_fraction_of_two_pos_mem`, `dyadic_support_fraction_reciprocalMass_diverges_or_gt_one`. Compared, runs 35544127144, 35624228171.

Theorem 6.109 (the Lean statement is stronger): `one_add_mul_card_le_two_mul_shifted_state`, `shifted_state_unbounded_of_infinite_support`, `exists_unbounded_shifted_odd_tail_nat_state_of_support_fraction`. Compared, runs 35544127144, 35624228171.

Theorem 6.110: `paper_zero_run_le_eps_logb`, `paper_zero_run_le_of_mem`. Compared, run 35643815458.

Proposition 6.111 (the Lean statement is stronger): `divisorSubpowerConst`, `card_divisors_pow_le_divisorSubpowerConst_pow_mul`, `card_divisors_le_divisorSubpowerConst_mul_rpow`.

Proposition 6.112: `primePowerLayer_comm`, `mixedPrimePowerLayerTwo_supportCoeffInt`, `mixedPrimePowerLayerTwo_twelve_fixture`. Compared, runs 35624228171, 35643815458.

Theorem 6.114: `localMersenneCrossingRanks`, `exists_first_localMersenne_crossing`.

Theorem 6.115: `integerHalfCarry_eq_scaled_residual_add_tail`, `mobiusCenteredHalfCarry_nonneg_of_supportSeries_lt_half`. Compared, run 35674034595.

Theorem 6.116: `halfGreedy_skipped_endpoint_trichotomy`.

Theorem 6.118: `mersenneGap_le`, `summable_mersenneGap_shift`, `mersenneGap_tail_le`, `tendsto_mersenneGap_tail_zero`. Compared, run 35643815458.

Theorem 6.120: `paper_dyadic_skip_test_iff`, `paper_sharp_skip_safe_lb3`, `paper_sharp_skip_safe_actual_tail`, `paper_three_channel_margin_identity`, `paper_sharp_weaker_than_dyadic`, `paper_sharp_strictly_weaker_realizable`, `paper_sharp_gives_three_over_three_t_sub_one`, `paper_two_channels_insufficient`, `paper_exact_mass_threshold`. Compared, run 35643815458.

Proposition 6.121: `unitNumerator_skipSafe`, `two_le_of_fatal`, `three_le_of_fatal_of_odd`, `unitNumerator_skipSafe_actualTail`. Compared, run 35682858662.

Theorem 6.122: `paper_two_sided_dyadic_bound`.

Proposition 6.123: `paper_finite_band_check_thirteen_to_thirty`, `paper_successor_remainders_fourteen_through_thirtyone`, `paper_universal_band_condition_unfolded`, `paper_universal_band_condition_would_close_half`.

Theorem 6.124: `paper_perturbed_order_preservation`, `paper_perturbed_prefixChoice_maximal`, `paper_perturbed_nextRemainder_three_branches`, `paper_perturbed_separation_global_maximality`, `paper_perturbed_two_stage_is_global_maximum`, `paper_perturbed_weak_cap_counterexample`, `prefixChoice_eq_below`.

Theorem 6.125: `ReverseCarryWord`, `overlappingReverseCarryWords_carryDifference_eq_twoPow_mul_odd`,

overlappingReverseCarryWords_twoPow_le_realBound,
overlappingMidpointReturns_twoPow_le_realBound.

Proposition 6.126 (the Lean statement is stronger): supportCoeff_insert_divisor,
supportCoeff_extend_true_eq_false_add_one_at_double,
supportCoeff_boundaryPair_unitDrop_at_double. *Compared, run 35643815458.*

Theorem 6.127: paper_bounded_double_or_recycle_countermodel,
paper_exists_seeded_bounded_double_or_recycle_model. *Compared, run 35643815458.*

Theorem 6.128: paper_exact_row_double_or_recycle, paper_finite_row_value_ne_half,
paper_exact_row_example_six_and_eleven. *Compared, run 35643815458.*

Proposition 6.129: paper_fractional_mass_bound_not_necessary,
paper_fractional_mass_bound_suffices_for_sharp_capacity. *Compared, run 35643815458.*

Theorem 6.131: paper_skipped_core_recycling_witness_bounded,
paper_returning_endpoint_may_fail_to_grow,
paper_exists_seeded_bounded_double_or_recycle_model,
paper_protected_row_crossing_beyond_cutoff, paper_protected_row_endpoint_growth.
Compared, runs 35643815458, 35674034595.

Theorem 6.133: paper_critical_crossing_support_is_greedy_prefix. *Compared, run 35674034595.*

Theorem 6.136: paper_relationInvariant_channels_rank_le_one,
paper_relationInvariant_channels_det_eq_zero. *Compared, run 35643815458.*

Theorem 6.139: paper_final_middle_cell_ne_neg_three,
paper_final_middle_cell_remaining_negative_values,
paper_mobiusCenteredHalfCarry_add_two,
paper_cofiniteRightTail_ne_zero_centeredEndpoint.

Proposition (The unsafe middle range is exactly three integers):
paper_unsafe_middle_range_is_three_integers.

Proposition (Exact Lebesgue-measure dichotomy):
paper_volume_supportedMersenneAchievementSet_dichotomy. *Compared, run 35643815458.*

Proposition 8.1: paper_one_orbit_stability, approx_orbit_induction,
tailGreedyRemainder_mersenne. *Compared, run 35674034595.*

Lemma 8.2: paper_no_ties, paper_no_ties_take, paper_no_ties_skip,
irrational_mersenneTail. *Compared, runs 35643815458, 35674034595.*

Lemma 8.4 (the Lean statement is stronger):
paper_forced_greedy_unique_support_and_criterion, paper_forced_greedy_low_ranks.
Compared, runs 35643815458, 35674034595.

Lemma 8.5: paper_parity_excludes_finite_support. *Compared, run 35643815458.*

Theorem 8.6: paper_witness_exclusion. *Compared, run 35643815458.*

Corollary 8.7: paper_half_lcm_horizon. *Compared, run 35643815458.*

Lemma 8.8: paper_mod_twelve_filter. *Compared, run 35643815458.*

Theorem 8.9: paper_rung_finite_decision. *Compared, run 35674034595.*

Theorem (A: a sufficient lower bound at every reset): paper_theoremA_half_membership,
paperResetSqrtEscape_iff_square, paper_theoremA_crossing_bound_square_le,
paper_theoremA_right_branch_forces_small_deviation,
half_mem_mersenneAchievementSet_of_resetSqrtEscape.

Theorem (C ; the one-sided finite decision boundary) (the Lean statement is stronger):
paper_one_sided_finite_decision_boundary. *Compared, run 35674034595.*

Proposition 10.4: paper_centred_completion_of_fixed_precision. *Compared, run 35643815458.*

Proposition 10.5: `paper_upper_branch_needs_no_exceptional_cell`.

Theorem 10.6: `paper_final_middle_cell_at_least_neg_two`.

Corollary 10.7: `paper_final_middle_cell_remaining_cells`.

Proposition 10.9 (the Lean statement is stronger):

`paper_pulse_family_no_autonomous_decoder`,
`paper_balanced_pulse_fanout_is_radius_succ`,
`paper_balanced_pulse_fanout_unbounded_corrected`,
`paper_pulse_family_finite_state_card`, `balancedPulseCoeff_injective`,
`balancedPulse_weighted_pair`. Compared, runs 35624228171, 35643815458.

Proposition 10.10: `paper_mobius_support_overshoots_half`,

`paper_first_positiveMobius_tail_term`. Compared, run 35643815458.

Proposition 10.11: `paper_finite_support_and_onesided_certificate`. Compared, run 35674034595.

Proposition 10.12: `paper_carry_survivor_extinction`, `paper_periodLcm_is_prefix_lcm`. Compared, run 35674034595.

Lemma 10.13: `paper_scalar_localization`, `paper_scalar_localization_size_bound`,
`paper_scalar_localization_zero_degenerate`. Compared, run 35643815458.

Corollary 10.14: `paper_mersenne_height`. Compared, run 35643815458.

Proposition 10.16: `paper_critical_band_index_collapse`.

Lemma 11.2 (the Lean statement is stronger): `theta_eq_tsum_divisorResidue`,
`theta_eq_tsum_geometricForm`, `theta_eq_divisorResidueSum`, `theta_eq_geometricForm`,
`residue_condition_iff`, `card_divisors_sub_one`, `iLeast_mem_Icc`, `dvd_add_iLeast`,
`not_dvd_of_lt_iLeast`, `iLeast_congr`, `mCount_eq_zero_of_lt_iLeast`,
`geometric_term_eq_zero_of_lt_iLeast`, `theta_eq_Psi`, `Psi_eq_of_residues_eq`,
`theta_one_one`, `theta_one_three`, `residue_cutoff_reading_fails`. Compared, run 35643815458.

Lemma 11.3: `paper_square_depth_terminal_bound`, `square_depth_witness`. Compared, run 35674034595.

Proposition 11.4: `paper_squarefree_support_engine_ceiling`. Compared, run 35624228171.

Proposition 12.3: `paper_cpqs_equiv`. Compared, run 35674034595.

Proposition 12.4: `paper_terminal_strip_equiv`, `paper_relaxed_constant_six_every_depth`. Compared, run 35674034595.

Theorem 13.1: `arithmetic_logarithmic_counterexample`,
`no_absolute_dyadic_kappaOne_constant`. Compared, run 35624228171.

Corollary 13.2: `sup_condExceedProb_le_paperCoverCost`,
`exists_support_paperCoverCost_ge`, `no_absolute_paperCoverCost_kappaOne_constant`.

Proposition 13.3 (the Lean statement is stronger): `logarithmic_initial_interval`. Compared, run 35624228171.

1.3 A finite estimate on common observation scales

For $1 < B \leq 2$, positive integers L, d, M , and an integer $R \geq 0$, put

$$w_{B,d}(n) = \frac{B^{n \bmod d}}{B^d - 1}, \quad \mathcal{D}_{L;R,M} F = \frac{1}{M} \sum_{j=R}^{R+M-1} \frac{1}{2^j} \sum_{m=1}^{2^j} F(Lm).$$

The finite estimate

$$\mathcal{D}_{L;R,M} w_{B,d} \leq \frac{1 + 4L/M}{d(B - 1)} \quad (12)$$

equivalently bounds $(B-1)\mathcal{D}_{L,R,M}w_{B,d}$ by $(1+4L/M)/d$. This normalized bound is uniform as $B \downarrow 1$; the unnormalized right side grows like $(B-1)^{-1}$. The ratio L/M measures the cost of incomplete modular periods.

To prove it, fix $T = 2^j$ and average over $1 \leq m \leq T$. When $d \leq LT$, the complete orbit has length d/g , with $g = (L, d)$, and total weight $1/(B^g - 1)$. Complete cycles and one remaining piece give

$$\frac{1}{T} \sum_{m=1}^T w_{B,d}(Lm) \leq \frac{g}{d(B^g - 1)} + \frac{1}{T(B^g - 1)} \leq \frac{1}{d(B-1)} + \frac{1}{T(B-1)}.$$

Across dyadic lengths satisfying $d \leq L2^j$, the reciprocal-length errors sum to at most $2L/[d(B-1)]$. When $d > 2LT$, there is no wrap and $2Lm \leq d-1$; hence $\sum_{i=0}^{d-1} B^i \geq dB^{(d-1)/2} \geq dB^{Lm}$. Each atom is then at most $1/[d(B-1)]$. Finally, at most one dyadic length satisfies $LT < d \leq 2LT$. For that length the geometric sum and convexity give

$$\frac{1}{T} \sum_{m=1}^T w_{B,d}(Lm) = \frac{B^L}{T(B^L - 1)} \frac{B^{LT} - 1}{B^d - 1} \leq \frac{LB^L}{d(B^L - 1)} \leq \frac{2L}{d(B-1)}.$$

Summing the main terms and these two error budgets proves (12). The release snapshot contains a proof body for the full estimate in [the finite dyadic averaging estimate](#). Its complete-cycle and no-wrap scalar ingredients are in [the complete-period and no-wrap estimates](#). These are source locators, not a fresh compilation or axiom-audit receipt; the argument displayed here is an ordinary proof. Nonnegative interchange permits summation against any coefficients c_d with $\sum_d c_d/d < \infty$.

1.4 Positive divisor majorants

For a finite set $F \subseteq \mathbb{N}_{>0}$, put $f_F(n) = \#\{a \in F : a \mid n\}$. The next theorem bounds a fractional power of this count by a nonnegative sum over divisors.

Theorem 1.2 (A summable family of divisor majorants). *For each $j \geq 1$, let $F_j \subseteq \mathbb{N}_{>0}$ be finite, let $0 < \alpha_j \leq 1$, and let $c_{j,d} \geq 0$ satisfy*

$$f_{F_j}(n)^{\alpha_j} \leq \sum_{d \mid n} c_{j,d} \quad (n \geq 1).$$

Set $C_j = \sum_{d \geq 1} c_{j,d}/d$. If

$$\sum_{j \geq 1} \frac{C_j 2^{j\alpha_j}}{2^{\alpha_j} - 1} < \infty, \tag{13}$$

then $X_A(b)$ is irrational for every infinite $A \subseteq \bigcup_j F_j$ and every integer $b \geq 2$.

Lean: [strengthenedPositiveCoverClaim](#).

Every finite set admits such a majorant: take $\alpha_j = 1$ and $c_{j,d} = \mathbf{1}_{F_j}(d)$. The restrictive part is summability over the whole family, not the existence of a majorant for an individual set. For example, $F_j = \{4^j\}$ with these choices contributes 2^{-j} to (13), so the theorem applies. Full support cannot admit a cover satisfying the theorem, because the proof gives arbitrarily small positive displacements and the full-support displacement is at least $1/2$.

Proof. Write $B_j = 2^{\alpha_j}$ and

$$U_j(N) = \sum_{r \geq 1} 2^{-r} f_{F_j}(N+r), \quad V_j(N) = \sum_{d \geq 1} c_{j,d} w_{B_j,d}(N).$$

The choice $B_j = 2^{\alpha_j}$ matches the decay after taking a fractional power. Subadditivity and the divisor majorant give

$$U_j(N)^{\alpha_j} \leq \sum_{r \geq 1} B_j^{-r} f_{F_j}(N+r)^{\alpha_j} \leq \sum_{d \geq 1} c_{j,d} \sum_{\substack{r \geq 1 \\ d|N+r}} B_j^{-r} = V_j(N).$$

The last identity sums a geometric progression in each residue class. The geometric-series identity also gives

$$\Delta_{2,F_j}(N) = U_j(N) - X_{F_j}(2) \leq U_j(N).$$

Consequently, once the first J covering sets have zero displacement, nonnegativity bounds the displacement of their union by $\sum_{j>J} U_j(N)$.

Fix $\varepsilon > 0$, set $t_j = \varepsilon 2^{-j}$, and choose J with

$$K_J := \sum_{j>J} \frac{C_j t_j^{-\alpha_j}}{B_j - 1} < \frac{1}{4}.$$

This is possible since $\varepsilon^{-\alpha_j} \leq \max(1, \varepsilon^{-1})$. Choose L divisible by every member of the first J covering sets. Equation (12) bounds the finite mean of $S_J(N) := \sum_{j>J} t_j^{-\alpha_j} V_j(N)$ by $(1 + 4L/M)K_J < 1/2$ whenever $M \geq 4L$. One sample therefore has $S_J(N) < 1$, forcing $U_j(N) < t_j$ for every $j > J$. Every exponent in the first J finite sets divides L , so those terms have zero displacement. Consequently,

$$0 < \Delta_{2,A}(N) \leq \sum_{j>J} U_j(N) \leq \varepsilon.$$

The fixed rational lattice excludes rationality at base two. For $0 \leq r < d$, the atom $(b^r - 1)/(b^d - 1)$ is nonincreasing in $b > 1$. For $r > 0$, cancellation of $b - 1$ writes it as $A(b)/(A(b) + C(b))$, with $A(b) = \sum_{i<r} b^i$ and $C(b) = \sum_{r \leq j < d} b^j$. Its derivative has the sign of $A'C - AC' = \sum_{i<r \leq j < d} (i-j)b^{i+j-1} < 0$. Thus $\Delta_{b,A}(N) \leq \Delta_{2,A}(N)$ for real $b \geq 2$; for integer b the rational-lattice argument gives irrationality. This ordinary argument is sharper than the factor-two bound in [the scalar comparison between bases](#). This locator is not a fresh Lean verification. Any prescribed positive divisor can be included in L , so the witnesses can also be required to be arbitrarily large. $\square$

The same proof permits any positive weights η_j with $\sum_j \eta_j = 1$: replace $2^{j\alpha_j}$ in (13) by $\eta_j^{-\alpha_j}$ and take $t_j = \varepsilon \eta_j$.

1.5 Combining the two support criteria

Separate small-displacement witnesses need not occur at the same index. For example, a sequence small only at even indices and one small only at odd indices need never have a small sum. The useful feature of (12) is its uniformity in the moving modulus: the positive-cover argument can use the exact observation window selected by the weighted proof.

Theorem 1.3 (mixed weighted and cover supports). *Let $E, V \subseteq \mathbb{N}_{>0}$. Suppose E has finite weighted mass (1) at $b = 2$ for a finite nonempty prime set $\mathcal{P}$. Suppose also that $V \subseteq \bigcup_j F_j$ for sets and majorants satisfying Theorem 1.2, with (13) or its positive-weight variant. Then $X_A(b)$ is irrational for every infinite $A \subseteq E \cup V$ and every integer $b \geq 2$.*

Lean: `mixedSupportClaim, arbitraryWeightMixedSupport_allBase_hereditary`.

Proof. Fix $\varepsilon > 0$ and an integer $N_0 \geq 1$, and put $\rho = \varepsilon/3$. Use the cover notation B_j, U_j, V_j above, with $\eta_j = 2^{-j}$ in the case (13). Set $t_j = \rho\eta_j$ and choose J so that

$$K_J = \sum_{j>J} \frac{C_j t_j^{-\alpha_j}}{B_j - 1} < \frac{1}{16}.$$

Choose a finite $F \subseteq E$ whose complementary weighted mass is $\kappa < \rho/16$. Let L be a positive common multiple of N_0 , all members of F , and all members of the first J covering sets. For large H , choose the modulus and averaging length from the base-two weighted proof:

$$Q = L \prod_{p \in \mathcal{P}} p^{\lfloor \log_p H \rfloor}, \quad G = \lfloor H / \max \mathcal{P} \rfloor, \quad M = \lfloor 2^{G/2} \rfloor.$$

The finite estimate (11), with complementary weighted mass κ , remains valid for this L : its proof requires only that every member of F divide L . It also applies to finite or empty E , since positivity was used only after the averaging estimate. Thus

$$\mathcal{D}_{Q;M,M}(\Delta_{2,E}/\rho) < \frac{1}{8}$$

for sufficiently large H . For the same finite distribution, (12) and nonnegative interchange give

$$\mathcal{D}_{Q;M,M} S_J \leq (1 + 4Q/M) K_J < \frac{1}{8}, \quad S_J = \sum_{j>J} t_j^{-\alpha_j} V_j,$$

because $Q/M \rightarrow 0$. Hence some sample $N = Qm$ satisfies $\Delta_{2,E}(N)/\rho + S_J(N) < 1$. At this index the weighted displacement is below ρ and $U_j(N) < t_j$ for every $j > J$. The earlier terms vanish because every exponent in their finite sets divides L , and hence divides N . Therefore

$$\Delta_{2,A}(N) \leq \Delta_{2,E}(N) + \Delta_{2,V}(N) < 2\rho < \varepsilon, \quad N \geq Q \geq L \geq N_0.$$

Overlaps between the supports only improve the inequality. Infinitude of A makes its displacement positive. The atom comparison used in the cover proof transfers arbitrarily small displacements to every integer base; the fixed rational lattice then excludes rationality. $\square$

The strengthened cover conclusion is kernel-checked as `strengthenedPositiveCoverClaim`, and the common-witness conclusion is kernel-checked as `mixedSupportClaim`. Their public dependency closure was replayed under Lean 4.29.1. The arbitrary-positive-weight cover and mixed-support conclusions are also kernel-checked as `arbitraryWeightPositiveCover_allBase_hereditary` and `arbitraryWeightMixedSupport_allBase_hereditary`; the latter gives the theorem's hereditary all-base conclusion for the arbitrary-weight branch. This theorem does not establish strict enlargement over both individual classes. The weighted non-cover host has a separate source body in `the weighted support without a strengthened cover`; the proposed cover-only host is not established here. It also leaves universal irrationality and the quantitative and prime-power thinning claims open here. [Math]

With arbitrary positive cover weights, the class of subsets of such mixed hosts is closed under finite unions and finite changes. For weighted supports use the union of the two finite prime sets: the prime part grows and $h/(2^h - 1)$ decreases. For two covers, interleave their finite sets with weights $\eta_j/2$ and $\theta_j/2$; the total cost grows by at most two, since $2^{\alpha_j} \leq 2$. Subsets inherit the same hosts, and finite sets have finite weighted mass. This argument uses the positive-weight variant; it does not silently reindex the dyadic weights in (13). Countable unions require a tail budget. Every prime singleton is admitted, but for the full prime support $\mathcal{D}$ one has $\Delta_{2,p}(N) > 1/3$ for every $N \geq 1$: indeed $\sum_{r \geq 1} 2^{-r} \omega(N+r) \geq 1$, whereas $X_{\mathcal{D}}(2) \leq \sum_{a \geq 2} (2^a - 1)^{-1} < 2/3$. Here $\omega(n)$ is the number of distinct prime divisors of n .

The elementary independence boundary. For the divisor counts c_A and a fixed $a \in A$, the identity $c_A(an) - c_A(n) = 1 - \mathbf{1}_{a|n}$ holds for every $n \geq 1$ exactly when a is coprime to every other member of A . Every summand $\mathbf{1}_{t|an} - \mathbf{1}_{t|n}$ is nonnegative; $t = a$ supplies the right side, and coprimality makes all other terms zero. If $g = (a, t) > 1$ for some $t \neq a$, take $n = t/g$ to obtain an additional positive term. For distinct odd primes p, q , the covariance of $\mathbf{1}_{2p|n}$ and $\mathbf{1}_{2q|n}$ over a common period is $1/(2pq) - 1/(4pq) = 1/(4pq)$. Thus dilating a prime support already destroys the independence used in the prime-incidence argument. This brief ordinary calculation marks a boundary of that argument, not an irrationality criterion.

1.6 The universal problem and the half-value question

For an integer $n \geq 1$ set $x_n := 1/(2^n - 1)$, and for $A \subseteq \{1, 2, 3, \dots\}$ write $x_A := \sum_{n \in A} x_n$; the sum always converges, since $x_n = O(2^{-n})$. Erdős's Problem #257 asks a single universally quantified question about this family.

Definition 1.4 (Universal #257, written (U)). Is $x_A = \sum_{n \in A} 1/(2^n - 1)$ irrational for every infinite $A \subseteq \mathbb{N}_{>0}$? **OPEN.** [Open]

Definition 1.5 (The half-value question, written (H)). Let the *Mersenne achievement set* be $A := \{y \in \mathbb{R} : \exists A \subseteq \mathbb{N}, 0 \notin A, y = \sum_{n \in A} x_n\}$ ([LEAN SOURCE](#)). Is $1/2 \in A$? **OPEN**. [Open]

The two are joined by one elementary and completely one-sided implication. Every $2^n - 1$ is odd, so every *finite* subset sum of $\{x_n\}$ has odd reduced denominator and cannot equal $1/2$ ([LEAN SOURCE](#)). Hence a witness to $1/2 \in A$ is necessarily an infinite A with $x_A = 1/2$ rational, which refutes (U) outright. Conversely, (U) implies $1/2 \notin A$ as one instance among uncountably many, and a proof of $1/2 \notin A$ leaves (U) entirely open. Thus (H) is one candidate counterexample to (U), not a restatement of it. Section 2.5 explains the corresponding difference between a finite certificate of nonmembership and an unbounded survival condition. [Math]

The reciprocal-summable and finite-prime weighted theorems, together with the checked support theorems below, leave (U) and (H) open. The later countermodels identify the particular arithmetic information lost by individual approaches; they do not prove that all such failures have one common cause.

1.7 Known support theorems

Full-support irrationality is Erdős's theorem [1]. The pairwise-coprime case is Erdős's theorem of 1968, where he also stated that the coprimality condition can be removed [2, p. 222]. Luca and Tachiya proved irrationality for every purely periodic integer weight that is not identically zero [3, Theorem A, p. 139]; the purely periodic, eventually periodic, residue-class, odd-support and periodic-weight rows below follow from their theorem, and odd support is their Example 2 [3, p. 140]. The lcm-gap supports, including factorials and powers of two, have summable reciprocals, so they fall under the reciprocal-summable theorem stated by Erdős. The Lean declarations in the table are formal proofs of these statements.

The following table includes two proof mechanisms: block certificates for the divisor counts $c_A(n)$, and rational-approximation estimates for a growing least-common-multiple gap. The latter need not be presented as instances of the former. The notation $sc_A(n)$ in later source excerpts means the same function as $c_A(n)$. All bases in the table are integers $b \geq 2$. For the lcm-gap row write $L_k = \text{lcm}(a_0, \dots, a_{k-1})$, with $L_0 = 1$.

Support class	Exact hypotheses	Site [Lean]
Full support $A = \mathbb{N}_{>0}$	$b \geq 2$	LEAN SOURCE
Multiples $A = d\mathbb{N}_{>0}$	$b \geq 2, d \geq 1$	LEAN SOURCE
Purely periodic	$m \geq 1$, m -periodic, contains a positive element	LEAN SOURCE
Eventually periodic	$m \geq 1$, m -periodic from N_0 , A infinite	LEAN SOURCE
Residue class	any $m \geq 1$, any c	LEAN SOURCE

Support class	Exact hypotheses	Site [Lean]
Odd support (density 1/2)	$b \geq 2$	LEAN SOURCE
Growing lcm gap	a_k strictly increasing and positive; $a_k - L_k \rightarrow \infty$	LEAN SOURCE
Factorials; powers of two	instances of the lcm-gap theorem	LEAN SOURCE , LEAN SOURCE
Pairwise coprime (Erdős 1968)	A infinite, pairwise coprime, and $\sum_{a \in A} 1/a < \infty$	LEAN SOURCE
Signed periodic weights	periodic integer weights; their divisor-transformed coefficients are nonnegative and frequently nonzero	LEAN SOURCE

The lcm-gap hypothesis does imply reciprocal summability. Eventually $a_{k+2} > \text{lcm}(a_k, a_{k+1})$. If $a_{k+1} \geq 2a_k$, monotonicity gives $a_{k+2} > 2a_k$. Otherwise $a_k \nmid a_{k+1}$, so $\text{lcm}(a_k, a_{k+1}) \geq 2a_{k+1} > 2a_k$. Thus each of the even- and odd-indexed subsequences eventually grows geometrically, proving $\sum_k 1/a_k < \infty$. This explains the inclusion claimed above, rather than using “lacunary” as a substitute for a growth hypothesis. For eventually periodic weights, the periodic tail must not be identically zero: an eventually zero weight sequence gives only a finite rational sum.

Full-support irrationality for every integer base $b \geq 2$ is classical (Erdős, 1948). The contribution recorded here is its Lean formalisation and the explicitly delimited structured-support developments, not a new extension from base two to arbitrary integer bases. Each support family retains its stated hypotheses; none establishes irrationality for every infinite support. The $b = 2$ case of the first row recovers the Erdős–Borwein constant $E = \sum_{n \geq 1} 1/(2^n - 1)$ as a one-line corollary of that classical theorem. Separately, the independence and tail estimates used by a particular certificate construction must be checked for that support. Barrier 2.7 discusses those construction requirements; the averaging proof above does not assume pairwise coprimality or independence of divisor events.

A separate conditional result allows common factors confined to divisors of one fixed integer Q , while the remaining factors must be pairwise coprime and have summable reciprocals ([LEAN SOURCE](#)). It also assumes that for every $K \geq 1$ there is an $N \geq 0$ such that

$$2^K \mid \sum_{r=1}^K c_A(N+r)2^{K-r}, \quad \sum_{r \geq 1} c_A(N+K+r)2^{-r} \leq 16.$$

Thus relaxing coprimality does not remove the simultaneous divisibility and tail-bound requirement. This additional hypothesis is not proved here, and the source file does not establish a dichotomy for arbitrary supports. [Lean][Open]

The record also gives a logarithmic bound on runs of zero divisor counts under rationality: for each $\varepsilon > 0$, a suitable constant B bounds the length by $\varepsilon \log_2 N + B$ beyond the stated threshold ([LEAN SOURCE](#)). For these supports there is already a simpler unconditional bound. If $a_0 = \min A$, every a_0 consecutive positive integers contain a multiple of a_0 , where c_A is positive. Hence every zero run has length at most $a_0 - 1$, whether or not $X_A(2)$ is rational. In particular the logarithmic bound alone does not distinguish rational-valued supports. The separate unbounded-tail result [LEAN SOURCE](#) also requires interpretation: an unbounded sequence can still satisfy $u(N) = o(2^N)$. [Lean][Math]

Remark 1.6 (What is and is not covered in the universal direction). The pinned Lean corpus contains no theorem for $A = \text{the primes}$, but this is a formalisation boundary, not a mathematical open case: Tao and Teräväinen prove the base-2 prime-support value irrational by first identifying it with $\sum_{n \geq 1} \omega(n)2^{-n}$. Duverney and Tachiya likewise settle squarefree support at every power-of-two base, and the perfect i th powers at every integer base (Corollary 1.2 and Example 1.1, p. 4 of the author preprint). Neither result is formalised here. The statement of Problem #257 has been formalised before, as a conjecture with an unfilled proof, in the *Formal Conjectures* [collection](#); that file also proves the Lambert identity for the full-support value and states its irrationality, the 1948 theorem of Erdős, with proof sorry. It is statement-level prior art, and no theorem recorded here is derived from it. The corpus still has no general theorem for an arbitrary positive-density nonperiodic or divisor-dense support, and the universal statement remains open.

Periodic weights: the two hypothesis sets. The periodic sequence theorem belongs to Luca and Tachiya [13]; their later author-written account states it as Theorem A [3, p. 139]. In the displayed Lean row above, nonnegativity and frequent nonvanishing are assumptions on $c_w(n) = \sum_{d|n} w(d)$, not on $w(n)$ itself. The formal row must therefore be read as a separately specified instance, not as an exact transcription of every clause of the published theorem.

A separate denominator-growth criterion. Erdős's growth theorem [21], as restated and recovered by Barreto, Kang, Kim, Kovač and Zhang [22, Remark 4(4), p. 5], gives irrationality of $\sum_n 1/a_n$ for an increasing integer sequence with $a_n \geq n^{1+\tau}$ for some $\tau > 0$ and $\limsup_n a_n^{1/2^n} = \infty$. For $a_n = b^{c_n} - 1$, where c_n are increasing positive exponents and $b \geq 2$ is an integer, the growth hypothesis becomes

$$\limsup_{n \rightarrow \infty} \frac{c_n}{2^n} = \infty.$$

Indeed $c_n \geq n$, so the polynomial lower bound holds eventually, and $\log_b(b^{c_n} - 1) = c_n + O(1)$. An eventual lower bound suffices: deleting finitely many terms subtracts a rational number and changes 2^n only by a fixed factor after reindexing.

This is not the reciprocal-summability condition of the reciprocal-support theorem in Section 1.1. The squares $c_n = n^2$ have convergent reciprocal sum but $c_n/2^n \rightarrow 0$. In the other direction, set $n_1 = 1$ and $n_{k+1} = n_k + 4^{n_k}$, and let block k consist of all integers from 4^{n_k} to $2 \cdot 4^{n_k} - 1$. The blocks are successive disjoint increasing blocks, and their union has increasing enumeration satisfying $c_{n_k} = 4^{n_k}$. Thus $c_{n_k}/2^{n_k} = 2^{n_k} \rightarrow \infty$,

whereas each block has reciprocal sum at least $1/2$. Its reciprocal sum therefore diverges. These examples separate these two hypotheses, not every support criterion in this paper. [Cited][Math]

1.8 Geometry and rational-target membership

Theorem 1.7 (Achievement-set geometry). *A is compact, closed, perfect, totally disconnected and nowhere dense, and $\text{volume}(A) = 1$. Thus its measure is positive although it contains no interval. Its convex hull is $[0, E]$, where $E = \sum_{n \geq 1} w_n$. The positive-index digit coding onto A is injective: each achievable real has exactly one support. achievement-set (uniform) [Lean]*

Sources: [LEAN SOURCE](#), [LEAN SOURCE](#), [LEAN SOURCE](#), [LEAN SOURCE](#), [LEAN SOURCE](#), [LEAN SOURCE](#).

Lean: [paper_achievement_geometry](#).

Theorem 1.8 (Support-restricted refinement). *Use zero-based indices in this statement: coordinate $j \in \mathbb{N}$ carries weight w_{j+1} . For $J \subseteq \mathbb{N}$, consider the sums that use only coordinates in J . If $\mathbb{N} \setminus J$ is finite, this achievement set has measure $2^{-|\mathbb{N} \setminus J|}$; if infinitely many coordinates are omitted, its measure is zero. Injectivity survives every restriction; perfectness is proved when J is infinite. No perfectness claim is made for finite J , whose coding range is finite. achievement-set (uniform) [Lean]*

Sources: [LEAN SOURCE](#), [LEAN SOURCE](#), [LEAN SOURCE](#), [LEAN SOURCE](#).

Hausdorff dimension of a restricted set. For example, retaining the positive exponents $2, 4, 6, \dots$ gives a set of measure zero but Hausdorff dimension $1/2$. More generally, let $J \subseteq \{1, 2, \dots\}$ and put

$$S_J = \left\{ \sum_{n \in J} \varepsilon_n w_n : \varepsilon_n \in \{0, 1\} \right\}, \quad s_N = \#(J \cap \{1, \dots, N\}), \quad d = \liminf_{N \rightarrow \infty} s_N / N.$$

Then $\dim_H S_J = d$. Here is a proof, separate from the supplied formal measure statements. If J is finite, both sides are zero, so suppose J is infinite.

Write $w_n = 2^{-n} + \delta_n$, where $\delta_n = \sum_{r \geq 2} 2^{-rn}$. A geometric sum gives

$$\sum_{k > n} \delta_k = \sum_{r \geq 2} \frac{2^{-rn}}{2^r - 1} \leq \delta_n / 3 < \delta_n.$$

The unique Mersenne digit coding therefore defines a map from S_J onto the binary digit-restricted set $B_J = \{\sum_{n \in J} \varepsilon_n 2^{-n}\}$ by keeping the same digits. This map is 1-Lipschitz. In fact, at the first differing digit n , the binary difference is nonnegative after orienting the larger digit first, and the Mersenne difference exceeds it by at least $\delta_n - \sum_{k > n} \delta_k > 0$.

The 2^{s_N} Mersenne prefix intervals have length at most $R_N \leq 2^{1-N}$. For any $t > d$, a subsequence with $s_N / N < (d + t)/2$ makes their total t -power length tend to zero. Hence $\dim_H S_J \leq d$. For the reverse inequality, give the allowed binary digits independent probabilities $1/2$. Each dyadic interval of length 2^{-N} has measure at most 2^{-s_N} ; endpoints have measure zero because J is infinite. An interval I with $2^{-N-1} < |I| \leq 2^{-N}$

meets at most three such dyadic intervals. For $0 < s < d$ and all sufficiently large N , its measure is therefore at most $3 \cdot 2^{-sN} \leq 3 \cdot 2^s |I|^s$. Summing this inequality over any sufficiently fine cover of B_J gives a positive lower bound for the sum of its s -power lengths. Thus $\dim_H B_J \geq s$; letting $s \uparrow d$ proves $\dim_H B_J \geq d$. A Lipschitz map cannot increase Hausdorff dimension, since it sends a cover to one of no larger diameters. The map above now yields $\dim_H S_J \geq d$, as required. For $d = 0$ the lower bound is automatic. [Math]

Remark 1.9 (Achievement-set terminology and source boundary). We use the modern term *achievement set* for the subsum set that [Takeya's 1914 note](#) describes as a set of partial sums. For an infinite summable sequence of positive terms, Takeya proves perfectness, and nowhere density under an all-index strict-tail condition; [Kovač-Tao, Remark 4.1](#) verifies that strict-tail condition for the fixed-base Mersenne weights. Nitecki proves that a subsum set whose terms all exceed their tails is a Cantor set of measure $\lim_n 2^n R_n$, and credits this to Hornich [4, Theorem 4(1), p. 9]; this is the classical source for the topology and measure used here. The local volume dichotomy, injectivity, and greedy-membership statements above and below are Lean-checked claims of this corpus. The sparse power-series criteria of [Kaneko-Suzuki-Tachiya](#) (Theorems 1–3, pp. 3–5 of arXiv v1) assume sparse coefficient supports; their relation to the divisor-incidence coefficients of these subseries is recorded after Theorem 1.1. [Cited]

Theorem 1.10 (Membership equals greedy survival; the fatal-gap dichotomy). *For a real target $x \geq 0$, let $r_n(x)$ be the remainder after the greedy rule has processed weights $w_1, \dots, w_n$, and let $R_n = \sum_{j>n} w_j$, with $r_0(x) = x$ and $R_0 = E$. Then*

$$x \in A \iff x \geq 0 \text{ and } r_n(x) \leq R_n \text{ for every } n \geq 0.$$

For every $x \in [0, E]$, nonmembership is equivalent to a finite strict gap between the two next-prefix intervals; see Observation 9.44. greedy recurrence (uniform) [Lean]

Sources: [LEAN SOURCE](#); [LEAN SOURCE](#); [LEAN SOURCE](#); [LEAN SOURCE](#); witness type [LEAN SOURCE](#). The general-target statement is checked by [LEAN SOURCE](#).

Two boundary degeneracies are excluded at every rank, unconditionally: $r_{k-1}(1/2) \neq w_k$ by the odd-denominator argument above, and $r_{k-1}(1/2) \neq R_k$ because the tail differs from E by a rational while E is irrational. Both greedy decisions are therefore strict at every rank. [Math]

The supplied computation record reports the following finite checks. The underlying tests use exact integer arithmetic; decimal margins are summaries of those data, not proofs of an asymptotic estimate. The scan to row 200,000 and the long grouped computation were not rerun for this revision. The quotient calculation through row 2500 was recomputed with exact integers, as specified below. Reset margins are measured by $\log_2 |\text{rem}(r+1) - 2^{r+1}| - (r+5)/2$. The separate real-greedy scan uses $\log_2(4^k(r_{k-1}(1/2) - w_k))$ at a take and $\log_2(4^k(R_k - r_{k-1}(1/2)))$ at a safe skip. These are different normalisations, not one common margin statistic. Truncation rungs $J = 3, \dots, 22$ all *proved* to survive (not merely computed). Seam orbit certified to row 200,000: branch counts R:M:U = 100197:49899:49898, the target inequality failing at exactly one reset row, $r = 7$; maximal pure-R run 19, attained at row 158,096.

Independently, rows $6 \leq n \leq 2500$: 1209 resets, zero classification anomalies, minimum in-scope margin +1.1119 bits at row 14 growing to ≈ 1246 bits by row 2500 (mean 627), never re-approaching the threshold.

Sharp tail margin, ranks 2 to 3000: no fatal skip, take fraction 0.4992, tightest take-margin +0.0340 bits at rank 7 (the near-tie $1/126 > 1/127$). In the stated units this is $\log_2(4^7(1/126 - 1/127)) = 0.034035\dots$. It is the smallest take-margin reported in this finite scan. The recorded finite greedy prefix through $m = 200,000$ satisfies $0 \leq 1/2 - \sum_{d \in G \cap [1, m]} w_d < 2^{-199999}$. The sum here is finite; this statement does not evaluate the infinite greedy sum.

A half-trapping macro receipt reaches depth 13,548,057 (labelled in its own source as a private exact finite computation receipt, not publication authority). Kernel-checked corroboration of the low rows: [LEAN SOURCE](#) verifies rows 13–30 by decide. [Cert]

Remark 1.11 (Literature, exactly). E is irrational [1]. Zudilin proves the irrationality measure bound $\mu(E) \leq 2.46497868\dots$ [8, Theorem 1, p. 154], correcting an earlier computation (see the remark at the end of his Section 3, p. 159). Transcendence of E is unknown [6, p. 2], and the Erdős–Graham conjecture that every such subset sum is irrational, i.e. (U), is itself open. For the binary digits of E , Erdős’s proof produces arbitrarily long runs of zeros. Crandall asked whether the block 11 occurs infinitely often in the binary expansion of E , and Campbell proves that it does [9, Theorem 1, p. 12]. [Cited]

The cited digit results establish the occurrence of particular patterns. An irrationality-measure estimate gives a different kind of information: a run of L equal digits after position N approximates E by a dyadic rational with denominator at most 2^N , giving $L < (\mu(E) - 1 + \varepsilon)N$ for all sufficiently large N . The reset argument considered here needs a bound for runs in its own inhomogeneous recurrence, not simply in the binary digits of E . No implication from the cited digit results to that bound is proved here. Such a bound is an input to this particular approach to (H), not a necessary ingredient of every possible proof of membership.

II. Limitations and extensions of recorded methods

2 Limitations of the recorded methods

2.1 What the obstructions do and do not show

This section studies the scope of particular reformulations and certificate mechanisms. Several share the greedy representation, but no proof identifies a single necessary obstruction to all methods. The statements below are useful insofar as they specify which hypotheses fail, which implications are equivalences, and which parameters must be uniform. The first structural fact is uniqueness at a fixed target and depth.

Observation 2.1 (Uniqueness within the fixed-target coding). Each Mersenne weight exceeds the sum of all subsequent weights. Consequently the digit coding is injective (Theorem 1.7): a represented target has only one support, recovered by the greedy rule. For the integer weights, both the scale and the cutoff matter. If $2 \leq k \leq M$ and $q_{M,d} = \lfloor 2^M / (2^d - 1) \rfloor$ for $2 \leq d \leq k$, then

$$q_{M,d} - \sum_{e=d+1}^k q_{M,e} \geq 2^{M-k}.$$

Indeed, with $w_d = (2^d - 1)^{-1}$ and $R_k = \sum_{e>k} w_e$,

$$w_d - \sum_{e=d+1}^k w_e = w_d - R_d + R_k > R_k > 2^{-k}.$$

Rounding the head down loses less than one, while rounding the later terms down can only increase the difference. The resulting integer is greater than $2^{M-k} - 1$, proving the bound. Comparing two words at their first differing rank gives the same separation for their sums. The rows with $M = 2s$, $k = s - 1$ and $s \geq 3$ therefore have gap at least 2^{s+1} ; the full rows with $k = M$ have gap at least 1. The first bound must not be applied after enlarging the cutoff: at $M = 6$, adjoining rank 3 introduces the subset sum $q_{6,3} = 9 < 16 = 2^{3+1}$.

When a finite construction is proved to encode a fixed greedy target, its choices are constrained by that support. Injectivity alone does not identify an arbitrary finite approximation with a real greedy prefix. Nor does it exclude averaging over indices, moduli, nearby targets or auxiliary constructions. *meta* (uniform) [Lean][Math]

Sources: [LEAN SOURCE](#), [LEAN SOURCE](#), [LEAN SOURCE](#).

The six limitations below require separate arguments. Fixed-target uniqueness explains canonicalisation (2.2); the recorded identities determine which sufficient conditions are equivalent to their goal (2.3). Measure and category (2.6), observation scales (2.4), finite versus infinite membership (2.5), and the hypotheses of the support certificate (2.7) do not follow from injectivity alone.

2.2 Uniqueness of the specified greedy words

Proposition 2.2 (Uniqueness under the stated finite conditions). *Put $G = \text{greedyMersenneSupport}(1/2)$, $w_n = (2^n - 1)^{-1}$ and $R_d = \sum_{n>d} w_n$. The following are three separate uniqueness statements.*

- (i) *If $D \subseteq \{1, \dots, d\}$ and $X_D(2) \leq 1/2 \leq X_D(2) + R_d$, then $D = G \cap \{1, \dots, d\}$.*
- (ii) *If $c \geq 4$, $D \subseteq \{2, \dots, c-1\}$ and $0 < 1/2 - X_D(2) < w_c$, then $D = G \cap \{1, \dots, c-1\}$.*
- (iii) *Let positive integer weights $v_1, \dots, v_k$ satisfy $v_i \geq g + \sum_{j>i} v_j$ for a positive integer g . For a nonnegative integer target T , let y^* be the word obtained by visiting these weights in order and taking each when it fits. An admissible Boolean word y has $0 \leq T - \sum_i y_i v_i < g$ if and only if $y = y^*$ and $0 \leq T - \sum_i y_i^* v_i < g$.*

Part (iii) gives uniqueness under the small-remainder condition, not existence. Theorem 6.100 gives a one-weight counterexample to dropping that condition. straddle-prefix ; integer-greedy (uniform) [Lean]

Sources: [LEAN SOURCE](#); [LEAN SOURCE](#); [LEAN SOURCE](#).

Proof. For (i), truncate D at each successive rank. The target remains in the corresponding prefix interval because all discarded terms are bounded by the complete tail. The two child intervals are disjoint: the lower one forces a greedy skip and the upper one forces a greedy take. Induction identifies every digit. Part (ii) implies (i) at depth $c-1$, since $w_c < R_{c-1}$. For (iii), consider the first disagreement with the integer greedy word. Taking a rejected weight would exceed T ; rejecting an accepted weight leaves a deficit at least g , even after taking every later weight. The converse is immediate. $\square$

Excludes at fixed depth. Under the three canonicalisation hypotheses just stated, a proof cannot gain freedom by choosing a different admissible word for the same target and depth. This says nothing about families satisfying different approximation conditions, or about averaging over depths or auxiliary data.

What the finite tests show. Three different calculations are recorded here. For unrestricted prefix residues, the proposed $2^{t/2}$ exclusion radius is reported to fail in 74% of the tested ranks $t \in [8, 26]$. Here the finite test uses the truncated, scaled fractional parts

$$b_d(t) = \left\lfloor 2^t \left\{ \frac{2^t}{2^d - 1} \right\} \right\rfloor, \quad \{y\} = y - \lfloor y \rfloor.$$

At $t = 15$, the values for $d = 2, 3, 5, 6, 10$ are respectively 21845, 4681, 1057, 4161, 1025. Their sum is $32769 = 2^{15} + 1$, so the sum has circular distance one from 0 modulo 2^{15} . This is a statement about these integer truncations, not the exact fractional-part sum. A heuristic count of near-hits has size $2^{t/6}$; that estimate is not a theorem about the greedy support. [Cert]

Long prescribed digit strings also occur for unrestricted finite words: the recorded examples at $t = 20$ and $t = 22$ have respectively eighteen leading ones and twenty leading zeros. A separate construction supplies about $t/12$ independently adjustable bits, with the reported finite check extending to $t = 601$. These examples disprove the proposed static exclusions, not an assertion about the actual greedy word. [Cert]

The integer-window search reported no counterexamples for $t \leq 22$ or at $t = 40, 42, 45$ over the tested cut positions. That absence is a finite observation. Proposition 2.2 explains why certain *specified* approximation conditions leave only the greedy word; it does not identify all three searches with one another or rule out averaging over other data. [Cert]

Remark 2.3 (A recorded hazard). The natural sufficient condition via fractional split-mass is *not* necessary on real crossing cores: the fixture $D = \{2, 3\}$, $c = 5$ is an explicit counterexample (LEAN SOURCE). The example refutes necessity of that sufficient condition. It does not refute using the condition on a class of cores for which it can be proved. [Lean]

2.3 Equivalent formulations of half-membership

The six conditions listed below are equivalent to half-membership. Their quantifier patterns can look less restrictive than a single compatible infinite construction, but the equivalences must be checked mathematically. They follow here from a carry identity and an unconditional tail estimate. No claim is made that every conceivable sufficient condition is equivalent to the goal.

Lemma 2.4 (Carry and residual value). *Let $A \subseteq \mathbb{N}$ with $1 \notin A$, put $\delta := 1/2 - x_A$, and let $\text{ihc}(A, N)$ denote the integer half-carry at level N . Then*

$$\text{ihc}(A, N) = 2^{N+1} \delta + T(N+1), \quad 0 \leq T(N+1) \leq 2\sqrt{N+1} + 4,$$

where $T(m) = \sum_{r \geq 1} c_A(m+r)2^{-r}$. half-carry recurrence (uniform) [Lean]

Sources: LEAN SOURCE; LEAN SOURCE; LEAN SOURCE.

Proposition 2.5 (A cofinal carry bound). *If $|\text{ihc}(A, N)| \leq C\sqrt{N+1} + C'$ for arbitrarily large N , with fixed constants C, C' , then $\delta = 0$. For the particular greedy support $A = G$ of $1/2$, a one-sided upper bound $\text{ihc}(G, N) \leq C\sqrt{N+1} + C'$ already suffices.*

Lean: half_of_cofinal_absolute_carry, greedy_half_of_cofinal_upper_carry.

Proof. Lemma 2.4 gives

$$2^{N+1}|\delta| \leq |\text{ihc}(A, N)| + T(N+1).$$

The right side is $O(\sqrt{N+1})$ along the stated sequence, so division by 2^{N+1} yields $\delta = 0$. For $A = G$, the greedy partial sums never exceed $1/2$, hence $\delta \geq 0$ (LEAN SOURCE). Since $T(N+1) \geq 0$, the one-sided upper bound gives the same conclusion. Conversely, $\delta = 0$ makes the carry equal to $T(N+1)$, so the bound holds with $C = 2$ and $C' = 4$. This converse does not assert the bound for arbitrary prescribed constants. For a prescribed A , the conclusion is $x_A = 1/2$; it is not merely membership witnessed by some other support. For $A = G$, this equality is equivalent to $1/2 \in A$ by LEAN SOURCE. $\square$

Lemma 2.6 (The terminal bound at square depths). *Suppose $x_A = 1/2$ and $1 \notin A$. Then for every $k \geq 1$,*

$$\text{ihc}(A, k^2 - 1) = T(k^2) \leq 2k + 4 = B(k^2).$$

half-carry recurrence (cofinal) [Math]

Lean: [square_depth_witness](#).

Proof. By Lemma 2.4 with $\delta = 0$, $\text{ihc}(A, k^2 - 1) = T(k^2) \leq 2\sqrt{k^2} + 4$. At a perfect square the real and truncated square roots agree exactly, $\sqrt{k^2} = k = \lfloor \sqrt{k^2} \rfloor$, so the real envelope coincides with the discrete strip $B(n) = 2\lfloor \sqrt{n} \rfloor + 4$ ([LEAN SOURCE](#)) with no slack required. $\square$

At the depths $M = k^2$, the real square root and the integer square root agree. Thus the reverse implication for the terminal bound uses the original constant 4, not a relaxed constant 6. Truncating an achieving support at $M = k^2$ leaves the carry at that depth unchanged, by [LEAN SOURCE](#). An achieving half-support excludes rank one because $w_1 = 1 > 1/2$.

Proposition 2.7 (Six equivalent membership conditions). *Each of the following is logically equivalent to $1/2 \in A$, not strictly weaker:*

- (a) $C_G(N) \leq 2\sqrt{N} + 4$ for every $N \geq 0$ ([LEAN SOURCE](#); [LEAN SOURCE](#));
- (b) there are arbitrarily large positive exponents outside G ([LEAN SOURCE](#));
- (c) at arbitrarily large depths n , some $D \subseteq \{2, \dots, n\}$ satisfies $Q(D, n) = 2^{n-1} - 1$ ([LEAN SOURCE](#)); no agreement between the sets D at different depths is required;
- (d) at arbitrarily large M , $\text{ihc}(G, M) \leq B(M + 1)$ ([LEAN SOURCE](#));
- (e) at arbitrarily large positive depths M , some $D \subseteq \{2, \dots, M\}$ satisfies $|\text{ihc}(D, M - 1)| \leq B(M)$ ([LEAN SOURCE](#));
- (f) the half-target greedy remainder never exceeds the full remaining tail, so no finite fatal gap exists ([LEAN SOURCE](#)); the seven equivalent forms in the later classification express this same condition.

Forward directions are supplied formal (e.g. [LEAN SOURCE](#), [LEAN SOURCE](#), [LEAN SOURCE](#)). For (d) and (e), the reverse directions use the square depths in Lemma 2.6. All six equivalences, including these reverse directions, are now Lean-checked in [LEAN SOURCE](#). For (a), membership gives $C_G(N) \leq 2\sqrt{N+1}+3 \leq 2\sqrt{N}+4$ when $N \geq 1$, and $C_G(0) = 0$. The forward implication still needs the greedy inequality $x_G \leq 1/2$. half-carry recurrence ; exact-row (cofinal) [Lean][Math]

Scope of the equivalences. These conditions are equivalent to membership, not strictly weaker substitutes. Equivalence does not make a formulation useless as an invariant, but it does mean that proving any of them is the same problem as (H). The equivalent conditions remain unproved. In particular, mutually *incompatible* witnesses at different endpoints do not buy a weaker sufficient condition, and the Π_2^0 shape of these hypotheses is not evidence of extra room.

Scope limit, preserved exactly. Item (a) does *not* extend to every A with $1 \notin A$. Dropping the greedy-specific fact $x_G \leq 1/2$ yields only the one-sided statement $\text{hbound} \iff (1/2 - x_A) \leq 0$. Thus the conclusion for a general support is a one-sided inequality, not equality with $1/2$.

2.4 The size of the required error bounds

One sufficient half-membership argument requires the reset inequality

$$|\text{rem}(r + 1) - 2^{r+1}| > 2^{(r+5)/2} \quad (r \geq 10 \text{ an upper or middle reset}).$$

For the normalised deviation $\lambda_{r+1} = (\text{rem}(r+1) - 2^{r+1})/2^{r+1}$, this is $|\lambda_{r+1}| > 2^{(3-r)/2}$. A near-integer estimate for a weighted divisor sum is a proposed way to obtain information about this inequality, not an equivalent condition established here. Section 11.1 displays the exact recurrence and the additional, changing skip-set term. Each obstruction below has its own hypotheses; none excludes every method at the stated scale.

Class (a): fixed precision or bounded state. The following three statements concern abstract recurrence models. Their conclusions apply to the Mersenne problem only within the specified models; they do not use its full divisor-count constraints.

Proposition 2.8 (Limits of the specified local summaries). (i) After L common steps the endpoint residue mod 2^L of an affine binary orbit is independent of the initial carry: $u(L) - v(L) = 2^L(u_0 - v_0)$ ([LEAN SOURCE](#)). (ii) Fix $m \geq 2$ and put $h = \lfloor (m+1)/2 \rfloor$. Consider the coefficient sequences

$$c_r(m) = h - r, \quad c_r(m+1) = 2r, \quad c_r(j) = 0 \ (j \neq m, m+1), \quad 0 \leq r \leq h.$$

Their binary sums are all $h2^{-m}$, and their scaled tails before position m agree, but their scaled tails immediately after m are r . Thus any exact label-and-decoder system recovering those latter tails must have at least $h+1 = \lfloor (m+1)/2 \rfloor + 1$ labels. A state determined only by the common preceding history, with no new input, cannot distinguish them ([LEAN SOURCE](#)). (iii) For any starting carry and any finite list of prescribed valuations and odd unit residues modulo 2^u , with $u \geq 1$, there are integer input coefficients with those data whose successor carries lie in their prescribed centred intervals. The unrestricted higher bits of the input coefficients may be chosen separately at each step, as in Proposition 10.4 ([LEAN SOURCE](#)), so these fixed-precision symbols alone do not exclude a completion in that specified family. This says nothing about extra constraints imposed by divisor counts from one common support. carry (uniform) [[Lean](#)]

Lean: `affineBinaryOrbit_mod_twoPow_eq`, `balancedPulse_weighted_pair`, `balancedPulse_endpoint_fanout`, `balancedPulse_label_card_lower_bound`, and 2 further declarations in the [coverage section](#).

For an integer recurrence $J(M+1) = 2J(M) - \delta(M+1)$, reduction modulo 2 gives $J(M+1) \equiv \delta(M+1) \pmod{2}$. For the particular recurrence $K(M+1) = 2K(M) - (\tau(M+1) - 1)$, with $K(2) = 1$, this gives, for $M \geq 2$, that $K(M)$ is even exactly when M is a square: divisors pair off except at a square. This is a genuine restriction at a specified index, but it is already encoded in the known forcing term. It supplies no additional run-length estimate by itself. [Math]

Class (b): global averages and Diophantine data.

Proposition 2.9 (Weighted denominator budget). Let $n \geq 2$, let $\text{Skip}_n \subseteq \{2, \dots, n-1\}$, and let D_n be the reduced denominator of the associated finite sum. Since D_n divides the product of its Mersenne denominators,

$$\log_2 D_n \leq \sum_{d \in \text{Skip}_n} \log_2(2^d - 1) \leq \sum_{d \in \text{Skip}_n} d \leq \frac{n(n-1)}{2} - 1.$$

The middle inequality is strict when the skip set is nonempty; for an empty skip set both sums are zero. This is an upper bound, not a quadratic asymptotic for D_n . Combined with the cited fixed

irrationality-measure exponent, it does not supply the required $2^{-3n/2}$ -scale estimate. Even an asymptotic for the exponent sum would need additional information about the skip distribution. diophantine (cofinal) [Math]

Lean: `skipSum_den_dvd_prod, weighted_denominator_budget`.

Scope of these obstructions. The common pre-index history in Proposition 2.8(ii) does not determine the post-index tail in its countermodel family. This does not exclude a state using additional arithmetic input from the actual support. The displayed denominator budget combined with the cited fixed irrationality-measure exponent does not yield the required run bound. These calculations do not exclude averaging over shifts, residues, or observation scales, nor a sharper analysis of the actual skip distribution.

Prime steps and bounded windows. In the universal recurrence, a prime step subtracts 1, so it sends $K(M)$ to $2K(M) - 1$; it is not exact doubling. Observation 10.1 gives an explicit frozen-row drift with values between 2 and 6 at the eleven indices 607, ..., 617, including the three prime indices 607, 613, and 617. This finite example shows that prime steps can occur within a small-value window. It neither proves statistical independence nor rules out every argument using prime-counting information. [Cert]

Remark 2.10 (Status of the exhaustion claim). The preceding counterexamples concern the specified finite summaries and estimates. They do not classify all available tools, even within the broader categories of local information and global averages. No exhaustion claim is used in a subsequent deduction.

2.5 What finite certificates decide

Theorem 2.11 (One-sidedness). *Non-membership of $1/2$ in A has an effectively checkable finite-certificate formulation, hence a Σ_1^0 formulation. Membership has the complementary Π_1^0 formulation. Survival through a tested finite depth alone does not establish membership; a uniform theorem or inductive invariant could. The arithmetical-hierarchy form by itself proves neither undecidability nor the absence of finite proofs. meta (uniform) [Lean][Math] The linked declarations establish the fatal-gap equivalences. The arithmetical-hierarchy classification also uses the effective tail estimate explained below; it is an ordinary computability deduction, not a separately replayed Lean theorem.*

Lean: `paper_one_sidedness, computablePred_halfFatalCertificateCode, half_not_mem_iff_exists_halfFatalCertificateCode, half_mem_iff_forall_not_halfFatalCertificateCode`, and 8 further declarations in the [coverage section](#).

Proof. The dichotomy $1/2 \in A$ or $\exists$ a fatal gap is unconditional (LEAN SOURCE), and the two alternatives are exactly complementary (LEAN SOURCE), with LEAN SOURCE a finite witness type. In the other direction, certified death implies non-membership (LEAN SOURCE, over LEAN SOURCE; $3/4$ is certified dead at level 1, lookahead 0, by LEAN SOURCE), whereas the absence of a certificate through depth D alone gives no conclusion about later depths. To make the search effective, note that a greedy remainder at a finite rank is rational. The tail R_n has rational upper bounds obtained by summing through a cutoff M and bounding the remaining sum by $2^{-M} + (2/3)4^{-M}$. These upper bounds converge to R_n . Thus the strict inequality $r_n(1/2) > R_n$, when true, is witnessed by some finite cutoff and a comparison of rational numbers. Enumerating ranks and cutoffs gives the

asserted finite-certificate search. Membership requires the survival inequality at every rank (Theorem 1.10). Checking a finite initial segment alone does not supply the later inequalities; a separate uniform argument or induction would be needed. $\square$

Finite search versus infinite survival. A discovered fatal-gap certificate proves non-membership. Further surviving prefixes give larger verified ranges, but do not establish membership without an argument covering all ranks. This distinction is unchanged whether one increases the geometric truncation depth, the number of quotient rows, or the depth of a grouped computation. The different finite parameters should not be reported as a single verification depth.

2.6 Why measure and category do not decide a rational target

Remark 2.12 (Achievement-set geometry does not decide rational-target membership). The achievement set has positive measure and is nowhere dense (Theorem 1.7). Neither fact alone decides membership of a specified rational. Injectivity fixes the support of an achievable value, but it does not exclude averaging over indices, shifts, residues, observation scales or auxiliary constructions. The recorded calculations obstruct the particular models and estimates analysed here; no impossibility theorem for all probabilistic, measure-theoretic or averaging methods is claimed. *achievement-set* (uniform) [Math]

The geometry itself is Lean-checked in Theorem 1.7: A has Lebesgue measure exactly 1 inside an ambient span of $E \approx 1.6067$, a relative density of about 62%, and is simultaneously nowhere dense, perfect and totally disconnected. Those facts may hold together. They do not yield competing theorems about specified rationals. Showing that A is “small” cannot decide (U) by measure: it is 62% of its ambient interval.

Recorded probabilistic heuristic. The supplied record reports an expected count of about 2×10^{-4} beyond $n = 31$ in a Bernoulli(1/2) digit model. This number is not a probability assigned to membership of the fixed point 1/2. Such an interpretation would require a specified probability space, a definition of the failure events, and a link between those events and the deterministic greedy orbit. None is provided here. Section 11.1 separates a reproducible finite phase experiment from its unproved probabilistic application. [Open]

2.7 Hypotheses of the CRT block-certificate argument

Observation 2.13 (Two inputs to the CRT certificate argument). The CRT block-certificate arguments considered here require both compatible divisibility prescriptions and a tail estimate strong enough for the height inequality $q(C + N + L + 2) < b^L$. Pairwise coprimality is one way to obtain the first input; reciprocal summability is one way to obtain the second. This observation concerns those certificate arguments only. It does not describe all theorems in Section 1: the reciprocal-summable and weighted averaging proofs do not require pairwise coprimality. *block-certificate hypotheses* (uniform) [Lean][Math]

What remains to check. One cannot apply a conditional CRT certificate theorem to arbitrary infinite A without establishing its hypotheses. Shared factors impose compatibility conditions on the prescribed residues; they do not make every CRT construction impossible. The tail budget in the cited proof is used at `CertificateKernel.lean:9586-`

-9598. Deleting that assumption does not complete the proof: one needs a replacement estimate for that step. This leaves open different constructions and different averaging arguments.

Explains only a method boundary. The primes are infinite and pairwise coprime, so they satisfy the first input. The identity [LEAN SOURCE](#) makes the dilation algebra exact there; yet this certificate theorem does not apply because $\sum_p 1/p$ diverges, so the stated reciprocal-sum hypothesis fails. This explains the absence of a *Lean theorem from this engine*; it does not make the base-2 value open, since Tao–Teräväinen settle it by a different analytic method. The sunflower generalisation inherits both axes and therefore does not recover that external theorem.

2.8 The distinct estimates still required

The half-target analysis involves a geometrically weighted divisor sum $\sum_{i \leq L} (\tau(M+i) - 1)2^{-i}$ and possible near-integer estimates. One universal-support certificate instead asks for an opening divisible block together with a controlled middle window. These inputs are related through incidence identities, but the displayed arguments do not prove them equivalent for all supports. They should therefore be retained as separate obligations. Digit-pattern occurrence results do not by themselves supply either required uniform estimate.

3 Unproved inputs for further arguments

The following proposals isolate inputs not supplied by the preceding results. Some are sufficient conditions, some are equivalent formulations, and one concerns reproducing an already known prime-support theorem. They are not an exhaustive list of possible approaches. The relevant question in each case is what must be proved, and whether the cited obstruction applies to that particular statement.

A claimed quantitative support condition

The absent analytic supplement claims a second criterion admitting reciprocal-divergent supports. Put $H_A(x) = \sum_{a \in A, a \leq x} 1/a$, $T_0 = 2$, $T_{j+1} = 2^{T_j}$, and $\ell(x) = \min\{j : x \leq T_j\}$. The condition $H_A(x) = o(\ell(x))$ is claimed to imply all-base irrationality, with no positive density assumption on A . Neither a proof of this implication nor the claimed sharper rational-phase lower bound is present in this checkout.

Two further analytic tests

Small displacements at a real base do not by themselves give the fixed rational lattice used in the integer-base proof. For $\beta = (3 + \sqrt{5})/2$, let $u = \beta^{-1}$ and $H = \{2^k m : k \geq 2, m \text{ odd}, m \leq 2^{2^k}\}$. The finite prefix identity and an algebraic norm show that the condition

$$\Delta_{\beta,H}(N) \left(1 + \sum_{\substack{a \in H \\ a \leq N}} \sum_{j=1}^{\lfloor N/a \rfloor} u^{N-ja} \right) \rightarrow 0$$

along some sequence of positive integers $N \rightarrow \infty$ would exclude $X_A(\beta)$ from $\mathbb{Q}(\sqrt{5})$ for every infinite $A \subseteq H$. The second factor controls the algebraic conjugate of the finite prefix; smallness of the displacement alone does not control this product. The product estimate remains unproved.

To justify the conditional implication, suppose instead that $X_A(\beta) \in \mathbb{Q}(\sqrt{5})$, and choose an integer $q \geq 1$ with $qX_A(\beta) \in \mathbb{Z}[\beta]$. The finite-prefix identity gives, for $N \geq 1$,

$$Y_N := q\Delta_{\beta,A}(N) = q(\beta^N - 1)X_A(\beta) - q \sum_{\substack{a \in A \\ a \leq N}} \sum_{j=1}^{\lfloor N/a \rfloor} \beta^{N-ja} \in \mathbb{Z}[\beta].$$

The infinite support makes $Y_N > 0$. Since $\beta^2 - 3\beta + 1 = 0$, conjugation σ sends β to $u = 3 - \beta \in (0, 1)$. For $m, n \in \mathbb{Z}$ the norm of $m + n\beta$ is $m^2 + 3mn + n^2$, so the nonzero Y_N has nonzero integer norm. Moreover,

$$|\sigma(Y_N)| \leq q \left(|\sigma(X_A(\beta))| + \sum_{\substack{a \in H \\ a \leq N}} \sum_{j=1}^{\lfloor N/a \rfloor} u^{N-ja} \right).$$

Using $0 < \Delta_{\beta,A}(N) \leq \Delta_{\beta,H}(N)$, we obtain

$$1 \leq |Y_N \sigma(Y_N)| \leq q^2 \max\{1, |\sigma(X_A(\beta))|\} \Delta_{\beta,H}(N) \left(1 + \sum_{\substack{a \in H \\ a \leq N}} \sum_{j=1}^{\lfloor N/a \rfloor} u^{N-ja} \right).$$

The asserted limit would contradict this inequality. The multiplier q and the conjugate of the putative value are fixed, independently of N . This proves only the implication from the displayed product estimate, not that estimate. [Math]

A different extension requires squarefree remote-tail control and a weighted displacement to be small at the same indices. CRT establishes compatibility of the finite opening divisibility conditions, but not this simultaneous smallness. Neither argument is claimed to establish its required analytic hypothesis.

3.1 A lower bound at reset ranks

What it needs. Either of the first two sufficient conditions below. Their different scales do not establish an implication between them. The sign condition (iii) is a separate question; no implication from it to half-membership is proved here.

- (i) *Reset-deviation bound.* For every upper or middle reset row $r \geq 10$, require $|\text{rem}(r+1) - 2^{r+1}| > 2^{(r+5)/2}$. Section 12.5 states the exact conditional implication to membership. A run-length estimate is a possible way to approach this bound, not an established exact reformulation of it.
- (ii) *Lean-facing form, already proved as a conditional implication.* `SeamUpperResetDyadicBandEscape`: for every $d \geq 13$ with `successorCarries` and every $j \leq d$, the reset charge $4 \cdot \text{overshoot} + \text{abovePulse}$ avoids the linear-width band $(2^{d-j+1} - 2(d+j), 2^{d-j+1}]$.

- (iii) *Sign law alone.* At every reset, M-resets have $\text{dev} > 0$ and U-resets have $\text{dev} < 0$. This specifies the sign but not a quantitative distance from zero. Neither an implication to half-membership nor a strict comparison with the first two conditions is established here.

Sources. [LEAN SOURCE](#) with conditional implication [LEAN SOURCE](#); quantifier collapse [LEAN SOURCE](#), which reduces the $\forall j \in [0, d]$ band check to one nearest-boundary index per row; finite verification [LEAN SOURCE](#); alternative implication [LEAN SOURCE](#) with conditional implication at :167 and base case [LEAN SOURCE](#).

Logical strength. The reset bound is a sufficient condition for half-membership in the cited argument. No converse is established here. The fact that it is quantitative, while membership is qualitative, does not prove that the condition is strictly stronger: that would require a separate implication or counterexample. The carry identity of Proposition 2.5 does not itself furnish this lower bound on the reset deviation.

What the finite checks establish. The reported computation checks the reset inequality over its stated range, with margins of 1.1119 bits at row 14 and about 1246 bits at row 2500; the reported mean is 627 bits. These are finite observations, not a proof of an asymptotic lower bound or of the inequality at every later reset. [Cert]

Uniqueness in Section 2.2 does not prevent this argument: the proposed bound concerns the actual greedy sequence. It does not require alternative supports for the same target.

Run lengths. The record reports a maximum right-branch run of length 19 at row 158,096, far below the proposed threshold $(r - 3)/2 = 79,046.5$ there. Observing larger runs at larger ranks does not refute the existence of some uniform finite bound. Neither a uniform bound nor the weaker asymptotic assertion $L_r = o(r)$ is proved by these data. The stated implication needs the required inequality beyond the verified initial ranks. [Cert]

The missing step is the quantitative lower bound itself. The finite observations do not supply the growing-precision estimate discussed in Section 2.4. [Open]

3.2 Increasing the depths of exact finite sums

What it needs. Rule out the recycle branch returning to a bounded endpoint forever. Any one of:

- (a) SkippedCoreCriticalQuotientSupply: for every below-half core D bounded by $[2, c]$ with genuine crossing deficit, $2^{(2c-2)-1} \leq Q(\text{insert } c \text{ } D) (2c - 2)$;
- (b) a direct proof that the crossing rank c produced by the recycle branch cannot repeat the value 4, or any bounded value, infinitely often;
- (c) a growth law for c as a function of the incoming endpoint n .

By Proposition 2.2(ii), form (a) need only be checked along the *one* canonical orbit `halfGreedyPrefixSupport`, not searched over D ; and by the recorded sharp-capacity biconditional it is exactly the $c - 2$ bit capacity test $SD1(2c - 2) < 2^{c-2}$. The fixture $D = \{2, 3\}$, $c = 5$ shows that the fractional-mass sufficient condition is not necessary. It does not rule out proving that condition for a suitably restricted family.

Sources. [LEAN SOURCE](#); the formal falsifier [LEAN SOURCE](#) (the model $n \mapsto (n = 6)$ satisfies the seed and the exact transition shape yet is not cofinal, so the schema plus an

endpoint-six seed provably does *not* give cofinality); conditional strict progress [LEAN SOURCE](#); unconditional recycling required input [LEAN SOURCE](#); hypothesis [LEAN SOURCE](#); capacity iff [LEAN SOURCE](#); seed fixture $\{2, 3, 6\}$ at [LEAN SOURCE](#).

The finite-row test and its unique candidate. The real-prefix argument in Proposition 2.2(ii) does not by itself identify an exact-row witness. The integer weights have a separate, decisive property: for fixed $n \geq 2$,

$$q_{n,d} = \left\lfloor \frac{2^n}{2^d - 1} \right\rfloor \quad (2 \leq d \leq n)$$

form a decreasing sequence in which each weight exceeds the sum of all later weights by at least 1. Apply the integer greedy rule with capacity $2^{n-1} - 1$: visit $d = 2, \dots, n$ and subtract $q_{n,d}$ when it fits. By Theorem 6.100 and Propositions 6.101–6.102, any exact-row witness has zero remainder and must be this greedy word. Thus there is at most one witness at a fixed depth, and its existence is tested by checking whether the final greedy remainder is zero. An exhaustive search over 2^{n-1} subsets is unnecessary.

This integer greedy word should not be silently identified with the real greedy prefix $G \cap [2, n]$: the two rules use different weights. The separation theorem proves uniqueness within the integer problem, not that identification.

Finite decidability and the cofinal requirement. The test decides the exact-row predicate at any specified depth. The claim that successful depths are arbitrarily large is still equivalent to half-membership, by Proposition 2.7(c). Neither uniqueness nor finite decidability supplies that cofinal assertion. The double-or-recycle argument must prove genuine progress in the depths, not rely on a purported freedom to choose many witnesses at one depth.

3.3 The known prime-support theorem

The base-2 prime-support value is not an open target:

$$\sum_{p \text{ prime}} \frac{1}{2^p - 1} = \sum_{n \geq 1} \frac{\omega(n)}{2^n}$$

by expanding each reciprocal Mersenne term geometrically and collecting divisors. Tao and Teräväinen prove the right-hand series irrational in Theorem 1.3 of [Quantitative correlations and some problems on prime factors of consecutive integers](#). Their proof is analytic and is cited here; it is not formalised in the pinned Lean corpus.

The local pairwise-coprime certificate theorem still does not cover the primes because it assumes summability of $\sum_p 1/p$. That is a boundary of this engine, not a boundary of the mathematics. The previously proposed windowed-tail repair is therefore withdrawn as a research contribution. It may remain useful as an audit question about whether the existing Lean engine can reproduce the known theorem, but reproducing a cited result is not new progress on #257. After Theorem 1.3 the authors state the prime-support extension to every integer base $b \geq 2$ and, separately, the full prime-power result at base 2, leaving those modifications to the reader. These statements should not be combined into an attributed all-base theorem for arbitrary prime-power subsets.

3.4 Equivalent membership conditions

Proposition 2.7 lists the equivalent conditions: a pointwise upper bound on C_G , infinitely many skipped ranks, exact quotient sums at unbounded depths, and the two cofinal carry bounds. For the carry conditions, the residual identity and its tail estimates are the relevant inputs: [LEAN SOURCE](#), [LEAN SOURCE](#), [LEAN SOURCE](#), and, for the terminal-strip conditions, evaluation of the square root at a perfect square. The one-sided scope limitation in Proposition 2.7 still applies to general supports.

Purpose of the equivalences. These results identify the exact logical content of the proposed hypotheses. They do not settle half-membership, nor do they make an equivalent formulation useless: a different expression of the same condition may expose an invariant or a more convenient induction. The formalisation task is separate from the unproved arithmetic assertion on either side.

3.5 Why zero windows do not split the support class

One proposed split asks whether every infinite support satisfies

c_A has arbitrarily late super-logarithmic zero windows
or a weighted block certificate exists at every precision.

For the present divisor counts, the first alternative never occurs. Writing $a_0 = \min A$, every run of zeros has length at most $a_0 - 1$, since each a_0 consecutive integers contain a multiple of a_0 . The logarithmic estimate [LEAN SOURCE](#) therefore does not split the support class into two useful cases. The lcm-gap theorem [LEAN SOURCE](#) remains an independent irrationality result, not an instance of this zero-window alternative. The proposed assertion consequently requires block certificates for every infinite support: for each required precision, one must construct an opening block with $2^K \mid \sum_{r=1}^K c_A(N+r)2^{K-r}$, a bounded middle contribution C , and $q(C+N+L+2) < 2^L$.

This reformulation does not remove the CRT compatibility and tail-estimate requirements of Section 2.7. It asks for a new construction that meets them for general supports. Nor does it identify that construction with the reset estimate used for the half-target: the two arguments use related divisor-count identities, but no equivalence of their missing hypotheses is proved here. [Open]

3.6 Comparison of the remaining inputs

Route 4 packages recorded equivalences; it does not establish their unproved existence hypotheses. Route 3 must be compared with the external prime-support theorem of Tao and Teräväinen [7, Theorem 1.3]; reproducing it in this framework requires its own analytic input. Route 2 offers a different finite search space under the conditions stated there. Routes 1 and 5 share some incidence identities, but this is not a proof that every missing hypothesis is equivalent. The appropriate comparison is the input each route needs, its quantifiers, and what a successful estimate would imply.

4 Reading the detailed record

The opening note has already defined the evidence bands, scale tags, coordinates, and catalogue terminology. The longer record has two different reading modes, and they should not be confused.

Read linearly. Sections 2 and 3 form one argument: the first identifies the recurring obstructions, and the second isolates the routes those obstructions do not exclude. Read them in order before selecting a new attack.

Use as reference. The mathematical-ingredient catalogue may be entered by coordinate: the certificate method and universal-direction record; the greedy and achievement-set geometry; the integer-quotient model, quotient identity, and dynamics; the half-carry and Boolean–Möbius coordinates; and the support rigidity results. The later implication tables record what chains with what, while the near-miss index gives the exact mismatch at each open obligation.

Scale and coordinate discipline. The direct certified instances remain fixed or bounded; rungs $J \leq 22$, seam rows $\leq 200,000$, ranks ≤ 3000 , and macro depth $\leq 13,548,057$; whereas the open obligations are cofinal. Theorem 2.11 shows that extending a bounded range does not establish a cofinal statement. Coordinate tags are equally literal: an obstruction in *half-carry recurrence*, *quotient-row*, or *diophantine* does not transfer to another representation without a proved map. Coordinate-free barriers are marked explicitly.

III. Detailed statement and arithmetic catalogue

5 The series, the greedy rule and finite identities

5.1 The object

In this part, write $x_n = w_n = (2^n - 1)^{-1}$ for every positive integer n , and $x_A = X_A(2) = \sum_{n \in A} x_n$ for $A \subseteq \mathbb{N}_{>0}$. These series converge by comparison with $\sum_n 2^{-n}$. Exponent 1 contributes the rational number $x_1 = 1$, so deleting or adjoining it does not affect irrationality. It cannot occur in a representation of $1/2$; those calculations therefore start at exponent 2.

Definition 5.1 (Universal Problem #257). **Is $x_A = \sum_{n \in A} 1/(2^n - 1)$ irrational for every infinite $A \subseteq \mathbb{N}_{>0}$?** The quantifier ranges over all infinite sets of positive exponents. It is **OPEN**. [Open]

The support theorems prove irrationality for specified infinite families. The following question instead asks whether one particular rational number is represented. A positive answer would disprove the universal assertion; a negative answer would settle only that target. The later analysis also treats $1/21$.

Definition 5.2 (The distinguished half-value question). Define the *Mersenne achievement set* $\mathcal{A} := \{y \in \mathbb{R} : \exists A \subseteq \mathbb{N}, 0 \notin A, y = \sum_{n \in A} x_n\}$, the set of reals realizable as a subset sum of the positive-index Mersenne weights $\{x_n\}_{n \geq 1}$. **Is $1/2 \in \mathcal{A}$?** This is **OPEN**. [Open]

Observation 5.3 (Why the half-value question controls universal #257). Any finite subset sum of $\{x_n\}$ has odd reduced denominator, since every $2^n - 1$ is odd; hence no finite A can sum to $1/2$ ([LEAN SOURCE](#)). Consequently, if $1/2 \in \mathcal{A}$, the witnessing support A is necessarily infinite, and $x_A = 1/2$ is rational. Such an A is then a single infinite set for which the universal-#257 sum is rational – an outright refutation of universal #257. Conversely, universal #257 being true forces $1/2 \notin \mathcal{A}$ as one instance among uncountably many. The half-value question is therefore one concrete test of the negation, not a restatement of the universal problem. Excluding $1/2$ would not exclude every other rational value, even though the support theorems already settle many infinite families. Representing $1/2$ would give an infinite-support counterexample to the universal assertion. (n/a) [Math]

5.2 What is proved

5.2.1 Full-support irrationality, for every base

Theorem 5.4 (Full-support irrationality, unconditional, every base $b \geq 2$). *For every integer $b \geq 2$,*

$$\sum_{k=0}^{\infty} \frac{1}{b^{k+1} - 1} \text{ is irrational.}$$

Hypotheses: $b \in \mathbb{N}, b \geq 2$. **Conclusion:** unconditional irrationality of the full-support Erdős–Borwein-type series at every base, not merely base 2. full-support (uniform) [Lean]
 Source: [LEAN SOURCE](#) ($'b : \mathbb{N}, hb : 2 \leq b'$), recorded via a certificate machine (bounded Bertrand/CRT first-block frame, middle-window divisor-pair average with pigeonhole selection, explicit parameter closure). The base-2 instance recovering the classical Erdős–Borwein constant's irrationality is [LEAN SOURCE](#), a one-line corollary.

This formalizes the positive-integer-base scope of Erdős's classical 1948 full-support theorem, which already covered integer bases of absolute value greater than one. The base-two Erdős–Borwein constant is a particular case; the Lean theorem is recorded with no open hypotheses.

5.2.2 Structured infinite supports

Beyond full support ($A = \{1, 2, 3, \dots\}$, i.e. $m = 1$ below), the following named infinite families are proved irrational for every integer base $b \geq 2$. The purely periodic, eventually periodic, residue-class and odd supports follow from the periodic theorem of Luca and Tachiya [3, Theorem A, p. 139], and the declarations below are formal proofs of these cases:

Theorem 5.5 (Purely periodic support). For integers $b \geq 2$ and $m \geq 1$, and every m -periodic $A \subseteq \mathbb{N}$ (i.e. $n + m \in A \Leftrightarrow n \in A$ for all n) containing a positive element, the support series $\sum_{a \in A, a \geq 1} (b^a - 1)^{-1}$ is irrational. **Hypotheses:** $b \geq 2, m \geq 1, A$ m -periodic, $\exists a > 0, a \in A$. periodic-support (uniform) [Lean]

Source: [LEAN SOURCE](#).

Theorem 5.6 (Eventually periodic support). For integers $b \geq 2$ and $m \geq 1$, if $A \subseteq \mathbb{N}$ is infinite and its membership is m -periodic from some threshold N_0 on (i.e. $n + m \in A \Leftrightarrow n \in A$ for all $n \geq N_0$), then $\sum_{a \in A, a \geq 1} (b^a - 1)^{-1}$ is irrational, by transferring irrationality across the finite symmetric difference from the shifted purely periodic set $A_{\text{pure}} := \{n : n + mN_0 \in A\}$. **Hypotheses:** $b \geq 2, m \geq 1, N_0 \in \mathbb{N}$, eventual m -periodicity from N_0 , A infinite. eventually-periodic-support (uniform) [Lean]

Source: [LEAN SOURCE](#).

Theorem 5.7 (Residue-class support). For integers $b \geq 2, m \geq 1$ and c , the series $\sum_{n \geq 1, n \equiv c \pmod{m}} (b^n - 1)^{-1}$ is irrational. residue-class-support (uniform) [Lean]

Source: [LEAN SOURCE](#) (specializes purely periodic support).

Lean: `irrational_residueClass_positive_support`.

Theorem 5.8 (Odd support). For every integer $b \geq 2$, the series $\sum_{n \geq 1, n \text{ odd}} (b^n - 1)^{-1}$ is irrational. This density-1/2 support is the case treated explicitly by Luca and Tachiya [3, Example 2, p. 140]. odd-support (uniform) [Lean]

Source: [LEAN SOURCE](#) (specializes residue-class support at $m = 2, c = 1$).

A support representing $1/2$ cannot be eventually periodic, by these irrationality theorems. This is a restriction on a hypothetical representation, not a proof that the actual greedy support is nonperiodic: if the greedy construction fails, it eventually selects every exponent. The periodic-support results therefore decide neither alternative for $1/2 \in A$.

5.2.3 Topology and measure of the achievement set

Theorem 5.9 (Achievement-set topology and measure). *A (Definition 5.2) is compact, closed, perfect, totally disconnected, and nowhere dense; its Lebesgue measure is exactly 1: $\text{volume}(A) = 1$. **Hypotheses:** none. **Conclusion:** A is a Cantor set of positive measure, often called a fat Cantor set. Its measure is 1, not the length of its ambient interval $[0, E]$, where $E \approx 1.6067$. achievement-set-topology (n/a) [Lean]*

Source: [LEAN SOURCE](#), [LEAN SOURCE](#), [LEAN SOURCE](#), [LEAN SOURCE](#), [LEAN SOURCE](#), [LEAN SOURCE](#). Strict separation and summability give the compactness, unique coding, and Cantor topology, and Hornich's strict-tail theorem, as proved by Nitecki [4, Theorem 4(1), p. 9], gives the measure as $\lim_N 2^N R_N$ with $R_N = \sum_{n \geq N} x_n$: each level- N cylinder has length R_N , and the 2^N disjoint cylinders have total length $2^N R_N$. The Mersenne-specific input is $2^N R_N \rightarrow 1$. Other weight sequences require their own tail asymptotic.

Use: Theorem 5.10 below (compactness is exactly what powers every "limit of achieved points is achieved" argument used downstream, including the seam-limit route of Part 2).

Lean: [paper_achievement_geometry](#), [mersenne_topology_quantitative](#).

Theorem 5.10 (Membership equals greedy survival at every level). *For a real x , let $r_n(x)$ be its greedy remainder after rank n , and let $R_n = \sum_{j \geq n} w_j$, as in the initial notation. Then*

$$x \in A \iff x \geq 0 \text{ and } r_n(x) \leq R_n \text{ for every } n \geq 0.$$

The inequality says that the unselected tail has enough total mass at every rank. A failure at any one rank excludes membership; the inequalities at all ranks give a representation. greedy-survival (uniform) [Lean]

Source: [LEAN SOURCE](#) (forward direction [LEAN SOURCE](#), reverse [LEAN SOURCE](#)).

5.3 The greedy rule for one half

Fix the target $1/2$ and the weights $x_n = 1/(2^n - 1)$. Here $x_1 = 1$, so exponent one cannot occur in a representation of $1/2$ by positive summands; the half-target discussion starts at exponent two. (The reciprocal formula is undefined only at exponent zero.) The greedy rule is deterministic. Strict tail domination makes it recover the support of any represented target (Theorem 5.10):

Definition 5.11 (The half-value greedy orbit). Let $\rho_k = r_{k-1}(1/2)$ be the residual entering rank k , so $\rho_2 = 1/2$. In this subsection $x_k = w_k$ and $T_{k+1} = R_k = \sum_{j \geq k} w_j$. At rank $k \geq 2$:

- **Take** iff $\rho_k \geq x_k$; if taken, $\rho_{k+1} := \rho_k - x_k$.
- **Skip** otherwise; $\rho_{k+1} := \rho_k$.

A skip at rank k is:

- **safe** iff $\rho_k \leq T_{k+1} := \sum_{j \geq k} x_j$ (the total remaining mass does not yet rule out a representation; this inequality alone does not guarantee one);
- **fatal** iff $\rho_k \in (T_{k+1}, x_k)$ (the residual is too large for the remaining tail to ever reach, yet too small to have been taken at k – an impossible gap).

For example, ranks 2 and 3 are selected, leaving $1/2 - 1/3 - 1/7 = 1/42$. Ranks 4 and 5 are skipped, while rank 6 is selected and leaves $1/126$. These decisions illustrate the rule; they do not establish its infinite survival.

By Theorem 5.10, $1/2 \in A$ iff no rank of this orbit is ever fatal. Two boundary degeneracies are excluded unconditionally and independently of any open hypothesis:

Remark 5.12 (No-ties: both greedy-orbit boundaries are strict). (a) $\rho_k \neq x_k$ at any rank: a finite Mersenne sum has odd reduced denominator (each $2^n - 1$ is odd), so the residual $1/2$ minus a finite prefix cannot equal the odd-denominator number x_k , sharpening the general denominator-parity fact used in Observation 5.3. (b) $\rho_k \neq T_{k+1}$ at any rank: the tail $T_{k+1} = E - 1 - \sum_{2 \leq j \leq k} x_j$, where $E := \sum_{n \geq 1} x_n$ is the Erdős–Borwein constant, is irrational by Erdős’s own 1948 theorem (the base-2 case of Theorem 5.4), while ρ_k is always rational; equality is impossible. **Conclusion:** both greedy-orbit boundaries are provably strict at every rank, so the orbit’s decisions stabilize after finitely many strict inequalities at any fixed truncation depth. *no-ties* (uniform) [Math]

5.4 Integer quotients and their remainder identity

The quotient rows are finite integer approximations to the real greedy rule, not identical finite stages of that rule. Their normalised weights tend to the Mersenne weights and their targets tend to $1/2$. Proposition 8.1 proves agreement at each fixed prefix once the row index is sufficiently large; it does not assert simultaneous agreement at all ranks. Fix a row $n \geq 6$. The integer quotient weights at row n are

$$w(n, d) := \left\lfloor \frac{4^n}{2^d - 1} \right\rfloor, \quad d = 2, \dots, n-1,$$

and the integer target is $T_n := 2^{2n-1} - 2^n$. The integer greedy process takes ranks $d = 2, \dots, n-1$ in ascending order against T_n exactly as in Definition 5.11, producing a take set $D_n \subseteq \{2, \dots, n-1\}$ and skip set $\text{Skip}_n := \{2, \dots, n-1\} \setminus D_n$. Define the *integer remainder* $\text{rem}(n) := T_n - \sum_{d \in D_n} w(n, d)$ and the *scaled remainder deviation* $\Delta_n := \text{rem}(n) - 2^n$.

Definition 5.13 (The full-support quotient recurrence). For $M \geq 2$, define

$$K(M) := 2^{M-1} - \sum_{d=2}^M \left\lfloor \frac{2^M}{2^d - 1} \right\rfloor.$$

This depends only on M , not on the selected support in a quotient row. The floor-quotient recurrence gives

$$K(2) = 1, \quad K(M+1) = 2K(M) - (\tau(M+1) - 1).$$

Here $\tau(n)$ counts the positive divisors of n . To see the correction term, double the quotients for $2 \leq d \leq M$ and add one for each such d dividing $M+1$. The new endpoint $d = M+1$ contributes one more quotient. Together these are exactly the $\tau(M+1) - 1$ divisors at least two. This is an identity for every $M \geq 2$, not merely a checked finite recurrence. *universal-recursion* (uniform) [Math][Cert]

Theorem 5.14 (Decomposition of the integer remainder). *For every quotient row $n \geq 6$,*

$$\Delta_n = K(2n) + \sum_{d \in \text{Skip}_n} \left\lfloor \frac{4^n}{2^d - 1} \right\rfloor. \quad (\text{I})$$

Here the take and skip sets, capacity and remainder are those defined at the start of this subsection. The identity separates the contribution of the full exponent range from that of the omitted ranks. It makes no claim that there are few omitted ranks; controlling their sum is a separate part of any application. quotient-row ; integer-greedy (Mersenne weights, binary digits of K) (fixed) [Math][Cert]
Proof. For $d = n$ the quotient is $2^n + 1$, and for $n < d \leq 2n$ it is 2^{2n-d} . Hence

$$\sum_{d=n}^{2n} \left\lfloor \frac{4^n}{2^d - 1} \right\rfloor = 2^n + 1 + \sum_{j=0}^{n-1} 2^j = 2^{n+1}.$$

Substitute this in the definition of $K(2n)$, split $\{2, \dots, n-1\} = D_n \sqcup \text{Skip}_n$, and use $\Delta_n = 2^{2n-1} - 2^{n+1} - \sum_{d \in D_n} w(n, d)$. This gives (I). **Certification:** verified rows $6 \leq n \leq 200$, zero mismatches, reported as reproduced in two separate computations (§5.7).

Use: (I) permits substitution of information about $K(2n)$. A resulting estimate for Δ_n must also retain or bound the skip-set sum; information about K alone need not control their cancellation. Lean: [paper_master_identity_floors](#).

5.5 A real-valued form of the quotient identity

Theorem 5.15 (Real (non-integer) form of the quotient identity). *Let $x_d := 1/(2^d - 1)$ and define the constant*

$$C := \sum_{d \geq 2} x_d - \frac{1}{2} = E - \frac{3}{2} = 0.1066951524152917 \dots,$$

where E is the Erdős–Borwein constant. Then, for every quotient row $n \geq 6$,

$$\Delta_n = 4^n \left(\sum_{d \in \text{Skip}_n} x_d - C \right) + \eta_n, \quad |\eta_n| < 2n + 2. \quad (\text{II})$$

Proof. Splitting the full series at n gives

$$\eta_n = 4^n R_{n-1} - 2^{n+1} + \sum_{d \in D_n} \left\{ \frac{4^n}{2^d - 1} \right\}.$$

The tail bounds $2^{-(n-1)} < R_{n-1} \leq 2^{-(n-1)} + \frac{2}{3}4^{-(n-1)}$ and $|D_n| \leq n - 2$ show $0 < \eta_n < n + 2/3 < 2n + 2$.

Use of the error bound. For $H > 2n+2$, the condition $|\sum_{d \in \text{Skip}_n} x_d - C| > (H+2n+2)/4^n$ is sufficient for $|\Delta_n| > H$. Conversely, $|\Delta_n| > H$ implies $|\sum_{d \in \text{Skip}_n} x_d - C| > (H - 2n - 2)/4^n$. Thus a square-root-exponential deviation corresponds to a separation of order $2^{-3n/2}$, with the displayed additive error retained. This is not an exact equivalence after simply discarding that error. real approximation and integer remainder (uniform) [Math]

Note: $C = E - 3/2$ is Mersenne-specific, but the shape “scaled deviation = $4^n \cdot (\text{finite skip-sum} - \text{target constant}) + O(n)$ ” is a template, not yet matched to any analogous constant on the #249 side.

Lean: [paper_real_quotient_core](#), [paper_real_quotient_margins](#), [row_constant_eq_tail](#), [mersenne_constant_decimal](#).

5.6 Dynamics

Theorem 5.16 (Exact recurrences for the three branches). *Fix $n \geq 5$. Among the quotient sums over subsets of $\{2, \dots, n-1\}$, let D_n give the largest sum at most T_n and let B_n give the smallest sum strictly greater than T_n . Thus D_n is the greedy support already defined. Put*

$$r = \text{rem}(n), \quad o = Q(B_n, 2n) - T_n,$$

and define the two correction terms by

$$p^- = \sum_{d \in D_n} (2\mathbf{1}_{d|2n+1} + \mathbf{1}_{d|2n+2}), \quad p^+ = \sum_{d \in B_n} (2\mathbf{1}_{d|2n+1} + \mathbf{1}_{d|2n+2}).$$

Both lie in $[0, 2(n-2)]$, since $d \geq 2$ cannot divide both consecutive integers. The next remainder is

$$\text{rem}(n+1) = \begin{cases} 2^{n+1} - 4o - p^+, & \text{if } 4o + p^+ \leq 2^{n+1} \quad (\text{U}), \\ 4r + 2^{n+1} - p^-, & \text{if not U and } 4r + 2^{n+1} - p^- < 2^{n+2} + 4 \quad (\text{M}), \\ 4r - 2^{n+1} - p^- - 4, & \text{otherwise} \quad (\text{R}). \end{cases}$$

Consequently, for $\lambda_n = (\text{rem}(n) - 2^n)/2^n$,

$$\lambda_{n+1} = \begin{cases} -(4o + p^+)/2^{n+1}, & U, \\ 2\lambda_n + 2 - p^-/2^{n+1}, & M, \\ 2\lambda_n - (p^- + 4)/2^{n+1}, & R. \end{cases} \quad (\text{III})$$

integer quotient recurrences(uniform)[Math]

Lean: [paper_dynamics](#), [paper_rowLower_existsUnique](#), [paper_rowUpper_existsUnique](#), [paper_greedySupport_isRowLower](#), and 11 further declarations in the [coverage section](#).

Proof. Applying the one-step quotient identity twice gives

$$q(2n+2, d) = 4q(2n, d) + 2\mathbf{1}_{d|2n+1} + \mathbf{1}_{d|2n+2}.$$

The new capacity is $T_{n+1} = 4T_n + 2^{n+1}$, and the new terminal weight, at exponent n , is $q(2n+2, n) = 2^{n+2} + 4$. Write $g = 2^{n+1}$. The old quotient sums are separated by at least g , so the corrected sums remain separated by at least $4g - 2(n-2) > 2g + 4$. This exceeds the new terminal weight. The upper word exists: already the terms at $d = 2, 3, 4$ have sum greater than T_n for $n \geq 5$, since $1/3 + 1/7 + 1/15 = 19/35 > 1/2$ and their total floor error is less than 3. Consequently, adding the terminal weight to a lower old word cannot pass the next corrected old word. Moreover, the increment $T_{n+1} - 4T_n = g$ is smaller than this separation, so no old word beyond the adjacent upper one can fit. It is therefore enough to compare the two adjacent old words and the terminal weight. The upper word fits precisely under the condition U; its remainder is then too small to take the terminal weight. Otherwise the lower word is used, and comparison with the terminal weight gives M or R. Substitution yields the displayed formulas. The supplied release gives the corresponding [abstract adjacent-cut identity](#) and [its identification with the concrete greedy row](#). Both are listed as `release_only` in the supplied index; these links identify the inspected source, not a new kernel replay. $\square$

Only the M and R formulas have the asserted small correction to a fixed affine map. To rewrite U in that form, let $\gamma_n = r + o - 2^{n+1} \geq 0$ be the excess of the actual adjacent spacing over its lower bound. Then

$$\lambda_{n+1} = 2\lambda_n - 2 - \frac{4\gamma_n + p^+}{2^{n+1}} \quad (\text{U}).$$

There is no $O(n)$ bound on γ_n supplied by the separation inequality. At $n = 13$, exact calculation gives

$$r = 15147, \quad o = 3997, \quad p^+ = 4, \quad \gamma_{13} = 2760, \quad \text{rem}(14) = 392.$$

Thus $\lambda_{14} - 2\lambda_{13} + 2 = -2761/4096$. In particular, the proposed uniform error bound $|\delta_n| \leq (2n+6)2^{-n}$ in a three-branch signed-digit formula is false already at this row. The adjacent-spacing term cannot be absorbed into the divisor correction.

Iterating the R formula is still legitimate along a run of right branches. Its accumulated correction and endpoint indices must be retained in a reset estimate. These exact recurrences do not identify the proposed ceiling $L_r < (r-3)/2$ with a converse to the conditional reset square-root estimate.

5.7 Reported finite checks

The following values and finite ranges are reported in the supplied record for the quotient-greedy sequence of Section 5.4. They provide test cases for the formulas above. The low-row remainders at $n = 14, 15, 32$ were reproduced by exact integer arithmetic in this revision's accompanying checks. The larger computation ranges remain reports from the supplied record, not new computations in this revision.

Observation 5.17 (Exact low-row values). At row $n = 14$: $\text{rem}(14) = 392$, with take set $D_{14} = \{2, 3, 6, 7\}$. At row $n = 15$: $\text{rem}(15) = 34333$. At row $n = 32$: $\text{rem}(32) = 3865046005$ and $\Delta_{32} = -429921291$. *quotient row, exact integer* (fixed) [Cert]

Source: the linked proof [LEAN SOURCE](#) uses exact successor remainders at rows 14–31 to verify the band-escape condition for reset indices $13 \leq d \leq 30$. The two ranges differ by one: the remainder used for reset d is at row $d + 1$. The source uses `decide+kernel`; its Lean build was not replayed for this revision.

Observation 5.18 (Certified computation depth, all four coordinates). The quotient identity (Theorem 5.14) is verified rows $6 \leq n \leq 200$ with zero mismatches, reported as reproduced in two separate computations, and the quotient-greedy orbit itself is extended and certified to row 2500 (and separately, in a longer single run, to row 200,000 – §SE-7 of the underlying record, 207.7s runtime, branch counts R:M:U = 100197:49899:49898, TARGET failing at exactly one reset row $\leq 200,000$, namely $r = 7$). *multiple: quotient row, binary digit stream of C, greedy set G* (fixed) [Cert]

Observation 5.19 (The four coordinate depths, side by side). Four related computations are reported, with different targets or integer representations, at the following ranges: truncation rung $J = 3, \dots, 22$ (all 20 rungs proved to survive, not merely computed – Part 3 of the truncation-rung record); quotient row $6 \leq n \leq 2500$ (Theorem 5.14 and the branch data underlying Theorem 5.16); integer margin scan to $m_c = 5 \times 10^5$ (advisory

Type B scan, not yet promoted to a certified theorem); sharp tail margin, ranks 2 through 3000 (no fatal skip at any rank; take fraction 0.4992; tightest certified take-margin 0.0340 bits at rank 7, the near-tie $1/126 \geq 1/127$). *all four coordinates* (fixed) [Cert]

Use: truncated weights define a different greedy problem, and the quotient-row construction uses integer weights rather than the original real weights. These are not automatically four readings of the same decision stream. Their relevance to full-weight membership comes from the approximation, limit, or identification theorems cited for each construction. A certificate at one range does not transfer to another without such a theorem, and these reports do not decide the universal problem.

The large computation ranges in this section are inherited from their identified sources. The source tags retain the distinction between formal proofs and reported finite checks; neither should be read as a new replay of those computations.

6 Detailed results and their hypotheses

This section catalogues machinery for Erdős #257 (is $\sum_{n \in A} 1/(2^n - 1)$ irrational for every infinite $A \subseteq \mathbb{N}_{\geq 1}$?) drawn from the half-membership route (is $1/2$ in the Mersenne achievement set?) and the support-rigidity / full-support route (irrationality of $\sum_{a \in A} 1/(2^a - 1)$ for named families A). Every entry states exact hypotheses, exact conclusion, coordinate, scale, evidence band, and Lean source location. None of these results decides #257. Entries are grouped into **conditional implications** (with any unproved hypothesis identified) and **equivalences and exact identities** (including structural properties, definitions and exact finite values). The scale labels distinguish a single instance, a bounded range, unbounded depths and a uniform assertion. They are descriptions of quantifiers, not a total ordering of the strength of the results.

6.1 Conditional membership tests

The finite tests below use $g_n = w_n - R_n$. When written as functions, $R(n) = R_n$, $g(n) = g_n$, $r(x, n) = r_n(x)$ and $C(A, N) = C_A(N)$. For a finite set D , $V(D) = X_D(2)$. For $n \geq 2$, the integer remainder $s(n) = \text{rem}(n)$ uses weights $q(2n, d)$, $2 \leq d < n$, and target $T_n = 2^{2n-1} - 2^n$; it is not the real remainder $r_n(1/2)$. In the branch conditions below, $D_s, B_s, o_s, p_s^-, p_s^+$ denote the lower and upper supports, overshoot and correction terms of Theorem 5.16, with the row added as a subscript. The labels U, M, R mean precisely its three arithmetic cases. Finally, $F(k, J) = F_k(J)$ denotes the expression in Theorem 6.13.

Theorem 6.1 (Infinitely many greedy skips force membership). *Fix $x \geq 0$. If $r_n(x) > R_n$ at some rank n , every later rank is selected and $r_{n+k}(x) > R_{n+k}$ for every $k \geq 0$. Consequently, infinitely many omitted positive ranks imply $x \in A$.*

[Lean](cofinal)greedy recurrence [LEAN SOURCE](#)

Lean: [fatal_absorbing](#).

Proof. The identity $R_n = w_{n+1} + R_{n+1}$ gives $r_n(x) > w_{n+1}$, so the rule takes rank $n + 1$ and leaves $r_{n+1}(x) = r_n(x) - w_{n+1} > R_{n+1}$. Induction proves persistence. Infinitely many skips therefore force $0 \leq r_n(x) \leq R_n$ at every rank; since $R_n \rightarrow 0$, the greedy sum equals x . $\square$

Theorem 6.2 (The quotient remainders converge to the greedy deficit). *Let G be the real greedy support for $1/2$, and let D_s be the integer-greedy support at row s . Then*

$$X_{D_s}(2) \rightarrow X_G(2), \quad \frac{\text{rem}(s)}{4^s} \rightarrow \delta := \frac{1}{2} - X_G(2) \geq 0.$$

In particular, the following conditions are equivalent: $1/2 \in A$; the full sequence $\text{rem}(s)/4^s$ tends to zero; and there exist $s_j \rightarrow \infty$ with $\text{rem}(s_j)/4^{s_j} \rightarrow 0$.

[Math](cofinal)integer quotients

Lean: [paper_seam_limit_unconditional](#), [tendsto_seamGreedyFiniteValue_greedyHalfTargetValue](#), [tendsto_seamGreedyNormalizedRemainder](#), [eventually_seamSupport_agrees](#), and 3 further declarations in the [coverage section](#).

Proof. For each fixed d , the scaled weight $4^{-s}q(2s, d)$ tends to w_d , and the scaled target tends to $1/2$. The finite-prefix stability argument of Proposition 8.1 therefore gives $D_s \cap \{2, \dots, K\} = G \cap \{2, \dots, K\}$ for all sufficiently large s , for each fixed K . At such rows,

$$|X_{D_s}(2) - X_G(2)| \leq R_K.$$

First letting $s \rightarrow \infty$ and then $K \rightarrow \infty$ proves convergence of the values. This uses summability, not an unproved uniform rate of prefix agreement.

For $s \geq 5$, put $\Phi_s = \sum_{d \in D_s} \{4^s / (2^d - 1)\}$, so $0 \leq \Phi_s < s - 2$. The quotient identity gives exactly

$$X_{D_s}(2) = \frac{1}{2} - 2^{-s} - \frac{\text{rem}(s)}{4^s} + \frac{\Phi_s}{4^s}.$$

The two error terms tend to zero, proving the asserted limit. The real greedy sums never exceed $1/2$, so $\delta \geq 0$. By strict tail separation, a support representing $1/2$ must agree with the greedy rule, as in Theorem 6.35. Thus membership is equivalent to $\delta = 0$, which is also equivalent to a zero limit along any one sequence $s_j \rightarrow \infty$. $\square$

The source theorem [LEAN SOURCE](#) proves the original forward implication from cofinal vanishing to membership, using $0 \leq 1/2 - X_{D_s}(2) \leq 2^{-s} + \text{rem}(s)/4^s$ and closedness: [LEAN SOURCE](#). The full-limit identification and converse above are ordinary arguments, not a claim about that declaration's formal statement. Convergence itself is automatic; proving that its limit is zero is exactly the unresolved membership question.

Theorem 6.3 (Two-channel and dyadic cap sufficiency). *Suppose that for every $n \geq 0$ at which the next weight is skipped,*

$$w_{n+1} > r_n(1/2) \implies r_n(1/2) \leq 2^{-(n+1)} + \frac{1}{3} 4^{-(n+1)}.$$

Then $1/2 \in A$. The stronger bound $r_n(1/2) \leq 2^{-(n+1)}$ at the same skipped ranks also suffices. The first bound retains two positive geometric terms of the full tail; it is larger, and hence less restrictive, than the second. Both are sufficient tests on the specified greedy orbit. The result does not assert that either bound holds at all its skipped ranks.

[Lean](uniform)two-channel-dyadic-cap [LEAN SOURCE](#) [LEAN SOURCE](#)

Theorem 6.4 (Second-channel phase separation sufficiency). Define $P_n = 4^n(2r_n(1/2) - 2^{-n})$. If

$$\frac{1}{6} + \frac{37}{56} 2^{-n} \leq \left| P_n - \frac{1}{3} \right| \quad \text{for every } n \geq 1,$$

then $1/2 \in A$. These are exact rational inequalities because the target and every finite greedy remainder are rational. The cited finite calculation establishes them for $1 \leq n \leq 6$; it therefore suffices to prove them for every $n \geq 7$.

Unlike Theorem 6.3, this hypothesis concerns every positive rank, not only skipped ranks. It excludes an explicit interval around $1/3$. For example, $P_n \notin (0, 1)$ implies the displayed inequality when $n \geq 2$, but no such avoidance theorem for the whole orbit is proved here.

[Lean](uniform)second-channel-phase [LEAN SOURCE](#) [LEAN SOURCE](#)

Theorem 6.5 (Straddle-prefix closed-set criterion). Suppose that for every $d \geq 0$ there is a finite support $D_d \subseteq \{1, \dots, d\}$ such that

$$X_{D_d}(2) \leq t \leq X_{D_d}(2) + R_d.$$

Then $t \in A$. The supports need not agree at different depths: compactness, together with $R_d \rightarrow 0$, supplies a limiting support. Conversely, truncating a support representing t supplies such finite supports. Thus this condition is an exact finite-approximation formulation of membership, not a criterion already known to hold for an unresolved target.

[Lean](uniform)dedekind-cut [LEAN SOURCE](#)

Lean: `straddle_all_depths_iff_mem, straddle_limiting_support_inputs`.

Theorem 6.6 (Keeping the largest skipped rank beyond two thirds). Let $D_s \subseteq \{2, \dots, s-1\}$ be the support of the integer-greedy row defined in Section 5.4. Assume the following implication for every $s \geq 14$ and every largest omitted rank $d = \max(\{2, \dots, s-1\} \setminus D_s)$:

$$2s < 3d \implies 2(s+1) < 3d \quad \text{or} \quad s \notin D_{s+1}.$$

Then $1/2 \in A$. The hypothesis says that either the same omitted rank remains beyond two thirds of the next row, or the next row omits its terminal rank. The latter alternative is exactly an upper or middle transition in the preceding classification.

The cited finite calculation establishes the initial inequality at row 14. The displayed implication then preserves the inequality by induction and supplies omitted ranks tending to infinity. It is an unproved condition on every later row; the verified base case alone does not establish it.

[Lean](uniform)integer quotients [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: `half_mem_mersenneAchievementSet_of_largestSkipLateStepSocket, largestSkipLateAt_fourteen, seamGreedy_terminal_false_iff_upperOrMiddle`.

Theorem 6.7 (Two sufficient bounds at middle transitions). Suppose that at every middle transition with row $s \geq 13$,

$$|D_s| + p_s^- + 5 < 4 \operatorname{rem}(s).$$

Then $1/2 \in A$. The stronger condition $\operatorname{rem}(s) \geq s$ at all such rows also suffices.

[Lean](uniform)integer quotients [LEAN SOURCE](#) [LEAN SOURCE](#)

Indeed, $|D_s| \leq s - 2$ and $p_s^- \leq 2(s - 2)$ give $|D_s| + p_s^- + 5 \leq 3s - 1 < 4s$. This explains the implication between the two hypotheses. The implication from the first bound to membership is supplied by the cited source. Neither all-row hypothesis is established by the finite computations.

Theorem 6.8 (The tail inequality at a final middle transition). *Suppose row $D \geq 13$ is middle and all rows after it are right. Let $F_D = D_D \cup \{D\}$ and put*

$$C_D = 4 \operatorname{rem}(D) - p_D^- - 4, \quad \Theta_D = \sum_{j \geq 1} c_{F_D}(2D + 2 + j)2^{-j}.$$

Then $C_D < \Theta_D$. Since $0 \leq \Theta_D \leq |F_D|$, an estimate in the opposite direction would exclude this scenario. No such reverse estimate is assumed or proved here. Remark 10.8 states the corresponding sufficient hypothesis with the value -3 excepted.

[Lean](uniform)integer quotients [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: `middleProducer_allRight_forces_carry_lt_tail`, `middleProducer_allRight_forces_rational_skip`, `producerCarry_insert_seamBelowSupport_eq_middleCoordinate`, `binaryCoeffTail_nonneg`, and 1 further declaration in the [coverage section](#).

Theorem 6.9 (A conditional two-sided dyadic bound). *Assume the following two conditions for every row $s \geq 5$:*

$$\begin{aligned} M \text{ at } s &\implies 4 \operatorname{rem}(s) - p_s^- - 4 \notin \{-3, -2, -1\}, \\ R \text{ at } s \text{ and } o_s \leq 2^s &\implies 4o_s + p_s^+ \leq 2^{s+2}. \end{aligned}$$

Then at every row $s \geq 5$,

$$\operatorname{rem}(s) \leq 2^s \quad \text{or} \quad o_s \leq 2^s.$$

The first hypothesis concerns only middle transitions. The second concerns only right transitions at which the old overshoot is at most 2^s . Neither condition is asserted here for the whole sequence.

[Lean](uniform)integer quotients [LEAN SOURCE](#) [LEAN SOURCE](#)

Proof. At row 5, $D_5 = \{2, 4\}$ and $\operatorname{rem}(5) = 71$, whereas $B_5 = \{2, 3\}$ and $o_5 = 7 \leq 32$. For the induction step use Theorem 5.16. The upper branch gives $\operatorname{rem}(s + 1) \leq 2^{s+1}$. On a middle branch put $c = 4 \operatorname{rem}(s) - p_s^- - 4$. If $c \leq -4$, the new remainder is $2^{s+1} + c + 4 \leq 2^{s+1}$. Otherwise the hypothesis forces $c \geq 0$. Adding the terminal weight to the corrected lower word then exceeds the target by $2^{s+1} - c > 0$, so the new minimal overshoot is at most 2^{s+1} .

On a right branch, an old remainder at most 2^s gives a new remainder at most 2^{s+1} by the recurrence. In the other case the inductive hypothesis gives $o_s \leq 2^s$; the corrected upper word exceeds the new target by $4o_s + p_s^+ - 2^{s+1}$, which is positive and at most 2^{s+1} by the second assumption. $\square$

This conclusion is a disjunction. It does not by itself give an upper bound for the remainder on rows where only the overshoot is small, and therefore is not by itself a membership criterion.

Theorem 6.10 (Upper-reset dyadic-band escape, checked for $13 \leq d \leq 30$). `SeamUpperResetDyadicBandEscape` requires the following at every actual upper reset $d \geq 13$: for every $0 \leq j \leq d$, the reset charge avoids a linear-width band immediately below the dyadic power 2^{d-j+1} ($2^{d-j+1} < \text{resetCharge}$ or $\text{resetCharge} + 2(d+j) \leq 2^{d-j+1}$). Granted this, $1/2 \in A$. The linked proof verifies this condition for $13 \leq d \leq 30$, using exact successor remainders at rows 14–31 (for example, $\text{rem}(14) = 392$ and $\text{rem}(31) = 4187487147$). This finite verification does not supply the hypothesis for every $d \geq 13$.

[Lean](uniform)dyadic-boundary [LEAN SOURCE](#) [LEAN SOURCE](#)

Theorem 6.11 (Möbius-centred carry nonnegativity below $1/2$). If $1 \notin A$ and $X_A(2) < 1/2$, then $C_A(N) \geq 0$ for every N . Indeed,

$$\text{ihc}(A, N) = 2^{N+1}(1/2 - X_A(2)) + \sum_{j>N+1} c_A(j)2^{N+1-j} > 0.$$

The carry is an integer, so it is at least 1; subtracting 1 gives the asserted nonnegativity of the centred carry. The strict inequality and integrality are both needed in this argument. The statement is about the divisor counts c_A , not a new irrationality criterion.

[Lean](uniform)divisor counts and finite sums [LEAN SOURCE](#) [LEAN SOURCE](#)

Theorem 6.12 (A square-root bound on the greedy carry). For the greedy support $G = G_{1/2}$, the centred carry is nonnegative. If in addition

$$C_G(N) \leq 2\sqrt{N} + 4 \quad \text{for every } N \geq 0,$$

then G has infinitely many skipped indices and $X_G(2) = 1/2$. The unresolved part is the upper bound along this particular orbit. The carry identity explains its strength: a positive gap $1/2 - X_G(2)$ would contribute a term of order 2^N , which cannot satisfy a square-root bound. This is distinct from the reset-deviation hypothesis of Theorem 6.10; no equivalence between the two hypotheses is asserted.

[Lean](uniform)greedy recurrence [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: `greedy_mobiusCenteredHalfCarry_nonneg`, `greedy_half_infinite_of_mobiusCenteredHalfCarry_sqrtBound`, `infinite_support_half_of_mobiusCenteredHalfCarry_sqrtBound`.

Theorem 6.13 (Equivalent sign and vanishing conditions). For the prefix $D = G \cap \{1, \dots, k\}$ define

$$F_k(J) = \sum_{i=1}^J c_D(k+1+i)2^{J-i} - 2^J c_D(k).$$

This is a finite integer expression: it compares the next J divisor counts of the fixed prefix with its centred carry. Let s_n denote the integer greedy remainder at depth $2n$, with weights $q(2n, d)$ for $2 \leq d < n$ and target $2^{2n-1} - 2^n$. The following conditions each imply $1/2 \in A$:

- (i) For every skipped rank $n \geq 3$, $F_{n-1}(n) \geq 0$.
- (ii) For every skipped rank $n \geq 3$ at which the real and integer greedy words agree on $\{2, \dots, n-1\}$, $s_n = 0$.

(iii) For every skipped rank $n \geq 3$, $B(2n) < s_n$.

Conditions (i) and (ii) are equivalent. Condition (iii) implies them; no converse or strict separation is established here. These hypotheses must hold at every indicated skipped rank, not merely through a finite computed range.

[Lean](uniform)integer quotients [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

Theorem 6.14 (Erdős–Borwein full-support irrationality, unconditional). *For every integer $b \geq 2$, the sum $\sum_{n \geq 1} (b^n - 1)^{-1}$ is irrational. This is Erdős’s theorem [1]; at base 2 the sum is the Erdős–Borwein constant. The result concerns full support, not all its infinite subsets.*

[Lean](uniform)divisor counts and finite sums [LEAN SOURCE](#) [LEAN SOURCE](#)

Theorem 6.15 (Pairwise-coprime support irrationality, Erdős 1968). *For every integer $b \geq 2$ and every infinite pairwise-coprime support $A \subseteq \mathbb{N}_{>0}$ with summable reciprocals, $\sum_{a \in A} 1/(b^a - 1)$ is irrational.*

[Lean](uniform)divisor counts and finite sums [LEAN SOURCE](#)

Theorem 6.16 (Irrationality from divisible coefficient blocks). *Let $b \geq 2$ be an integer and let $c : \mathbb{N} \rightarrow \mathbb{N}$ satisfy $c(n) \leq n$ for every n . Suppose that for every integer $q \geq 1$ there are nonnegative integers N, K, L, C , with $K \leq L$, such that*

$$\begin{aligned} b^r &\mid c(N+r) \quad (1 \leq r \leq K), \\ \sum_{r=K+1}^L c(N+r)b^{L-r} &\leq C, \\ c(N+L+1+t) &> 0 \quad \text{for some integer } t \geq 0, \\ q(C+N+L+2) &< b^L. \end{aligned}$$

Then $\sum_{n \geq 1} c(n)b^{-n}$ is irrational.

[Lean](uniform)binary digits [LEAN SOURCE](#)

Proof. Write $x = \sum_{n \geq 1} c(n)b^{-n}$. Multiplication by b^N makes the terms through N integral. The first condition also makes the terms at $N+1, \dots, N+K$ integral. Subtracting their sum leaves, for some integer z ,

$$0 < b^N x - z = \sum_{r > K} \frac{c(N+r)}{b^r} \leq \frac{C+N+L+2}{b^L} < \frac{1}{q}.$$

Positivity follows from the specified nonzero coefficient beyond L . The middle sum contributes at most C/b^L , while $c(N+r) \leq N+r$ gives

$$\sum_{r > L} \frac{c(N+r)}{b^r} \leq b^{-L} \left(\frac{N+L}{b-1} + \frac{b}{(b-1)^2} \right) \leq \frac{N+L+2}{b^L}.$$

If $x = p/q$ were rational, the positive difference $b^N x - z$ could not be smaller than $1/q$. □

A certificate at one precision is not enough: the same coefficient sequence must satisfy these conditions for every q . The positive-tail condition is essential; without it the zero sequence would satisfy the other conditions at every precision. The full-support theorem supplies such certificates for $c = \tau$. Their existence for $c = \varphi$ is not proved here. This is a limitation of this particular sufficient criterion, not a comparison of every approach to Problem 249.

Theorem 6.17 (Irrationality from a gap beyond the preceding least common multiple). *For an integer base $b \geq 2$ and strictly increasing support $a : \mathbb{N} \rightarrow \mathbb{N}$ with $a(0) \geq 1$: if $a(k) - \text{lcm}(a(0), \dots, a(k-1)) \rightarrow \infty$, then $\sum_k' 1/(b^{a(k)} - 1)$ is irrational. The base may vary over the integers, but the denominators remain $b^{a(k)} - 1$. Applying the argument to another denominator sequence would require its own divisibility and tail estimates.*

[Lean](uniform)lcm-gap [LEAN SOURCE](#)

Theorem 6.18 (Factorial-support and 2^k -support instances). *For every integer $b \geq 2$, both*

$$\sum_{k \geq 0} \frac{1}{b^{(k+1)!} - 1} \quad \text{and} \quad \sum_{k \geq 0} \frac{1}{b^{2^k} - 1}$$

are irrational. These are instances of the preceding theorem. For the factorial support the preceding least common multiple is $k!$ when $k \geq 1$, so the gap is $k k!$. For the powers of two it is 2^{k-1} , so the gap is 2^{k-1} . Both tend to infinity. The conclusion concerns these two supports, not every infinite support.

[Lean](uniform)lcm-gap [LEAN SOURCE](#) [LEAN SOURCE](#)

Theorem 6.19 (Multiples-support irrationality via dilation). *For integers $b \geq 2$ and $d \geq 1$,*

$$X_{d\mathbb{N}_{>0}}(b) = \sum_{k \geq 1} \frac{1}{(b^d)^k - 1}$$

is irrational. This is the full-support result (Theorem 6.14) at the integer base b^d ; no new support argument is required.

[Lean](uniform)dilation [LEAN SOURCE](#) [LEAN SOURCE](#)

Luca and Tachiya proved irrationality for every purely periodic integer weight that is not identically zero [3, Theorem A, p. 139]. The periodic, eventually periodic, residue-class and odd-support statements below follow from their theorem; the declarations are formal proofs of these cases, and the mechanism described is that of the certificate proof.

Theorem 6.20 (Periodic-support irrationality). *Let $b \geq 2$ and $m \geq 1$ be integers. Suppose $A \subseteq \mathbb{N}_{>0}$ is nonempty and satisfies $n \in A$ if and only if $n + m \in A$ for every $n \geq 1$. Then $X_A(b)$ is irrational.*

The indicator of A is a nonzero purely periodic integer weight, so this follows from Luca and Tachiya's theorem stated above. The linked declaration supplies a separate formal proof. At $m = 1$ the only nonempty periodic support is full support; residue classes and unions of residue classes give the other immediate examples.

[Lean](uniform)periodic-sieve [LEAN SOURCE](#)

Theorem 6.21 (Eventually-periodic support irrationality). *Let $b \geq 2$ and $m \geq 1$ be integers. An infinite support whose membership is m -periodic from some threshold N_0 onward has irrational $X_A(b)$.*

[Lean](uniform)periodic-sieve [LEAN SOURCE](#)

Theorem 6.22 (Residue-class and odd-support irrationality). *For integers $b \geq 2$, $m \geq 1$ and a residue c , $\sum_{\substack{n \geq 1 \\ n \equiv c \pmod{m}}} (b^n - 1)^{-1}$ is irrational. Specializing $m = 2, c = 1$: $\sum_{n \text{ odd}} 1/(b^n - 1)$ is irrational for every $b \geq 2$, the case treated explicitly in [3, Example 2, p. 140].*

[Lean](uniform)periodic-sieve [LEAN SOURCE](#) [LEAN SOURCE](#)

Comparison with the known theorem. Luca and Tachiya proved irrationality for every nonzero purely periodic integer weight, of any signs [3, Theorem A, p. 139], so the terminating branch below occurs only for the zero weight. The theorem records the dichotomy reached by the certificate method, together with its one-sided closures.

Theorem 6.23 (A formal dichotomy for signed periodic weights). *Let $b \geq 2$ and $m \geq 1$ be integers, and let $w : \mathbb{N}_{>0} \rightarrow \mathbb{Z}$ be m -periodic. Put*

$$x = \sum_{a \geq 1} \frac{w(a)}{b^a - 1}, \quad c_w(n) = \sum_{d|n} w(d).$$

The cited formal argument gives the dichotomy that x is irrational or $b^k x \in \mathbb{Z}$ for some integer $k \geq 0$. It excludes the latter alternative if c_w has one sign throughout and is nonzero at arbitrarily large indices. These sign assumptions concern c_w , not w . As noted above, Luca and Tachiya's theorem already excludes the terminating alternative for every nonzero periodic w , without either sign restriction. The zero weight gives $x = 0$.

[Lean](uniform)signed-divisor-calculus [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

For the next two results write $M_r = 2^r - 1$ and, for a positive squarefree integer r , set

$$A_r = \sum_{d|r} \mu(d) \frac{r}{d} \frac{M_r}{M_d} \in \mathbb{Z}, \quad B(r) = \frac{A_r}{M_r} = \sum_{d|r} \frac{\mu(d)(r/d)}{2^d - 1}.$$

Every quotient in A_r is integral because $d \mid r$. The source names A_r the Möbius numerator; $B(r)$ is a finite rational sum, not the full-support Lambert series.

Theorem 6.24 (Denominators of finite Mersenne sums). *Let $t \geq 1$ and $h \geq 1$ be integers, and let $r \geq 1$ be squarefree with every prime factor at most t . Let P be a set of prime divisors of r such that $t < 2p$ for each $p \in P$, and put $C = \prod_{p \in P} (2^p - 1)$. Then*

$$\frac{C}{\gcd(C, h)} \mid \text{den}(hB(r)).$$

In particular, C divides this reduced denominator when $\gcd(C, h) = 1$. The factors in C are pairwise coprime, since $\gcd(2^p - 1, 2^q - 1) = 2^{\gcd(p, q)} - 1 = 1$ for distinct primes. The assertion concerns this finite signed sum; it does not replace the approximation hypotheses in Theorem 6.14.

[Lean](uniform)cyclotomic [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: [paper_mersenne_channel_survival](#), [paper_mersenne_channel_survival_of_coprime_scale](#), [paper_channel_factor_gcd_eq_one](#), [paper_channel_factors_pairwise_coprime](#), and 5 further declarations in the [coverage section](#).

Theorem 6.25 (Exponential growth of the reduced denominators). *For an integer $t \geq 5$, put*

$$H_t = \text{lcm}(1, \dots, t), \quad r_t = \prod_{p \leq t} p, \quad h_t = H_t / r_t, \quad D_t = \text{den}(h_t B(r_t)),$$

where the product is over primes. Then $D_t \geq 2^{t/2}$. Its exact denominator is

$$D_t = \frac{2^{r_t} - 1}{\gcd(2^{r_t} - 1, h_t A_{r_t})}.$$

[Lean][Math](uniform)cyclotomic [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: [lcmHeight_scaledMobiusShadow_den_lower_bound](#), [upperHalfMersenneProduct_lower_bound](#), [lcmHeight_scaledMobiusShadow_den_exact](#).

Proof. Apply Bertrand's postulate [24, §§4–5, pp. 197–198] to $m = \lceil t/2 \rceil \geq 3$. There is a prime $m < p < 2m$, hence $p \leq t$ and $p - 1 \geq t/2$. Every prime divisor ℓ of $2^p - 1$ has multiplicative order p for 2 modulo ℓ . Since p is odd, $2p \mid \ell - 1$; thus $\ell > t$. All prime factors of h_t are at most t , so Theorem 6.24 gives

$$2^{t/2} \leq 2^{p-1} < 2^p - 1 \leq D_t.$$

The exact expression for D_t is the usual reduction of the fraction $h_t A_{r_t} / (2^{r_t} - 1)$. $\square$

The first source directly uses natural-number division $t/2$, that is, $\lfloor t/2 \rfloor$. The ceiling choice in the ordinary argument supplies the displayed real-exponent bound also when t is odd. The second source gives an equivalent exact denominator formula in its numerator notation. Neither denominator growth nor that identity supplies an approximation-error estimate.

6.2 Exact identities and reductions

Observation 6.26 (A finite half-value approximation and the first correction tail). The finite prefix has value

$$\frac{1}{2} - X_{\{2,3,6,7\}}(2) = \frac{1}{16002}.$$

The excess of the full tail after rank 1 over its dyadic value is

$$R_1 - \frac{1}{2} = \sum_{n \geq 2} \left(\frac{1}{2^n - 1} - 2^{-n} \right) \leq \frac{1}{12} + \frac{1}{28} = \frac{5}{42}.$$

Indeed, expand each summand as $4^{-n} + \sum_{j \geq 3} 2^{-jn}$; for $n \geq 2$, the second sum is at most $2 \cdot 8^{-n}$. Thus $\frac{1}{2} - (R_1 - \frac{1}{2}) \geq \frac{8}{21} > \frac{3}{8}$. The first identity is an exact rational check, whereas the two tail inequalities use a convergent-series estimate.

[Lean](fixed)exact sums and tail estimates [LEAN SOURCE](#) [LEAN SOURCE](#)

Observation 6.27 (Period-4 sign weight vanishes on a residue class). Let $w(n)$ have successive values $1, 0, -1, 0$, repeated with period 4. For every $n \equiv 3 \pmod{4}$,

$$c_w(n) = \sum_{d|n} w(d) = 0.$$

Indeed, n is odd, and paired divisors $d, n/d$ have opposite values of w because their product is 3 modulo 4. There are no fixed pairs: such an n is not a square. Thus an argument requiring a nonzero coefficient at every chosen residue must check that requirement; periodicity alone does not provide it. This cancellation does not conflict with irrationality of the full weighted series, which follows from Luca and Tachiya's theorem [3, Theorem A, p. 139].

[Lean](fixed)signed-divisor-calculus [LEAN SOURCE](#)

Theorem 6.28 (The Lambert-series identity for the Möbius function). *The absolutely convergent signed series satisfies*

$$\sum_{d \geq 1} \frac{\mu(d)}{2^d - 1} = \frac{1}{2}.$$

This classical identity is also recorded in [6, Example 1.1, p. 4]. Expand each denominator geometrically and interchange the absolutely convergent sums. The coefficient of 2^{-n} is $\sum_{d|n} \mu(d)$, equal to 1 for $n = 1$ and 0 otherwise. This proves the identity. Its coefficients are $-1, 0, 1$, not indicators of a support; the identity therefore does not represent $1/2$ as a subseries with all coefficients in $\{0, 1\}$.

[Cited](fixed)divisor counts and finite sums [LEAN SOURCE](#)

Corollary 6.29 (Negative-Möbius Boolean support overshoots $1/2$). Let $N = \{d \geq 2 : \mu(d) = -1\}$. Isolating the $d = 1$ term in Theorem 6.28 gives

$$X_N(2) = \frac{1}{2} + \sum_{\substack{d \geq 2 \\ \mu(d)=1}} \frac{1}{2^d - 1} \geq \frac{1}{2} + \frac{1}{63} > \frac{1}{2},$$

since $\mu(6) = 1$. This rules out the particular candidate N . It does not rule out other infinite supports representing $1/2$.

[Lean](fixed)divisor counts and finite sums [LEAN SOURCE](#) [LEAN SOURCE](#)

Observation 6.30 (Dyadic safety need not survive a take). For a residual $p/(2D)$, taking rank b subtracts w_b and gives

$$\frac{p}{2D} - \frac{1}{2^b - 1} = \frac{p(2^b - 1) - 2D}{2D(2^b - 1)}.$$

The residual $7/34$ is at most 2^{-2} . Taking rank 3 leaves $15/238$, which is strictly between 2^{-4} and $w_4 = 1/15$. Thus rank 4 is skipped but fails the dyadic safety test. The displayed numerator has increased from 7 to 15; that increase does not preserve safety.

The same failure occurs with an initial denominator containing all Mersenne denominators through rank 7. With $D = \text{lcm}(3, 7, 15, 31, 63, 127) = 1240155$, one has

$$\frac{12149}{2D} - \frac{1}{255} = \frac{41179}{2 \cdot 21082635}.$$

The initial residual is at most 2^{-7} . After this rank-8 take, ranks 9 and 10 are skipped, but the new residual exceeds 2^{-10} . The exact numerator of that dyadic excess is $2^9 \cdot 41179 - 21082635 = 1013 > 0$. These are counterexamples to the stated local safety implication, not assertions that either state occurs on the real half-greedy orbit.

[Cert](fixed)integer quotients [LEAN SOURCE](#) [LEAN SOURCE](#)

Theorem 6.31 (The half-skip dichotomy via Erdős–Borwein irrationality). *The half target satisfies*

$$\frac{1}{2} \in A \iff \mathbb{N}_{>0} \setminus G \text{ is infinite,}$$

where G is the greedy support for $1/2$. In the forward direction, finitely many skipped exponents would express the full Mersenne sum as $1/2$ plus a finite rational sum, contradicting Erdős’s full-support irrationality theorem. The reverse implication uses the greedy tail criterion: once a remainder exceeds the whole available tail, every later exponent is selected. Infinitely many skips exclude that failure. The equivalence does not establish that infinitely many skips actually occur.

[Lean](cofinal)greedy recurrence [LEAN SOURCE](#) [LEAN SOURCE](#)

Theorem 6.32 (Equivalent descriptions of half-membership). *Each of the following is equivalent to $1/2 \in A$:*

- (1) *The integer-greedy sequence is not eventually always on branch R.*
- (2) *There are arbitrarily large $s \geq 5$ with $s \notin D_{s+1}$.*
- (3) *Upper or middle transitions occur at arbitrarily large rows.*
- (4) *Some sequence $s_j \rightarrow \infty$ satisfies $s_j \notin D_{s_j+1}$ for every j .*
- (5) *There are rows $s_j \rightarrow \infty$ and omitted ranks $d_j \in \{2, \dots, s_j - 1\} \setminus D_{s_j}$ with $d_j \rightarrow \infty$.*
- (6) *The positive ranks omitted by the real half-greedy rule form an infinite set.*
- (7) *The real half-greedy rule has no last omitted positive rank.*

[Lean](cofinal)integer quotients [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: `half_mem_mersenneAchievementSet_iff_not_seamGreedyEventuallyRight`, `half_mem_mersenneAchievementSet_iff_unboundedTerminalFalse`, `half_mem_mersenneAchievementSet_iff_unboundedUpperOrMiddle`, `half_mem_mersenneAchievementSet_iff_cofinalTerminalFalse`, and 3 further declarations in the [coverage section](#).

Items (1)–(4) differ only by the branch classification and the choice of a sequence of rows. Item (5) requires the omitted ranks themselves to tend to infinity, not merely their containing rows. The cited identification with the real greedy support supplies (6)–(7). These equivalences do not prove that any of the conditions holds.

Lemma 6.33 (The values on either side of an eventual right continuation). *If the seam eventually always extends “true” (right branch) from some row S on with a fixed lower prefix u , the resulting cofinite-support value stays strictly below $1/2$ (`prefix_add_mersenneTail_lt_half_of_eventually_right`); the matching alternative “upper competitor” word gives a strict excess above $1/2$ (`half_lt_upper_competitor_of_eventually_right`). Neither of these two*

cofinite continuations represents $1/2$. This does not rule out an eventually-right integer orbit; it describes the two values in that case. The inequalities are used in Theorem 6.44.

[Lean](cofinal)integer quotients [LEAN SOURCE](#) [LEAN SOURCE](#)

Lemma 6.34 (Each Mersenne weight exceeds its remaining tail). For $n \geq 1$, let $w_n = (2^n - 1)^{-1}$ and $R_n = \sum_{j>n} w_j$. Then

$$R_n = w_{n+1} + R_{n+1}, \quad 2^{-n} < R_n \leq 2w_{n+1} < w_n.$$

For the upper bound, compare each tail term with $2^{1-j}w_{n+1}$ at index $n+j$, $j \geq 1$, and sum the geometric series. The last inequality follows by comparing $2/(2^{n+1} - 1)$ with $1/(2^n - 1)$. This strict term-versus-tail inequality is the separation hypothesis used in the greedy arguments.

It is not available for the totient weights $\varphi(n)/2^n$: the terms at indices 4 and 5 already sum to $2/16 + 4/32 = 1/4 = \varphi(3)/2^3$, and the remaining tail is positive. A general strict-tail argument therefore cannot be transferred to that sequence without a different hypothesis or proof.

[Lean](uniform)divisor counts and finite sums [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: `mersenneTail_eq_weight_add`, `halfTwoChannelCap_lt_mersenneTail`, `mersenneTail_le_two_mul_weight`, `two_mul_mersenneWeight_succ_lt`, and 1 further declaration in the [coverage section](#).

Theorem 6.35 (Greedy survival and membership). With the greedy remainder $r_n(x)$ and complete tail R_n ,

$$x \in A \iff x \geq 0 \text{ and } r_n(x) \leq R_n \text{ for every } n \geq 0.$$

If all inequalities hold, the nonnegative remainders tend to zero because $R_n \rightarrow 0$, so the greedy partial sums converge to x . Conversely, strict tail domination $w_n > R_n$ forces the greedy choices in any representation and therefore gives every inequality. The argument applies to positive summable weights satisfying that separation hypothesis; it is not a consequence of summability alone.

[Lean](uniform)greedy recurrence [LEAN SOURCE](#)

Lemma 6.36 (The two next-prefix intervals and their gap). Let $d \geq 0$ be an integer, let $u \subseteq \{1, \dots, d\}$, and suppose the target t lies in the interval $[X_u(2), X_u(2) + R_d]$. At depth $d+1$, either t lies in the lower interval

$$[X_u(2), X_u(2) + R_{d+1}],$$

or in the upper interval

$$[X_u(2) + w_{d+1}, X_u(2) + w_{d+1} + R_{d+1}],$$

or in the open gap between them. These alternatives are disjoint because $R_{d+1} < w_{d+1}$, and exhaustive because $R_d = w_{d+1} + R_{d+1}$. For $t = 1/2$, Lemma 6.38 also excludes the interval endpoints, so all relevant comparisons are strict. Endpoint exclusion is not what makes the two child intervals disjoint; the strict-tail inequality does that.

[Lean](uniform)dedekind-cut [LEAN SOURCE](#) [LEAN SOURCE](#)

Lemma 6.37 (A greedy gap excludes every representation). *Let $d \geq 0$ be an integer and $u \subseteq \{1, \dots, d\}$ a finite prefix. If*

$$X_u(2) + R_{d+1} < t < X_u(2) + w_{d+1},$$

then no support agreeing with u through rank d represents t : omitting $d + 1$ leaves value at most the lower endpoint, while including it gives value at least the upper endpoint. Other length- d prefixes have disjoint containing intervals by the first-difference argument using $w_n > R_n$. Since the displayed gap lies inside the interval for u , none of those prefixes can represent t either. Thus a certified strict gap excludes every representation, not just one proposed continuation.

[Lean](uniform)dedekind-cut [LEAN SOURCE](#)

Lean: `fatal_gap_excludes_every_representation, fatal_gap_endpoint_bounds, depth_prefix_interval_disjoint, fatal_gap_within_prefix_interval`.

Lemma 6.38 (No finite Mersenne sum equals one half). *For every finite $u \subseteq \mathbb{N}_{>0}$, $X_u(2) \neq 1/2$: its reduced denominator is odd. Also $X_u(2) + R_d \neq 1/2$ for every $d \geq 0$, since R_d is the irrational full Mersenne sum minus a finite rational sum. These observations exclude equality at the finite-prefix and complete-tail endpoints of Lemma 6.36.*

[Lean](uniform)parity-irrationality [LEAN SOURCE](#) [LEAN SOURCE](#)

Lemma 6.39 (Straddle words are canonical: they agree with the greedy prefix). *Let $u \subseteq \{1, \dots, d\}$ satisfy $X_u(2) \leq 1/2 \leq X_u(2) + R_d$. Then*

$$u = G \cap \{1, \dots, d\},$$

where G is the real greedy support for $1/2$. At the first disagreement, if u takes rank k and the greedy rule skips it, then $X_u(2) > 1/2$. In the opposite case, the greedy remainder before k is at least w_k , whereas the remaining contribution allowed by u , including R_d , is at most $R_k < w_k$. Both cases contradict the displayed interval. Thus there is at most one such word at each depth; the lemma does not assert its existence at every depth.

[Lean](uniform)greedy recurrence [LEAN SOURCE](#) [LEAN SOURCE](#)

Theorem 6.40 (Last-skip iff local fatality: a pointwise criterion). *Keep G for the selected half-greedy support, $r_M(1/2)$ for its remainder and R_M for the complete tail after rank M . A positive rank M is the last omitted rank if and only if $M \notin G$ and $r_M(1/2) > R_M$. Consequently,*

$$\frac{1}{2} \in A \iff r_M(1/2) \leq R_M \text{ for every positive } M \notin G.$$

Thus it suffices to check the tail inequality at the ranks actually omitted by the greedy rule. This is an equivalence on that fixed orbit; it does not establish the inequality at its unboundedly many possible omitted ranks.

[Lean](uniform)greedy recurrence [LEAN SOURCE](#) [LEAN SOURCE](#)

Lemma 6.41 (Seam upper-or-middle classification). *At a row $s \geq 5$, the terminal rank is omitted in the next row, $s \notin D_{s+1}$, if and only if either*

$$4o_s + p_s^+ \leq 2^{s+1},$$

or this inequality fails and

$$4 \operatorname{rem}(s) + 2^{s+1} - p_s^- < 2^{s+2} + 4.$$

These are exactly branches U and M of Theorem 5.16. On either branch the new terminal weight does not fit; on R it does. This is an exact classification, not a conjecture inferred from observed branch labels.

[Lean](uniform)integer quotients [LEAN SOURCE](#)

Lemma 6.42 (The exact gap at the largest omitted rank). *Let $2 \leq d < s$ and $2s < 3d$. Write $W_s(E) = \sum_{e \in E} \lfloor 4^s / (2^e - 1) \rfloor$ for a finite set $E \subseteq \{2, \dots, s-1\}$. For $u \subseteq \{2, \dots, d-1\}$, put $E_- = u \cup \{d+1, \dots, s-1\}$ and $E_+ = u \cup \{d\}$. Then*

$$3W_s(E_-) + 3 \cdot 2^{s+1} + 2 \cdot 4^{s-d} + 4 = 3W_s(E_+).$$

The correction is independent of the common prefix u . For the integer-greedy rows $s \geq 5$, a right transition preserves the largest omitted rank, whereas an upper or middle transition makes s the largest omitted rank in the next row.

[Lean](uniform)integer quotients [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: [largest_false_rank_algebra](#), [paper_largest_false_rank_algebra](#).

Proof. For $d \leq e < s$, the inequality $2s < 3d$ gives

$$\left\lfloor \frac{4^s}{2^e - 1} \right\rfloor = 2^{2s-e} + 4^{s-e}.$$

Subtract the sum over $d < e < s$ from the term at d . The powers of 2 leave 2^{s+1} ; the powers of 4 leave $(2 \cdot 4^{s-d} + 4)/3$. The common prefix u cancels. The branch statements follow from the support updates in Theorem 5.16: the right branch appends s to the lower word, and the other two branches omit s . $\square$

For example, $s = 5$, $d = 4$ and $u = \{2\}$ give $W_s(E_-) = 341$ and $W_s(E_+) = 409$; the displayed identity is $3 \cdot 341 + 204 = 3 \cdot 409$. Both the factor 3 on the right and the restriction $2s < 3d$ are essential to this formula.

Theorem 6.43 (Reduction to the nearest dyadic boundary). *Let d, E be nonnegative integers with $E \leq 2^{d+1}$, and let j_* be the largest $j \in \{0, \dots, d\}$ for which $E \leq 2^{d-j+1}$. The set is nonempty, so j_* is well defined, including when $E = 0$. Then*

$$\begin{aligned} \forall j \in \{0, \dots, d\}, \quad 2^{d-j+1} < E \text{ or } E + 2(d+j) \leq 2^{d-j+1} \\ \iff E + 2(d+j_*) \leq 2^{d-j_*+1}. \end{aligned}$$

For $j > j_*$ the first alternative holds by maximality. For $j \leq j_*$, the left side of the required inequality increases with j and its right side decreases, so the condition at j_* implies all the others. This is an elementary reduction from $d+1$ inequalities to one. Specialized to the seam reset charge (where the actual upper-reset condition gives $E \leq 2^{d+1}$, as explained in Proposition 6.80), the reduced hypothesis `SeamUpperResetCriticalBandEscape` is proved logically equivalent to Theorem 6.10's band-avoidance hypothesis.

[Lean](uniform)dyadic-boundary [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: `exists_criticalDyadicBandIndex`, `dyadicBandEscape_iff_exists_critical`, `seamUpperResetCriticalBandEscape_iff`.

Theorem 6.44 (Excluding minus three at a final middle transition). *Suppose that row $D \geq 13$ is a middle transition in Theorem 5.16 and every transition at a row $s \geq D + 1$ is right. Then*

$$C_D := 4 \operatorname{rem}(D) - p_D^- - 4 \neq -3.$$

This excludes one particular value under the stated tail assumption; it neither excludes every final middle transition nor proves the all-middle-row hypothesis of Theorem 6.9.

[Lean](uniform)integer quotients [LEAN SOURCE](#) [LEAN SOURCE](#)

To see the contradiction, complete D_D by all ranks greater than D . The resulting support $B = D_D \cup \{D + 1, D + 2, \dots\}$ has $X_B(2) < 1/2$ under the all-right assumption. The exact carry identity in the cited proof gives $C_B(2D + 1) = C_D + 3$. If $C_D = -3$, this is zero. The recurrence $C_B(N + 1) = 2C_B(N) + 1 - c_B(N + 2)$ and nonnegativity from Theorem 6.11 first force $C_B(2D + 2) = 0$: indeed $2D + 3 \in B$, so $c_B(2D + 3) \geq 1$. Next $D + 2, 2D + 4 \in B$ imply $c_B(2D + 4) \geq 2$, giving $C_B(2D + 3) < 0$, a contradiction. The stronger conclusion $C_D \geq -2$ is proved in Theorem 10.6.

Lemma 6.45 (The signed position of a skipped greedy prefix). *Let $s \geq 5$ be omitted by the real greedy support G , and put*

$$H_s = G \cap \{2, \dots, s - 1\}, \quad f_s = \sum_{d \in H_s} \left\lfloor \frac{4^s}{2^d - 1} \right\rfloor - T_s.$$

Exactly one of the following holds:

$$\begin{aligned} f_s < 0 : & \quad H_s = D_s, \quad f_s = -\operatorname{rem}(s), \quad \operatorname{rem}(s) \geq 1; \\ f_s = 0 : & \quad H_s = D_s, \quad \operatorname{rem}(s) = 0; \\ f_s > 0 : & \quad H_s = B_s, \quad f_s = o_s. \end{aligned}$$

Thus the real greedy prefix is one of the two adjacent integer words. This classifies its signed distance f_s from T_s ; it does not exclude the negative case. Nor does it identify f_s with the different coordinate $C_s = 4\operatorname{rem}(s) - p_s^- - 4$ used in the final-middle argument.

[Lean](uniform)integer quotients [LEAN SOURCE](#)

Lemma 6.46 (Spacing of reverse-carry words). *For $i = 1, 2$, let integer sequences a_i, b_i, u_i satisfy*

$$b_i(m) + 2u_i(m) = a_i(m) + u_i(m + 1).$$

Let $k, L \geq 0$ be integers. Suppose that at index k the coefficients agree and $b_1(k) - b_2(k) = 1$. Suppose also that both coefficients and bits agree at indices $k + 1, \dots, k + L$. Then

$$u_1(k + L + 1) - u_2(k + L + 1) = 2^L(2(u_1(k) - u_2(k)) + 1).$$

The factor in parentheses is an odd integer. If the two terminal carries have absolute values at most B_1, B_2 , respectively, then $2^L \leq B_1 + B_2$. For a common bound B the conclusion is $2^L \leq 2B$,

not $2^L \leq B$. The formula follows by subtracting the two recurrences: the initial bit difference gives the odd factor, and each subsequent agreement doubles it. An application must supply the stated agreements and terminal bounds.

[Lean](uniform)binary digits [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: [paper_reverse_carry_word](#), [reverse_carry_word_common_bound_sharp](#).

Lemma 6.47 (Linear functionals factoring through one value). *Let V be a vector space over $\mathbb{Q}$, let $\text{ev} : V \rightarrow \mathbb{Q}$ be linear, and choose $e \in V$ with $\text{ev}(e) = 1$. Suppose that each linear functional $\ell_j : V \rightarrow \mathbb{Q}$ vanishes on $\ker(\text{ev})$. Then $v - \text{ev}(v)e \in \ker(\text{ev})$ gives $\ell_j(v) = \ell_j(e)\text{ev}(v)$. Consequently any finite evaluation matrix $(\ell_j(v_i))_{i,j}$ is an outer product and has rank at most one. Every square minor of order at least two therefore vanishes. This excludes determinant arguments formed from these particular functionals, not determinant methods with additional independent information.*

[Lean](uniform)linear-algebra [LEAN SOURCE](#)

Lean: [linearDescender_eq_smul_eval](#), [relationInvariantLinearChannels_det_eq_zero](#).

Theorem 6.48 (Two-thirds band: exact localisation of post-take skip-unsafety). *Write a positive residual as $1/R$. A skipped rank k passes the sufficient dyadic test precisely when $R \geq 2^k$. Suppose a weight at rank b is taken without exhausting the residual: with $q = 2^b - 1$, assume $0 < R < q$. The new reciprocal residual is $Rq/(q - R)$. If the next take is at rank $c \geq b + 2$, put $m = 2^{c-1}$. The last skipped rank is dyadically unsafe exactly when*

$$m - 1 < \frac{Rq}{q - R} < m, \quad \text{equivalently} \quad \frac{q(m - 1)}{q + m - 1} < R < \frac{qm}{q + m}.$$

Clearing the positive denominators proves the equivalence. The interval has width $q^2/((q + m)(q + m - 1))$.

For a single skipped rank, $c = b + 2$ and $m = 2q + 2$. The width is then $q^2/((3q + 1)(3q + 2)) < 1/9$, and the interval lies inside $2q < 3R < 2q + 2/3$. In this single-skip case, the cited corollaries exclude integral R and show that a reduced pre-take residual $p/(2D)$ with p, D, q odd can be unsafe only if $p \geq 7$. The latter conclusion uses the divisibility by 4 of $6D - 2pq$ under the unsafe-band inequalities.

These are local transition statements, not bounds on the actual half-greedy orbit. Dyadic safety is only sufficient for the current tail test. The odd-coprime data $(p, D, b) = (17, 41, 3)$ give an unsafe single-skip example; no conclusion that the half-greedy orbit avoids such data follows from the band calculation.

[Lean](uniform)rational-band [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: [paper_two_thirds_band](#).

Theorem 6.49 (A weaker sufficient test against the remaining tail). *Let k, u, L be positive integers, and put $a = 2L - (2^k - 1)u$. Suppose $a > 0$, equivalently that $\rho = u/(2L) < w_k$ and the greedy rule skips weight w_k . The dyadic sufficient test $\rho \leq 2^{-k}$ is equivalent to $u \leq a$. The weaker sufficient condition*

$$2u \leq 3a$$

ensures $\rho < R_k$, by comparison with $2^{-k} + (3 \cdot 4^k)^{-1} + (7 \cdot 8^k)^{-1} < R_k$. This conclusion excludes a fatal tail-mass deficit at the current step; it does not assert that ρ is representable by the remaining weights, or that all future greedy steps survive.

The containment of these sufficient conditions is strict even for valid rational data: $(k, u, L, a) = (2, 7, 13, 5)$ gives $1/4 < 7/26 < R_2$. In contrast, the scalar pair $(u, a) = (3, 2)$ in the linked inequality lemma does not arise from integral L under the present parity relation. Unit numerators are nonfatal at such a skipped step, since then $a \geq 1$. Conversely, a fatal step requires $3a < 2u$, hence $u \geq 2$, or $u \geq 3$ when u is odd.

[Lean](uniform)divisor counts and finite sums [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)
Lean: [paper_sharp_fatal_gap](#).

Lemma 6.50 (Summability of one gap length per level). *For $n \geq 1$, put $g_n = w_n - R_n > 0$. The sum of one gap length per level satisfies, for $N \geq 0$,*

$$\sum_{n \geq N} g_n \leq \frac{2}{9} 4^{-N} + \frac{3}{7} 8^{-N}.$$

This follows by summing the per-level upper bounds geometrically. It is not the measure of the union of all gaps: level n has 2^{n-1} translated gaps of length g_n . The corresponding weighted sum and its geometric interpretation are given after Theorem 6.118. Neither estimate decides membership of a specified point such as $1/2$.

[Lean](uniform)divisor counts and finite sums [LEAN SOURCE](#) [LEAN SOURCE](#)
Lean: [mersenneGap_pos](#), [summable_mersenneGap_succ](#), [mersenneGap_tail_le](#), [tendsto_mersenneGap_tail_zero](#).

Lemma 6.51 (The effect of adding one divisor). *Let $N \geq 0$ and let finite supports D_0, D_1 differ only by $D_1 = D_0 \cup \{N+1\}$, with $N+1 \notin D_0$. For every positive integer m ,*

$$c_{D_1}(m) - c_{D_0}(m) = \mathbf{1}_{N+1|m}.$$

In particular, the difference is 1 at $m = 2(N+1)$. This follows directly from the definition of a divisor count. Its use in another coefficient sequence requires proving that sequence has the same support-incidence description.

[Lean](uniform)divisor-incidence [LEAN SOURCE](#)
Lean: [supportCoeff_insert_eq_add_indicator](#), [supportCoeff_extend_true_eq_false_add_one_at_double](#).

Theorem 6.52 (Uniqueness of an integer recurrence with a vanishing scaled limit). *For any nonnegative-integer coefficient sequence $c : \mathbb{N} \rightarrow \mathbb{N}$ with $c(n) \leq n$: the binary coefficient series $X_c = \sum'_{n \geq 1} c(n)/2^n$ is rational iff there exists a positive integer multiplier v and an integer orbit $u : \mathbb{N} \rightarrow \mathbb{Z}$ satisfying the exact carry recurrence $u(N+1) = 2u(N) - v \cdot c(N+1)$ together with the condition $u(N)/2^N \rightarrow 0$. Every integer sequence satisfying both conditions obeys $u(N) = v \cdot T_c(N)$ exactly, where $T_c(N) = \sum_{j \geq 1} c(N+j)/2^j$ is the scaled tail. For each fixed v there is at most one such sequence. Positivity of the orbit alone is deliberately not used as an equivalent criterion: a homogeneous 2^N -scaled perturbation can be added to any orbit without breaking the recurrence, so the limit condition cannot be dropped. The choices $c = c_A$ for Problem 257 and $c = \varphi$ for Problem 249 both satisfy $0 \leq c(n) \leq n$. The telescoping argument is given again, with an example, in Theorem 6.105.*

[Lean](uniform)binary digits [LEAN SOURCE](#) [LEAN SOURCE](#)

Lemma 6.53 (Changing finitely many support elements). *Let $b \geq 2$ be an integer, and let $A, B \subseteq \mathbb{N}_{\geq 1}$. If $X_A(b)$ is irrational and A, B have finite symmetric difference, then $X_B(b)$ is irrational too: the two sums differ by a finite sum of rational numbers. This observation supplies the finite modifications in Theorem 6.21.*

[Lean](uniform)tail-transfer [LEAN SOURCE](#)

Lean: `irrational_erdosSupportSeries_of_tail, irrational_erdosSupportSeries_tail_of_irrational`.

Lemma 6.54 (An integer test for the interval between the dyadic and Mersenne weights). *For positive integers p, L and $n \geq 0$, write the residual as $p/(2L)$ and put $E = 2^n p - L$. Clearing positive denominators gives*

$$\frac{p}{2L} \leq \frac{1}{2^{n+1}} \iff E \leq 0, \quad \frac{1}{2^{n+1}} < \frac{p}{2L} < \frac{1}{2^{n+1} - 1} \iff 0 < E < \frac{p}{2}.$$

The first test concerns the dyadic bound, not the greedy selection threshold itself. The second identifies precisely the interval in which a Mersenne weight is skipped but the dyadic safety test fails.

[Lean](uniform)integer quotients [LEAN SOURCE](#) [LEAN SOURCE](#)

Lemma 6.55 (The odd denominator survives dyadic subtraction). *Let r/D be a reduced fraction with $D > 0$ odd, let $p \in \mathbb{Z}$, and let $c \geq 0$ be an integer. The reduced denominator of $p/2^c - r/D$ is divisible by D and divides $2^c D$. Indeed, its unreduced numerator $pD - 2^c r$ is coprime to D , since $\gcd(r, D) = \gcd(2^c, D) = 1$. Reduction can therefore remove only powers of 2 from the displayed denominator. No property of Mersenne weights is used.*

[Lean](uniform)denominator-algebra [LEAN SOURCE](#)

Lemma 6.56 (A guaranteed divisor of a reduced denominator). *Let $a \in \mathbb{Z}$, $D \in \mathbb{N}_{>0}$ and $m, C, h \in \mathbb{N}$. Write $\text{den}(x)$ for the positive reduced denominator of a rational number x . If $m \mid D$ and $\gcd(m, |a|) = 1$, then $m \mid \text{den}(a/D)$. If $C \mid D$ and $\gcd(C, |a|) = 1$, then*

$$\frac{C}{\gcd(C, h)} \mid \text{den}(ha/D).$$

Thus scaling may remove part of the guaranteed divisor. Indeed, reduction divides D by $\gcd(D, |a|)$, so no prime power in m can be lost. After scaling, a prime $p \mid C$ can lose at most $v_p(h)$ powers from the guaranteed divisor. This is a lower bound on the reduced denominator, not an equality: for $D = 12$, $C = 6$, $a = 1$, $h = 2$, it guarantees the divisor 3, while the reduced denominator is 6. The lemma feeds Theorem 6.24 directly.

[Lean](uniform)denominator-algebra [LEAN SOURCE](#) [LEAN SOURCE](#)

Lemma 6.57 (Commuting prime-power differences). *For a function $g : \mathbb{N}_{>0} \rightarrow \mathbb{Z}$, a positive integer p and an integer $e \geq 1$, define*

$$(\Delta_{p,e}g)(n) = g(p^e n) - g(p^{e-1}n).$$

These operators commute. If $p \neq q$ are primes, $e, f \geq 1$, $A \subseteq \mathbb{N}_{>0}$, $n \geq 1$ and $\gcd(n, pq) = 1$, then

$$(\Delta_{q,f}\Delta_{p,e}c_A)(n) = \sum_{d|n} \mathbf{1}_A(p^e q^f d).$$

Thus the mixed difference counts precisely the support elements whose p - and q -adic exponents are e and f in this divisor sum.

[Lean](uniform)p-adic [LEAN SOURCE](#) [LEAN SOURCE](#)

Proof. Expanding either order gives the same four terms:

$$\begin{aligned} &g(p^e q^f n) - g(p^{e-1} q^f n) \\ &\quad - g(p^e q^{f-1} n) + g(p^{e-1} q^{f-1} n). \end{aligned}$$

This commutation needs no primality or coprimality. For $g = c_A$, write each divisor of $p^e q^f n$ uniquely as $p^i q^j d$, with $d | n$, $0 \leq i \leq e$ and $0 \leq j \leq f$. The two differences cancel all terms except $i = e, j = f$. $\square$

For $A = \{12\}$, $p = 2$, $e = 2$, $q = 3$, $f = 1$, $n = 1$, the four counts are 1, 0, 0, 0, so the mixed difference is 1. This coefficient identity alone does not transfer rationality through the subsequence operation $g(n) \mapsto g(pn)$.

Definition 6.58 (The Mersenne achievement set and the support-series object). For a support $A \subseteq \mathbb{N}_{\geq 1}$, its Mersenne subsum is

$$x_A = \sum_{n \geq 1} \frac{\mathbf{1}_A(n)}{2^n - 1}, \quad \mathcal{A} = \{x_A : A \subseteq \mathbb{N}_{\geq 1}\}.$$

Finite, empty, and infinite supports are all allowed in $\mathcal{A}$. More generally, for an integer base $b \geq 2$, $X_A(b) = \sum_{n \geq 1} \mathbf{1}_A(n)/(b^n - 1)$ and $c_A(n) = \sum_{d|n} \mathbf{1}_A(d) = (\mathbf{1}_A * \mathbf{1})(n)$ for $n \geq 1$, where $*$ denotes Dirichlet convolution. Full support gives the series in Theorem 6.14. The linked definitions implement these positive-index sums.

[Lean](n/a)binary digits [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

Proposition 6.59 (Achievement-set topology: compact, closed, perfect, measure exactly one). $\mathcal{A}$ is compact (continuous image of the binary-sequence Cantor space $\mathbb{N} \rightarrow \text{Fin } 2$ under the product topology, via `positiveMersenneDigitValue`), hence closed; it is also perfect, totally disconnected, and nowhere dense, with Lebesgue measure exactly 1. The compactness/closedness argument (binary coding $\rightarrow$ Cantor space $\rightarrow$ continuous image) is a fully generic technique for characterizing the achievement set of any absolutely convergent digit-weighted series, not specific to Mersenne denominators; reusable for a $\varphi(n)/2^n$ subsum set after checking summability. No separation of successive weights is needed for compactness or closedness; the stronger topological conclusions require their own hypotheses. Closedness alone is what powers every “limit of a sequence of achieved points is achieved” argument in this catalogue (e.g. Theorem 6.2, Theorem 6.5).

[Lean](n/a)binary digits [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: [paper_achievement_set_topology](#).

Theorem 6.60 (Nonmembership of one half and a finite fatal gap). *The value $1/2$ is not in A if and only if there are an integer $d \geq 0$ and a finite set $u \subseteq \{1, \dots, d\}$ such that*

$$X_u(2) + R_{d+1} < \frac{1}{2} < X_u(2) + w_{d+1}.$$

One implication is Lemma 6.37. For the other, follow the greedy prefixes: nonmembership forces a first failure of the tail inequality, hence one of these gaps. The endpoint equalities are excluded by Lemma 6.38.

A finite prefix and strict gap provide a nonmembership witness. The tail inequality can be certified by rational truncation bounds as in Theorem 2.11. Membership, in contrast, asserts that no such witness exists; failure to find one in a finite search is not a proof. The reduction uses $w_n > R_n$, not compactness alone. The totient weights fail this hypothesis, as the example after Lemma 6.34 shows.

[Lean](n/a)dedekind-cut [LEAN SOURCE](#) [LEAN SOURCE](#)

Definition 6.61 (Separated integer values and bounded perturbations). Let α be a set, let $s : \alpha \rightarrow \mathbb{N}$ be injective, and let $p : \alpha \rightarrow \mathbb{N}$. Suppose integers $g \geq 1$ and $B \geq 0$ satisfy

$$s(x) < s(y) \implies s(x) + g \leq s(y), \quad 0 \leq p(x) \leq B < 3g.$$

The updated values are $t(x) = 4s(x) + p(x)$. For a capacity $C \geq 0$, an adjacent pair consists of $x_-, x_+ \in \alpha$ with $s(x_-)$ the largest value at most C and $s(x_+)$ the smallest value greater than C . Existence of both values is part of the assumption; it is not guaranteed for an arbitrary family. Set

$$r = C - s(x_-), \quad o = s(x_+) - C, \quad p_- = p(x_-), \quad p_+ = p(x_+).$$

The inequality $B < 3g$ preserves the order of the updated values. The next result uses the stronger bound $B < g$ to ensure that the selected updated value is admissible at capacity $4C + g$. The linked structures encode these data.

[Lean](n/a)abstract-perturbed-greedy [LEAN SOURCE](#) [LEAN SOURCE](#)

Theorem 6.62 (Perturbed-family maximality and the three-branch recurrence). *Use the family and adjacent pair of Definition 6.61, and assume in addition that $B < g$. At capacity $C' = 4C + g$, the largest admissible updated value is $t(x_+)$ if $4o + p_+ \leq g$, and $t(x_-)$ otherwise. After this choice, apply the take-if-possible rule to an additional weight $W = 2g + 4$. The resulting remainder is*

$$\begin{cases} g - 4o - p_+, & 4o + p_+ \leq g, \\ 4r + g - p_-, & 4o + p_+ > g \text{ and } 4r + g - p_- < W, \\ 4r - g - p_- - 4, & 4o + p_+ > g \text{ and } 4r + g - p_- \geq W. \end{cases}$$

[Lean](n/a)abstract-perturbed-greedy [LEAN SOURCE](#) [LEAN SOURCE](#)

Proof. The lower candidate satisfies $t(x_-) \leq 4C + B < C'$, and $t(x_+) \leq C'$ is exactly the test $4o + p_+ \leq g$. Order preservation excludes earlier values from improving the selected one. Any value after x_+ has $s(x) \geq s(x_+) + g > C + g$, and hence $t(x) > C'$. Thus no other candidate is missed.

On the first branch the remainder is at most $g < W$, so the extra weight is not taken. The other two branches compare the lower candidate's remainder $4r + g - p_-$ with W . All displayed differences are nonnegative on their respective branches. $\square$

The theorem describes a two-stage rule, not automatically the maximum of $\{t(x) + \varepsilon W : x \in \alpha, \varepsilon \in \{0, 1\}\}$ below C' . For example, take old values 4, 5, 7, $g = 1$, $B = 0$, $p = 0$ and $C = 6$. Then $C' = 25$ and $W = 6$. The rule chooses 20 and cannot add 6, but $16 + 6 = 22$ is a larger admissible extended value. A sufficient extra hypothesis for global maximality is $4g - B \geq W$: every earlier updated value plus W is then at most the selected updated value. In the Mersenne application $g = 2^{s+1}$, $B = 2(s - 2)$ and $s \geq 5$, so this inequality holds. Thus the concrete recurrence is unaffected, but an application to another weight system must check this separation as well.

The hypothesis $B < g$ is additional; it does not follow from $B < 3g$ in Definition 6.61. The weaker bound alone does not ensure admissibility of the chosen value. Take $g = 2$, $B = 3$, old values 0, 2, 4, perturbations 0, 3, 0, and $C = 2$. The rule selects the old value 2, whose update is $11 > C' = 10$, although the updated value 0 is admissible. This example explains why the stronger hypothesis must accompany the maximality statement.

6.3 Consequence theorems

This subsection concerns finite sets D whose floor-quotient sums nearly represent $1/2$. Its main implication is elementary to state: if such finite sums approximate $1/2$ at unbounded depths with an error tending to zero, compactness supplies a representing support. The difficult input is the existence of the finite sets, not the final compactness argument. Definitions and conditional implications are separated below. A hypothesis marked [Open] is not established merely because the implication using it has a formal proof.

Finite quotients and their remainders

Definition 6.63 (An exact finite quotient sum). For an integer $n \geq 1$, an exact row at depth n is a set $D \subseteq \{2, \dots, n\}$ satisfying

$$Q(D, n) = \sum_{d \in D} \left\lfloor \frac{2^n}{2^d - 1} \right\rfloor = 2^{n-1} - 1.$$

The word “exact” refers to this integer equality, not to $X_D(2) = 1/2$, which is impossible for finite D . The formal definition also allows $n = 0$ using truncated natural-number subtraction. That degenerate case is immaterial to existence at unbounded depths.

[Lean] (n/a) *divisor counts and finite sums* [LEAN SOURCE](#)

Definition 6.64 (Floor quotients and the signed endpoint expression).

$$q(M, d) := \left\lfloor \frac{2^M}{2^d - 1} \right\rfloor, \quad Q(D, M) := \sum_{d \in D} q(M, d),$$

The signed endpoint expression has three arguments:

$$H(D, k, n) = 2S(D, k, n-1) + 1 - c_D(n), \quad c_D(n) = \#\{d \in D : d \mid n\}.$$

The suffix in this formula is defined next. The quotients are integers; this final expression is allowed to be negative.

[Lean] (n/a) *divisor counts and finite sums* [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

Definition 6.65 (The remaining binary suffix).

$$S(D, k, M) := \max\{2^{M-k} - Q(D, M) - 1, 0\} \quad (M \geq k),$$

the nonnegative remainder after subtracting the quotient contributions from the non-terminating binary expansion of 2^{-k} through place M . The formal definition uses truncated natural-number subtraction; in the applications where the displayed difference is nonnegative, the maximum may be omitted. The formal convention for $M < k$ also truncates $M - k$ to zero, a case not used in these capacity estimates.

[Lean] (n/a) *binary digits* [LEAN SOURCE](#)

Definition 6.66 (Exact quotient sums at unbounded depths). The cofinal exact-row condition says that for every integer N there are $n \geq \max\{N, 1\}$ and $D \subseteq \{2, \dots, n\}$ with

$$Q(D, n) = 2^{n-1} - 1.$$

No agreement is required between supports chosen at different depths. This removes a compatibility requirement from the data to be supplied, not from the underlying membership question: Proposition 2.7 identifies this cofinal existence statement with half-membership. A failed construction of compatible rows therefore does not refute the condition, and the condition itself is not proved here.

[Open] (cofinal) *divisor counts and finite sums* [LEAN SOURCE](#)

Two implications yielding half-membership

Theorem 6.67 (Compactness from exact finite sums). *Suppose the cofinal exact-row condition in Definition 6.66 holds. Then $1/2 \in \mathcal{A}$. For each N , choose an exact row D_N at depth $n_N \geq \max\{N, 1\}$. Proposition 6.98 gives*

$$\left| X_{D_N}(2) - \frac{1}{2} \right| \leq \frac{n_N + 1}{2^{n_N}} \rightarrow 0.$$

Every $X_{D_N}(2)$ belongs to the closed set $\mathcal{A}$, so its limit does too. No finite support represents $1/2$; the representing support is therefore infinite. The proof does not require the chosen rows to be nested. The unproved input is their existence at unbounded depths.

[Lean] (cofinal) *topological-achievement-set* [LEAN SOURCE](#)

Lean: `CofinalExactLocalMersenneHalfRows`, `abs_exactLocalMersenneRowValue_sub_half_1`, `half_mem_mersenneAchievementSet_of_cofinalExactLocalRows`, `finite_boolSupport_ne_half`.

Theorem 6.68 (Infinitely many greedy skips suffice). *Suppose that the greedy remainder for $1/2$ satisfies*

$$\forall N \exists c \geq \max\{N, 4\} : \quad 0 < r_{c-1}(1/2) < w_c.$$

Then exact quotient rows exist at unbounded depths, and $1/2 \in \mathcal{A}$. A single such skip at c gives an exact row at depth $2c - 2$ without an additional capacity hypothesis. Taking arbitrarily large c therefore supplies the rows needed for the compactness implication in Theorem 6.67. The existence of arbitrarily late positive skips is the hypothesis, not an established property of the greedy sequence. It is equivalent to half-membership and remains unproved here.

[Lean] (cofinal) greedy recurrence [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: `CofinalPositiveHalfGreedySkips`, `exactLocalMersenneHalfRow_of_positiveHalfGreedySkip`, `cofinalExactLocalMersenneHalfRows_of_positiveHalfGreedySkips`, `half_mem_mersenneAchievementSet_of_positiveHalfGreedySkips`, and 1 further declaration in the [coverage section](#).

Compatible finite approximations

Theorem 6.69 (A compatible family of finite supports). *Let $b_{n,d} \in \{0, 1\}$ satisfy $b_{n+1,d} = b_{n,d}$ whenever $2d \leq n$. Define*

$$E_n = \{d : 2 \leq d \leq n, b_{n,d} = 1\}, \quad D_n = E_n \cap \{2, \dots, \lfloor n/2 \rfloor\}.$$

The four conditions on the finite rows, required for every $n \geq 2$, are

$$\begin{aligned} 2^{\max\{c_{D_n}(n)-1, 0\}} - 1 &\leq S(D_n, 1, n-1), \\ \sum_{\substack{d \in E_n \\ d > \lfloor n/2 \rfloor}} 2^{n-d} &= H(D_n, 1, n), \\ H(D_n, 1, n) &< 2^{n-\lfloor n/2 \rfloor}, \\ Q(E_n, n) &= 2^{n-1} - 1. \end{aligned}$$

The first condition implies that the integer in the next two lines is nonnegative. The second specifies the binary value of the upper half of the row, the third bounds that value by its available number of bits, and the fourth is the exact quotient equation. The maximum in the first line records natural-number subtraction at the case $c_{D_n}(n) = 0$.

These are the finite hypotheses of the compatible-limit construction. The bit condition fixes coordinate d from row $2d$ onward. Unlike the independent finite supports in Definition 6.66, the rows therefore have a prescribed common limit. This extra structure is not, by itself, a proof of strictness between the corresponding existence statements.

[Open] (uniform) binary digits [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: `paper_compatible_bit_stable`, `paper_compatible_finite_row_conditions`, `paper_compatible_first_condition_gives_nonneg`, `paper_compatible_rows_agree_with_limit`.

A quotient bound at a crossing

Theorem 6.70 (A sufficient quotient bound at a crossing). *Consider the following condition on finite sets D and integers $c \geq 4$:*

$$\begin{aligned} D \subseteq \{2, \dots, c-1\}, \quad X_D(2) < \frac{1}{2} < X_D(2) + w_c \\ \implies Q(D \cup \{c\}, 2c-2) \geq 2^{2c-3}. \end{aligned}$$

It requires the quotient inequality at every strict crossing of one half by an added weight. Theorem 6.93 identifies this inequality with the binary bound needed in the finite construction. If the condition holds for all such D, c , the induction below produces exact sums at unbounded depths and hence an infinite support of value $1/2$. The universal crossing condition itself is unproved.

[Open] (uniform) binary digits [LEAN SOURCE](#)

Lean: `SkippedCoreCriticalQuotientSupply, cofinalExactLocalMersenneHalfRows_of_criticalQuotientSupply, half_mem_mersenneAchievementSet_of_criticalQuotientSupply, finite_boolSupport_ne_half.`

Theorem 6.71 (Induction from the depth-six example). *An already-formalised induction from the endpoint-six seed (Definition 6.97) consumes Theorem 6.70 at every step: each ProtectedExactLocalMersenneRow either doubles below half (unconditional) or recycles at its first crossing rank $e > \text{cutoff}$, giving endpoint $2e - 2 > \text{previous endpoint}$; protection (endpoint $< 2 \cdot \text{cutoff}$, new ranks $> \text{cutoff}$) is exactly what converts the non-growing recycle endpoint of the bare dichotomy (Theorem 6.127) into strict progress. The below-half branch never fires twice from the seed arithmetic, so the supply is needed at essentially every step. Ends literally at CofinalExactLocalMersenneHalfRows.*

[Lean] (uniform) binary digits [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

Theorem 6.72 (The quotient condition on greedy prefixes). *Let $D_c = G \cap \{2, \dots, c-1\}$. The condition on the actual greedy prefixes is*

$$\begin{aligned} c \geq 4, \quad c \notin G \\ \implies Q(D_c \cup \{c\}, 2c-2) \geq 2^{2c-3}. \end{aligned}$$

It asks for the quotient bound at every skipped rank of the specified half-greedy orbit. This restricts the crossing test of Theorem 6.70 to its actual greedy prefixes. The finite uniqueness results justify that restriction where the cited induction uses it; they do not prove the displayed inequality.

[Open] (uniform) greedy recurrence [LEAN SOURCE](#)

Lean: `HalfGreedySkippedCriticalQuotientSupply, skippedCoreCriticalQuotientSupply_if_halfGreedySkipped.`

Theorem 6.73 (Two consecutive skips). *If $c \geq 6$ and the real half-greedy rule skips both c and $c+1$, put $D = G \cap \{2, \dots, c-1\}$. Then*

$$S(D, 1, 2c-3) < 2^{c-3}.$$

The one-step quotient recurrence gives $S(D, 1, 2c - 2) < 2^{c-2}$, so Theorem 6.75 supplies an exact row at depth $2c - 2$ whose new ranks are all greater than c . The hypothesis concerns this pair of skipped ranks; no unbounded sequence of such pairs is proved here. For the stated precritical-suffix condition at every skipped rank, the remaining tests are the skip-then-take cases, with ranks 4 and 5 handled separately in the cited proof.

[Lean] (uniform) greedy recurrence [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

Proof. The two skips imply $0 < \rho = 1/2 - X_D(2) < w_{c+1}$. Put $M = 2c - 3$. The fractional-part error in $Q(D, M)$ is less than $c - 2$, because $D \subseteq \{2, \dots, c - 1\}$. Therefore

$$S(D, 1, M) < \frac{2^{2c-3}}{2^{c+1} - 1} + c - 3 < 2^{c-4} + c - 2 \leq 2^{c-3},$$

where the last inequality holds for $c \geq 6$. The suffix is nonnegative because $X_D(2) < 1/2$, so its next-step identity is $S(D, 1, M + 1) = 2S(D, 1, M) + 1 - c_D(M + 1)$. Integrality gives the required strict bound at $M + 1$. $\square$

Theorem 6.74 (A later skip after a selected block). *Suppose the real half-greedy rule skips rank c , takes $c + 1, \dots, c + t - 1$, and skips $c + t$. If*

$$0 < t \leq c - 3, \quad c - 2 \leq 2^{c-t-3},$$

then the precritical suffix bound at c holds. It yields sharp capacity and hence an exact row at depth $2c - 2$. The case $t = 1$ includes Theorem 6.73; the proof also permits longer selected blocks. The arithmetic condition is precisely $t \leq c - 3 - \lceil \log_2(c - 2) \rceil$. Its role in the proof is to make the dyadic allowance 2^{c-t-3} cover the bound $|D| \leq c - 2$ for the earlier selected support.

To obtain cofinal exact rows by this result, the stated gap condition must hold at cofinally many skipped ranks c , not just at one rank or throughout a finite sample. A finite empirical skip frequency alone gives no such pointwise gap bound.

[Lean] (uniform) greedy recurrence [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

Filling the remaining binary positions

Theorem 6.75 (Filling the remaining binary positions). *For any $c \geq 4$ and any finite $D \subseteq [2, c]$ with $\text{value}(D) < 1/2$ and the sharp capacity*

$$S(D, 1, 2c - 2) < 2^{c-2},$$

there is an exact row E at endpoint $2c - 2$ with $D \subseteq E$ and every new rank strictly above c . The implication requires the displayed capacity bound, but neither a crossing condition nor a specified real greedy prefix. The crossing hypothesis used in Theorem 6.70 is one way to seek that input, not an additional premise of this filling result. No strict comparison between the corresponding cofinal existence statements is asserted.

[Lean] (uniform) binary digits [LEAN SOURCE](#)

Theorem 6.76 (An exact sum at depth $2c - 2$). *Under the hypotheses of Theorem 6.75, there is a set $E \subseteq \{2, \dots, 2c - 2\}$ with*

$$Q(E, 2c - 2) = 2^{2c-3} - 1.$$

This conclusion forgets the additional support-extension information in that theorem. It is the form used when only the existence of an exact row is needed.

[Lean] (uniform) binary digits [LEAN SOURCE](#)

Proposition 6.77 (Quotients in the upper half of the index range). *For integers $d \geq 2$ and $d \leq M < 2d$,*

$$q(M, d) = \left\lfloor \frac{2^M}{2^d - 1} \right\rfloor = 2^{M-d}.$$

Indeed, the quotient is $2^{M-d} + 2^{M-d}/(2^d - 1)$, whose second term is strictly between zero and one. In Theorem 6.75, the available ranks $d = c + 1, \dots, 2c - 2$ therefore supply the binary weights $2^{c-3}, \dots, 1$. They represent every integer from 0 to $2^{c-2} - 1$, which explains both the capacity bound and the choice of terminal depth.

[Lean] (uniform) binary digits [LEAN SOURCE](#)

Lean: `localMersenneQuotient_eq_two_pow_sub_of_half_lt, exists_boolean_word_of_lt_two_pow`.

Theorem 6.78 (A finite sum from a skipped prefix). *Let $c \geq 4$ and $D \subseteq \{2, \dots, c - 1\}$ satisfy*

$$0 < \frac{1}{2} - X_D(2) < \frac{1}{2^c - 1}.$$

Then there is E with $D \subseteq E \subseteq \{2, \dots, 2c - 2\}$ and $Q(E, 2c - 2) = 2^{2c-3} - 1$. No separate sharp-capacity assumption is required. The binary completion in the linked proof uses ranks $c, \dots, 2c - 2$; it may therefore insert c . Theorem 6.75 instead assumes the sharper $(c - 2)$ -bit bound so that every added rank is strictly greater than c .

[Lean] (bounded) binary digits [LEAN SOURCE](#)

Lean: `paper_exact_row_from_skipped_prefix`.

For a finite $F \subseteq \{2, \dots, n\}$ with $X_F(2) > 1/2$, take c to be its first crossing rank and $D = F \cap \{2, \dots, c - 1\}$. Then $c \geq 4$ because $1/3 + 1/7 < 1/2$, and the displayed hypotheses hold: equality at the lower endpoint is excluded by odd denominators. This is the first-crossing application in `exists_skippedCoreExactRow_of_value_above` and `exactLocalMersenneHalfRow_of_first_localMersenne_crossing` (`BooleanMobiusExactRowCrossing.lean:155, 191`).

Reducing the dyadic-boundary checks

Theorem 6.79 (Testing one nearest dyadic boundary). *Let d, E be nonnegative integers with $E \leq 2^{d+1}$, and let*

$$j_* := \max\{0 \leq j \leq d : E \leq 2^{d-j+1}\}.$$

The boundary 2^{d-j_+1} is the smallest power at least E among the finite list $2, 4, \dots, 2^{d+1}$; in particular $j_* = d$ when $E \leq 2$. The following conditions are equivalent:*

$$\text{for every } 0 \leq j \leq d, \quad 2^{d-j+1} < E \text{ or } E + 2(d+j) \leq 2^{d-j+1}; \\ E + 2(d+j_*) \leq 2^{d-j_*+1}.$$

The first condition is `DyadicBandEscape(d, E)` and j_* satisfies `CriticalDyadicBandIndex(d, E, j_*)`.

[Lean] (uniform) dyadic-boundary [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: [paper_critical_dyadic_band_index_unique](#), [paper_critical_dyadic_boundary_is_smallest](#), [paper_critical_dyadic_band_index_eq_top_of_le_two](#), [paper_dyadic_band_escape_iff_single_test](#).

Proof. For $j > j_*$ the boundary is below E ; for $j < j_*$ the boundary is larger and the required width $2(d+j)$ is smaller. This proves sufficiency of the single inequality; necessity is its instance at j_* . $\square$

For example, $d = 5, E = 20$ gives $j_* = 1$ and $20 + 12 = 32$, so equality at the allowed edge satisfies every band condition.

At an actual upper reset, Proposition 6.80 supplies $E \leq 2^{d+1}$. The equivalence replaces all its local band tests by one test. The arithmetic inequality at every required reset remains an unproved hypothesis of the membership implication.

Proposition 6.80 (An upper bound for the reset expression). *Let $d \geq 5$ be an actual upper-reset row. Write E_d for its reset charge, the sum of four times the adjacent upper overshoot and the corresponding nonnegative correction term. The upper-reset identity is*

$$\text{rem}(d+1) + E_d = 2^{d+1}.$$

The nonnegativity of the successor remainder therefore gives $E_d \leq 2^{d+1}$, the side condition needed for Theorem 6.79. This argument uses the upper-reset branch assumption. It does not require the separate, conditional two-sided bound of Theorem 6.9, and does not assert the charge bound at arbitrary rows.

[Lean] (uniform) dyadic-boundary [LEAN SOURCE](#)

Lean: [seamUpperBranch_remainder_add_resetCharge_eq](#), [seamUpperResetCharge_le](#).

Nonnegativity of the omitted-tail margin

The next three statements compare the conditions of Theorem 6.13. The first two are equivalent; the third is sufficient for them.

Definition 6.81 (A nonnegativity condition at skipped ranks).

$$\forall \text{ skipped rank } n \geq 3, \quad 0 \leq F(n-1, n).$$

This condition is not established here. Its usefulness is the following implication: assuming it, `half_mem_mersenneAchievementSet_of_skippedFullShellNonnegative` proves $1/2 \in A$.

[Open] (uniform) frozen-margin [LEAN SOURCE](#) [LEAN SOURCE](#)

Theorem 6.82 (An equivalent vanishing condition).

$\forall$ skipped rank $n \geq 3$ whose actual word equals the seam-greedy word,
the integer remainder at n is zero.

This condition is equivalent to Definition 6.81. At an actual skipped rank, the source proves that $F(n-1, n) < 0$ holds exactly when the real prefix agrees with the integer greedy word and its integer remainder is positive. On agreement, $F(n-1, n)$ is the negative of that nonnegative remainder. Thus nonnegativity of the margin at every skip is equivalent to vanishing of the remainder at every aligned skip. The equivalence is proved; neither condition is established for all required ranks. Its formal statement is `skippedSeamAlignmentZero_iff_skippedFullShellNonnegative`.

[Open] (uniform) frozen-margin [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: `HalfGreedySkippedSeamAlignmentZero`, `HalfGreedySkippedFullShellNonnegative`, `skipped_fullShell_neg_iff_alignment_and_seamRemainder_pos`, `greedyHalfFrozenMargin_fullShell_eq_neg_seamRemainder_of_alignment`, and 1 further declaration in the [coverage section](#).

Theorem 6.83 (A sufficient lower bound for the integer remainder).

$\forall$ skipped rank $n \geq 3$, $B(2n) < \text{rem}(n)$,

a sufficient condition for the nonnegativity condition above. Indeed, a negative margin at an actual skip would force $1 \leq \text{rem}(n) \leq B(2n)$, contradicting this bound. This proves the implication, not a converse or strict separation of the conditions.

[Open] (uniform) frozen-margin [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: `paper_seam_escape_forces_remainder_band`, `paper_seam_escape_implies_full_shell_nonnegative`, `paper_seam_escape_implies_half_membership`.

Square-root bounds for half-carries

Theorem 6.84 (A square-root carry bound implies half-membership). Write G for the greedy support and $C_G(N)$ for its centred carry. If

$$\forall N, \quad C_G(N) \leq 2\sqrt{N} + 4,$$

then the greedy selected support G is infinite and $X_G(2) = 1/2$. Nonnegativity of C_G is unconditional (see Theorem 6.11 and the greedy case above); only the upper $2\sqrt{N} + 4$ bound remains open. This is a different hypothesis from the dyadic-band condition and the largest-skip condition; logical independence is not asserted.

[Open] (cofinal) sqrt-carry-growth [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: `greedy_mobiusCenteredHalfCarry_nonneg`, `greedy_half_infinite_of_mobiusCenteredHalfCarry_upperBound`.

Theorem 6.85 (Terminal bounds at unbounded depths). Suppose that for every $N \geq 0$ there are $M \geq \max\{N, 1\}$ and $D \subseteq \{2, \dots, M\}$ such that

$$|\text{ihc}(D, M-1)| \leq B(M) = 2\lfloor \sqrt{M} \rfloor + 4.$$

Then some infinite $A \subseteq \mathbb{N}_{>0}$ satisfies $X_A(2) = 1/2$. No compatibility between different finite supports or bound on their earlier carries is required.

These are not the exact-row conditions of Definition 6.66. For such a finite support,

$$\text{ihc}(D, M-1) = 2^{M-1} - Q(D, M),$$

so an exact row has carry 1, whereas the displayed bound allows several positive and negative values. At $M = 6$, for example, $D = \{2, 3\}$ has $Q(D, 6) = 30$ and carry 2: it satisfies the terminal bound but is not an exact row. Both cofinal existence statements nevertheless imply, and are implied by, half-membership, as explained in Proposition 2.7. The difference between their finite data must not be confused with a strict logical weakening of the membership problem.

[Open] (cofinal) sqrt-carry-growth [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: [paper_terminal_strip_forces_half_membership](#), [paper_integerHalfCarry_eq_two_pow_sub_localPrefixQuotient](#), [paper_exact_row_integerHalfCarry_eq_one](#), [paper_terminal_strip_witness_six](#), and 1 further declaration in the [coverage section](#).

Theorem 6.86 (Cofinal returns of the greedy carry). Suppose that for every $N \geq 0$ there is $M \geq N$ such that

$$\text{ihc}(G, M) \leq B(M+1) = 2\lfloor \sqrt{M+1} \rfloor + 4.$$

Then G is infinite and $X_G(2) = 1/2$. Unlike Theorem 6.84, this hypothesis bounds the actual greedy carry only at unboundedly many indices. It still concerns the same fixed support G , not independently chosen finite supports. The last source below compares this carry with the carry obtained by fixing a finite prefix. Its comparison includes a contribution from later omitted ranks. No equivalence with the finite-support hypothesis of Theorem 6.75 follows without controlling that additional contribution.

[Open] (cofinal) sqrt-carry-growth [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

A margin at an unspecified later horizon

Theorem 6.87 (An eventual nonnegative margin suffices). Fix a positive depth k and put $D = G \cap \{2, \dots, k\}$. The following conditions are equivalent:

$$r_k(1/2) < 2^{-(k+1)} \iff F_k(J) \geq 0 \text{ for some } J \geq 0.$$

Here $F_k(J)$ is the integer expression defined in Theorem 6.13. Its normalised value is

$$2^{-J}F_k(J) = \sum_{i=1}^J c_D(k+1+i)2^{-i} - C_D(k).$$

This normalised expression, not necessarily $F_k(J)$ itself, is nondecreasing in J . Its limit is $1 - 2^{k+1}r_k(1/2)$, by the finite-support Lambert identity. This is positive exactly under the strict dyadic inequality on the left. At positive k , oddness of the reduced excess numerator excludes equality at zero.

The result supplies some finite horizon, not the horizon $J = c - 3$ required in the crossing application. More quantitatively, write $\eta = 1 - 2^{k+1}r_k(1/2) > 0$. Since $\sum_{r \geq 1} c_D(m+r)2^{-r} \leq m+2$, the omitted tail shows that $(k+J+3)2^{-J} < \eta$ suffices for $F_k(J) > 0$. The value η is rational for the finite support D , so this is an effective sufficient test. Proving it at the prescribed horizon is a separate arithmetic obligation.

[Lean] (uniform) frozen-margin [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: [paper_halfGreedyPrefixSupport_eq_greedy_inter_Icc](#), [paper_eventual_nonnegative_margin_equivalence](#), [paper_frozen_margin_normalised_value](#), [paper_frozen_margin_normalised_monotone](#), and 7 further declarations in the [coverage section](#).

Equivalent conditions for infinitely many greedy skips

Theorem 6.88 (Half-membership and infinitely many skips).

$$(1/2 : \mathbb{R}) \in A \iff (\text{greedyMersenneSkippedSupport}(1/2)).\text{Infinite}.$$

Odd-denominator parity supplies the positivity clause in Theorem 6.68 automatically. Thus its cofinal-positive-skip hypothesis is exactly the assertion that the greedy skipped support is infinite, which the displayed theorem identifies with half-membership. See Observation 6.132 for the reason this restatement does not itself prove membership.

[Lean] (cofinal) topological-achievement-set [LEAN SOURCE](#) [LEAN SOURCE](#)

6.4 Invariants

This subsection collects identities and bounds used in the preceding criteria. Each statement retains its own hypotheses; a bound for one branch is not an invariant of the entire recurrence.

One-step quotient identities

Theorem 6.89 (The next floor quotient). For $d \geq 2$ and $M \geq 0$, the floor quotient satisfies

$$q(M+1, d) = 2q(M, d) + \mathbf{1}_{d|M+1}.$$

Indeed, $q(M, d) = \sum_{j=1}^{\lfloor M/d \rfloor} 2^{M-jd}$. Increasing M doubles each existing term and adds 1 exactly when $M+1$ is divisible by d . Thus a quotient doubles or doubles plus one; it does not remain unchanged except when both values are zero. Summing over a fixed finite support gives the corresponding update with added term $c_D(M+1)$.

[Lean] (uniform) binary digits [LEAN SOURCE](#)

Lean: [paper_next_floor_quotient](#), [paper_next_floor_quotient_no_fixed_point](#), [paper_next_quotient_sum](#), [paper_floor_quotient_geometric_sum](#).

Theorem 6.90 (The next quotient sum). For a fixed finite set $D \subseteq \{2, 3, \dots\}$ and $M \geq 0$,

$$Q(D, M+1) = 2Q(D, M) + c_D(M+1).$$

This is the sum of Theorem 6.89 over $d \in D$. Keeping D fixed is essential: changing the support between steps adds a separate difference of quotient sums. The exclusion of $d = 1$ is also essential for the displayed correction term, since $q(M, 1) = 2^M$ doubles without an added unit.

[Lean] (uniform) binary digits [LEAN SOURCE](#)

Theorem 6.91 (The signed endpoint recurrence). *The signed next-step expression is $H(D, k, n) = 2S(D, k, n - 1) + 1 - c_D(n)$. The last term counts the selected exponents dividing the new endpoint. This is the same recurrence shape as the generic tempered-orbit recurrence $u(N + 1) = 2u(N) - v \cdot c(N + 1)$ from the public `GenericTailOrbitRigidity.lean` T7 criterion (Theorem 6.105 below), specialised to Mersenne local repair. A matching recurrence alone does not imply rationality: Theorem 6.105 also requires one fixed coefficient sequence, integer states and a vanishing scaled limit. Those additional conditions must be checked before that criterion can be applied to these finite-row quantities.*

[Lean] (uniform) binary digits [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: [paper_signed_endpoint_recurrence](#), [paper_endpoint_term_counts_divisors](#), [paper_signed_endpoint_defect_succ](#), [paper_repair_integer_eq_endpoint_defect](#).

Bounds for the remaining binary positions

Definition 6.92 (The sharper binary bound). Let $c \geq 4$, $D \subseteq \{2, \dots, c - 1\}$, and $X_D(2) < 1/2$. The sharper bound is the condition

$$S(D, 1, 2c - 2) < 2^{c-2}.$$

It permits a binary completion using ranks $c + 1, \dots, 2c - 2$, with rank c omitted, as explained below.

[Lean] (n/a) binary digits [LEAN SOURCE](#)

Theorem 6.93 (Equivalent forms of the sharper bound). *Under the hypotheses of Definition 6.92:*

$$S(D, 1, 2c - 2) < 2^{c-2} \iff 2^{(2c-2)-1} \leq Q(\text{insert } c \text{ } D, 2c - 2).$$

The threshold on the right is one greater than the exact-row target $2^{2c-3} - 1$. Thus this is a test for omitting rank c , not for including it in the completed row.

[Lean] (uniform) binary digits [LEAN SOURCE](#)

Proof. Put $M = 2c - 2$. Since $X_D(2) < 1/2$, one has $Q(D, M) < 2^{M-1}$, and the integer remainder is $S(D, 1, M) = 2^{M-1} - Q(D, M) - 1 \geq 0$ without truncation. The geometric quotient identity gives $q(M, d) = 2^{M-d}$ for $c \leq d \leq M$, since $2d > M$. In particular $q(M, c) = 2^{c-2}$, so

$$S(D, 1, M) < 2^{c-2} \iff Q(D \cup \{c\}, M) \geq 2^{M-1}.$$

For the completion, use the ranks $c + 1, \dots, M$ instead. Their quotient weights are $2^{c-3}, \dots, 1$ and represent every integer from 0 to $2^{c-2} - 1$. A subset of them therefore has quotient sum $S(D, 1, M)$ and completes D to the exact target $2^{M-1} - 1$ while omitting c . These finite identities prove the equivalence and explain the sharper bound. They do not assert that a completed row stays below one half in real value. $\square$

For example, $D = \{2, 3\}$ and $c = 5$ give $M = 8$, $Q(D, 8) = 121$ and remainder $6 < 8$. Inserting rank 5 adds 8, giving 129 rather than the exact target 127. The correct completion adds ranks 6 and 7, whose quotients are 4 and 2. [Math]

Theorem 6.94 (An unconditional bound with one extra bit). *Unconditionally, for every $c \geq 4$ and every below-half core $D \subseteq [2, c)$ with deficit $< w(c)$:*

$$S(D, 1, 2c - 2) < 2^{c-1}.$$

The bound holds uniformly in c , but is twice the threshold needed in Theorem 6.93 to apply Theorem 6.75. Before using $|D| \leq c - 2 \leq 2^{c-2}$, the proof gives the sharper additive estimate $S(D, 1, 2c - 2) < 2^{c-2} + |D|$. Thus the critical-capacity inequality would follow by excluding the integer band $[2^{c-2}, 2^{c-2} + c - 3]$, which contains $c - 2$ integers. This is the same band shape as Theorem 6.79's dyadic-band condition. Its width is linear in c , whereas the square-root reset condition discussed in Section 12.5 has width $2^{(r+5)/2}$ in a different parameter r . Comparing these widths alone proves no implication between the two hypotheses.

[Lean] (uniform) binary digits [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: [paper_unconditional_bound_one_extra_bit](#), [paper_sharper_additive_estimate](#), [paper_capacity_band_exclusion](#).

Doubling the endpoint

Theorem 6.95 (The binary bound after doubling). *Let $n \geq 6$ and $D \subseteq \{2, \dots, n\}$ satisfy $2 \in D$, $Q(D, n) = 2^{n-1} - 1$, and $X_D(2) < 1/2$. Then*

$$S(D, 1, 2n - 1) < 2^{n-1}.$$

The missing quotient can therefore be filled using only the ranks $n + 1, \dots, 2n - 1$, without changing D .

[Lean] (uniform) binary digits [LEAN SOURCE](#)

Lean: [localBinarySuffix_two_mul_sub_one_lt_upperWindow_of_exact_below](#), [exists_exactRowStrictUpperExtension_two_mul_sub_one_of_exact_below](#).

Proof. Write $F_M = \sum_{d \in D} \{2^M / (2^d - 1)\}$ and $P = 2^{n-1}$. The exact row at n gives $2^n(1/2 - X_D(2)) = 1 - F_n$. At $M = 2n - 1$ this becomes

$$S(D, 1, 2n - 1) = P(1 - F_n) + F_{2n-1} - 1.$$

The below-half assumption makes this an ordinary integer difference, so no truncation is needed. Since $2 \in D$, one has $F_n \geq 1/3$; also $F_{2n-1} \leq |D| \leq n - 1$. Therefore

$$S(D, 1, 2n - 1) \leq \frac{2}{3}P + n - 2 < P,$$

using $3(n - 2) < 2^{n-1}$ for $n \geq 6$. The ranks $n + 1, \dots, 2n - 1$ supply exactly $n - 1$ binary digits, with total capacity $2^{n-1} - 1$. The binary expansion of the nonnegative missing quotient gives the required extension. $\square$

For example, $n = 6$ and $D = \{2, 3, 6\}$ give $X_D(2) = 31/63$, $Q(D, 11) = 1006$ and missing quotient $1023 - 1006 = 17$. The ranks 7 and 11 supply $16 + 1$. The resulting exact row $E = \{2, 3, 6, 7, 11\}$ has $X_E(2) - 1/2 = 13955/32756094 > 0$. Thus the extension is exact at its integer target, but need not remain below half in real value.

Proposition 6.96 (Every exact sum contains exponent two). *Let $n \geq 3$ and let $D \subseteq \{2, \dots, n\}$ satisfy $Q(D, n) = 2^{n-1} - 1$. Then $2 \in D$. Thus the rank-two hypothesis in Theorem 6.95 is automatic for its exact rows. The interval restriction on D is part of the assertion, not an assumption about arbitrary finite quotient sums.*

[Lean] (uniform) binary digits [LEAN SOURCE](#)

The depth-six example

Definition 6.97 (An exact sum at depth six). $\text{exactRowSixSupport} := \{2, 3, 6\}$ satisfies

$$Q(\{2, 3, 6\}, 6) = 2^5 - 1 = 31, \quad V(\{2, 3, 6\}) < 1/2,$$

hence $\text{ExactLocalMersenneHalfRow}(6)$. Closed numeric fixture, norm_num-checked. This is the chosen seed for the $n \geq 6$ transition theorem; the induction starts at this depth. It is not the first exact quotient row without that restriction: at $n = 2$, $D = \{2\}$ already gives $Q(D, 2) = 1 = 2^{2-1} - 1$.

[Cert] (fixed) divisor counts and finite sums [LEAN SOURCE](#) [LEAN SOURCE](#)

Approximation error and finite denominators

Proposition 6.98 (Error in the finite subseries value). *An exact row at endpoint n has real value within $O(n/2^n)$ of $1/2$: the quantitative bound $|y_n - 1/2| \leq (n+1)/2^n$ consumed by Theorem 6.67 to turn a sequence of exact rows tending to depth infinity into a sequence of values tending to $1/2$.*

[Lean] (uniform) divisor counts and finite sums [LEAN SOURCE](#)

Proposition 6.99 (No finite support has value one half). *A finite sum of reciprocals of odd integers has odd denominator in lowest terms, so it cannot equal $1/2$. Applied to the denominators $2^a - 1$, this proves that a support representing $1/2$ must be infinite, as used in Theorem 6.67.*

[Cited] (uniform) p-adic [LEAN SOURCE](#)

Lean: [paper_finite_sum_inv_odd_den_odd](#), [paper_finite_sum_inv_odd_ne_half](#), [paper_finiteErdosSum_den_odd](#), [paper_finite_support_series_ne_half](#), and 1 further declaration in the [coverage section](#).

Uniqueness for gap-dominated finite weights

Theorem 6.100 (Small remainder forces the greedy word). *Let $g \geq 1$ and let $w_1, \dots, w_m$ be nonnegative integer weights satisfying*

$$w_i \geq g + \sum_{j>i} w_j \quad (1 \leq i \leq m).$$

Fix an integer capacity $C \geq 0$. The greedy word γ is obtained by taking w_i exactly when it does not exceed the remaining capacity. For an admissible Boolean word ε , meaning $\sum_i \varepsilon_i w_i \leq C$, one has

$$C - \sum_i \varepsilon_i w_i < g \quad \Longleftrightarrow \quad \varepsilon = \gamma \text{ and } C - \sum_i \gamma_i w_i < g.$$

Indeed, two different words have sums separated by at least g , as seen at their first differing index; the greedy sum is the largest admissible sum. Thus there is at most one admissible word with remainder below g , and any such word is greedy. The greedy word alone need not have remainder below g : for weights (5), gap $g = 2$ and capacity $C = 3$, its remainder is 3.

[Lean] (uniform) greedy recurrence [LEAN SOURCE](#)

Proposition 6.101 (Gap domination of the integer weights). *For integers $1 \leq d \leq R \leq M$, the quotient weights satisfy*

$$\left\lfloor \frac{2^M}{2^d - 1} \right\rfloor \geq 2^{M-R} + \sum_{j=d+1}^R \left\lfloor \frac{2^M}{2^j - 1} \right\rfloor.$$

Thus the finite weight list at ranks $2, \dots, R$ satisfies the hypothesis of Theorem 6.100 with $g = 2^{M-R}$. At $M = 2R - 1$ this gap is 2^{R-1} , and at $M = 2R$ it is 2^R . The inequality is unconditional; it does not assert that the greedy remainder is smaller than the gap.

[Lean] (uniform) greedy recurrence [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

Proposition 6.102 (Uniqueness of the finite greedy representation). *Let $1 \leq M$, $0 \leq R \leq M$, and $0 \leq a < 2^{M-R}$. If*

$$\sum_{d=2}^R \varepsilon_d \left\lfloor \frac{2^M}{2^d - 1} \right\rfloor + a = 2^{M-1} - 1, \quad \varepsilon_d \in \{0, 1\},$$

then ε is the greedy word for capacity $2^{M-1} - 1$ and a is its remainder. This follows by applying Theorem 6.100 with the preceding gap inequality. The statement identifies any such representation; it does not prove that a representation with $a < 2^{M-R}$ exists. The degenerate formal case $M = 0$ has empty word and zero target, using truncated natural-number subtraction.

[Lean] (uniform) greedy recurrence [LEAN SOURCE](#)

A geometric form of the quotient condition

Theorem 6.103 (A division-free form of the binary bound). *For $M \geq 0$ and $d \geq 2$,*

$$\left\lfloor \frac{2^M}{2^d - 1} \right\rfloor = \sum_{j=1}^{\lfloor M/d \rfloor} 2^{M-jd},$$

where an empty sum is zero. To see this, write $M = qd + r$ with $0 \leq r < d$ and expand the finite geometric sum. The remaining fraction is $2^r / (2^d - 1)$, which lies strictly between 0 and 1. Substitution in Theorem 6.93 expresses its quotient condition as a finite sum of powers of 2. This is an exact rewriting, not a weaker hypothesis or a new existence result.

[Lean] (uniform) geometric-series [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: `localMersenneGeometricQuotient`, `localMersenneQuotient_eq_geometric`, `localBinarySuffix_two_mul_sub_two_lt_criticalCapacity_iff_geometric`, `localBinarySuffix_two_mul_sub_two_lt_criticalCapacity_iff_geometricCore`.

Bounds for a general coefficient sequence

Theorem 6.104 (An upper bound for a binary coefficient tail). *Let $c : \mathbb{N} \rightarrow \mathbb{N}$ satisfy $c(n) \leq n$ for every n . Then*

$$\sum_{r \geq 1} c(N+r)2^{-r} \leq N+2 \quad (N \geq 0).$$

Indeed, $c(N+r) \leq N+r$, while $\sum_{r \geq 1} 2^{-r} = 1$ and $\sum_{r \geq 1} r2^{-r} = 2$. Thus the scaled tail is $O(N)$, and in particular $o(2^N)$, as required in Theorem 6.105. The coefficient bound also holds for Euler's totient function.

[Cited] (uniform) binary digits [LEAN SOURCE](#)

Lean: `binaryCoeffTail_le, binaryCoeffTail_div_pow_tendsto_zero`.

Theorem 6.105 (Rationality through an integer recurrence). *Let $c : \mathbb{N} \rightarrow \mathbb{N}$ satisfy $c(n) \leq n$ for every n . Then $\sum_{n \geq 1} c(n)2^{-n}$ is rational if and only if there exist a positive integer v and a sequence $u : \mathbb{N} \rightarrow \mathbb{Z}$ such that*

$$u(N+1) = 2u(N) - vc(N+1) \quad (N \geq 0), \quad \frac{u(N)}{2^N} \rightarrow 0.$$

The limit condition is $u(N) = o(2^N)$; it does not assume that u is bounded. No divisor-count hypothesis is imposed on c . Thus the ordinary statement also applies to $c = \varphi$; the identification with the separate formal #249 development, including its indexing conventions, is not asserted here.

[Lean] (uniform) binary digits [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: `binaryCoeffSeries, IsTemperedBinaryOrbit, not_irrational_binaryCoeffSeries_iff_exists_temperedBinaryOrbit`.

The equivalence follows by telescoping the recurrence:

$$\frac{u(N)}{2^N} = u(0) - v \sum_{n=1}^N \frac{c(n)}{2^n}.$$

The limit condition gives the rational value $u(0)/v$. Conversely, if the series equals p/v , define $u(N) = 2^N p - v \sum_{n=1}^N 2^{N-n} c(n)$. This is an integer and equals $v \sum_{r \geq 1} c(N+r)2^{-r}$, so the preceding tail bound proves the limit. This argument also explains why an arbitrary integer solution of the recurrence is insufficient: its homogeneous part can grow like 2^N .

For example, the finite support $A = \{2\}$ has $c_A(n) = \mathbf{1}_{2|n}$ and $X_A(2) = 1/3$. With $v = 3$, the sequence $u(N) = 1$ for even N and $u(N) = 2$ for odd N satisfies the recurrence. The criterion therefore includes finite supports; it does not provide the infinite support needed to contradict Problem 257.

Theorem 6.106 (Möbius inversion of the divisor counts). *For a set $A \subseteq \mathbb{N}$ and each positive integer n ,*

$$\sum_{d|n} \mu(d) c_A(n/d) = \mathbf{1}_A(n),$$

where μ is the Möbius function and $c_A(n)$ counts the positive elements of A dividing n . Thus the divisor counts determine the positive support exactly. In particular their Möbius transform takes only the values 0 and 1. The identity does not depend on a base.

[Lean] (uniform) mobius-inversion [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: `moebius_mul_supportCoeffAF`, `mobius_supportCoeff_eq_one_iff`, `mobius_supportCoeff_boolean`.

Theorem 6.107 (The support series as a coefficient series). For $A \subseteq \mathbb{N}_{>0}$,

$$X_A(2) = \sum_{a \in A} \frac{1}{2^a - 1} = \sum_{n \geq 1} \frac{c_A(n)}{2^n}.$$

Expanding each denominator as a geometric series and interchanging nonnegative sums gives the identity. Since $c_A(n) \leq \tau(n) \leq n$, Theorem 6.105 applies. Together with Theorem 6.106, it expresses rationality through an integer recurrence whose coefficients all come from the same support.

[Lean] (uniform) divisor counts and finite sums [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: `erdosSupportSeries_two_eq_binaryCoeffSeries`, `supportCoeff_le_card_divisors`, `supportCoeff_le_self`, `erdosSupportSeries_rational_iff_exists_temperedCarry`.

Restrictions on a support with rational value

The next results impose separate restrictions on rational-valued supports. Their sources are `RationalSupportCarrySkeleton.lean`, `SublogDivisorCoverage.lean` and `MaximalOmegaLayer.lean`. The related strict-gap result is Proposition 6.99.

Theorem 6.108 (A restriction on dyadic rational values). If an infinite support $A \subseteq \mathbb{N}_{>0}$ has $X_A(2) = p/2^c$ for integers p and $c \geq 0$, then $\sum_{a \in A} 1/a$ either diverges or converges to a value greater than 1. The cited proof averages the shifted integer recurrence and uses a common multiple of two distinct support elements. Thus a convergent reciprocal sum of at most 1 is excluded. This is a separate necessary condition: the reciprocal-summable criterion proved earlier already excludes every support with a convergent reciprocal sum, not just those whose sum is at most 1.

[Lean] (uniform) cesaro-tail [LEAN SOURCE](#) [LEAN SOURCE](#)

Theorem 6.109 (Unboundedness of a positive shifted recurrence). Let $A \subseteq \mathbb{N}_{>0}$ be infinite, fix a shift $c \geq 0$ and a positive integer v , and let u be a positive integer sequence satisfying

$$u(n+1) + vc_A(c+n+1) = 2u(n) \quad (n \geq 0).$$

Then u is unbounded. Indeed, under a proposed bound $u(n) \leq B$, choose $2B+1$ distinct support elements and a common multiple larger than c . At the corresponding index the divisor count is at least $2B+1$, contradicting the recurrence. This does not conflict with Theorem 6.105, whose condition is $u(n)/2^n \rightarrow 0$, not boundedness. An unbounded sequence can still satisfy that limit.

[Lean] (uniform) common-multiple-forcing [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

Theorem 6.110 (Bounds for intervals with zero divisor counts). *Suppose that $X_A(2) = p/(2^c v)$ for an infinite positive support A and an odd positive integer v . For every $\varepsilon > 0$ there is B such that, for $N \geq 1$, a run of h zero divisor counts starting after $c + N$ satisfies $h \leq \varepsilon \log_2 N + B$. In fact this conclusion holds without rationality: fix any $a \in A$. Every a consecutive positive integers include a multiple of a , where c_A is positive. Hence $h \leq a - 1$, so one can take $B = a - 1$. The zero-run conclusion therefore imposes no additional restriction on a fixed nonempty support, regardless of the tail estimates used in the linked proof.*

[Lean] (uniform) divisor-envelope [LEAN SOURCE](#)

Lean: [paper_zero_run_le_eps_logb](#), [paper_zero_run_le_of_mem](#).

Proposition 6.111 (A subpower bound for divisor counts). *For positive integers n, k , $\tau(n)^k \leq (k^{2^k})^k \cdot n$, or equivalently $\tau(n) \leq k^{2^k} n^{1/k}$. The constant absorbs the finitely many primes below 2^k ; for larger primes, $(v+1)^k \leq p^v$ controls each factor of the divisor product. This estimate is independent of the support problem. The elementary zero-run bound above already follows from a single positive support element, without this estimate or a recurrence.*

[Lean] (uniform) divisor-envelope [LEAN SOURCE](#)

Lean: [divisorSubpowerConst](#), [card_divisors_pow_le_divisorSubpowerConst_pow_mul](#), [card_divisors_le_divisorSubpowerConst_mul_rpow](#).

Proposition 6.112 (Two distinct prime-power differences commute). *Lemma 6.57 gives the four-term expansion and its divisor-count interpretation. Commutation is elementary for arbitrary positive multipliers. The extraction formula uses distinct primes, positive exponents e, f and $\gcd(n, pq) = 1$. The example $A = \{12\}$ shows the extracted coefficient explicitly. No rationality statement about multiplicative subsequences follows from commutation alone.*

[Lean] (uniform) mobius-inversion [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: [primePowerLayer_comm](#), [mixedPrimePowerLayerTwo_supportCoeffInt](#), [mixedPrimePowerLayerTwo_twelve_fixture](#).

The first crossing of the half-value

Definition 6.113 (The first crossing rank).

$$\text{localMersenneCrossingRanks}(E) := E.\text{filter}(c \mapsto \frac{1}{2} < V(E.\text{filter}(\cdot \leq c))),$$

the ranks at which E 's inclusive ordered prefix is already above one half.

[Lean] (n/a) binary digits [LEAN SOURCE](#)

Theorem 6.114 (Existence of a first crossing). *Let $E \subseteq \{2, 3, \dots\}$ be finite and suppose $X_E(2) > 1/2$. There is a least $c \in E$ such that*

$$X_{E \cap [2, c)}(2) < \frac{1}{2} < X_{E \cap [2, c]}(2), \quad c \geq 4.$$

Positivity of the summands makes the partial sums increasing. Their odd reduced denominators exclude equality with $1/2$, so the first crossing is strict on both sides. Finally $w_2 + w_3 = 10/21 < 1/2$, which excludes a crossing before rank 4. This is the finite crossing data used in the later exact-row constructions; the lower bound on c and strictness are specific to these weights.

[Lean] (uniform) binary digits [LEAN SOURCE](#)

Lean: [localMersenneCrossingRanks](#), [exists_first_localMersenne_crossing](#).

Centred carries at skipped ranks

Theorem 6.115 (Nonnegative centred carries below the half-value). *If $1 \notin A$ and $X_A(2) < 1/2$, then*

$$C_A(N) = \text{ihc}(A, N) - 1 \geq 0 \quad (N \geq 0).$$

Indeed, Lemma 2.4 gives

$$\text{ihc}(A, N) = 2^{N+1}(1/2 - X_A(2)) + \sum_{r \geq 1} c_A(N+1+r)2^{-r} > 0.$$

The first term is strictly positive and the tail is nonnegative. Since the half-carry is an integer, it is at least 1; subtracting 1 proves the centred bound. Nonnegativity alone, without this strictness and integrality step, would only give $C_A(N) \geq -1$. The argument supplies the lower bound used with the conditional upper bound in Theorem 6.84.

[Lean] (uniform) mobius-centred-carry [LEAN SOURCE](#) [LEAN SOURCE](#)

Theorem 6.116 (Three possibilities at a skipped endpoint). *For a rank $s \geq 5$ omitted by the real greedy support, let H_s and f_s be as in Lemma 6.45. That lemma identifies the actual prefix as D_s when $f_s \leq 0$ and as B_s when $f_s > 0$, with the exact remainder or overshoot in each case. To obtain the nonnegative-margin condition in Definition 6.81, the negative case would still have to be excluded at every required skipped rank. The trichotomy itself does not exclude it. In particular, the value -3 excluded for C_D under the all-right-tail hypothesis of Theorem 6.44 is not an exclusion for f_s : the coordinates and the hypotheses are different.*

[Lean] (uniform) frozen-margin [LEAN SOURCE](#)

Gap lengths and the measure of their union

Definition 6.117 (The gap after a Mersenne weight). For $n \geq 1$, let $R_n = \sum_{j>n} w_j$ and $g_n = w_n - R_n > 0$. The quantity g_n is the length of each gap created at level n , not the sum of all gap lengths created at that level.

[Lean] (n/a) divisor counts and finite sums [LEAN SOURCE](#) [LEAN SOURCE](#)

Theorem 6.118 (The unweighted sum of remaining gap lengths). *For every $N \geq 0$, the unweighted sum of gap lengths satisfies*

$$\sum_{n>N} g_n \leq \frac{2}{9} 4^{-N} + \frac{3}{7} 8^{-N},$$

which tends to zero. The estimate follows by summing the per-level bound; it counts one length per level. It does not, by itself, bound the measure of the union of all gaps, because there are 2^{n-1} disjoint translated gaps at level n .

[Lean] (uniform) summed-gap-mass [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

The measure of the union of gaps created after level N is instead

$$\sum_{n>N} 2^{n-1} g_n = 2^N R_N - 1 \leq \frac{2}{3} 2^{-N} \quad (N \geq 1).$$

Indeed, the finite sum from $N + 1$ to M telescopes to $2^N R_N - 2^M R_M$, and $2^M R_M \rightarrow 1$. The final bound uses $R_N \leq 2^{-N} + (2/3)4^{-N}$. Disjointness follows from the strict-tail construction of the achievement set. This is an ordinary geometric deduction, not the statement of the preceding linked gap-sum lemma. A small union measure still gives no answer about a prescribed point. [Math]

A sufficient test for a safe skip

Definition 6.119 (Three geometric terms in the tail lower bound). `mersenneTailLB3` $(k) := 1/2^k + 1/(3 \cdot 4^k) + 1/(7 \cdot 8^k) < R(k)$; a strict rational lower bound obtained by truncating the Lambert-type series at three terms.

[Lean] (n/a) `lambert-bound` [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

Theorem 6.120 (A sufficient inequality for a safe skip). *Let $k, u, L \geq 1$ be integers, set $a = 2L - (2^k - 1)u$, and suppose $a > 0$. For the skipped rational remainder $\rho = u/(2L)$, $\rho \leq 2^{-k}$ is equivalent to $u \leq a$, whereas*

$$2u \leq 3a \implies \rho < R_k.$$

The latter implication uses the three-term lower bound above. More explicitly, with $t = 2^k \geq 2$,

$$\left(\frac{1}{t} + \frac{1}{3t^2} + \frac{1}{7t^3} \right) - \frac{3}{3t-1} = \frac{2t-3}{21t^3(3t-1)} > 0.$$

The inequality $2u \leq 3a$ gives $\rho \leq 3/(3t-1)$, proving the claim. The first two terms alone would not prove this comparison. This is a sufficient exclusion of the current tail-mass deficit, not an exact test for membership in the remaining achievement set.

The sufficient condition is weaker than $u \leq a$. A realizable example is $(k, u, L, a) = (2, 7, 13, 5)$; it passes $2u \leq 3a$ but not $u \leq a$. The linked scalar example $(u, a) = (3, 2)$ proves strict containment of the inequalities, but is not integral rational data for this substitution. For comparison, the exact mass threshold is $a/u \geq R_k^{-1} - (2^k - 1)$; its right-hand side lies strictly between 0 and $2/3$. None of these comparisons asserts that the half-greedy orbit satisfies the sufficient inequality at every skipped rank.

[Lean] (uniform) `lambert-bound` [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: [paper_dyadic_skip_test_iff](#), [paper_sharp_skip_safe_lb3](#), [paper_sharp_skip_safe_actual_tail](#), [paper_three_channel_margin_identity](#), and 5 further declarations in the [coverage section](#).

Proposition 6.121 (Two unconditional safety cases). *Under the positive-integer and skipped-step hypotheses of Theorem 6.120, $u = 1$ implies $a \geq 1$, so the current remainder is less than R_k . Conversely, a fatal tail-mass deficit forces $3a < 2u$, hence $u \geq 2$, or $u \geq 3$ when u is odd. These are statements about the current step, not infinite survival. The exact declaration names are linked below.*

[Lean] (uniform) `lambert-bound` [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

Two-sided bounds near dyadic endpoints

Theorem 6.122 (Two-sided dyadic bounds). *Under the local hypothesis excluding the specified three middle cells and imposing the right-pulse bound, induction (base case at row 5 verified by decide) propagates it to the universal two-sided bound*

$$\forall s \geq 5, \quad \min(\text{rem}(s), \text{overshoot}(s)) \leq 2^s.$$

This conditional bound says that at least one of the remainder and the adjacent overshoot is at most 2^s . It does not assert that both are at most 2^s ; that stronger assertion would replace the minimum by a maximum. It is not needed for the charge bound at an actual upper reset: Proposition 6.80 obtains that bound directly from the upper-reset identity. Applying the inductive argument to another recurrence would require its own transition and separation estimates.

[Lean] (uniform) dyadic-scale [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: [paper_two_sided_dyadic_bound](#).

Proposition 6.123 (A finite band check for $13 \leq d \leq 30$). *For every actual upper-reset index $13 \leq d \leq 30$ and every $0 \leq j \leq d$, the linked certificate verifies*

$$2^{d-j+1} < \text{resetCharge} \quad \text{or} \quad \text{resetCharge} + 2(d+j) \leq 2^{d-j+1}.$$

It computes the successor remainder at each row $d+1$, from $\text{rem}(14) = 392$ through $\text{rem}(31) = 4187487147$. This proves the required condition on a finite range, not the universal condition for all $d \geq 13$. A proof for unbounded reset indices must supply a further argument; repeatedly extending the table alone does not establish that statement.

[Cert] (fixed) dyadic-boundary [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: [paper_finite_band_check_thirteen_to_thirty](#), [paper_successor_remainders_fourteen_through_thirtyone](#), [paper_universal_band_condition_unfolded](#), [paper_universal_band_condition_would_close_half](#).

General identities for perturbed greedy recurrences

Theorem 6.124 (A general perturbed greedy recurrence). *Let the old integer values be separated by at least $g \geq 1$, and let the update be $t(x) = 4s(x) + p(x)$ with $0 \leq p(x) \leq B < g$. Suppose x_- and x_+ are the adjacent old values on either side of capacity C . Theorem 6.62 gives the largest admissible updated value at capacity $4C + g$ and the exact three-branch remainder after testing the extra weight $2g + 4$.*

The underlying structure in Definition 6.61 assumes only $B < 3g$; that suffices for order preservation, but the cited maximality application has the additional hypothesis $B < g$. The counterexample after Theorem 6.62 shows what can fail without it. For a maximum over both choices of the extra weight, rather than the stated two-stage rule, one also needs a separation argument; $4g - B \geq 2g + 4$ suffices, as shown after that theorem.

[Lean] (uniform) abstract-perturbed-greedy [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: [paper_perturbed_order_preservation](#), [paper_perturbed_prefixChoice_maximal](#), [paper_perturbed_nextRemainder_three_branches](#), [paper_perturbed_separation_global_maximality](#), and 3 further declarations in the [coverage section](#).

Theorem 6.125 (Spacing of compatible reverse-carry words). *Consider two integer carry recurrences $b_i(m) + 2u_i(m) = a_i(m) + u_i(m + 1)$, $i = 1, 2$. If $a_1(k) = a_2(k)$ and the output bits at k are 1 and 0, then the carry difference at $k + 1$ is odd. If the coefficients and bits agree for the following L positions, the difference at $k + L + 1$ is 2^L times that odd integer. Consequently absolute terminal bounds B_1, B_2 give $2^L \leq B_1 + B_2$. Lemma 6.46 states the exact identity and proves it by subtraction. The linked formal statements encode these agreements; they do not supply terminal bounds for an unrelated digit system.*

[Lean] (uniform) reverse-carry-word [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

Proposition 6.126 (Changing one support bit at a doubled rank). *Let two supports agree except that the second includes $N + 1$ while the first does not. At the argument $2(N + 1)$ their divisor counts differ by exactly 1, because $N + 1$ divides $2(N + 1)$. More generally the difference at a positive argument m is $\mathbf{1}_{N+1|m}$. The linked statements apply this identity to the finite supports in their hypotheses. No conclusion about a different coefficient sequence follows without identifying its own support change.*

[Lean] (uniform) half-divisor [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: `supportCoeff_insert_divisor, supportCoeff_extend_true_eq_false_add_one_at_double, supportCoeff_boundaryPair_unitDrop_at_double`.

6.5 Obstructions and countermodels

The results below test particular proposed implications. A countermodel to an abstract transition rule shows that the rule needs an additional hypothesis; it does not disprove a more structured arithmetic argument. Each entry states the information retained by its model and the conclusion that this information fails to imply.

Theorem 6.127 (A bounded model of the doubling-or-return alternative). *The dichotomy $\text{ExactLocalMersenneHalfRow}(2n-1) \vee \exists c, 4 \leq c \leq n \wedge \text{ExactLocalMersenneHalfRow}(2c-2)$ (proved for $n \geq 6$ at Theorem 6.128 below) is not by itself enough for cofinality. Countermodel: $\text{boundedDoubleOrRecycleModel}(n) := (n = 6)$ satisfies exactly the same two-branch transition shape (seed at 6, and `boundedDoubleOrRecycleModel_transition` reproduces the $\vee$ shape by always taking the recycle branch with $c = 4$, conclusion back at $2 \cdot 4 - 2 = 6$), yet*

$\neg(\forall N, \exists n \geq N, \text{boundedDoubleOrRecycleModel}(n))$ (not cofinal ; only ever true at $n = 6$).

Packaged existentially as `exists_seeded_bounded_double_or_recycle_model`. The double-or-recycle transition shape, even together with an endpoint-six seed, is logically insufficient to conclude cofinal exact rows. A genuinely new progress input (strict endpoint growth, or an independent cofinality argument) is required. The countermodel retains only the seed and this transition rule. It therefore tests those premises, not additional arithmetic information that a more structured induction may use.

[Lean] (fixed) meta-logical [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: `paper_bounded_double_or_recycle_countermodel, paper_exists_seeded_bounded_double_or_recycle_model`.

Theorem 6.128 (Doubling or returning to an earlier depth). *For $n \geq 6$, $\text{ExactLocalMersenneHalfRow}(n)$:*

$$\text{ExactLocalMersenneHalfRow}(2n - 1) \vee \exists c, 4 \leq c \leq n \wedge \text{ExactLocalMersenneHalfRow}(2c - 2).$$

Every exact row of depth at least 6 has at least one of the two asserted continuations. The proof separates the below-half case from the above-half case, but the two existential conclusions need not be exclusive. The returned depth $2c - 2$ is not claimed to exceed n . The finite Mersenne value cannot equal $1/2$, by Observation 6.138 (`finiteErdosSum_den_odd`); this says nothing about equality of the returned and original depths. For example, at $n = 6$, the support $\{2, 3, 6\}$ is an exact row, and so is $\{2, 3, 6, 7, 11\}$ at depth 11. The second conclusion also holds with $c = 4$, since $2c - 2 = 6$. Thus both conclusions hold in this example, and the second allows no depth increase. What is missing, per Theorem 6.127, is not another dichotomy but strict endpoint progress in the recycle branch, or a proof the below-half branch recurs; re-deriving this disjunction without an additional progress argument would not establish cofinality.

[Lean] (bounded) divisor counts and finite sums [LEAN SOURCE](#)

Lean: [paper_exact_row_double_or_recycle](#), [paper_finite_row_value_ne_half](#), [paper_exact_row_example_six_and_eleven](#).

Proposition 6.129 (A sufficient fractional-mass bound need not hold). *Take $D = \{2, 3\}$ and $c = 5$. Direct calculation gives*

$$X_D(2) = \frac{10}{21} < \frac{1}{2} < \frac{331}{651} = X_{D \cup \{5\}}(2), \quad S(D, 1, 8) = 6 < 8.$$

However, the combined fractional mass is

$$\sum_{d \in D \cup \{5\}} \frac{2^8 \bmod (2^d - 1)}{2^d - 1} = \frac{757}{651} > 1.$$

The linked `norm_num` proof checks this example. Thus the combined fractional-mass bound by 1, which suffices for the sharp capacity estimate, is not necessary for that estimate. Thus the fractional-mass bound cannot hold at every real crossing core. A proof using it would need to handle the exceptional cores separately; the example does not rule out all uses of fractional-mass estimates. Theorem 6.133 identifies the support at a critical crossing with a real greedy prefix, but does not prove the missing bound on those prefixes.

[Cert] (fixed) divisor counts and finite sums [LEAN SOURCE](#)

Lean: [paper_fractional_mass_bound_not_necessary](#), [paper_fractional_mass_bound_suffices_for_sharp_capacity](#).

Observation 6.130 (A doubled finite sum can exceed one half). Theorem 6.95's doubling extension is not preserved iteratively: below-halfness is not preserved under one application of doubling, and demonstrably fails at the very first step. From the identity $\text{value}(E) - \frac{1}{2} = (\text{localFractionMass}(E, 2n-1) - 1)/2^{2n-1}$, at the seed $D = \{2, 3, 6\}$ (Definition 6.97) with $M = 11$ the core alone contributes

$$\frac{2^{11} \bmod 3}{3} + \frac{2^{11} \bmod 7}{7} + \frac{2^{11} \bmod 63}{63} = \frac{2}{3} + \frac{4}{7} + \frac{32}{63} \approx 1.746 > 1,$$

so the doubled row is strictly above half. Iterating the cheap (doubling) arm is therefore impossible from the only recorded seed; the chain is forced into the capacity-gated recycle arm at step one. What would close the doubling route instead: cofinally many BELOW-HALF exact rows, $\forall N \exists n \geq N \exists D$ exact at n with $\text{localFractionMass}(D, n) < 1$; a single such row at each of cofinally many n would make doubling unnecessary; a self-reproducing one would close everything.

[Cert] (fixed) *divisor counts and finite sums* [LEAN SOURCE](#)

Theorem 6.131 (The returning endpoint need not be larger). *Theorem 6.78* (`exists_skippedCoreExactRow_of_value_above`) is unconditional, but its witness is $\exists c \leq n$, not $\exists c$ large: $2c-2$ may be $\leq n$, so the endpoint need not grow. *Theorem 6.127's* `exists_seeded_bounded_double_or_recycle_model` is the explicit falsifier of the naive hope that growth comes for free: `boundedDoubleOrRecycleModel(n) := (n = 6)` satisfies the same transition schema plus a seed and is not cofinal. Growth is recovered only inside `ProtectedExactLocalMersenneRow` (*Theorem 6.71*), whose invariants `endpoint < 2 · cutoff` and `new_above_cutoff` force $c > \text{cutoff}$ hence $2c - 2 > \text{endpoint}$; and maintaining those invariants is precisely what needs the strict-upper (sharp capacity) fill of *Theorem 6.75* rather than this general recycling theorem. [Lean] (bounded) *binary digits* [LEAN SOURCE](#)

Lean: `paper_skipped_core_recycling_witness_bounded`, `paper_returning_endpoint_may_fail_to_grow`, `paper_exists_seeded_bounded_double_or_recycle_model`, `paper_protected_row_crossing_beyond_cutoff`, and 1 further declaration in the [coverage section](#).

Observation 6.132 (The skip condition is equivalent to membership). The positivity conjunct of *Theorem 6.68's* hypothesis, $0 < r(1/2, c-1)$, is unconditionally dischargeable (`localMersennePrefixValue_halfGreedy_lt_half`, itself just the odd-denominator parity fact of *Observation 6.138*, together with `greedyMersenneRemainderRat_eq_sub_finiteErdosSum`). Hence

$\text{CofinalPositiveHalfGreedySkips} \iff (\text{greedyMersenneSkippedSupport}(1/2)).\text{Infinite}$,

which *Theorem 6.88* (A6) proves equivalent to $1/2 \in A$. Proving *Theorem 6.68's* hypothesis would therefore settle the half-target question. This equivalence does not by itself make the hypothesis easier, but it also does not preclude a useful argument in greedy coordinates. *Theorem 6.75* offers a different sufficient condition whose finite supports need not be greedy prefixes.

[Lean] (cofinal) *greedy recurrence* [LEAN SOURCE](#)

Theorem 6.133 (Uniqueness at a critical crossing). *For $c \geq 4$, D bounded $[2, c)$, below-half, with genuine crossing deficit $\frac{1}{2} - \text{value}(D) < w(c)$:*

$$D = \text{halfGreedyPrefixSupport}(c - 1).$$

*Thus $D = G \cap \{1, \dots, c - 1\}$: at a critical crossing the support is fixed by c . The proof uses the strict inequality between each Mersenne weight and the sum of all later weights, as in *Theorem 6.100*. Consequently, the support quantifier in *Theorem 6.70* does not allow arbitrary choices of D . This reduces that hypothesis to a statement about the actual greedy sequence. It does not rule out proving a new bound for that sequence.*

[Lean] (uniform) *greedy recurrence* [LEAN SOURCE](#)

Lean: [paper_critical_crossing_support_is_greedy_prefix](#).

Observation 6.134 (Different requirements on finite witnesses). Theorem 6.85 (the cofinal terminal-strip condition) tolerates any support inside the depth- M window but demands the carry inside a $\sim 2\sqrt{M}$ strip; Theorem 6.75 tolerates a carry up to 2^{c-2} (exponentially looser) but demands the support be confined to ranks $< c$, because the exact fill encodes the residue in the pure upper window where the Mersenne quotient is exactly 2^{M-d} (Proposition 6.77). The two conditions therefore differ in the permitted support as well as the size of the carry. If the square-root witnesses could also be chosen with support in $[2, c)$ at depth $2c - 3$, then $2\sqrt{2c} + 4 < 2^{c-2}$ for $c \geq 8$ would give the numerical bound needed for the fill. The required support restriction is not proved here. Comparing the displayed bounds alone does not order the two existence hypotheses.

[Lean] (cofinal) *sqrt-carry-growth* [LEAN SOURCE](#)

Observation 6.135 (An unsafe example in the two-thirds interval). Read in full alongside HalfGreedyFatalGap for scope contrast: writing a greedy residual as $\text{rem} = 1/R$, a skip at rank k is dyadically safe iff $R \geq 2^k$. Two unconditional corollaries hold (an integral reciprocal R is never unsafe, `not_twoThirdsBand_of_int`; an unsafe run with odd p, D, q forces $p \geq 7$, `seven_le_of_intBand_odd`) ; but the file's own docstring is explicit that dyadic safety is sufficient, not necessary, for the true greedy process to survive (the true test uses the actual tail $\sum_{j>k} \text{weight}(j)$, which exceeds 2^{-k}), and nothing in this file asserts the actual half-greedy orbit avoids the band: $(p, D, b) = (17, 41, 3)$ is an explicit odd-coprime example that IS dyadically unsafe. This sharpens directly into Theorem 6.120's $2u \leq 3a$ criterion, which narrows exactly this gap for the single-skip case ; but the general question of whether the fatal region overlaps a reachable band under the actual greedy trajectory is settled by neither file.

[Lean] (n/a) *rational-band* [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

Theorem 6.136 (Vanishing of the specified linear-channel determinant). *Let V be a vector space over $\mathbb{Q}$, let $e : V \rightarrow \mathbb{Q}$ be linear, and let $(\ell_j)_{j \in I}$ be a finite family of linear functionals vanishing on $\ker e$. For any vectors $(v_i)_{i \in I'}$, the matrix $(\ell_j(v_i))_{i,j \in I'}$ has rank at most one, so every square minor of size at least two vanishes. Indeed, the functionals descend to $V / \ker e$, which has dimension at most one. This elementary linear-algebra argument applies at every matrix size. It does not cover additional functionals that fail to vanish on $\ker e$.*

[Lean] (uniform) *linear-algebra* [LEAN SOURCE](#)

Lean: [paper_relationInvariant_channels_rank_le_one](#), [paper_relationInvariant_channels_det_eq_zero](#).

Observation 6.137 (The additive error in the binary bound). Theorem 6.94's 2^{c-1} bound and Theorem 6.93's sharp 2^{c-2} requirement differ by one bit, and the proof of the loose bound shows exactly where: it derives $A < 2^{c-2} + |D|$ and discards $|D| \leq c - 2 \leq 2^{c-2}$. Under this additive upper bound, the remaining assertion is that the integer suffix avoids $[2^{c-2}, 2^{c-2} + c - 3]$. This identifies the unresolved interval; it does not prescribe how to exclude it. A sharper estimate, an arithmetic restriction on the suffix, or another argument could in principle establish the required inequality.

[Lean] (uniform) *binary digits* [LEAN SOURCE](#)

Observation 6.138 (Source status of finite-denominator parity). For finite $D \subseteq \mathbb{N}_{>0}$, each denominator $2^d - 1$ is odd. The sum $V(D)$ can therefore be written over an odd common denominator; its reduced denominator, being a divisor of that denominator, is odd. In particular, $V(D) \neq 1/2$. This elementary argument supplies the parity fact used in Proposition 6.129 and Proposition 6.99. The linked declaration is retained for formal-source comparison; no new kernel replay is asserted.

[Cited] (uniform) *p*-adic [LEAN SOURCE](#)

Theorem 6.139 (Excluding the cell with value minus three). *At a middle row $D \geq 13$ followed only by right transitions, $4\text{rem}(D) - p_D^- - 4 \neq -3$, by Theorem 6.44. Its proof uses the nonnegative centred carry for the completed support, not a finite search. The values $-2, -1$ remain among the three exceptional negative cells. Excluding them under this extra tail assumption would still not exclude nonnegative values of the coordinate, or establish the all-middle-row and right-branch hypotheses in Theorem 6.9. The complete remaining tail inequality is stated in Remark 10.8.*

[Lean] (uniform) *mobius-centred-carry* [LEAN SOURCE](#) [LEAN SOURCE](#) [LEAN SOURCE](#)

Lean: [paper_final_middle_cell_ne_neg_three](#), [paper_final_middle_cell_remaining_negative_values](#), [paper_mobiusCenteredHalfCarry_add_two](#), [paper_cofiniteRightTail_ne_zero_centeredEndpoint](#).

6.6 Further finite and conditional results

This subsection records further finite identities and conditional membership results from the supplied source files. They concern the middle transition, finite families with a common prefix, and rational targets other than $1/2$. Their hypotheses are stated separately from the existence assertions still needed to apply them.

6.6.1 The integer update in a middle interval

The following identities explain the two-sided bound and the reset-band condition used above. They also distinguish the signed middle coordinate $C_s = 4\text{rem}(s) - p_s^- - 4$ from the nonnegative label in the next estimate.

Proposition 6.140 (The unsafe middle range is exactly three integers). *For $s \geq 5$, write $R = \text{rem}(s)$ and $P = p_s^-$. Since $4R - P - 4$ is an integer,*

$$\neg(4R - P - 4 \leq -4 \vee 0 \leq 4R - P - 4) \iff 4R - P - 4 \in \{-3, -2, -1\}$$

Lean: [paper_unsafe_middle_range_is_three_integers](#).

([LEAN SOURCE](#), [Lean], (uniform), *divisor counts and finite sums*, proved by omega from the integer definitions). The complement of the two safe integer ranges is precisely $\{-3, -2, -1\}$. The conclusion is only a classification; excluding any of these values requires an additional argument.

For a row $s \geq 5$, set

$$E_s = \sum_{d \in D_s} \left\{ \frac{4^s}{2^d - 1} \right\}, \quad Z_s = \text{rem}(s) - E_s, \quad c = 4\text{rem}(s) - p_s^- = C_s + 4.$$

Here braces denote fractional part. Assume $c \geq 0$, as in the linked statements below. In particular, the exceptional signed values $C_s = -3, -2, -1$ correspond to $c = 1, 2, 3$, not to negative labels. The floor-quotient recurrence gives $p_s^- \leq 4E_s$, and hence

$$4Z_s \leq c.$$

This is [LEAN SOURCE](#) ([Lean]). On a middle transition $D_{s+1} = D_s$. Writing $\rho = 1/2 - X_{D_s}(2)$, the exact quotient identity therefore gives

$$4^{s+1}\rho = 2^{s+2} + 4Z_s \leq 2^{s+2} + c,$$

as in [LEAN SOURCE](#) ([Lean]).

For $C_s = -3$, the correct label is $c = 1$. This estimate does have a local consequence: since $4^{s+1}R_s > 2^{s+2} + 4/3$, it gives $\rho < R_s$. Thus the completed support $D_s \cup \{s+1, s+2, \dots\}$ has value greater than $1/2$. It does not prove that every later greedy step survives. For $s \geq 13$, under the additional all-right-tail hypothesis, the same completed support would have value less than $1/2$, giving the contradiction in [LEAN SOURCE](#). That tail hypothesis must not be dropped. For $c = 2, 3$, the displayed upper bound alone no longer gives $\rho < R_s$.

A parallel iff for the right branch: an overshoot witness $\leq 2^s$ produces a pulse leak above 2^{s+2} exactly when the overshoot sits within an explicit distance k of 2^s with $4k < \text{abovePulse}$ ([LEAN SOURCE](#), [Lean]). A new one-coordinate normal form, `seamPureHalfPrefixRemainders` $s := \text{integerGreedyRemainder}(\text{seamWeights } s)(2^{2s-1})$ ([LEAN SOURCE](#)), has an exact case-split formula ([LEAN SOURCE](#), [Lean]) whose equivalence with the two-sided bound is stated explicitly in [LEAN SOURCE](#) ([Lean]).

Observation 6.141 (The reset-band hypothesis has a certified base case). `SeamUpperResetDyadicBandEscape` (O5's second hypothesis, discussed after §12.5 above) is reduced by three new lemmas to a quarter-scale overshoot exclusion: the iff [LEAN SOURCE](#) ([Lean]) rewrites the reset-charge band as a successor-remainder avoidance; the identity [LEAN SOURCE](#) ([Lean]) rescales it by a factor of 4; and [LEAN SOURCE](#) ([Lean]) derives the band from a quarter-scale overshoot exclusion plus the pulse bound $\text{abovePulse} \leq 2(d-2)$. Row $d = 13$, the first row of the late regime, is then proved unconditionally to escape *every* band simultaneously: its successor remainder is exactly 392 ([LEAN SOURCE](#), [Lean], [Cert] for the numeral 392). This is one certified row of a hypothesis that is $\forall d \geq 13$; it is evidence the base case is not itself an obstruction, nothing more.

6.6.2 Finite approximations without compatibility

For a finite support $D \subseteq \{2, \dots, M\}$, the integer

$$K_D(M) = 2^{M-1} - \sum_{j=2}^M 2^{M-j} \#\{d \in D : d \mid j\}$$

is the half-carry at the terminal rank. It is $\text{ihc}(D, M-1)$ in the notation used earlier. The condition below asks for $M_j \rightarrow \infty$ and finite sets D_j with $|K_{D_j}(M_j)|/2^{M_j} \rightarrow 0$. It does not require the sets to agree or bound the carries at earlier ranks. The estimate

$$\left| \sum_{d \in D_j} \frac{1}{2^d - 1} - \frac{1}{2} \right| \leq \frac{|K_{D_j}(M_j)| + 2\sqrt{M_j} + 4}{2^{M_j}}$$

explains the hypothesis: it makes the finite values converge to $1/2$. Compactness then gives membership, but not the finite approximants needed to apply the theorem. A square-root terminal bound is sufficient and stronger as a bound on these particular witnesses; no claim of a strict logical separation between the corresponding existence statements follows.

Here is the finite family used in the cited construction. Put $B(m) = 2\lfloor \sqrt{m} \rfloor + 4$. A depth- M family with a shared prefix through $K \leq M$ consists of sets $A_1, \dots, A_{B(M)} \subseteq \{2, \dots, M\}$ such that

$$1 \leq K_{A_k}(m) \leq B(m) \quad (1 \leq m \leq M), \quad K_{A_k}(M) = k,$$

all A_k agree through K , and one integer $E \geq B(M)$ satisfies

$$\sum_{j=K+1}^M 1_{A_k}(j) 2^{M-j} + k = E \quad (1 \leq k \leq B(M)).$$

The suffixes therefore have consecutive binary values $E - k$. The following statements explain how these particular families supply finite approximations; they do not construct them at arbitrary depths.

Proposition 6.142. *Every shared-prefix family just defined at depth N contains a finite support A with $|K_A(N)| \leq B(N)$ ([LEAN SOURCE](#), [Lean], support rigidity, half-carry recurrence), and at a feedback row this survives both outputs of the corpus's total feedback theorem; full-cylinder advance or a localized one-hole seam, which can delete at most one of carries 3, 4 ([LEAN SOURCE](#), [Lean]).*

Lean: `paper_shared_prefix_family_contains_strip_witness`, `paper_shared_prefix_family_strip_witness_after_feedback_all_depths`, `paper_feedback_row_total_dichotomy`, `halfTerminalOnlyStripWitness_of_feedbackAdvance`, and 3 further declarations in the [coverage section](#).

Four concrete certified instances follow at depths 51–54 ([LEAN SOURCE](#), [LEAN SOURCE](#), [Lean], [Cert] for the underlying stage), which is a finite witness chain, not a cofinal one.

Observation 6.143 (Shared-prefix families at unbounded depths). The conditional theorem

$$\left(\begin{array}{l} \text{such shared-prefix families exist} \\ \text{at arbitrarily large depths } M \end{array} \right) \Rightarrow \exists A, A.\text{Infinite} \wedge X_A(2) = 1/2$$

is proved unconditionally ([LEAN SOURCE](#), [Lean]), with a positive-support refinement using the same odd-denominator lemma [LEAN SOURCE](#) already used for O2 above ([LEAN SOURCE](#), [Lean]). The hypothesis requires these finite families at arbitrarily large depths and is [Open]. The terminal bound then implies half-membership. This is a sufficient condition; no converse or strict comparison with other sufficient conditions is established here.

6.6.3 Removing one case from a sufficient tail condition

Five declarations extend the already-cited `SeamMiddleProducerSqrtEscape` hypothesis with a reduction in the cases that remain to be checked, using the proved exclusion of the -3 cell. No strict separation of the existence conditions is inferred from dropping that case.

Definition 6.144. `SeamMiddleProducerTailEscape` is the exact hypothesis that at every genuine middle transition ($s \geq 13$, no successor carry, below the terminal-weight threshold) the terminal-augmented binary coefficient tail is strictly below the signed carry ([LEAN SOURCE](#)). `SeamMiddleProducerTailEscapeExceptNegThree` is the same hypothesis restricted to cells $\neq -3$ ([LEAN SOURCE](#)).

The reduction `TailEscape` $\Rightarrow$ `ExceptNegThree` is trivial by hypothesis-dropping ([LEAN SOURCE](#), [Lean]); the substance is in the converse direction of the implication. The reduced hypothesis, with cell -3 discharged by the already-proved [LEAN SOURCE](#), already suffices to close HALF:

$$\text{SeamMiddleProducerTailEscapeExceptNegThree} \Rightarrow (1/2 : \mathbb{R}) \in A$$

([LEAN SOURCE](#), [Lean]), with the unreduced hypothesis and the already-cited square-root majorant both routing through it ([LEAN SOURCE](#), [Lean]). So the chain now reads `SqrtEscape` $\Rightarrow$ `TailEscape` $\Rightarrow$ `ExceptNegThree` $\Rightarrow$ HALF, a chain of sufficient implications. The displayed arrows do not prove that any implication is strict; all three antecedents remain [Open].

6.6.4 An eventual nonnegative margin

The finite-horizon condition can be stated without a separate private predicate. Recall the integer margin $F_k(J)$ from Theorem 6.13 and the equivalence in [LEAN SOURCE](#).

Definition 6.145 (A nonnegative margin by the next depth). At every skipped step $k+1$, require

$$\exists J \in \{0, \dots, k+1\}, \quad F_k(J) \geq 0.$$

This is the hypothesis encoded by [LEAN SOURCE](#).

It implies $1/2 \in A$, as stated in [LEAN SOURCE](#) and deduced through [LEAN SOURCE](#) ([Lean]). Indeed, the finite coefficient window is at most its infinite tail, so $F_k(J) \geq 0$ implies $r_k(1/2) \leq 2^{-(k+1)}$ at that skip. This is below R_{k+1} . At $k = 0$ the remainder is exactly $1/2$ and the same inequality holds; the strict version in Theorem 6.87 requires positive k .

The bound $J \leq k+1$ is *not used* in this implication. Allowing any finite J at each skip already suffices; requiring first passage by $k+1$ is a stronger finite-horizon demand, not a necessary ingredient of the membership proof. Neither all-skip existence condition is established here, and no equivalence between the bounded and unbounded horizons is asserted. [Open]

6.6.5 Rewriting the signed carry recurrence

Numeral-adjacent boundary words have more structure than the previously used coefficient bound. If the lower suffix numeral is one larger than the upper suffix numeral, there is a unique carry pivot d : rank d enters the lower support, every rank e with $d < e \leq N$ flips from lower-unselected to upper-selected, and all ranks below d agree (LEAN SOURCE, [Lean], (uniform), support rigidity, half-carry recurrence). Hence for every positive coefficient row $m \geq 1$ the boundary discrepancy is the exact divisor-incidence identity

$$c_{\text{low}}(m) + \#\{e : d < e \leq N, e \mid m\} = c_{\text{up}}(m) + \mathbf{1}_{d \mid m},$$

not merely an inequality (LEAN SOURCE, [Lean]).

The stage-level theorem consumes adjacent profiled prefixes directly (LEAN SOURCE, [Lean]). This is a local normal form for the two-sheet boundary. It does not produce adjacent stages cofinally and therefore does not close HALF or universal #257.

6.6.6 Further restrictions on represented values

This module works in the new `ErdosProblems.Erdos257` namespace and states in its own header that its results “concern achievement-set geometry and do not settle the universal irrationality problem.” For an arbitrary support restriction $J \subseteq \mathbb{N}$, define `supportedMersenneAchievementSet J` as the range of Mersenne digit values supported on J (LEAN SOURCE). It is compact (LEAN SOURCE, [Lean]), nowhere dense for every J (LEAN SOURCE, [Lean]), and perfect; compact with no isolated points; whenever J is infinite (LEAN SOURCE, [Lean]), via a Cantor-space no-isolated-point argument on the coding space itself (LEAN SOURCE, [Lean]).

Proposition 6.146 (Exact Lebesgue-measure dichotomy). *Either $J = F^c$ for a finite F , and $\text{vol}(\text{supportedMersenneAchievementSet } J) = 2^{-|F|}$ exactly, or J^c is infinite and the volume is exactly 0.*

Lean: `paper_volume_supportedMersenneAchievementSet_dichotomy`.

(LEAN SOURCE, [Lean], (uniform), divisor counts and finite sums). The finite case is an exact doubling induction over inserted coordinates (LEAN SOURCE, [Lean], using disjointness of the two coordinate-split faces, LEAN SOURCE, [Lean], which is exactly where unique binary coding is spent). None of this module touches irrationality of any subseries value; it is a self-contained measure-and-topology classification of achievement sets, orthogonal to O1–O5. The dichotomy is the case of retained coordinates of Hornich’s strict-tail measure formula [4, Theorem 4(1), p. 9]: the retained Mersenne weights still exceed their tails, and the Mersenne tail estimate supplies the values $2^{-|F|}$ and 0.

6.6.7 Conditional counterexamples at rational targets

This short module restates U257 in the new namespace (LEAN SOURCE) and gives it a problem-centric interface to the legacy half-terminal required input already used for O2 above: any finite supports with vanishing scaled terminal carry witness gives an infinite

A with $X_A(2) = 1/2$ ([LEAN SOURCE](#), [Lean]), hence refutes U257 outright ([LEAN SOURCE](#), [Lean]). The module's own docstring states plainly that the required input sequence itself "has only been checked to a large finite cutoff": this is a restatement of the already-[Open] O2 required input obligation in a new namespace, not a new result, and it duplicates §O2's own logic (any HALF witness is $\neg$ U257) rather than extending it.

The module now also proves an unconditional finite-support obstruction at the second rational target $1/21$:

$$F \subseteq \{n : n \geq 2\} \implies \sum_{n \in F} \frac{1}{2^n - 1} \neq \frac{1}{21}$$

for every finite F ([LEAN SOURCE](#), [Lean], (uniform), *finite-lcm*). The denominator-order identity forces $\text{lcm}(F) = 6$, after which the eight subsets of $\{2, 3, 6\}$ are exact arithmetic. This removes finite termination for that target; it does not establish achievement-set membership.

6.6.8 Multiplicity of the stated two–three representations

This module records the exact arithmetic of the primitive cone $2p + 3q = n$, $p, q > 0$, $\gcd(p, q) = 1$ ([LEAN SOURCE](#)). It proves that every $n \geq 11$ has such a representation ([LEAN SOURCE](#), [Lean], (uniform), *finite-lcm*), while rank 10 has none ([LEAN SOURCE](#), [Lean]). Rank 11 already has two distinct witnesses ([LEAN SOURCE](#), [Lean]), and so does every rank $10k$ with $k \geq 2$ ([LEAN SOURCE](#), [Lean], (uniform), *finite-lcm*). These theorems are deliberately recorded as route diagnostics rather than as a counterexample: the file does not prove that the associated integer-multiplicity expansion of $1/21$ Booleanises. The recurring multiplicities are precisely the unresolved collision data, not zero–one support digits.

6.6.9 Integer-quotient tests for one twenty-first

The later quotient-greedy analysis removes the diffuse alternatives. Let `TwentyOneFatalAlignedBranch` denote the explicit branch consisting of a fatal greedy gap, finite skipped support (hence cofinite selection), eventual alignment between the quotient rows and the rational greedy prefix, and eventual visits to every doubling block. Then

$$\frac{1}{21} \in \mathcal{A} \iff \neg \text{TwentyOneFatalAlignedBranch}.$$

[LEAN SOURCE](#) [Lean] (uniform) *greedy recurrence*.

On the fatal branch the canonical quotient remainder s_R is eventually strictly above the closed capacity 2^R ([LEAN SOURCE](#)), and thereafter the support appends $R + 1$ while s_R follows one explicit affine recurrence ([LEAN SOURCE](#)). Conversely, closed rows $s_R \leq 2^R$ along any unbounded sequence decay after normalisation and put $1/21$ in the achievement set ([LEAN SOURCE](#); [LEAN SOURCE](#)).

This is the exact current contribution: membership has been reduced to excluding one named permanent affine-supercapacity regime. No theorem excludes that regime, forces unbounded closed returns, or proves membership. Combined with the finite-support obstruction above, membership would provide an infinite-support rational counterexample to universal #257; its present status remains [Open].

IV. Scales, conditional interfaces and unresolved inputs

7 Finite ranges and unbounded parameters

Every claim in this programme carries a *scale* tag alongside its evidence band. The two axes are independent and both matter to a reader deciding whether a result can be used as a premise. Evidence labels record the type of support supplied with a claim. [Lean] refers to the cited formal source and its recorded build status, not a new kernel run in this revision. A newly cited release file marked `release_only` in the packet is described as [Lean source]; that label does not promote it to the packet's checked-build class. [Cert] denotes an exact finite computation; the surrounding text distinguishes fresh checks from inherited reports. [Math] denotes an ordinary mathematical argument, not independent human certification. [Cited] invokes the stated external result, and [Open] marks an unproved hypothesis.

The scale records the quantified range. A (fixed) claim concerns a specified value; (bounded) denotes either a finite tested range or a witness bounded in terms of the input. A (uniform) assertion applies to every index in its stated range, *with its hypotheses retained*; it need not be unconditional. A (cofinal) assertion has the form $\forall N \exists n \geq N, P(n)$. Compatibility between its witnesses is an additional requirement only when explicitly imposed. Definitions and scale-free statements are marked (n/a).

The important distinction is between obtaining a witness in a finite range and proving the quantified assertion used by a membership test. For example, half-membership follows from unbounded successful exact rows or from infinitely many greedy skips. A different sufficient condition asks for a bound at every reset beyond its threshold. These statements have different quantifiers, even when the same transition formulas occur in their proofs. Each row must retain its predicate, quantified variable, and every premise. An unconditional eventual assertion $\forall n \geq N_0, P(n)$ does give cofinal witnesses to that same predicate, by taking $n = \max(N, N_0)$. A uniform conditional implication $H(n) \Rightarrow C(n)$ supplies no witness until $H(n)$ is established, and a checked finite band supplies no assertion beyond its endpoints. These are different logical shapes and are recorded separately below.

Table 2 separates established finite or quantified results from unproved membership hypotheses and the proved implications that use them. Its lower part therefore contains proved theorems as well as open assumptions. The scale alone determines neither evidence nor truth status; the stated predicate and premises do.

Result	Scale	Quantifier prefix	Coordinate
<i>Fixed ; one concrete witness or numeral</i>			
Seed for continuation from $n \geq 6$: support $\{2, 3, 6\}$ satisfies $Q = 2^5 - 1$. (fixed) LEAN SOURCE [Cert]		$n = 6$	<i>divisor counts and finite sums</i>
One-point counter-model: the double-or-recycle transition shape alone, even seeded, is not cofinal. LEAN SOURCE [Lean]	(fixed)	$n = 6$ (free-standing model)	<i>meta-logical</i>
The combined-residue-mass ≤ 1 sufficient test for sharp capacity is not necessary: fixture $D = \{2, 3\}$, $c = 5$. LEAN SOURCE [Cert]	(fixed)	$D = \{2, 3\}$, $c = 5$	<i>divisor counts and finite sums</i>
Exact residual fixture: $\frac{1}{2} - W(\{2, 3, 6, 7\}) = \frac{1}{16002}$. LEAN SOURCE [Cert]	(fixed)	none (numeric instance)	<i>regression-fixture</i>
Closed-form Möbius–Lambert value $\sum_{d \geq 1} \mu(d)/(2^d - 1) = 1/2$ (imported identity). LEAN SOURCE [Cited]	(fixed)	none (closed-form value)	<i>divisor counts and finite sums</i>
The one natural negative-Möbius candidate support overshoots: $1/2 < \sum_{d \in N} 1/(2^d - 1)$ for $N = \{d \geq 2 : \mu(d) = -1\}$. LEAN SOURCE [Lean]	(fixed)	one fixed candidate support N	<i>divisor counts and finite sums</i>
<i>Bounded ; a witness exists, controlled by the input</i>			
From an exact row, either repeated doubling is available or an exact row is obtained at some $2c - 2$ with $c \leq n$; these alternatives need not be disjoint. LEAN SOURCE [Lean]	(bounded)	$\forall n \geq 6$, given a row at n : $(\dots) \vee \exists c \leq n(\dots)$	<i>divisor counts and finite sums</i>
Every above-half bounded support recycles into a genuine exact row at some $2c - 2$ with $c \leq n$. LEAN SOURCE [Lean]	(bounded)	$E \subseteq [2, n]$, $X_E(2) > 1/2$: $\exists c \leq n$	<i>binary digits</i>
(#249 cross-argument check) $\sum \varphi(n)/2^n$ matches no rational whose denominator divides $2^{14}(2^h - 1)$, for some $1 \leq h \leq 16$. LEAN SOURCE [Cert]	(bounded)	$1 \leq h \leq 16$ (explicit bound)	<i>binary digits</i>
Truncation-rung ladder: HalfRung(J) proved for every $3 \leq J \leq 22$ via the finite decision procedure (§8, Theorem 8.9). (source: erdos257_truncation_rung_ladder_2026_07_24.md, §5) [Math] + [Cert]	(bounded)	$3 \leq J \leq 22$ (finite table, exhaustive)	<i>truncated greedy sums</i>
<i>Uniform assertions and conditional implications; premises retained</i>			
Master achievement/greedy-survival criterion. LEAN SOURCE [Lean]	(uniform)	$\forall x \geq 0$, membership iff survival at every rank	<i>greedy recurrence</i>
Membership iff every actually-skipped rank survives. LEAN SOURCE [Lean]	(uniform)	$\forall M \in \text{skipped support}$	<i>greedy recurrence</i>
Every crossing core is forced to be the canonical rational half-greedy prefix. LEAN SOURCE [Lean]	(uniform)	$c \geq 4$, $D \subseteq [2, c]$, $0 < 1/2 - X_D(2) < w_c$	<i>greedy recurrence</i>
Largest-skip conditional implication used in the reset argument. LEAN SOURCE [Lean]	(uniform)	$\forall s \geq 14$, given the hypothesis	<i>integer quotients</i>

Result	Scale	Quantifier prefix	Coordinate
Band-avoidance certified at reset indices $13 \leq r \leq 30$, using successor remainders at indices $14, \dots, 31$. This is a finite range of the conditional hypothesis. LEAN SOURCE [Cert]	(bounded)	$13 \leq r \leq 30$, exhaustive	<i>integer quotients</i>
Critical-band-index reduction: a $(d+1)$ -way check collapses to one nearest-boundary check. LEAN SOURCE [Lean]	(uniform)	$\forall d, E$ with $E \leq 2^{d+1}$	<i>dyadic-boundary</i>
Parity obstruction: no finite support has value $1/2$. LEAN SOURCE [Lean]	(uniform)	$\forall A$ finite, $0 \notin A$	<i>denominator-parity</i>
Full-support irrationality at every integer base. LEAN SOURCE ; base two: LEAN SOURCE [Lean]	(uniform)	$\forall b \geq 2$	<i>divisor counts and finite sums</i>
Erdős 1968: any infinite, pairwise-coprime, reciprocal-summable index set gives an irrational support series. LEAN SOURCE [Lean]	(uniform)	$\forall b \geq 2 \forall A$ (infinite, coprime, summable)	<i>divisor counts and finite sums</i>
Generic block-certificate engine (problem-agnostic): weighted-coefficient series is irrational given a per-precision divisibility/window certificate. LEAN SOURCE [Lean]	(uniform)	$\forall b \geq 2 \forall c (\forall m, c(m) \leq m)$, given $\forall q \exists N, K, L, C$	<i>binary digits</i>
Finite greedy-prefix denominators are odd (source of the parity boundary exclusion, §8). LEAN SOURCE [Lean]	(uniform)	$\forall n$	<i>denominator-parity</i>
Unproved hypotheses and the proved implications that use them			
<i>Unbounded-witness and all-reset conditions have different quantifiers</i>			
The exact-row existence condition: exact rows recur at arbitrarily large endpoints, with no coherence required between witnesses. LEAN SOURCE [Open]	(cofinal)	$\forall N \exists n \geq N$, $\text{ExactLocalMersenneHalfRow}(n)$	<i>divisor counts and finite sums</i>
The preceding cofinal exact-row condition implies $1/2 \in A$ by compactness. LEAN SOURCE [Lean] (theorem proved; <i>hypothesis</i> open)	(cofinal)	given $\forall N \exists n \geq N(\dots)$	<i>topological-achievement-set</i>
An equivalent membership condition: the canonical rational half-greedy orbit itself has infinitely many positive skip events. LEAN SOURCE [Open]	(cofinal)	$\forall N \exists c \geq \max(N, 4)$, $0 < \text{rem}(c-1) < w_c$	<i>greedy recurrence</i>
Infinitely many positive greedy skips imply $1/2 \in A$. LEAN SOURCE [Lean] (theorem proved; hypothesis open)	(cofinal)	given $\forall N \exists c \geq N(\dots)$	<i>greedy recurrence</i>
Master half-branch dichotomy: membership iff the greedy skip set is infinite. LEAN SOURCE [Lean] (iff proved; both sides cofinal-shaped)	(cofinal)	iff $(\text{skippedSupport}).\text{Infinite}$	<i>greedy recurrence</i>

Result	Scale	Quantifier prefix	Coordinate
Seven-way classification hub: membership iff the seam word is not eventually right, iff unbounded terminal-false, iff cofinally many skipped ranks, <i>etc.</i> LEAN SOURCE [Lean]	(cofinal)	every right-hand side is $\exists^\infty$ / Unbounded-shaped	<i>integer quotients</i>
Ladder dichotomy: either infinitely many rungs survive (#257 false, explicit limit support) or an eventual bad-rank obstruction closes only this required input. (source: erdos257_truncation_rung_ladder_2026_07_24.md, Cor. 8) [Math]	(cofinal)	open branch (a): $\exists^\infty J$, HalfRung(J)	<i>truncated sums</i> <i>greedy</i>
Reset square-root escape, a sufficient all-reset hypothesis. The release source gives the conditional implication; the run-length comparison is discussed separately in Section 12.5. [Open]	(all resets)	$\forall r \geq 10$ that are resets, $ \text{rem}(r+1) - 2^{r+1} > 2^{(r+5)/2}$	<i>integer quotients</i>
Table 2: Quantifiers and evidence for the stated results. The labels separate finite tests, uniform implications, unbounded witness conditions and all-reset hypotheses. A theorem conditional on an unproved hypothesis is not a proof of that hypothesis.			

Remark 7.1 (Finite checks and quantified conclusions). The supplied audit examined eighteen possible implications from recorded results to the required unbounded-witness conditions. It reported no such implication. This is a report about those attempts, not a theorem excluding other arguments and not a fresh kernel replay.

For example, [LEAN SOURCE](#) checks reset indices 13 through 30; it has no conclusion beyond that range. Likewise, a bound depending on the depth need not produce a witness at a larger depth. Conversely, a proved eventual assertion of the *same* witness predicate immediately gives cofinal witnesses. The issue is the precise premise and conclusion, not the name of its scale or the tactic used in a finite proof.

8 Approximations to the greedy orbit

Four reported calculations

Fix the weights $x_n := 1/(2^n - 1)$, $n \geq 2$, and the target $1/2$. The single object under study throughout this programme is the greedy orbit for this target: a residual ρ starts at $1/2$, and at rank k the orbit *takes* iff $\rho \geq x_k$. A skip at rank k is *safe* iff $\rho \leq T_{k+1} := \sum_{j>k} x_j$ and *fatal* iff $\rho \in (T_{k+1}, x_k)$; the fatal interval is exactly the gap no later tail can repair. The four calculations below concern this target, but their finite parameters and certification records differ. Agreement with the real greedy prefix requires a convergence argument for the approximants used in each calculation; agreement over a tested range is not that argument.

- **Truncation rung J .** Replace the full weight x_n by the J -term truncation $w_n^{(J)} := \sum_{q=1}^J 2^{-qn}$ and ask whether some Boolean support hits $1/2$ exactly under the truncated weights. Certified for $3 \leq J \leq 22$ (Table 2, bounded row; Theorem 8.6–8.9 below).
- **Seam row s .** The base-4 seam recursion tracking the deviation $\Delta_s := \text{rem}(s) - 2^s$ of the integer greedy orbit from its dyadic target. Certified rows 6 through 2500 for the reset-crossing classification (§7, uniform rows; source: `erdos257_reset_crossing_unification_2026_07_24.md`, §8), and the underlying real orbit separately certified to row 200,000 (source: `erdos257_sqrt_escape_digit_reduction_2026_07_24.md`, §3): exactly one TARGET violation in that range, at $r = 7$.
- **Integer margin m_c .** An advisory Type B scan of the integer-greedy margin out to 5×10^5 , tracking how far the certified region sits from the danger threshold.
- **Sharp tail margin.** The rank-by-rank take/skip margin against the exact tail, checked for ranks 2 through 3000: 1497 takes against 1502 skips, zero fatal skips anywhere, minimum safe-skip margin +2.9922 bits at rank 5, and the smallest reported take-margin +0.0340 bits at rank 7, where the remainder $1/126$ exceeds the weight $1/127$. Here the margins are the base-two logarithms of the corresponding positive slack multiplied by 4^k (source: `erdos257_reset_crossing_unification_2026_07_24.md`, §8).

Proposition 8.1 (Stability of each fixed greedy prefix). *Let $t_j \rightarrow 1/2$ and, for each fixed $n \geq 2$, let $v_n^{(j)} \rightarrow x_n = (2^n - 1)^{-1}$, with $v_n^{(j)} > 0$. Apply the greedy rule with target t_j and*

weights $v_n^{(j)}$ in increasing order of n , through depths $m_j \rightarrow \infty$. For every fixed depth K , the decisions at ranks $2, \dots, K$ eventually agree with those of the real half-greedy rule. This assertion concerns finite prefixes, not survival at all ranks.

Lean: [paper_one_orbit_stability](#), [approx_orbit_induction](#), [tailGreedyRemainder_mersenne](#).

Proof. At rank 2, the target and the weight converge to their real counterparts. Their comparison is strict by Lemma 8.2, so its sign is eventually unchanged. Inductively, once all earlier decisions agree, the approximate remainder is the target minus a fixed finite sum of convergent weights. It therefore converges to the real remainder. The next comparison is again strict, so the next decision eventually agrees. A finite induction through K proves the claim. $\square$

Lemma 8.2 (No-ties lemma). *At every rank k of the full greedy orbit for target $1/2$, both defining comparisons are strict: $\rho \neq x_k$ and $\rho \neq T_{k+1}$.*

Lean: [paper_no_ties](#), [paper_no_ties_take](#), [paper_no_ties_skip](#), [irrational_mersenneTail](#).

Proof. The two strict comparisons have different reasons: parity excludes equality with a weight, and irrationality excludes equality with the complete tail.

Take boundary excluded (parity). Write a finite Mersenne sum as $p/q = \sum_{d \in F} x_d$ in lowest terms. Its denominator q is odd, because every $2^d - 1$ is odd (equivalently, [LEAN SOURCE](#), itself traced to [LEAN SOURCE](#)). Hence the residual after any finite prefix is $\rho = \frac{1}{2} - \frac{p}{q} = \frac{q-2p}{2q}$: its reduced denominator is *even*. Every weight $x_k = 1/(2^k - 1)$ has *odd* denominator. An even-denominator rational can never equal an odd-denominator rational, so $\rho = x_k$ is impossible at any rank; the take comparison is never a tie.

Skip-safety boundary excluded (irrationality). The tail past rank k is $T_{k+1} = E - 1 - \sum_{2 \leq j \leq k} x_j$, where $E = \sum_{n \geq 1} 1/(2^n - 1)$ is the Erdős–Borwein constant: a rational shift of E . Erdős’s own 1948 theorem gives E irrational ([LEAN SOURCE](#), instantiated at $b = 2$), so T_{k+1} is irrational for every k . The residual ρ arising from any finite Boolean prefix is rational. An irrational number never equals a rational one, so $\rho = T_{k+1}$ is impossible at any rank; the skip-safety comparison is never a tie either.

Both boundaries of the fatal interval (T_{k+1}, x_k) are therefore approached only strictly, at every rank, unconditionally. $\square$

Remark 8.3. The hypotheses of Proposition 8.1 can be checked for two concrete approximations. The truncated geometric weights $\sum_{q=1}^J 2^{-q^n}$ increase to x_n . For the integer quotient rows, division by 4^s gives weights $4^{-s} \lfloor 4^s / (2^n - 1) \rfloor \rightarrow x_n$ and target $4^{-s} (2^{2s-1} - 2^s) = 1/2 - 2^{-s} \rightarrow 1/2$. Hence both greedy rules agree with each fixed real prefix once their parameters are large enough. Convergent upper and lower bounds for R_k likewise determine the strict tail comparison at a fixed rank. No conclusion about the advisory margin scan follows without specifying its approximants and proving their convergence. Fixed-prefix agreement is not uniform agreement through a growing depth. For the quotient rows, however, summability bounds the unmatched tail by R_K and does give convergence of the coded values, as proved in Theorem 6.2. It identifies their limit as $X_G(2)$; it does not establish that this value is $1/2$.

Successive tail truncations

Finite truncations of the geometric expansion give rational weights and permit a finite decision procedure for each fixed truncation order J . The bounds below explain why only finitely many ranks need to be tested. Fix $J \geq 2$ and set

$$w_n^{(J)} := \sum_{q=1}^J 2^{-qn}, \quad T_{n+1}^{(J)} := \sum_{q=1}^J \frac{2^{-qn}}{2^q - 1},$$

$$\text{HalfRung}(J) : \iff \exists A \subseteq \{2, 3, \dots\}, \quad \sum_{n \in A} w_n^{(J)} = \frac{1}{2}.$$

The arguments below are ordinary mathematics. The accompanying exact integer computation rechecks the finite instances $3 \leq J \leq 22$; neither the arguments nor that computation are presented as a Lean verification.

Lemma 8.4 (Forced greedy). *For every $J \geq 2$: $w_n^{(J)} > T_{n+1}^{(J)}$ for all n (the $q = 1$ terms agree exactly, and every $q \geq 2$ tail term is strictly smaller than the corresponding weight term). Consequently the greedy support is the unique candidate support, and $\text{HalfRung}(J)$ holds iff the greedy orbit for $1/2$ under weights $w_n^{(J)}$ never lands in a fatal interval $(T_{n+1}^{(J)}, w_n^{(J)})$. Rank 1 is always a safe skip; ranks 2 and 3 are always takes.*

Lean: [paper_forced_greedy_unique_support_and_criterion](#), [paper_forced_greedy_low_ranks](#).

Lemma 8.5 (Parity forces infinite support). *For every $J \geq 2$, no finite $A \subseteq \{2, 3, \dots\}$ attains $\text{HalfRung}(J)$.*

Lean: [paper_parity_excludes_finite_support](#).

Proof. The empty set has value zero. For a nonempty finite set, put $m = \max A$. After multiplication by 2^m , the m th summand is odd and every earlier summand is even, whereas the target 2^{Jm-1} is even. This parity contradiction excludes finite support. $\square$

Theorem 8.6 (Witness exclusion). *Define the misalignment mass*

$$\mu_J(M) := \sum_{q=2}^J \frac{2^{M \bmod q}}{2^q - 1}.$$

Let $J \geq 3$, $n \geq 4$, and suppose some $M \in [n, 2n - 2]$ has $\mu_J(M) \leq \frac{11}{15}$. Then no Boolean prefix $D \subseteq \{2, \dots, n - 1\}$ satisfies

$$T_{n+1}^{(J)} < \frac{1}{2} - \sum_{d \in D} w_d^{(J)} < w_n^{(J)};$$

This excludes the specified gap for every Boolean prefix, not only the greedy prefix.

Lean: [paper_witness_exclusion](#).

Proof. Suppose such a prefix exists and write $\rho = 1/2 - \sum_{d \in D} w_d^{(J)}$. Since $T_{n+1}^{(J)} \geq 2^{-n}$ and $n \leq M \leq 2n - 2$,

$$0 < 2^M \rho - 2^{M-n} < \sum_{q=2}^J 2^{M-qn} \leq \frac{2^{M-2n}}{1-2^{-n}} \leq \frac{4}{15}.$$

Here $n \geq 4$ gives the last inequality. Split each scaled summand 2^{M-qn} according to whether $qd \leq M$. The terms with $qd \leq M$ are integers. The other terms have total

$$F = \sum_{d \in D} \sum_{\substack{1 \leq q \leq J \\ qd > M}} 2^{M-qn} \leq \sum_{q=2}^J \sum_{d > \lfloor M/q \rfloor} 2^{M-qn} = \mu_J(M) \leq \frac{11}{15}.$$

There is no $q = 1$ contribution to F , since $d < n \leq M$. Thus $2^M \rho + F - 2^{M-n}$ is an integer strictly between zero and one, a contradiction. The same calculation also applies when $J = 2$. $\square$

Corollary 8.7 (Half-LCM horizon). *Let $J \geq 2$ and $L_J := \text{lcm}(2, 3, \dots, J)$. For every $n \geq \max\{4, L_J/2 + 1\}$ there is an $M \in [n, 2n - 2]$ with $\mu_J(M) < 11/15$. Thus witness exclusion leaves only the finite window $[4, L_J/2]$ to check.*

Lean: [paper_half_lcm_horizon](#).

Proof. Let M be the least multiple of L_J with $M \geq n$. If $M = L_J$, then $M \leq 2n - 2$ is the assumed inequality. If $M = kL_J$ with $k \geq 2$, then $n \geq (k - 1)L_J + 1$ gives $2n - 2 \geq 2(k - 1)L_J \geq M$. At this witness,

$$\mu_J(M) = \sum_{q=2}^J \frac{1}{2^q - 1} < \sum_{q=2}^{\infty} \frac{4}{3 \cdot 2^q} = \frac{2}{3} < \frac{11}{15}.$$

The bound uses $2^q - 1 \geq 3 \cdot 2^{q-2}$, with strict inequality for $q > 2$, so no numerical approximation to the full series is needed. Apply Theorem 8.6; its proof also covers $J = 2$. Ranks 2 and 3 are taken, as already noted. $\square$

Lemma 8.8 (Mod-12 filter, $J \geq 7$). *For $J \geq 7$, the inequality $\mu_J(M) \leq 11/15$ implies $12 \mid M$. Thus only multiples of 12 need be tested as witnesses in each interval $[n, 2n - 2]$; no potentially uncovered rank is discarded.*

Lean: [paper_mod_twelve_filter](#).

Proof. If M is odd, the $q = 2, 3$ terms give at least $2/3 + 1/7 > 11/15$. If $M \equiv 2 \pmod{4}$, the $q = 2, 3, 4$ terms give at least $1/3 + 1/7 + 4/15 = 26/35 > 11/15$. Finally, if $4 \mid M$ but $3 \nmid M$, the $q = 2, 3, 4, 6$ terms give at least $1/3 + 2/7 + 1/15 + 4/63 = 236/315 > 11/15$. Each case contradicts the assumed upper bound. $\square$

Theorem 8.9 (Finite decision procedure). *Call $n \in [4, L_J/2]$ **bad** if no $M \in [n, 2n - 2]$ has $\mu_J(M) \leq \frac{11}{15}$, and set $B(J) := \max(\text{bad} \cup \{3\})$. Then $\text{HalfRung}(J)$ holds iff the greedy orbit for $1/2$ under weights $w_n^{(J)}$ survives every rank from 2 through $B(J)$. This is a finite exact decision procedure.*

Lean: [paper_rung_finite_decision](#).

Proof. Survival at all ranks plainly implies survival through $B(J)$. Conversely, every later rank at most $L_J/2$ has a witness by the definition of $B(J)$, and every larger rank has one by Corollary 8.7. Witness exclusion rules out a fatal gap at all these ranks. The strict-tail criterion then gives membership. For $J = 2$, the same proof works because $\mu_2(M) \leq 2/3$ for every M ; the window is empty and $B(2) = 3$. $\square$

Remark 8.10 (Verified truncations and the unbounded-index question). The finite decision procedure verifies $\text{HalfRung}(J)$ for every $3 \leq J \leq 22$. The supplied checker uses exact integers: it verifies witness intervals covering the ranks beyond $B(J)$, proves that the listed bad ranks have no witness, and tests the greedy remainders through $B(J)$. In particular, it reproduces $B(19) = 1008$ and $B(20) = B(21) = B(22) = 1530$. The certificate data and an independent verifier are in `audit/final-proof-reading/truncation/`. These are finite computer-assisted proofs using the ordinary lemmas above, not Lean proofs and not a test of the full weights through all ranks.

If A represents $1/2$ with the J -truncated weights, then

$$0 \leq X_A(2) - \frac{1}{2} \leq \sum_{q>J} \frac{2^{-2q}}{1 - 2^{-q}} \leq 2 \sum_{q>J} 4^{-q} = \frac{2}{3} 4^{-J}.$$

Consequently, such representations for an unbounded sequence of J would put $1/2$ in the closed set A . Its representing support would be infinite, since finite Mersenne sums have odd denominator (Lemma 8.2); Lemma 8.5 separately excludes finite support for each truncation. The unbounded-index assertion is still unproved. Twenty verified values of J do not establish it.

9 Which hypotheses remain unproved

This section compares the stated results with the remaining hypotheses. Parts I–II assembled the corpus; here every result in it is tested against three exact open obligations and rendered as a usable premise: what it yields, the exact mismatch against the obligation, a classified gap kind (`hypothesis_strength`, a different parameter range, quantifier_order, a different expression of the same condition, or multiple), and the exact auxiliary statement that would close it. Nothing here decides #249 or #257; the corpus's own moderator standard applies throughout: a conditional reduction is a result with a hypothesis, not a proof of that hypothesis, and a finite verified list is not a cofinal supply.

9.1 The three #257 obligations, stated exactly

Definition 9.1 (257-reset). For every upper or middle reset row $r \geq 10$, require $\Delta_{r+1}^2 > 2^{r+5}$, where $\Delta_{r+1} = \text{rem}(r+1) - 2^{r+1}$. The finite range $10 \leq r < 31$ must be covered as well as the tail. Section 12.5 gives the exact conditional implication from this hypothesis. A proposed bound on right-run lengths is a different statement until the pulse corrections and endpoint conventions are matched. Finite observations establish neither failure of every constant run bound nor logarithmic asymptotic growth.

Definition 9.2 (257-cofinal-rows). $\forall N \exists n \geq N$ with [LEAN SOURCE](#) n ; i.e. cofinally many finite prefixes whose local Mersenne quotient lands exactly on the half-value residue

class. Consumed unconditionally by [LEAN SOURCE](#). No coherence between the witnessing rows is demanded and none should be assumed: the rows may be, and in the corpus's own words, mutually incompatible ; only their endpoint lengths need tend to infinity. Requiring compatibility would add a condition on the witnesses. Whether that condition can be obtained from this existence statement needs a separate extraction argument; it is not a strict logical comparison merely because the witnesses are described differently.

Definition 9.3 (257-universal). The original question is the base-2 statement: $\sum_{n \in A} (2^n - 1)^{-1}$ is irrational for every infinite $A \subseteq \mathbb{N}_{\geq 1}$. Asking the same for every integer base $b \geq 2$ is a stronger all-base extension, not the definition of Problem #257 itself.

9.2 Reading a row

Each entry identifies the available result, the assumptions needed to apply it, and the further estimate still required. A source link supports the stated declaration, not every proposed consequence in the discussion. Ordinary deductions and finite computations are identified separately.

9.3 Reset estimates and their missing hypotheses

Observation 9.4 (Remainders along a right-branch run). [Lean] [LEAN SOURCE](#). For integers $s \geq 5$ and $k \geq 0$, the cited theorem assumes $\text{rem}(s) \leq 2^s$ and $2^k(2^s - \text{rem}(s)) \geq 2^s$ and gives an upper or middle transition at some $t \in [s, s+k]$. Its literal contrapositive concerns *all* $k+1$ transition indices in that closed interval.

For a run of exactly $k \geq 1$ right transitions, at indices $s, \dots, s+k-1$, a direct recurrence gives the sharper conventionally indexed statement. Put $D_j = 2^{s+j} - \text{rem}(s+j)$ and $p_j = p_{s+j}^- \geq 0$. The right-branch identity yields

$$D_{j+1} = 4D_j + p_j + 4, \quad 4^k D_0 < D_k \leq 2^{s+k}.$$

The strict inequality uses $k \geq 1$, and the last bound uses $\text{rem}(s+k) \geq 0$. Hence $D_0 < 2^{s-k}$. In particular, $D_0 \geq 1$ excludes $k \geq s$ consecutive right transitions. A stronger lower bound on D_0 shortens that run, but must be established at the starting row of the run. Applying this argument after a reset requires the post-reset index and the *deficit-side* condition; it does not prove the absolute reset bound on both signs.

Observation 9.5 (The affine excess recurrence). [Lean] [LEAN SOURCE](#), [LEAN SOURCE](#), [LEAN SOURCE](#) (recurrence [LEAN SOURCE](#)). For k right transitions from row $S \geq 5$, let $X_j = \text{rem}(S+j) - 2^{S+j}$ and $p_j = p_{S+j}^-$. The cited identity is

$$X_0 = \frac{X_k}{4^k} + \sum_{j=0}^{k-1} \frac{p_j + 4}{4^{j+1}}.$$

The source also gives $0 \leq p_j \leq 2(S+j-2)$. Thus the following consequences are ordinary deductions from the identity:

$$\sum_{j=0}^{k-1} \frac{p_j + 4}{4^{j+1}} \leq \sum_{j \geq 0} \frac{2(S+j)}{4^{j+1}} = \frac{2S}{3} + \frac{2}{9}.$$

If the terminal row satisfies $\text{rem}(S+k) < 2^{S+k+1}$, then

$$X_0 < 2^{S-k} + \frac{2S}{3} + \frac{2}{9}.$$

Only this terminal bound is needed; no bound at every intermediate row is required. If instead the largest omitted rank d at the terminal row satisfies $2(S+k) < 3d$, the next entry gives the different bound

$$X_0 < 2^{S-k} + \frac{2 \cdot 4^{S+k-d} + 4}{3 \cdot 4^k} + \frac{2S}{3} + \frac{2}{9}.$$

The extra term must be retained. Lateness of an omitted rank is not known to be equivalent to $\text{rem}(s) < 2^{s+1}$; the cited estimate alone does not imply that sharper ceiling. Neither upper bound supplies the missing lower bound on an actual reset deviation.

Observation 9.6 (A bound at later rows). [Lean] [LEAN SOURCE](#). For $s \geq 5$, if d is the largest omitted rank and $2s < 3d$, the cited theorem gives

$$\text{rem}(s) < 2^{s+1} + \frac{2 \cdot 4^{s-d} + 4}{3}.$$

Since $s-d < s/3$, its additive error is at most $(2 \cdot 2^{2s/3} + 4)/3$, hence $O(2^{2s/3})$. The undivided error $2 \cdot 4^{s-d} + 4$ need not be less than $2^{2s/3}$: already $s=7, d=5$ gives $36 > 2^{14/3}$. These numbers test the numerical estimate, not an asserted pair on the greedy orbit. The theorem supplies an upper bound only under the largest-omitted-rank and lateness assumptions. When it is used in the preceding recurrence, its additive error remains part of the result.

Observation 9.7 (Exclusion of equality at a branch boundary). [Lean] [LEAN SOURCE](#), with [LEAN SOURCE](#). The first cited theorem excludes equality between the corrected real-valued quotient expression and the terminal *take threshold*: such equality would give a finite Mersenne sum equal to $1/2$, contrary to denominator parity. The second classifies the signed integer position of a real greedy prefix at a skipped rank; its statement explicitly permits the zero-position case.

Neither statement identifies its boundary with $\text{rem}(r) = 2^r$. They therefore do not, by themselves, prove $|\text{rem}(r) - 2^r| \geq 1$. Integrality gives that inequality only after non-vanishing of this particular deviation has been proved. An estimate of the form $2^{r/2+o(r)}$ would not settle the required constant either: $2^{r/2}/r$ has that form but is smaller than $2^{(r+5)/2}$. The reset hypothesis requires its stated constant, strict inequality and post-reset index, not merely an exponent approaching $1/2$.

Observation 9.8 (A sharper test for a fatal greedy gap). [Lean] [LEAN SOURCE](#), [LEAN SOURCE](#), [LEAN SOURCE](#), [LEAN SOURCE](#). Let $k \geq 1$ be a skipped rank and write its positive rational remainder in lowest terms as $\rho = u/(2L) < w_k$, with u, L odd and positive. Set $a = 2L - (2^k - 1)u > 0$. The cited sufficient condition $2u \leq 3a$ guarantees $\rho < R_k$ and

improves the simpler condition $u \leq a$. Its contrapositive gives $3a < 2u$ at a fatal skip, and hence $u \geq 3$ when u is odd.

This is not the exact threshold. With $\theta_k = R_k^{-1} - (2^k - 1)$, direct rearrangement gives

$$\rho < R_k \iff u\theta_k < a.$$

The tail estimate gives $\theta_k < 2/3$ and $\theta_k \rightarrow 2/3$; $2u \leq 3a$ is therefore sufficient, not necessary. For example, at $k = 2$ the local rational state $\rho = 77/282$ has $u = 77$, $L = 141$, $a = 51$, so $2u > 3a$, yet $\sum_{j=3}^{12} (2^j - 1)^{-1} > 77/282$ proves $\rho < R_2$. This is a local state, not an asserted point of the half-greedy orbit. Improving a bounded tail ratio alone does not provide a lower bound growing exponentially in k on an arithmetic numerator; such a bound would require additional information about the actual prefix.

Observation 9.9 (Exclusion of a linear-width band at rows 13–30). [Lean] [LEAN SOURCE](#). **Conclusion.** At every upper reset row d , $13 \leq d \leq 30$, the reset charge $E_d = 4 \cdot \text{overshoot}_d + \text{abovePulse}_d$ avoids the linear-width band $(2^{d-j+1} - 2(d+j), 2^{d-j+1}]$ for every $j \leq d$; the finite segment of [LEAN SOURCE](#), which by [LEAN SOURCE](#) delivers the whole #257-false payoff. **Parameter range.** proved at 18 explicit rows by `interval_cases` plus `per-row decide+kernel`; the obligation needs all $d \geq 13$. The excluded band has width $O(d)$, whereas the reset condition uses a width of order $2^{(r+5)/2}$. This is a comparison of numerical scales, not a proof that one hypothesis on the orbit implies the other. **Missing step.** A structural lower bound on the distance from E_d to the nearest dyadic power at or above it, of width $2(d+j_0)$ at the single critical index j_0 ; already collapsed from $\forall j$ to one inequality by [LEAN SOURCE](#). Only $O(d)$ separation is required, not `sqrt-of-scale`.

Observation 9.10 (Middle predecessors and right-branch runs). [Lean] [LEAN SOURCE](#), [LEAN SOURCE](#), [LEAN SOURCE](#). **Conclusion.** Unconditional, uniform in $s \geq 5$: a middle (M) branch forces $s+1 \leq \text{rem}(s+1)$, using no lower bound on the source remainder; with the row-scale source bound this upgrades to $2^{s+1} + (s+1) \leq \text{rem}(s+1)$. The minimal-counterexample argument at line 3606 then shows any row whose remainder is smaller than its index must have an upper-reset (U) ancestor, since M-ancestors and R-runs are both excluded. **Scope of the hypothesis.** two of the three ancestor channels (M, R) are closed unconditionally at all scales; only the U channel remains, and it needs only the $O(d)$ -wide band escape of F5b above; `sqrt-of-dyadic-scale` escape is being demanded by the 257-reset statement where linear-width escape at U-resets alone suffices for this downstream route. **Missing step.** The single remaining channel; $\forall d \geq 13$ with successor carries, the reset charge avoids the width- $2(d+j_0)$ band at its critical dyadic index. Then `remainder_ge_row`, hence `SeamMiddleProducerRowEscape`, hence $1/2 \in A$, all already in Lean.

Observation 9.11 (A linear lower bound at middle transitions). [Lean] [LEAN SOURCE](#), [LEAN SOURCE](#) (via `.toCardEscape`, `.toTailEscape`, [LEAN SOURCE](#)). **Conclusion.** If at every late middle row ($s \geq 13$, no carry, middle branch) $s \leq \text{rem}(s)$, then $1/2 \in A$. Module `docstring`: “far weaker than an exponential bound, but stronger than the already exposed square-root-scale carry target.” **Scope of the hypothesis.** this hypothesis asks for a *linear-in-row* lower bound on the remainder itself; the obligation supplies an *exponential* lower bound on the deviation from 2^s at reset rows. For unrestricted numerical

values, a large absolute deviation from 2^s permits $\text{rem}(s) = 0$, which violates the linear lower bound. That observation is not a counterexample within the actual recurrence, so incomparability of the two orbit hypotheses is not established. **Missing step.** $\forall s \geq 13$ at a middle transition, $s \leq \text{rem}(s)$. The upper-reset band condition described above is one sufficient input to `remainder_ge_row`; no converse is asserted.

Observation 9.12 (A square-root lower bound in the row index). [Lean] [LEAN SOURCE](#), its implication to the exact tail bound (`.toTailEscape`, line 98), [LEAN SOURCE](#). **Conclusion.** If at every late middle row $2\sqrt{2s+2} + 4 < \text{producerCarry}(s)$, then $1/2 \in \mathcal{A}$; `producerCarry` equals the middle coordinate $4\text{rem} - \text{belowPulse} - 4$ ([LEAN SOURCE](#)), so the requirement is essentially $\text{rem}(s) \gtrsim (\sqrt{s} + \text{belowPulse})/4$. **Comparison.** this is “square root” of the *row index* s (a `supportCoeff` tail majorant), not square root of the *dyadic scale* $2^{(r+5)/2}$. The two `sqrt`’s in this record are exponentially far apart and easy to conflate; this hypothesis is row-index bound. The cited conditional theorem uses it to rule out a last non-right transition. An implication between the different orbit hypotheses would require a further argument. **Missing step.** $\forall s \geq 13$ at a middle transition, $4\text{rem}(s) - \text{belowPulse}(s) - 4 > 2\sqrt{2s+2} + 4$; a row-scale, not dyadic-scale, lower bound on the middle coordinate.

Observation 9.13 (The remaining interval of possible integer values). [Lean] [LEAN SOURCE](#), [LEAN SOURCE](#), [LEAN SOURCE](#), [LEAN SOURCE](#), [LEAN SOURCE](#). **Conclusion.** Under the hypothesis of a last non-right transition at row $D \geq 13$, the upper case is excluded. In the middle case put

$$X_{D+1} := \text{rem}(D+1) - 2^{D+1} = 4\text{rem}(D) - p_D^- = C_D + 4.$$

The cited bounds give $2 \leq X_{D+1} < 2\sqrt{2D+2} + 8$. This permits nonnegative C_D as well as the negative exceptional values $-2, -1$; the latter do not exhaust the remaining interval. **Scope of the hypothesis.** the residual bad set is an explicit $O(\sqrt{D})$ -wide integer window at each candidate last non-right transition row, not an interval at dyadic scale. The reset hypothesis would exclude this window. Its width is much smaller than $2^{(r+5)/2}$, but width alone does not measure the difficulty of proving that the actual sequence avoids it. **Missing step.** Exclude every integer

$$2 \leq X_{D+1} \leq \lceil 2\sqrt{2D+2} \rceil + 7$$

at a hypothetical last middle transition $D \geq 13$. The ceiling is necessary when converting the strict real upper bound to an integer interval. Excluding only $C_D = -2, -1$ does not suffice.

Observation 9.14 (Excluding three integer values). [Lean] [LEAN SOURCE](#), [LEAN SOURCE](#), [LEAN SOURCE](#). Theorem 6.9 states the two needed hypotheses and their row ranges explicitly. Excluding $-3, -2, -1$ is required only on middle branches; the overshoot inequality is required only on right branches where $o_s \leq 2^s$. The resulting invariant bounds at least one of the two distances from the target, not necessarily the lower remainder. The exclusion of -3 in the special last-middle-then-all-right configuration does not establish the unrestricted middle-branch hypothesis. Both universal assumptions remain to be proved for the actual sequence.

Observation 9.15 (An upper bound on the remainder). [Lean] [LEAN SOURCE](#), [LEAN SOURCE](#) and the cofinal implication [LEAN SOURCE](#). The estimate $\text{rem}(s) < 2^{s+1}$ for every $s \geq 6$ would give $\text{rem}(s)/4^s \rightarrow 0$ and hence $1/2 \in A$. Theorem 6.2 identifies the normalized limit unconditionally; requiring its value to be zero, even only along one unbounded sequence, is equivalent to membership. No strict separation of these conditions on the actual sequence is proved. The upper estimate is also the missing ceiling in the excess-side right-run argument. The existing bound on rows satisfying $2s < 3d$, where d is the largest omitted rank, still requires that row condition wherever it is applied. The finite checks do not supply it at every later row.

Observation 9.16 (Reported computation through row 200,000). [Cert] **Conclusion.** Exact integer arithmetic to row 200,000 ; branch counts $R : M : U = 100197 : 49899 : 49898$; the obligation *fails only at* $r = 7$; maximum R-run 19 at row 158,096 against the required threshold $\approx (r - 3)/2 \approx 79,000$ (margin $\sim 4000\times$). Independently: 1209 resets checked to row 2500, zero classification anomalies, exactly two crossing cells ever ($s=7, d=5$ and $s=10, d=7$), both below hypothesis scope. **Parameter range.** fixed-scale exhaustive verification versus a statement at every reset $r \geq 31$. This is the canonical difference between a finite and an unbounded parameter range this programme warns against ; a finite verified list is not a cofinal supply, however large the margin. **Missing step.** Nothing computational; only a theorem. The certificate's value is fixing the true constant ($r \geq 31$, sole violation $r = 7$) so the theorem's scope can be stated exactly.

Observation 9.17 (Structure of resets after long runs). [Math] **Conclusion.** At a dangerous reset r ($|\text{rem}(r+1) - 2^{r+1}| \leq 2^{(r+5)/2}$) whose preceding reset r_0 has pure R-run length $L' = r - r_0 - 1 > (r + 5)/4$, the deviation w_{r_0+1} is pinned to at most two adjacent integers determined by the divisor-pulse stream alone; a chain of n dangerous resets forces $n-1$ nested exact integer-coincidence towers. **Scope of the hypothesis.** applies only to dangerous resets preceded by a long pure R-run ($L' > (r + 5)/4$) ; an isolated dangerous reset after a short run is untouched. Pinning to two integers is not yet a contradiction: no theorem excludes those two values. Prose, not Lean. **Missing step.** (i) a congruence or pulse-stream obstruction showing the two pinned integer values are unattainable, and (ii) removal of the long-preceding-run hypothesis, or a proof that dangerous resets must enter chains to which the exclusion applies. Excluding the two pinned values addresses only resets satisfying the long-run hypothesis; covering all remaining resets is a separate requirement. Neither step alone proves a global exclusion.

Observation 9.18 (The sign condition and growth of the margin). [Cert] **Conclusion.** At all 1209 observed resets, $M \Rightarrow \text{deviation} > 0$ and $U \Rightarrow \text{deviation} < 0$, with zero exceptions; the margin over the $2^{(r+5)/2}$ threshold grows from 1.11 bits at $r = 14$ to ~ 1246 bits at $r = 2500$. The record isolates this as a separate, possibly more tractable target than full sqrt-escape. **Different restrictions.** The upper sign follows from the positive reset charge: $w_{r+1} = -(4o_r + p_r^+) < 0$. The middle sign is equivalent to $4\text{rem}(r) > p_r$ and needs further information about the orbit. Neither sign implies a square-root lower bound on the absolute deviation. **Missing step.** Prove that middle inequality and then a quantitative margin on each branch. The exact branch calculation is given in Section 11.1; it is not a substitute for the magnitude estimate.

Observation 9.19 (Comparison of finite truncation lengths). [Cert] The exact finite procedure in Theorem 8.9 verifies the twenty truncations $3 \leq J \leq 22$. The proof combines finite integer certificates with a witness bound valid beyond the finite search horizon; it does not infer an infinite statement from a long finite run. These certificates are reproduced in `audit/final-proof-reading/truncation/`. Representations for an unbounded sequence of J would yield a representation of $1/2$ for the full weights by the compactness estimate following that theorem. That unbounded-index premise remains unproved. The size of the least-common-multiple horizon affects this finite algorithm, but is not an impossibility theorem for other arguments.

9.4 Finite quotient sums and unbounded depths

Observation 9.20 (Exact rows from a sharp capacity bound). [Lean] [LEAN SOURCE](#), [LEAN SOURCE](#). **Conclusion.** For any $c \geq 4$ and any finite $D \subseteq [2, c)$ with $V D < 1/2$ and $S D 1 (2c-2) < 2^{c-2}$: an exact row E at endpoint $2c-2$ with $D \subseteq E$, every new rank strictly above c . No additional crossing condition or identification with a greedy prefix is needed. **Scope of the hypothesis.** needs a per- c input $S D 1 (2c-2) < 2^{c-2}$, and the corpus only ever supplies that input through the crossing-deficit route, which [LEAN SOURCE](#) then identifies the core with $D = \text{halfGreedyPrefixSupport}(c-1)$, under its critical-crossing premises. The fill implication does not require that identification; this difference in its formulation is not a proof of strictness between the associated cofinal existence statements. **Closes it (*)**: $\forall N \exists c \geq N, \exists \text{ finite } D \subseteq [2, c) \text{ with } V D < 1/2 \text{ and } S D 1 (2c-2) < 2^{c-2}$; equivalently $C \uparrow D (2c-3) < 2^{c-2}$. No coherence between the D 's, no requirement that D be a greedy prefix.

Observation 9.21 (A conditional induction for unbounded depths). [Lean] [LEAN SOURCE](#) $\rightarrow$ [LEAN SOURCE](#) $\rightarrow$ [LEAN SOURCE](#) ([LEAN SOURCE](#); seed line 1023; strict progress [LEAN SOURCE](#)). **Conclusion.** Under the sharp-capacity supply hypothesis below, the induction starts from the depth-six seed and constructs exact rows at unbounded depths. A protected row is either doubled while its real value stays below $1/2$, or extended using its first crossing rank e . In the latter case protection gives $2e - 2$ strictly larger than the old endpoint. **Scope of the hypothesis.** the recycle branch consumes `SkippedCoreCriticalQuotientSupply` whenever recycling is required. Protection (endpoint $< 2 \cdot \text{cutoff}$ and new ranks $> \text{cutoff}$) makes that extension strictly increase the endpoint. The induction establishes the conditional implication, not the supply hypothesis. **Missing step.** The sharp inequality at one crossing rank per step, canonical form [LEAN SOURCE](#): $\forall c \geq 4$ skipped by the rational half-greedy orbit, $2^{(2c-2)-1} \leq Q(\text{insert } c (\text{halfGreedyPrefixSupport}(c-1))) (2c-2)$.

Observation 9.22 (Unconditional capacity, off by one bit). [Lean] [LEAN SOURCE](#) (word form line 228). **Conclusion.** Unconditionally, $\forall c \geq 4$ and every below-half core $D \subseteq [2, c)$ with deficit $< w c$: $S D 1 (2c-2) < 2^{c-1}$. Uniform in c , no table, no case split. **Parameter range.** one bit short. The strict-upper fill needs $< 2^{c-2}$, this gives $< 2^{c-1}$. The proof derives $A < 2^{c-2} + |D|$ then discards $|D| \leq c - 2 \leq 2^{c-2}$; the loss is purely additive. The true unproved statement is only that A avoids the linear-width band $[2^{c-2}, 2^{c-2} + c - 3]$, width $c-2$ inside a range of size 2^{c-2} . **Missing step.** A linear-width dyadic-band anti-concentration: at cofinally many crossing ranks c , $S D 1 (2c-2) \notin [2^{c-2}, 2^{c-2} + c -$

3]. Same band shape as the earlier linear-band conditions. The reset condition uses a different coordinate and parameter, with width $2^{(r+5)/2}$; no implication between the hypotheses follows merely from this comparison of widths.

Observation 9.23 (Doubling extension, fails at the only recorded seed). [Lean] [LEAN SOURCE](#), line 296; rank-two discharge [LEAN SOURCE](#), line 108. **Conclusion.** $\forall n \geq 6$: a below-half exact row at n produces an exact row at $2n-1$, support-extending, all new ranks $> n$. The only unconditional endpoint-doubling required input in the corpus. **Scope of the hypothesis.** below-halfness is not preserved and demonstrably fails at the first step. For the seed $D = \{2, 3, 6\}$ at $M = 11$ the core alone contributes $2/3 + 4/7 + 32/63 \approx 1.746 > 1$, so the doubled row is strictly *above* half. Iterating the cheap arm is therefore impossible from the only recorded seed; the chain is forced into the capacity-gated recycle arm at step one. **Missing step.** Cofinally many below-half exact rows ; $\forall N \exists n \geq N \exists D$ exact at n with $\text{localFractionMass } D \ n < 1$. A single such row at each of cofinally many n makes doubling unnecessary; a self-reproducing one closes everything.

Observation 9.24 (Consecutive-skip precritical bound). [Lean] [LEAN SOURCE](#) (feeds [LEAN SOURCE](#), then lines 970/986). **Conclusion.** Unconditional: if $c \geq 6$ and both rank c and rank $c+1$ are skipped by the rational half-greedy orbit, then $S(\text{halfGreedyPrefixSupport}(c-1)) \ 1 \ (2c-3) < 2^{c-3}$, doubling into sharp capacity at c , hence an exact row at $2c-2$. No supply hypothesis anywhere. **Scope of the hypothesis.** needs two *consecutive* skipped ranks. The residual open set is exactly the skip-then-take rows: [LEAN SOURCE](#) proves the whole hypothesis equivalent to [LEAN SOURCE](#), with $c = 4, 5$ already discharged by decide. **Missing step.** Cofinally many $c \geq 6$ at which two consecutive ranks $c, c+1$ are both skipped by the rational half-greedy orbit ; or dually, the pre-take residual at cofinally many skipped-then-taken ranks. Either one closes the obligation with no further input.

Observation 9.25 (Bounded-gap precritical bound). [Lean] [LEAN SOURCE](#) (with [LEAN SOURCE](#), [LEAN SOURCE](#)). **Conclusion.** Unconditional and uniform in *both* parameters: skip at c , takes at $c+1, \dots, c+t-1$, skip at $c+t$, with $0 < t \leq c-3$ and $c-2 \leq 2^{c-t-3} \Rightarrow$ precritical suffix bound at $c \Rightarrow$ sharp capacity $\Rightarrow$ exact row at $2c-2$. Handles any skip gap $t \lesssim c-3-\log_2(c-2)$, not just $t=1$. **Scope of the hypothesis.** the only missing input is an orbit-level skip-gap bound ; that the next skipped rank after c occurs at $c+t$ with $c-2 \leq 2^{c-t-3}$. The proof is fully uniform in (c, t) , every constant explicit, no table. **Missing step.** For cofinally many skipped ranks c of the rational half-greedy orbit, the next skipped rank is at most $c+(c-3-\lceil \log_2(c-2) \rceil)$. The finite skip-density observation does not imply this pointwise gap bound at unbounded ranks. Establishing the latter would supply the hypothesis of the stated conditional argument.

Observation 9.26 (Frozen margin, ineffective horizon). [Lean] [LEAN SOURCE](#), [LEAN SOURCE](#) (margin def line 1021, recurrence line 1036, monotonicity line 1094). **Conclusion.** At every dyadically safe rank $k > 0$, $F \ k \ J \geq 0$ for *some* horizon J ; the finite coefficient window eventually covers the centred carry. Nonnegativity is monotone up in J . **Parameter range.** the horizon comes from a non-effective limit argument (`tendsto_finiteCoeff Window_atTop` plus an existence extraction), so J is unbounded. The exact-row route

needs the crossing by the specific horizon $J = c - 3$ (LEAN SOURCE, $k = c - 1$). All ingredients to make it effective already exist: LEAN SOURCE gives $\text{gap} = \text{tail}(k+1+J)/2^J$, and LEAN SOURCE gives $\text{tail}(m) \leq m + 2$. **Missing step.** $\text{Margin}(k, J) \geq 0$ as soon as $(k+J+3)/2^J < |\text{halfGreedyNextDyadicExcessNumerator } k| / \text{halfGreedyPrefixDenominator } k$. Setting $J = k - 2$ turns the hypothesis into one explicit dyadic-safety margin $|E_k|/D_k > (2k+1)/2^{k-2}$ at every skipped rank. Pure bookkeeping over existing lemmas. *Observation 9.27* (A nonnegative margin at an unspecified horizon). [Lean] LEAN SOURCE, LEAN SOURCE, LEAN SOURCE. **Open property and proved conditional implication:** HalfGreedySkippedFullShellNonnegative asks that every skipped rank $n \geq 3$ satisfy $0 \leq F(n-1)n$. Assuming this property, half_mem_mersenneAchievementSet_of_skippedFullShellNonnegative proves $1/2 \in A$. **Order of quantifiers.** exactly three doubling steps of horizon short. The exact-row hypothesis needs $\text{margin}(k, k-2) \geq 0$; this gives $\text{margin}(k, k+1) \geq 0$. By the succ-recurrence, $\text{margin}(k, k+1) = 8 \cdot \text{margin}(k, k-2) + 4 \text{sc}(2k) + 2 \text{sc}(2k+1) + \text{sc}(2k+2)$, sc a divisor-count $\text{supportCoeff} \leq \tau \leq 2\sqrt{2k+2}$. Since nonnegativity is monotone up in J , the cited statement is the weaker one and cannot supply the obligation as it stands. **Missing step.** A horizon-tightening ; at every skipped rank, $\text{margin}(k, k+1) \geq 4 \text{sc}(2k) + 2 \text{sc}(2k+1) + \text{sc}(2k+2)$: the full-shell margin must exceed an explicit $O(\sqrt{k})$ divisor-count quantity. Same $\sqrt{\cdot}$ -scale currency as the LPC-sqrt row above.

Observation 9.28 (Cofinal positive skips is the obligation). [Lean] LEAN SOURCE, LEAN SOURCE, line 60; A5 LEAN SOURCE, A6 (LEAN SOURCE); positivity discharge LEAN SOURCE. **Conclusion.** Cofinally many positive rational-greedy skips $\Rightarrow$ the existence of exact quotient sums at unbounded depths $\Rightarrow 1/2 \in A$. Advertised elsewhere as the shortest known path. **Scope of the hypothesis.** it is not weaker than the problem ; it is the problem. The positivity conjunct $0 < r(1/2)(c-1)$ is unconditionally dischargeable by the odd-denominator parity fact, so CofinalPositiveHalfGreedySkips $\Leftrightarrow$ (greedyMersenneSkippedSupport(1/2)).Infinite, which Theorem 6.88 proves equivalent to $1/2 \in A$. Thus proving this hypothesis would settle the half-target question itself. **What remains.** One must prove that this fixed greedy sequence has infinitely many skips. Equivalence neither supplies that proof nor rules out finding a useful invariant for the sequence. The condition in Theorem 6.75 instead permits finite supports that are not greedy prefixes.

Observation 9.29 (Terminal bounds and restrictions on finite support). [Lean] LEAN SOURCE, LEAN SOURCE, LEAN SOURCE. **Conclusion.** Cofinally many depths M carrying a finite normalized word a with $|\text{ihc}(\text{wordSupport } a)(M-1)| \leq B M \Rightarrow \exists A$ infinite with $X_A(2) = 1/2$. Same architecture as the obligation: cofinal, no coherence, mutually incompatible witnesses allowed.

Different restrictions. The terminal condition permits *any* support inside the depth- M window but demands the carry inside a $\sim 2\sqrt{M}$ strip; the obligation form $(\star)$ tolerates a carry up to $2^{M/2-1}$; exponentially looser ; but demands the support be confined to ranks $< M/2 + 1$, since the exact fill encodes the residue in the pure upper window where the Mersenne quotient is exactly 2^{M-d} . The gap is support *locality*, not carry size.

Missing step. A lower-half confinement upgrade of the strip witnesses ; $\forall N \exists c \geq N$ with a strip-scale witness whose support lies in $[2, c)$ at depth $2c-3$. Then $2\sqrt{2c} + 4 <$

2^{c-2} for $c \geq 8$ gives the required numerical capacity bound. The confinement of the support remains an additional, unproved condition; the comparison does not establish a simpler existence theorem.

Observation 9.30 (The live greedy carry and the carry of a fixed prefix). [Lean] [LEAN SOURCE](#), [LEAN SOURCE](#), line 157; [LEAN SOURCE](#), line 805 (nonnegativity from line 842 / [LEAN SOURCE](#)). **Conclusion.** Cofinal returns of the actual greedy carry to the square-root strip (C4c), or a pointwise $2\sqrt{N} + 4$ bound on the Möbius-centred carry (G3), each give an infinite support with value exactly $1/2$. **Comparison.** these bound the carry of the *infinite* greedy support `greedyMersenneSupport(1/2)`; sharp capacity is about the *frozen finite prefix* `halfGreedyPrefixSupport(c-1)`. The two carries agree only at index k (inside [LEAN SOURCE](#)), not at depth $2c-3$ where the fill needs it. The bridge is already recorded: [LEAN SOURCE](#): $\text{sharp capacity} \Leftrightarrow C G(2c-4) \leq \text{futureSkipCapacity } G c(c-3)$; transport costs exactly the future-skip capacity term. **Missing step.** Either a strip bound stated for the frozen prefix rather than the live orbit, or $\text{futureSkipCapacity } G c(c-3) \geq B(2c-3)$; holds as soon as one skip occurs early in $[c+1, 2c-4]$, since the capacity is dyadically weighted. Same input as the bounded-gap precritical bound above.

Observation 9.31 (Above-support recycling, no growth guarantee). [Lean] [LEAN SOURCE](#), [LEAN SOURCE](#), [LEAN SOURCE](#). **Conclusion.** Unconditional: every finite above-half support bounded by n recycles through its first crossing rank c into a genuine exact row at endpoint $2c-2$, with $4 \leq c \leq n$. No capacity hypothesis at all. **Order of quantifiers.** the witness is $\exists c \leq n$, not $\exists c$ large; $2c-2$ may be $\leq n$, so the endpoint need not grow. [LEAN SOURCE](#) is the explicit falsifier; `boundedDoubleOrRecycleModel n := (n=6)` satisfies the same transition schema plus a seed and is *not* cofinal. Growth is recovered only inside `ProtectedExactLocalMersenneRow`, whose invariants force $c > \text{cutoff}$, hence $2c-2 > \text{endpoint}$; and maintaining those invariants is precisely what needs the sharp-capacity fill rather than this general recycling. **Missing step.** A crossing-location lower bound; for the supports actually produced, the first crossing rank satisfies $2c-2 > n$. Equivalently, keep protection alive without sharp capacity, e.g. a fill placing all new ranks above c using only the recorded $c-1$ -bit capacity.

Observation 9.32 (Double-or-recycle dichotomy, proved insufficient). [Lean] [LEAN SOURCE](#), [LEAN SOURCE](#), [LEAN SOURCE](#). **Conclusion.** $\forall n \geq 6$, an exact row at n gives an exact row at $2n-1$ or at $2c-2$ for some $4 \leq c \leq n$; plus a formal proof that this shape, even seeded at endpoint six, is logically insufficient for cofinality. **Order of quantifiers.** the dichotomy is proved by splitting on the sign of the witness value (`finiteErdosSum_den_odd` excludes equality); the no-go pins the deficit precisely; what is missing is not another dichotomy but strict endpoint progress in the recycle branch, or a proof that the below-half branch recurs. Re-deriving this dichotomy without controlling repeated returns to bounded depths would leave the same gap. **Missing step.** Either (a) the recycle branch upgraded to $2c-2 > n$, or (b) cofinally many below-half exact rows so the doubling branch alone is cofinal. These are two sufficient ways to obtain progress; the counter-model does not rule out other assumptions that also force unbounded depths.

9.5 Universal support criteria and their limits

Observation 9.33 (A sufficient block-certificate condition). [Lean] [LEAN SOURCE](#) (carry variant [LEAN SOURCE](#)). **Conclusion.** $\forall b \geq 2, \forall A \subseteq \mathbb{N}$ (no infinitude, no structure hypothesis at all): if $\forall q > 0 \exists N, K, L, C$ with $K \leq L, b^r \mid c_A(N+r)$ for $r \in [1, K]$, $\sum_{r=K+1}^L c_A(N+r)b^{L-r} \leq C, \exists t, 0 < c_A(N+L+1+t)$, and $q(C+N+L+2) < b^L$; then $X_A(b)$ is irrational. The cited full-support, multiples, periodic, residue-class, pairwise-coprime and fixed-core constructions use this criterion. No such factorisation is asserted for the weighted support theorem in Section 1. **Scope of the hypothesis.** already universally quantified over A ; the universal quantifier is *not* the gap. The hypothesis `hcrt` is a cofinal required input supplied only for named classes; the theorem's own docstring: "certificate existence for f_A ... for arbitrary infinite A remains an open obligation." **Limit of this sufficient condition.** Supplying these certificates for every infinite support would imply the universal statement, but Proposition 11.4 rules out that supply for the two specified certificate schemes. The carry variant replaces digit-wise divisibility by aggregate divisibility $b^K \mid \sum_r c_A(N+r)b^{K-r}$; the squarefree obstruction applies to that variant too. Universal certificate supply is therefore not being posed as an unresolved equivalent formulation of Problem 257.

Observation 9.34 (Rational targets and infinitely many greedy skips). [Lean] [LEAN SOURCE](#) (forward [LEAN SOURCE](#); reverse [LEAN SOURCE](#)). For every rational $x \geq 0$,

$$x \in A \iff \text{the greedy algorithm for } x \text{ skips infinitely many ranks.}$$

The reverse implication holds for every real $x \geq 0$: a fatal remainder $r_n(x) > R_n$ forces every subsequent weight to be taken, because $R_n = w_{n+1} + R_{n+1}$. Hence infinitely many skips exclude a fatal remainder, and Theorem 1.10 gives membership. For the forward implication, membership identifies the greedy sum with x . Finitely many skips would make that sum E minus a finite rational sum, which is irrational by Erdős's theorem [1]. This contradicts rationality of x and proves the stated extension as an ordinary argument.

Infinitely many skips do not mean infinitely many selected ranks: $x = 1/3$ has the single-element support $\{2\}$ and skips every other rank. Thus the equivalence does not decide which rational targets are represented, or provide an infinite support with rational value. The supplied half-target declarations remain identified above; a new target-uniform Lean theorem is not claimed. [Math]

Observation 9.35 (Finite sums exclude every rational with even denominator). [Lean] [LEAN SOURCE](#); engine [LEAN SOURCE](#). **Conclusion.** $\forall A$ finite with $0 \notin A$: $X_A(2) \neq 1/2$; any support achieving exactly $1/2$ must be infinite. **Parameter range.** the headline is pinned to target $1/2$ while its engine `Odd((finiteErdosSum F 2).den)` is already fully general over finite supports. **Sufficient specialization:** verbatim re-derivation; $\forall t : \mathbb{Q}$ with even reduced denominator, $\forall A$ finite with $0 \notin A$, $X_A(2) \neq (t : \mathbb{R})$. No new mathematics. Thus a represented rational with even reduced denominator would refute universal #257 at $b = 2$. The converse is not established. If `FinVal` denotes the finite-support values, the exact unrestricted condition is $(\mathbb{Q} \cap A) \setminus \text{FinVal} \neq \emptyset$.

Observation 9.36 (The denominator in the reciprocal-mass estimate). [Lean] [LEAN SOURCE](#) (line 2116); generic- v sibling [LEAN SOURCE](#). **Conclusion.** For A infinite with $X_A(2) =$

$p/2^c$ (dyadic rational): $\neg \text{Summable}(1/a \text{ on } A) \vee 1 < \text{reciprocalMass } A$. **Comparison.** restricted to denominators 2^c ($v=1$) while the obligation needs every rational $p/(2^c v)$. This restriction is a call-site choice, not a proof constraint ; every supporting lemma carries a free $\{v : \mathbb{N}\}$ ($h v : 0 < v$), and the general- v conclusion is *already recorded* at line 1962 for arbitrary (v, h, L, F) . **Missing step.** Instantiate the recorded generic theorem at $F = \{a, b\}$, $L = \text{lcm } a b$, $h = \text{multiplicative order of } 2 \text{ mod } v$, giving $\rho(A) \geq \text{doublingWrapCount}(p, v, h)/h + 1/\text{lcm}(a, b)$ for every rational value. The one genuinely new step: for $v > 1$ the excess can be 0, so strictness needs $1 \leq \text{doublingWrapCount}(p, v, h)/h + \text{booleanCollisionSurplus } |F|/L$ (equivalently: the mean least residue of $p \cdot 2^n \text{ mod } v$ is $\geq v(1 - |F|/2L)$).

Observation 9.37 (Zero windows of a fixed support). [Lean] [LEAN SOURCE](#) (line 392). **Conclusion.** For any A with a positive element and $X_A(2) = p/(2^c v)$, $v > 0$ (any rational, not just dyadic): $\forall \varepsilon > 0 \exists B \geq 0 \forall N \geq 1 \forall h, \text{SupportCoeffZeroWindow } A (c+N) h \rightarrow h \leq \varepsilon \log_2 N + B$. The constant B may depend on the fixed support and the other parameters. More basically, if a_0 is any positive element of A , then each block of a_0 consecutive integers contains a multiple of a_0 . Every zero window of c_A therefore has length at most $a_0 - 1$, whether or not $X_A(2)$ is rational. For a fixed nonempty support, the displayed logarithmic bound thus gives no additional restriction. A proposed alternative between super-logarithmic zero windows and general block certificates cannot split this class: the first branch is already impossible. It would require the certificate branch for every support, which the obstruction in Proposition 11.4 rules out for the specified schemes.

Observation 9.38 (Prime support and the reciprocal-sum hypothesis). [Lean][Cited] [LEAN SOURCE](#) (direct certificate required input [LEAN SOURCE](#)). **Conclusion.** For every $b \geq 2$, an infinite pairwise-coprime support with summable reciprocal mass has irrational support series. **Method boundary:** the primes satisfy the other hypotheses but $\sum_p 1/p$ diverges, so this Lean required input does not apply. It is false, however, to infer that the base-2 prime-support value is open. Tao–Teräväinen prove

$$\sum_{p \text{ prime}} \frac{1}{2^p - 1} = \sum_{n \geq 1} \frac{\omega(n)}{2^n}$$

irrational in Theorem 1.3 of [their 2025 preprint](#). That result is cited, not formalised here. **What remains diagnostic:** replacing the global tail budget by a windowed one would ask whether this particular certificate engine can recover the known theorem. It is no longer an open mathematical target or a claimed new route for #257.

Observation 9.39 (The two hypotheses of the sunflower criterion). [Lean] [LEAN SOURCE](#) (selector def line 406; carry hypothesis line 381; bouquet structure line 33). **Conclusion.** For $A = (\text{finite exceptional set dividing } Q) \cup \{\text{core}_i \cdot \text{petal}_i\}$ with cores $\mid Q$, petals pairwise coprime and coprime to Q , $\sum 1/\text{petal}_i < \infty$: if $\text{SunflowerForcedSlotTailSelection } A$ then irrational.

Different restrictions. two gaps. (i) On the support: the bouquet class strictly generalises pairwise-coprime but still demands globally pairwise-coprime, summable petals ; it inherits NM-06's exclusion of the primes and adds nothing in the divisor-dense regime.

(ii) On the selector: SunflowerForcedSlotTailSelection is an unproved cofinal supply ; grep for “dichotomy” in the file returns only the namespace; despite the filename there is no dichotomy theorem there.

Missing step. For the selector, $\forall K > 0 \exists N, 2^K \mid$ the carried first block $\wedge$ binaryCoef fTail(c_A)($N+K$) ≤ 16 , from the bouquet’s persistent-reduced-modulus averaging ; the file states this is what that averaging argument is supposed to give. For the class: a bouquet-existence theorem $A.\text{Infinite} \rightarrow \text{Nonempty}(\text{bouquet structure})$ is *false* as stated (take A = all multiples of 6 ; no coprime petal decomposition). A theorem for every infinite support would therefore require a different structural hypothesis, not merely a proof of that false existence assertion.

Observation 9.40 (What eventual periodicity supplies). [Lean] [LEAN SOURCE](#) (engine line 11262, required input line 11072; finite-perturbation transfer lines 9139, 9148). **Conclusion.** $\forall b \geq 2 \forall m > 0 \forall N_0$: A infinite and m -periodic above $N_0 \Rightarrow$ irrational. Plus: the class of supports with irrational series is closed under finite symmetric difference in both directions. The statement follows from the periodic theorem of Luca and Tachiya [3, Theorem A, p. 139].

Scope of the hypothesis. The theorem allows every positive period and every finite starting threshold. It still requires an eventually periodic support, so it does not include all aperiodic supports covered by the other criteria in this record. Finite symmetric difference preserves eventual periodicity. Thus the finite-perturbation argument alone cannot extend this theorem to an aperiodic tail; a further hypothesis or argument is needed.

Missing step. A “locally periodic at cofinally many scales” version ; $\forall q > 0 \exists N \exists m \leq f(N)$ such that $A \cap [N, N+L]$ is m -periodic for L large enough that the periodic sieve manufactures one block certificate. The periodic required input already only needs periodicity across the certificate window (window-local), so this is the closest genuine widening ; but proving cofinally many periodic windows exist for arbitrary A is a new statement, not a re-run.

Observation 9.41 (The reduced denominator in a gap criterion). [Lean] [LEAN SOURCE](#) (instances [LEAN SOURCE](#), [LEAN SOURCE](#)). Let $A = \{a_0 < a_1 < \dots\} \subseteq \mathbb{N}_{>0}$, let $b \geq 2$ be an integer, and write the first k terms in lowest terms as

$$S_k = \sum_{j < k} \frac{1}{b^{a_j} - 1} = \frac{P_k}{Q_k}, \quad Q_k \geq 1.$$

A sufficient irrationality condition is

$$Q_k b^{-a_k} \rightarrow 0, \quad \text{equivalently} \quad a_k - \log_b Q_k \rightarrow +\infty.$$

The next omitted exponent is a_k , not the last exponent included in S_k . Since $a_{k+j} \geq a_k + j$, the positive omitted tail obeys

$$0 < X_A(b) - S_k \leq \left(\frac{b}{b-1} \right)^2 b^{-a_k}.$$

If $X_A(b) = p/q$ were rational, positivity would give $qQ_k(X_A(b) - S_k) \geq 1$. The displayed tail estimate contradicts this under the proposed condition. This is the denominator-times-error argument in [LEAN SOURCE](#), with all indices specified.

For $L_k = \text{lcm}(a_0, \dots, a_{k-1})$ one has $Q_k \mid b^{L_k} - 1$, hence $\log_b Q_k < L_k$. Thus the earlier condition $a_k - L_k \rightarrow \infty$ suffices. Using actual reduced denominators can improve this estimate, but requires an *upper* bound for Q_k relative to b^{a_k} . A surviving-divisor *lower* bound does not provide that upper bound. In particular, the divisibility statement in [LEAN SOURCE](#) and the lower bound in [LEAN SOURCE](#) do not alone establish the proposed condition for an arbitrary support. No such condition is proved here for the remaining dense supports. [Math]

Observation 9.42 (Rationality expressed by the integer carry). [Lean] [LEAN SOURCE](#) (line 376); T7 [LEAN SOURCE](#). **Conclusion.** $\forall A \subseteq \mathbb{N}$ (no hypothesis at all): $\text{HasRationalValue}(X_A(2)) \Leftrightarrow \exists q > 0 \exists U : \mathbb{N} \rightarrow \mathbb{Z}, \text{IsTemperedBinaryOrbit}(c_A) q U$. **Comparison.** already universal over A ; the exact coordinate in which the obligation should be stated, and the literal shared argument with #249 (instantiate T7 at $c := \text{Nat.totient}$). It does not refute the orbit: [LEAN SOURCE](#) pins the orbit uniquely to $u(N) = qT_c(N)$, and $T_c(N) \leq N+2$ confirms it is never ruled out on growth grounds for any c with $c(n) \leq n$. **Missing step.** A support-specific non-existence theorem; $\forall A$ infinite, $\forall q > 0, \neg \exists U$ with $\text{IsTemperedBinaryOrbit}(c_A) q U$. Because T7 is an iff and the orbit is unique, this is logically equivalent to the obligation; the row's value is naming the single object (the tempered carry orbit of c_A) whose non-existence is the entire content, for every A at once.

Observation 9.43 (Why unbounded carry alone is insufficient). [Lean] [LEAN SOURCE](#), [LEAN SOURCE](#), [LEAN SOURCE](#). Let $A \subseteq \mathbb{N}_{>0}$ be infinite, let $c \geq 0$ and $v \geq 1$ be integers, and let $u(n)$ be positive integers satisfying

$$u(n+1) + vc_A(c+n+1) = 2u(n) \quad (n \geq 0).$$

For a finite $F \subseteq A$ and a common multiple L of its elements with $L > c$, put $n = L - c - 1 \geq 0$. Directly from the recurrence,

$$2u(L - c - 1) = u(L - c) + vc_A(L) \geq 1 + v|F|.$$

This argument uses positivity of the next state, not an assumed tail representation or a limit condition. A general positive solution may also have an exponentially growing homogeneous part. The condition $L > c$ ensures that the coefficient row really is L ; truncated natural-number subtraction cannot replace it.

Finite subsets F of arbitrarily large size, followed by positive common multiples exceeding c , prove that u is unbounded. In the tempered case the separate identity $u(n) = v \sum_{r \geq 1} c_A(c+n+r)2^{-r}$ gives $u(n) \leq v(c+n+2)$. Unboundedness is consistent with this linear upper bound and with $u(n)/2^n \rightarrow 0$. A contradiction would require an incompatible upper estimate at the same chosen indices. [Math]

Observation 9.44 (A finite-gap criterion for every target in the convex hull). [Lean] [LEAN SOURCE](#), [LEAN SOURCE](#); general- A killer [LEAN SOURCE](#). For every real $x \in [0, E]$, not just rational targets, one has

$$x \notin A \iff \exists m \geq 1 \exists D \subseteq \{1, \dots, m-1\}, \quad X_D(2) + R_m < x < X_D(2) + w_m.$$

Here $R_m = \sum_{j>m} w_j$. The implication to the right has a short greedy proof. If $x \notin A$, choose the first m with $r_m(x) > R_m$ using Theorem 1.10. Since $r_0(x) = x \leq E$, we have $m \geq 1$ and $r_{m-1}(x) \leq R_{m-1} = w_m + R_m$. A take at m would leave a remainder at most R_m , so this step must be a skip. Thus $R_m < r_{m-1}(x) < w_m$, giving the displayed gap with the greedy prefix D .

Conversely, the displayed gap contains no achievable value. A support with prefix D either skips m , giving value at most $X_D(2) + R_m$, or takes m , giving value at least $X_D(2) + w_m$. A different prefix first differs at some $j < m$; the inequality $w_j > R_j$ orders its entire continuation interval strictly on one side of the interval with prefix D . Both gap endpoints themselves are achievable, by the supports $D \cup \{m+1, m+2, \dots\}$ and $D \cup \{m\}$. Consequently no parity argument or irrationality assertion about E is needed for this ordinary general-target dichotomy. The restriction to $[0, E]$ is essential to this formulation: targets outside the convex hull are already excluded and need not lie in an internal gap. The general-target equivalence and both endpoint constructions are Lean-checked in [LEAN SOURCE](#). [Math]

Observation 9.45 (Why the signed periodic theorem still requires periodicity). [Lean] [LEAN SOURCE](#), [LEAN SOURCE](#) (nonpos mirror line 14315). **Conclusion.** $\forall b \geq 2 \forall m > 0 \forall \mathbb{Z}$ -valued m -periodic weight w : $\text{irrational}(\sum w(a)/(b^a - 1)) \vee \exists k, z, b^k x = z$; the terminating branch closes unconditionally as soon as $\text{intWeightedCoeff } w$ is one-signed and frequently nonzero. Luca and Tachiya's theorem already gives irrationality for every nonzero periodic weight [3, Theorem A, p. 139].

Different restrictions. two mismatches. (i) periodicity of the weight is still required ; this is the signed lift of NM-08, not an escape from it. (ii) the one-sidedness clause is *free* for #257 (a 0/1 support indicator gives $c_A \geq 0$ automatically) ; the sign machinery buys nothing here; it was built for the signed #249 argument. The genuinely transferable content is the dichotomy shape itself: irrational-or- b -adically- terminating, terminating branch killable by a positive far tail.

Missing step. An aperiodic version of the dichotomy ; for arbitrary A , either $X_A(2)$ is irrational or 2^k -it is an integer. For arbitrary infinite A , exclusion of the terminating branch is itself an additional unproved assertion; finite-support odd-denominator parity does not pass automatically to an infinite limit. An aperiodic dichotomy together with a valid terminating-branch exclusion would close the obligation. Periodicity is used only to manufacture the full-block certificates ("periodicity alone manufactures the full-block certificates," per the file's own docstring), so removing it is again NM-01's missing required input, restated a third time.

9.6 Three combined conditional arguments

The first conditional implication also appears in the supplied release source, as identified below. The other two entries are ordinary mathematical arguments about what the stated hypotheses would imply. None establishes its unproved orbit hypothesis or decides #257.

Theorem 9.46 (A : a sufficient lower bound at every reset). *Assume that $|\text{rem}(r+1) - 2^{r+1}| > 2^{(r+5)/2}$ for every upper or middle reset $r \geq 10$. Then $1/2 \in A$.*

The supplied release source states this implication as [LEAN SOURCE](#). The packet classifies that file as `release_only`; its proof body was inspected for this exposition, without a new kernel replay. Here is the role of the constants. Put

$$B(r) = 2^{\lfloor (r+4)/2 \rfloor} + 2r + 3.$$

For $r \geq 10$, the source proves $B(r)^2 \leq 2^{r+5}$. A right branch at the first crossing of the largest omitted rank r would force $|\Delta_{r+1}| \leq B(r)$, by tracing the exact affine recurrence back to that reset. The assumed strict lower bound excludes this crossing. The largest-skip induction then gives half-membership. The hypothesis concerns every actual reset beyond the threshold; neither a long finite check nor an asserted run-length asymptotic supplies it.

Lean: [paper_theoremA_half_membership](#), [paperResetSqrtEscape_iff_square](#), [paper_theoremA_crossing_bound_square_le](#), [paper_theoremA_right_branch_forces_small_deviation](#), and 1 further declaration in the [coverage section](#).

Remark 9.47 (B: rigidity of dangerous resets). A dangerous reset (one within $2^{(r+5)/2}$ of the sqrt-escape threshold) that is preceded by a long pure R-run pins the deviation at the preceding reset to one of at most two adjacent pulse-determined integers, where one might expect a free real number: the affine excess recurrence (Observation 9.5) is exact, so once the run length and the pulse word are fixed, w_{r_0+1} has no remaining degree of freedom. Observation 12.9 states the hypotheses and the run-length threshold. Consequently a chain of n consecutive dangerous resets with long runs is a tower of $n-1$ nested exact integer conditions, none free to vary once the pulse stream is fixed. Failure of sqrt-escape therefore requires such an integer-coincidence tower. No probability model is supplied, so these conditions do not by themselves make dangerous resets unlikely; the open step is excluding the two pinned values. [Math]

Theorem 9.48 (C ; the one-sided finite decision boundary). The checked equivalence [LEAN SOURCE](#) separates two logically different kinds of evidence. A fatal greedy gap found at a finite rank is a finite certificate that $1/2 \notin A$. By contrast, membership requires survival at every rank; the corpus supplies no finite certificate that the orbit survives forever and no completion theorem turning a long surviving prefix into membership.

The reset-crossing hypotheses, truncation-rung ladder, and sharp-capacity inequalities feed sufficient conditions for membership. They are not proved equivalent to one another, nor is failure to find one at a given depth evidence of survival. Consequently these searches give at most a semi-decision procedure for nonmembership of the specific value $1/2$. Halting does not prove the universal statement in #257, and non-halting does not prove that $1/2$ is represented. In particular, neither truth nor falsity of the universal problem is shown semi-decidable here. Deeper negative search alone establishes nothing beyond the tested finite ranks; this is the quantifier boundary recorded by SE-7/RC-7.

Lean: [paper_one_sided_finite_decision_boundary](#).

9.7 Four limitations of the stated implications

The following conditions are equivalent to their target statements within the specified scope. An equivalence may still be useful by exposing a quantity that is easier to estimate. What it does not supply is a proof of that estimate, or a strict weakening of the logical obligation.

Remark 9.49 (hbound is equivalent to the conclusion, greedy support only). [LEAN SOURCE](#) and its packaged form [LEAN SOURCE](#) take a hypothesis hbound: $\forall N, C \wedge N \leq 2\sqrt{N}+4$. For the canonical greedy support $A = \text{greedyMersenneSupport}(1/2)$ specifically, hbound is logically *equivalent* to the theorem's own conclusion $X_A(2) = 1/2$; not a strictly weaker sufficient condition. This follows from the exact residual identity at line 842, the unconditional sqrt-envelope tail bound ([LEAN SOURCE](#)), the unconditional tail-nonnegativity lemma ([LEAN SOURCE](#)), and the greedy-specific fact $X_A(2) \leq 1/2$ (line 859). **The scope limit is exact and must not be widened by inference:** this equivalence does *not* extend to every A with $1 \notin A$; dropping the greedy-specific $\leq 1/2$ fact yields only $\text{hbound} \Leftrightarrow \delta \leq 0$ (a one-sided condition), not $\delta = 0$. For the sibling hypothesis [LEAN SOURCE](#), the former `Real.sqrt/Nat.sqrt` adapter gap is repaired by restricting cofinal witness depths to perfect squares. Any remaining gap lies in the required input hypothesis, not in rounding.

Remark 9.50 (What the terminal carry condition proves). [LEAN SOURCE](#) at the original strip constant is logically equivalent to $1/2 \in A$ after restricting cofinal witness depths to perfect squares and using the recorded tail bound. The square-root conversion is no longer an open adapter; the open content is the required input itself.

Remark 9.51 (Positive greedy skips and half-membership). Repeated here for emphasis from the 257-cofinal-rows entry above: once the positivity conjunct of [LEAN SOURCE](#) is recognised as unconditionally free (odd-numerator/even-subtrahend parity plus greedy nonnegativity, discharged at [LEAN SOURCE](#)), the hypothesis collapses to $(\text{greedyMersenneSkippedSupport}(1/2)).\text{Infinite}$, proved by A6 ([LEAN SOURCE](#)) to be *equivalent* to $1/2 \in A$. A proof or disproof of this one Prop settles #257-at-1/2 and vice versa; it is not a waypoint toward that settlement.

Remark 9.52 (Scope discipline for the equivalence class). Each equivalence uses its stated domain. The centred-carry upper bound concerns the actual greedy support G , not every support omitting 1. The cofinal terminal condition already works with the original constant 4, because perfect-square depths avoid a rounding loss. The positive-skip equivalence is stated here at the target $1/2$; extending it to another target requires checking the target-dependent arguments. None of these qualifications supplies the still-missing unbounded sequence of witnesses.

10 Hypothesis-specific obstructions

The examples in this section test particular proposed implications for the reset estimate. Some give counterexamples in abstract recurrence models; others compute an exact finite orbit or compare two quantitative bounds. An abstract counterexample is not a state of the actual greedy orbit unless that correspondence is established. Additional arithmetic hypotheses may therefore change the conclusion.

Throughout, $\text{rem}(s)$ is the quotient-row remainder, $\lambda_n := \Delta_n/2^n$ the normalized seam deviation of Theorem 5.16, and A the Mersenne achievement set. Reset $\sqrt{\text{escape}}$ from row 10 is the statement $|\text{rem}(r+1) - 2^{r+1}| > 2^{(r+5)/2}$ for every reset row $r \geq 10$ (RC-4, Part III). It is a sufficient input for the cited half-membership argument, not a necessary condition for every possible proof.

10.1 Four tested approaches to the reset estimate

Observation 10.1 (Prime steps can occur within a bounded drift window). **Route as conceived.** The universal recursion $K(M+1) = 2K(M) - (\tau(M+1) - 1)$ (Part I, SE-1) doubles at every step and subtracts the divisor-excess $\tau(M+1) - 1$. At a prime $M+1 = p$, $\tau(p) = 2$ so the subtracted term is 1; the naive hope is that primality of an index forces the step to be “almost pure doubling,” and that prime density (Chebyshev/PNT-scale) then bounds how long a run of near-pure-doubling steps can last, which would bound the danger run-length L_r in the right-branch recurrence in Theorem 5.16. The upper branch has a different, adjacent-spacing correction, so a run-length argument cannot use the discarded uniform-error formula across resets.

A finite test. Use the skip set of the quotient row $n = 607$ from Section 5.4, and define the frozen-row drift

$$J_{607}(M) = K(M) + \sum_{d \in \text{Skip}_{607}} \left\lfloor \frac{2^M}{2^d - 1} \right\rfloor.$$

Exact integer arithmetic gives, in order at $M = 607, \dots, 617$,

$$5, 5, 4, 4, 6, 2, 3, 4, 5, 2, 3.$$

There are eleven values and ten transitions. The indices 607, 613, and 617 are prime. For $M+1 \geq 607$ the recurrence is

$$J_{607}(M+1) = 2J_{607}(M) - (\tau(M+1) - 1 - c_{\text{Skip}_{607}}(M+1)).$$

At a prime index the last coefficient is 1. The composite-index coefficients can compensate for those prime steps, as the displayed values illustrate. This table is a fresh finite arithmetic check, not a replay of the longer historical run search.

Scope. Primality at some indices does not force this drift to leave a small interval immediately. The example supplies no universal run-length bound and no statistical conclusion about prime indices. It also does not exclude a proof that combines prime counts with quantitative information about the intervening divisor counts. *prime-density* [Cert]

Observation 10.2 (Parity and perfect squares in the recurrence). **Route as conceived.** $K(M)$ is even exactly when M is a perfect square (a true, provable fact). The hope was to use this parity law to forbid certain residues or run patterns in the digit stream of K , since parity constraints are a classical first attack on digit-run questions.

What the parity calculation says. The recurrence implies $J(M+1) \equiv \delta(M+1) \pmod{2}$ for every integer forcing sequence. For K , the additional arithmetic fact that $\tau(M)$ is odd exactly at squares gives the stated parity pattern. One should not call that pattern a property of every forcing sequence, or say that it excludes no residue: at each fixed index it does specify the parity.

Scope. Restating a parity consequence of a known recurrence does not add a new independent constraint. To turn it into a run-length bound would require further information about the sizes or arrangement of the forcing terms. No argument here rules out using parity together with such information. *parity-tautology*

Use of the observation. The parity law does constrain $K(M)$ at each specified index and is a useful check on any proposed run pattern. The missing implication is from this residue information to a quantitative bound on run lengths. No such implication is proved by the parity calculation alone. [Math]

Observation 10.3 (What the denominator estimate does not give). One proposed separation argument would prevent $\sum_{d \in \text{Skip}_n} x_d$ from approximating $C := E - \frac{3}{2} = 0.1066951524152917 \dots$ at scale $2^{-3n/2}$. An irrationality-measure bound is relevant only after relating the reduced denominator of that finite sum to n .

For $n \geq 2$ and $\text{Skip}_n \subseteq \{2, \dots, n-1\}$, its reduced denominator D_n satisfies

$$\log_2 D_n \leq \sum_{d \in \text{Skip}_n} \log_2(2^d - 1) \leq \sum_{d \in \text{Skip}_n} d \leq \frac{n(n-1)}{2} - 1.$$

The middle inequality is strict when the skip set is nonempty. This is an upper bound; it is not an asymptotic formula for D_n . In particular, cardinality alone does not justify an asymptotic of order $n^2/4$.

Zudilin proves $\mu(E) \leq 2.46497868 \dots$ [8, Theorem 1, p. 154]. Rational translation preserves the irrationality exponent, so for any fixed $\nu > 2.46497868 \dots$ there is a constant $c_\nu > 0$ such that every reduced rational p/q , with $q \geq 1$, satisfies $|C - p/q| \geq c_\nu q^{-\nu}$. Using only the displayed denominator estimate therefore gives

$$\left| C - \sum_{d \in \text{Skip}_n} x_d \right| \geq c_\nu 2^{-\nu(n(n-1)/2-1)}.$$

This guaranteed lower bound is eventually smaller than $2^{-3n/2}$, so it does not prove the desired separation. A sharper estimate for the actual reduced denominators, or an approximation result using additional structure of these numerators, could change the comparison. Neither is excluded by this calculation. *diophantine-approximation*[Math]

Retaining only the parity of an input coefficient leaves considerable freedom. In $e' = 2e + c$, every incoming integer e admits the odd choice $c = -1 - 2e$, which gives $e' = -1$. The next proposition uses the same centred-remainder construction when more unit bits and a valuation are prescribed.

Proposition 10.4 (Centred completion of fixed-precision 2-adic data). *Fix $u \geq 1$ and a finite list $(v_i, a_i)_{0 \leq i < m}$, where $v_i \geq 0$ are integers and the integers a_i are odd. For every initial integer e_0 , there are integers z_i, e_{i+1} such that*

$$e_{i+1} = 2e_i + 2^{v_i}(a_i + 2^u z_i), \quad |e_{i+1}| \leq 2^{v_i+u-1} \quad (0 \leq i < m).$$

Thus one may prescribe the valuation and u bits of the odd unit of each input coefficient while keeping every successor state in its stated centred interval. The initial state is arbitrary. *2-adic-valuation-unit*[Lean](finite words)

Lean: [paper_centred_completion_of_fixed_precision](#).

Proof. For a single step, put $M = 2^{v_i+u}$ and $R = M/2$, and choose

$$e_{i+1} = ((2e_i + 2^{v_i}a_i + R) \bmod M) - R.$$

The remainder convention $0 \leq x \bmod M < M$ gives $-R \leq e_{i+1} < R$. Also $e_{i+1} - 2e_i \equiv 2^{v_i}a_i \pmod{M}$, so this difference has the required form $2^{v_i}(a_i + 2^u z_i)$. Repeating the construction proves the finite-list statement. It uses no relation between the freely chosen quotients z_i at different steps. $\square$

The formal-source statements are [LEAN SOURCE](#), [LEAN SOURCE](#) and [LEAN SOURCE](#). They concern this unrestricted congruence model, not coefficients already fixed by a divisor-count sequence. Fixed-precision data alone cannot exclude every completion in this model. The proposition does not show that the completion comes from a Mersenne support, satisfies additional arithmetic relations, or belongs to the particular greedy orbit. In particular, it is not a general obstruction to every fixed-window argument for Problem 249 or 257.

These observations have separate scopes: a finite prime-window example, a parity consequence of a known recurrence, an insufficient denominator estimate, and a completion theorem for specified fixed-precision symbols. They do not exhaust the possible approaches to the reset estimate or force every proof to use digit-carry machinery. The next section studies that machinery because it gives explicit conditional implications, not because other methods have been ruled out.

10.2 Excluding particular final dyadic intervals

The two-sided bound $\min(\text{rem}(s), \text{overshoot}(s)) \leq 2^s$ for all $s \geq 5$ is conditional. It propagates by induction from the local hypothesis `SeamTwoSidedDyadicCellEscape`, which at every row excludes three specified integer values on the middle branch and imposes one further inequality on the right branch:

$$\begin{aligned} s \geq 5, \quad \text{no carry and a middle transition} &\implies C_s \notin \{-3, -2, -1\}; \\ s \geq 5, \quad \text{no carry and a right transition, } \text{overshoot}(s) \leq 2^s &\implies \text{charge}(s) \leq 2^{s+2}. \end{aligned}$$

Here $C_s = 4 \text{rem}(s) - p_s^- - 4$. On a middle transition, $\text{rem}(s+1) - 2^{s+1} = C_s + 4$; on a right transition, $\text{rem}(s+1) = C_s - 2^{s+1}$. Thus C_s must not be identified with the next row's centred remainder. [LEAN SOURCE](#)

Proposition 10.5 (The upper (carries) successor needs no exceptional-cell exclusion). *Let $s \geq 5$ and suppose the upper branch occurs at row s , that is, $(\text{seamAdjacentCut } s \text{ hs}).\text{successorCarries}$ holds. Then*

$$\text{rem}(s+1) \leq 2^{s+1}.$$

[LEAN SOURCE](#) Indeed, the reset identity expresses 2^{s+1} as $\text{rem}(s+1)$ plus a nonnegative reset charge: [LEAN SOURCE](#). Thus the upper branch needs no exceptional-value exclusion for the next-row two-sided invariant. This conclusion concerns that induction step, not the absence of upper transitions on the orbit. *quotient-branch-classification* [\[Lean\]\(n/a\)](#)

Lean: [paper_upper_branch_needs_no_exceptional_cell](#).

Theorem 10.6 ($C_D = -3$ is impossible at the final middle transition, $D \geq 13$). Consider a hypothetical final middle transition: a middle transition at row $D \geq 13$ ($\neg$ carries, and the middle-branch inequality $4\text{rem}(D) + \text{gap} - \text{belowPulse} < \text{terminalWeight}$ holds at D), followed by an all-right tail forever after ($\forall s \geq D+1, \text{seamGreedyWord}(s+1) = \text{seamGreedyWord}(s).\text{extend true}$). Under these assumptions,

$$\text{belowPulse}(D) + 2 \leq 4 \cdot \text{rem}(D),$$

i.e. $C_D \geq -2$; the cell $C_D = -3$ (and every more negative value) is excluded. [LEAN SOURCE](#)

Proof of the bound. The all-right-tail assumption puts the terminal-augmented finite prefix strictly above $1/2$ ([LEAN SOURCE](#)). Consequently the producer carry is below its complete incidence tail ([LEAN SOURCE](#)), or equivalently $Z_D < 1 + 1/(2^D - 1)$, where $Z_D = \text{rem}(D) - \text{seamWordFloorError}(D)$ ([LEAN SOURCE](#)). The lower bound uses the other endpoint: under the same tail assumption, the finite value at D plus the complete Mersenne tail remains below $1/2$. The exact scaled remainder identity therefore gives

$$Z_D > 4^D \sum_{e>D} \frac{1}{2^e - 1} - 2^D > \frac{1}{3}.$$

The last inequality follows by retaining the first two geometric channels of the complete Mersenne tail. Put $E_D = \text{seamWordFloorError}(D)$. Pulse absorption gives $\text{belowPulse}(D) \leq 4E_D$, so

$$4\text{rem}(D) - \text{belowPulse}(D) \geq 4(\text{rem}(D) - E_D) = 4Z_D > \frac{4}{3}.$$

The left side is an integer, hence is at least 2, as required. The upper threshold inequality alone would not give this lower bound.

Scope. This excludes $C_D = -3$ at a middle row $D \geq 13$ under the additional all-right-tail assumption. The general induction requires exclusion of all three cells at every middle row, not just under this extra hypothesis, together with its right-branch bound. The two statements share the same coordinate C_D but have different hypotheses. Theorem 10.6 is universal over rows satisfying the additional all-right-tail assumption; it does not assert that such a row exists. The induction hypothesis concerns every middle row, without that tail assumption. Excluding $C_D = -3$ in the former statement therefore does not supply the exclusion required in the latter. *final-required input; conclusion-excess* **Object or representation:** about the object under the stated hypothesis; a genuine arithmetic consequence of the all-right-tail assumption via the fatal-gap orbit, not a coordinate artifact. [Lean](n/a)

Lean: [paper_final_middle_cell_at_least_neg_two](#).

Corollary 10.7 (Only $C_D \in \{-2, -1\}$ remain among the exceptional dyadic cells). At a final middle row $D \geq 13$, Theorem 10.6 removes -3 from the exceptional set $\{-3, -2, -1\}$, leaving $-2, -1$ within that set. It does not assert $C_D \in \{-2, -1\}$: nonnegative values have not been excluded. Proposition 10.5 handles upper transitions in the two-sided induction; it supplies no additional restriction on a middle coordinate. For the all-middle-row induction, the exclusion of all three values still needs proof without an all-right-tail assumption, together with the separate right-branch inequality. *final-required input; conclusion-excess* [Lean]

Lean: [paper_final_middle_cell_remaining_cells](#).

Remark 10.8 (A sufficient tail inequality and its source). Let U_D be the finite set of exponents selected by the below word at row D , and put $F_D = U_D \cup \{D\}$. With $c_{F_D}(n) = \#\{d \in F_D : d \mid n\}$, define

$$\Theta_D := \sum_{j \geq 1} c_{F_D}(2D + 2 + j) 2^{-j}.$$

This is a tail sum, not a ratio. Since $0 \leq c_{F_D}(n) \leq |F_D|$, it satisfies $0 \leq \Theta_D \leq |F_D|$. The integer $C_D = 4 \operatorname{rem}(D) - \operatorname{belowPulse}(D) - 4$ is the producer carry in the cited source. Under the final-middle and all-right-tail assumptions, that source proves $C_D < \Theta_D$: [LEAN SOURCE](#).

An inequality in the opposite direction would exclude such a final middle transition. The following sufficient condition retains the exception already handled by Theorem 10.6:

$$(\text{TD}) \quad \forall D \geq 13, \quad (D \text{ is a middle row and } C_D \neq -3) \implies \Theta_D < C_D.$$

Because $\Theta_D \geq 0$, this condition in particular excludes $C_D = -2$ and $C_D = -1$. It requires an inequality at the other indicated middle rows as well, so it is not merely a reformulation of excluding two integer values. Nor does it supply the separate right-branch hypothesis used in the two-sided induction.

The supplied release already defines this condition as `SeamMiddleProducer-TailEscapeExceptNegThree` and proves half-membership from it: [LEAN SOURCE](#). The inequality itself is not established here. Its role is a conditional implication, not a newly identified necessary condition.

The two more concrete sufficient bounds in Theorem 6.7 imply the tail inequality even without the exception. In the source they are [LEAN SOURCE](#) and [LEAN SOURCE](#). The cardinality condition is

$$|U_D| + \operatorname{belowPulse}(D) + 5 < 4 \operatorname{rem}(D).$$

It gives $|F_D| = |U_D| + 1 < C_D$, hence $\Theta_D < C_D$. The simpler row condition $D \leq \operatorname{rem}(D)$ is stronger: $\operatorname{belowPulse}(D) \leq 2|U_D|$ and $|U_D| \leq D - 2$ give $|U_D| + \operatorname{belowPulse}(D) + 5 \leq 3D - 1 < 4D$. The corresponding half-membership implication is [LEAN SOURCE](#). No converse between these sufficient conditions is claimed. [Open](cofinal)

10.3 Related counterexamples and restrictions

The next results give counterexamples to particular summaries and estimates for integer carries. Their hypotheses matter: a countermodel for one summary is not a prohibition on every proof that uses carries or on every other input.

Proposition 10.9 (A finite-state restriction for the stated pulse family). *For a “balanced pulse” family at location $m \geq 2$ (radius $\rho = \lfloor (m+1)/2 \rfloor$, parameters $0 \leq r \leq \rho$, moving mass between positions m and $m+1$ with the weighted total $2c(m) + c(m+1) = 2\rho$ of the two*

coefficients fixed), if a predecessor state is constant across the whole family, then no function $\text{decode} : \text{State} \rightarrow \mathbb{N}$ can recover the parameter r from $\text{state}(r)$ for every r . The family has exactly $\rho + 1 = \lfloor (m+1)/2 \rfloor + 1$ members, so the fan-out is unbounded in m . More strongly, any finite set of states carrying an exact decoder for the family has at least $\rho + 1$ elements, unbounded in m . [LEAN SOURCE](#) (exact family size at [LEAN SOURCE](#)).

Scope. Excludes bounded-state encodings of pre- m history that must distinguish every member of the displayed balanced-pulse family, whether the coefficients are φ or a Möbius-support indicator. Applying it to #249 or #257 requires showing that the relevant orbit realises that family, and it does not rule out every finite-state proof strategy. **binary digits, generic** **Object or representation:** about a representation class, the encodings of history that separate the balanced-pulse family; it says nothing about C or A directly. [\[Lean\]\(n/a\)](#)

Lean: [paper_pulse_family_no_autonomous_decoder](#), [paper_balanced_pulse_fanout_is_radius_succ](#), [paper_balanced_pulse_fanout_unbounded_corrected](#), [paper_pulse_family_finite_state_card](#), and 2 further declarations in the [coverage section](#).

Proposition 10.10 (Möbius-support countermodel: the natural negative-sign candidate overshoots). *The signed Lambert identity $\sum_{d \geq 1} \mu(d)/(2^d - 1) = 1/2$ is exact. Writing $N := \{d : \mu(d) = -1\}$: $\sum_{d \in N} 1/(2^d - 1) = 1/2 + \sum_{d \in P} 1/(2^d - 1)$ where $P := \{d \geq 2 : \mu(d) = 1\}$, and quantitatively $1/2 + 1/63 \leq \sum_{d \in N} 1/(2^d - 1)$ (using the first positive tail term $d = 6$, $\mu(6) = 1$); the negative-Möbius support strictly overshoots $1/2$ by at least $1/63$. [LEAN SOURCE](#) (exact decomposition at [LEAN SOURCE](#)).*

Scope. Rules out exactly one natural candidate infinite Boolean support (the negative-Möbius set) as a witness for $1/2 \in A$. A route-sufficiency no-go only; it says nothing about whether some other infinite Boolean support sums to $1/2$. **divisor counts and finite sums** **Object or representation:** about the object; a genuine value inequality for one specific candidate set, not a coordinate artifact. Margin $1/63$ is the concrete number any repair attempt (adding/removing finitely many elements) must close or exceed. [\[Lean\]\(fixed\)](#)

Lean: [paper_mobius_support_overshoots_half](#), [paper_first_positiveMobius_tail_term](#).

Proposition 10.11 (Finite computation and half-membership). *No finite positive-index Boolean support has value exactly $1/2$: the reduced denominator of any finite Mersenne subset-sum is provably odd (each $2^n - 1$ is odd), while $1/2$ needs an even denominator. [LEAN SOURCE](#). Separately, `CertifiedGreedyMersenneDeath` is a finite-depth certificate of $x \notin A$, obtained by decidable tests on rational input. For example, the supplied source excludes $3/4$ at level 1 with lookahead 0: [LEAN SOURCE](#). Failure to find such a certificate at a given depth does not establish membership. It records survival of that finite test, not survival at every depth. [LEAN SOURCE](#).*

Scope. Any support representing $1/2$, if it exists, must be infinite. This follows from denominator parity, independently of the certificate search. The search supplies a different fact: a successful exclusion certificate proves nonmembership, whereas its absence in a bounded search does not prove membership or rationality. The same distinction must be checked separately for any other certificate family. **denominator parity; finite search** [\[Lean\]\(fixed\)](#)

Lean: [paper_finite_support_and_onesided_certificate](#).

Proposition 10.12 (Period exclusion for the totient series). *For $\sum_n \varphi(n)/2^n$, the cited finite test proves that $\text{totientTail}(N + h) - \text{totientTail}(N)$ is not an integer by excluding every*

possible integer state in a bounded range within finitely many steps. If the series were rational, some positive period h_0 would make these tail differences integral for every sufficiently large N . Telescoping would then give integrality also for every positive multiple mh_0 .

Consequently, the required certificate supply is: for every $h_0 \geq 1$ and every lower bound N_0 , there exist $m \geq 1$, $N \geq N_0$ and a finite test length K excluding integrality for (mh_0, N) . The unbounded choice of N is essential; a counterexample before the unknown eventual threshold does not exclude an eventual period. [LEAN SOURCE](#). It is sufficient instead to obtain such certificates for $h = \text{lcm}(1, \dots, t)$ with both t and N exceeding arbitrary prescribed lower bounds: [LEAN SOURCE](#). The supplied finite theorem excludes the tail differences for $1 \leq h \leq 16$ at $(N, L) = (14, 9)$, and hence excludes rational values whose reduced denominator divides $2^{14}(2^h - 1)$ for one of those h . It does not establish the unbounded certificate supply. [LEAN SOURCE](#).

Scope. The declarations in `CarrySurvivorExtinction.lean`, `AdjacentCarryTube.lean`, `AdjacentPhaseSeparation.lean` and `TotientCarryKernelRigidity.lean` used in this comparison concern the totient series of Problem 249. Their carry recurrences may suggest constructions for Problem 257, but applying them there requires a new argument for the divisor transform of its chosen support. The totient-specific theorems do not supply that argument. *integer carries; totient series* `[Lean](cofinal)`

Lean: [paper_carry_survivor_extinction](#), [paper_periodLcm_is_prefix_lcm](#).

10.4 The scalar-localisation height obstruction

Lemma 10.13 (Denominator complement survives scaling). *For $x : \mathbb{Q}$, $c : \mathbb{Z}$, $H : \mathbb{N}$: if $H \mid x.\text{den}$ and $(c \cdot x).\text{den} \mid H$; i.e. multiplying by the integer c shrinks the displayed denominator down into H ; then the **complementary** denominator factor $x.\text{den}/H$ divides c :*

$$H \mid x.\text{den} \wedge (c \cdot x).\text{den} \mid H \implies x.\text{den}/H \mid |c|.$$

[LEAN SOURCE](#). Equivalently, writing $D = x.\text{den}/H$, we have $D \mid c$, so $t = c/D$ is an integer and $Hcx = t x.\text{num}$. Here $H > 0$ follows from $H \mid x.\text{den}$, since the reduced denominator is positive. [LEAN SOURCE](#).

Scope. This is elementary rational arithmetic, independent of the problem. If $c \neq 0$, divisibility gives the size bound $|c| \geq x.\text{den}/H$. The nonzero condition is essential: $c = 0$ clears every denominator and satisfies the divisibility conclusion without any positive lower bound on $|c|$. The lemma therefore constrains denominator clearing by a nonzero bounded multiplier, not arbitrary multiplication. *rational-denominator; height* **Object or representation:** about a representation; pure denominator bookkeeping for $\mathbb{Q}$, with zero problem-specific content; the upstream primitive ([LEAN SOURCE](#)) is the domain-neutral extraction this lemma builds on. `[Lean](n/a)`

Lean: [paper_scalar_localization](#), [paper_scalar_localization_size_bound](#), [paper_scalar_localization_zero_degenerate](#).

Corollary 10.14 (Mersenne specialisation). *Let x be a positive rational number and let $r \geq 0$, $n \geq 1$ be integers. If $2^r \mid x.\text{num}.\text{natAbs}$ and $x < 2/(2^n - 1)$, then $2^r \cdot (2^n - 1) < 2 \cdot x.\text{den}$; a numerator 2-power lower bound plus a Mersenne-scale upper bound on x together force a denominator lower bound, without introducing a global prefix LCM. [LEAN SOURCE](#) (generic form [LEAN SOURCE](#)). To see the inequality directly, write $x = a/b$ in lowest terms with $a, b > 0$. The hypotheses give $a \geq 2^r$ and $a(2^n - 1) < 2b$. The conclusion follows by substitution. An*

application of Lemma 10.13 must in addition supply its denominator-divisibility assumptions. rational-denominator; height, Mersenne instance [Lean](n/a)

Lean: [paper_mersenne_height](#).

10.5 The final-skip band formula, and what it does not show

Definition 10.15 (Upper-reset dyadic band escape). At each upper-reset row $d \geq 13$, write $E_d = 4 \text{overshoot}(d) + \text{abovePulse}(d)$ for the reset expression in the preceding transition formulas. The condition is

$$\text{for every } 0 \leq j \leq d, \quad E_d > 2^{d-j+1} \quad \text{or} \quad E_d + 2(d+j) \leq 2^{d-j+1}.$$

Thus E_d must avoid the interval $(2^{d-j+1} - 2(d+j), 2^{d-j+1}]$ at every indicated scale. **LEAN SOURCE**. In words: at every upper-reset row $d \geq 13$, the reset charge must avoid a linear-width forbidden band immediately below every dyadic threshold 2^{d-j+1} , for every $j \leq d$ simultaneously. Granted this hypothesis, **LEAN SOURCE** gives $1/2 \in \mathcal{A}$ outright.

Proposition 10.16 (Quantifier collapse: one critical index suffices, not $d+1$). Let $d, E \in \mathbb{N}$ and assume $E \leq 2^{d+1}$. Then the purely combinatorial statement $\text{DyadicBandEscape}(d, E) \iff \exists j, \text{CriticalDyadicBandIndex}(d, E, j) \wedge E + 2(d+j) \leq 2^{d-j+1}$ collapses the $\forall j \in [0, d]$ band-avoidance condition (formally $d+1$ separate inequalities) to checking exactly one nearest-boundary index j . The range assumption ensures that a dyadic threshold lies at or above E . Choose the smallest such threshold: smaller thresholds are already below E , while larger thresholds have narrower forbidden bands. Without the range assumption all bands escape automatically when $E > 2^{d+1}$, but no critical index exists; $(d, E) = (0, 3)$ is the smallest example. Specialised to the concrete seam reset charge, $\text{SeamUpperResetCriticalBandEscape}$ is proved logically equivalent to Definition 10.15's hypothesis. **LEAN SOURCE** (seam specialisation **LEAN SOURCE**). Zero Mersenne/seam content in the core lemma ; pure (d, E, j) arithmetic over powers of 2. dyadic-boundary, generic [Lean](n/a)

Lean: [paper_critical_band_index_collapse](#).

Observation 10.17 (A verified finite range of reset indices). The linked certificate proves the condition in Definition 10.15 for each reset index $13 \leq d \leq 30$. It uses successor remainders at rows 14–31, including $\text{rem}(14) = 392$ and $\text{rem}(31) = 4,187,487,147$. Thus the last remainder row is 31, but the last reset index covered by this theorem is 30. **LEAN SOURCE**. [Cert](bounded)

Remark 10.18 (What this does *not* show ; read the quantifiers exactly). Definition 10.15 is a $\forall d \geq 13$ statement; Certificate 10.17 is a finite verification in the cited Lean source for $13 \leq d \leq 30$. **The finite certificate does not establish the universal hypothesis**. Nothing in the tree proves $\text{SeamUpperResetDyadicBandEscape}$ for all $d \geq 13$; this is exactly the unproved hypothesis named in Part I (SE-9) and Part III (RC-4/D1's LargestSkipLateStepSocket), stated here in its own coordinate.

The finite theorem supplies no reset index beyond 30. It therefore does not prove that the orbit avoids the band at unboundedly many reset indices, let alone at every reset. The universal band condition remains an unproved sufficient hypothesis.

For the branch convention of Theorem 5.16, direct integer recomputation at reset-test indices $6 \leq r \leq 2500$ gives 605 middle and 604 upper resets, hence 1209 resets. There are 1286 right transitions. The earlier reported count 1264 does not agree with this convention. The recomputation checks the sign inequalities at these resets, the square-root bound at resets $r \geq 10$, and the displayed band condition at upper resets $r \geq 13$, with no failures. This remains a finite ordinary computation, distinct from the cited Lean theorem for $13 \leq r \leq 30$ and from an assertion at every later reset. [Open](cofinal)

11 Unproved inputs for the recorded implications

The preceding parts prove support criteria, finite identities and conditional implications. They do not decide the universal problem or either rational target. This section states the remaining hypotheses for particular implications and compares them with the estimates already available. Ordinary arguments, reported finite computations and unproved proposals are distinguished by the source tags explained at the start of the paper. No single necessary obstruction is claimed to account for every unsuccessful estimate.

For example, the first two weighted divisor counts after $M = 3$ give $(\tau(4) - 1)/2 + (\tau(5) - 1)/4 = 5/4$. The finite weighted divisor sum below records the same calculation for any starting index and length. Its role in particular universal-support criteria is discussed separately; it is not a proved common necessary obstruction for the two problems.

Definition 11.1 (The short-window divisor phase). For $M \geq 1$ and $L \geq 1$ put

$$\Theta_L(M) := \sum_{i=1}^L (\tau(M+i) - 1) 2^{-i} \in \mathbb{Q}_{\geq 0},$$

where τ is the number-of-divisors function, so that $\tau(n) - 1 = \#\{d \geq 2 : d \mid n\}$ counts the divisors of n other than 1.

Iteration of the integer recurrence $K(M+1) = 2K(M) - (\tau(M+1) - 1)$ gives the exact identity

$$\Theta_L(M) = K(M) - 2^{-L}K(M+L) \quad (M \geq 2, L \geq 1).$$

Consequently $2^L \Theta_L(M)$ is an integer, and

$$\text{dist}(\Theta_L(M), \mathbb{Z}) = 2^{-L} \text{dist}(K(M+L), 2^L \mathbb{Z}).$$

A nonzero distance is at least 2^{-L} , but nonvanishing itself requires an argument. For instance $\Theta_1(3) = 1$ has distance zero. The identity concerns K ; the quotient deviation also contains the skip-set correction in Theorem 5.14. No positive lower bound for this phase is asserted to be necessary for the universal support problem.

11.1 Reset bounds and finite weighted divisor sums

The statement needed

The first two conditions below each have a conditional implication to half-membership. No implication between them is asserted. The sign condition is weaker than a quantitative margin in its conclusion, but its relation to either membership hypothesis must still be proved.

Form (i), the reset-deviation bound. Write $w_s = \text{rem}(s) - 2^s$. The actual hypothesis is

$$\forall r \geq 10 \text{ that are upper or middle resets : } |w_{r+1}| > 2^{(r+5)/2}.$$

Equivalently, $|\lambda_{r+1}| > 2^{(3-r)/2}$ for $\lambda_{r+1} = w_{r+1}/2^{r+1}$. This normalisation concerns distance of the deviation from zero, not distance of an unrelated phase from the nearest integer. The run-length estimate $L_r < (r-3)/2$ for resets $r \geq 31$ is a proposed comparison target. Its exact relation to the displayed bound requires the pulse terms and the initial range $10 \leq r < 31$; it is not stated as an equivalence here. A sublinear estimate gives an eventual ceiling, not automatically the required finite initial cases. The supplied computation reports a maximum run of length 19 through row 200,000, attained at row 158,096. That observation is finite evidence, not an asymptotic theorem.

Form (ii), the Lean-facing form already proved as a conditional implication.

$$\text{SeamUpperResetDyadicBandEscape} := \forall d \geq 13, (\text{seamAdjacentCut } d).\text{successorCarries} \Rightarrow \forall j \leq d,$$

$$2^{d-j+1} < E_d \quad \vee \quad E_d + 2(d+j) \leq 2^{d-j+1}, \quad E_d := 4 \cdot \text{overshoot}_d + \text{abovePulse}_d,$$

([LEAN SOURCE](#)), with the closed conditional implication [LEAN SOURCE](#), [Lean]. The $\forall j \leq d$ band check collapses to a single nearest-boundary index per row by [LEAN SOURCE](#), and rows $13 \leq d \leq 30$ are verified in the cited source by [LEAN SOURCE](#).

Form (iii), the sign of the reset deviation.

$$\forall \text{ reset rows } r : \quad r \text{ is an } M\text{-reset} \Rightarrow w_{r+1} > 0, \quad r \text{ is a } U\text{-reset} \Rightarrow w_{r+1} < 0.$$

The finite recomputation for this revision verifies these signs at all 1209 resets in rows 6–2500, with 605 middle and 604 upper resets. The upper sign follows directly from the exact branch identity below. The middle sign remains an unproved assertion about the orbit; neither sign alone gives the required lower bound on the magnitude.

Divisor residues and the changing skip set

Expanding each divisor count separates its residue conditions. This is a finite counting identity. It does not prescribe whether an eventual proof should be analytic, combinatorial, or use another method.

Lemma 11.2 (Divisor-residue form of the short-window phase). *For all $M, L \geq 1$,*

$$\Theta_L(M) = \sum_{d \geq 2} \sum_{\substack{1 \leq i \leq L \\ i \equiv -M \pmod{d}}} 2^{-i} = \sum_{d \geq 2} 2^{-i_d(M)} \cdot \frac{1 - 2^{-d m_d}}{1 - 2^{-d}},$$

where $i_d(M) \in [1, d]$ is the least $i \geq 1$ with $d \mid M + i$ (so $i_d(M)$ depends only on $M \bmod d$), $m_d := \#\{1 \leq i \leq L : d \mid M + i\}$, and terms with $i_d(M) > L$ are empty. In particular, for every integer $D \geq M + L$,

$$\Theta_L(M) = \sum_{d=2}^D \sum_{i=1}^L 2^{-i} \mathbf{1}_{d \mid M+i},$$

and the right-hand side, with an arbitrary integer $x \geq 0$ in place of M , depends on x only through its residues modulo the integers d with $2 \leq d \leq D$. The cutoff D has to be fixed before the argument varies: at $(M, L) = (1, 1)$ the cutoff is $M + L = 2$, the integers 1 and 3 have the same residue modulo every d with $2 \leq d \leq 2$, and yet $\Theta_1(1) = 1/2$ while $\Theta_1(3) = 1$.

Lean: `theta_eq_tsum_divisorResidue`, `theta_eq_tsum_geometricForm`, `theta_eq_divisorResidueSum`, `theta_eq_geometricForm`, and 13 further declarations in the [coverage section](#).

Proof. $\tau(n) - 1 = \sum_{d \geq 2} [d \mid n]$; exchange the order of summation and sum the geometric progression $i_d, i_d + d, i_d + 2d, \dots$ inside $[1, L]$. No $d > M + L$ divides any $M + i$ with $1 \leq i \leq L$, which gives the cutoff form, and each indicator $\mathbf{1}_{d \mid x+i}$ depends only on $x \bmod d$. $\square$

[Lean] ([LEAN SOURCE](#), cutoff form at [LEAN SOURCE](#)), and verified exactly as rationals at $(M, L) \in \{1000, 10001, 123456\} \times \{12, 20\}$ [Cert].

A finite divisor cutoff makes the averaging question precise. For integers $D \geq 2$ and $L \geq 1$, define

$$\Psi_{L,D}(x) = \sum_{d=2}^D \sum_{i=1}^L 2^{-i} \mathbf{1}_{d \mid x+i}, \quad Q_D = \text{lcm}(1, \dots, D).$$

This is the cutoff sum of Lemma 11.2. It is periodic in x with period Q_D , and $\Theta_L(M) = \Psi_{L,D}(M)$ whenever $M \geq 1$ and $D \geq M + L$. Averaging each indicator over a complete period gives

$$\frac{1}{Q_D} \sum_{x=0}^{Q_D-1} \Psi_{L,D}(x) = (1 - 2^{-L}) \sum_{d=2}^D \frac{1}{d}.$$

The mean diverges as $D \rightarrow \infty$, even for fixed L . Thus this family does not converge to a finite integrable observable whose limiting mean could simply be used in an equidistribution argument.

The same issue arises if one uses the profinite integers $\widehat{\mathbb{Z}}$, whose Haar probability measure is uniform on each finite residue quotient. The finite-cutoff functions are continuous there. Their increasing limit is infinite almost everywhere: the events $p \mid x + 1$ are independent over primes by the Chinese remainder theorem, have probabilities $1/p$, and occur infinitely often almost surely. Each contributes $1/2$ to the term at offset 1. The positive integer inputs M , for which $\Theta_L(M)$ is finite, do not justify treating this divergent limit as a continuous or integrable function. Only the finite period average above is used in the present argument.

There is a related residue expression in Theorem 5.15, but its index and divisor set differ:

$$\sum_{d \in D_n} \left\{ \frac{4^n}{2^d - 1} \right\} = \sum_{d \in D_n} \frac{2^{-i_d(2n)}}{1 - 2^{-d}}.$$

Indeed $i_d(2n) = d - (2n \bmod d)$, with value d when the remainder is zero. The argument is $2n$, not n , because the quotient scale is $4^n = 2^{2n}$. For example, at $n = 7, d = 3$ the fractional part is $4/7$; using $i_3(7)$ would instead give $2/7$. This sum uses the selected finite set D_n , whereas Θ_L uses all divisors with offsets truncated at L . Their common geometric summand does not make the two sums equal.

To keep the changing support visible, put $S_n = \sum_{d \in \text{Skip}_n} \lfloor 4^n / (2^d - 1) \rfloor$. The master identity $\Delta_n = K(2n) + S_n$ and the recurrence for K give, for $n \geq 6$ and $\ell \geq 1$,

$$2^{-2\ell} \Delta_{n+\ell} = \Delta_n - \Theta_{2\ell}(2n) + 2^{-2\ell} S_{n+\ell} - S_n.$$

A phase estimate would have to control this additional term and identify the required reset rows before implying the reset bound. No such implication is proved here. The scales also differ from a logarithmic-window assertion: ℓ quotient transitions use 2ℓ binary steps starting at $M = 2n$. If ℓ is proportional to n , the phase length is comparable to M , not to $\log M$. A theorem for logarithmic windows would need a separate argument connecting its range to the reset condition.

Finite residue expansions, divisor-sum estimates and probabilistic local estimates are possible sources of further information. None is an input already supplied here. An application must state the divisor cutoff, offset length, error at the required precision, and control on the changing skip-set sum, rather than appeal to the name of a method or to equidistribution of a different statistic.

The specific obstacle

What the cited short-interval theorems control. Matomäki and Radziwiłł prove that a real-valued multiplicative function bounded by 1 has, in almost all windows $[x, x + h]$ with $h \rightarrow \infty$, essentially its long average [10, Theorem 1, pp. 1–2]. Mangerel extends such bounds to a class of divisor-bounded multiplicative functions [11, Theorem 1.7, p. 7, and Corollary 1.8, pp. 7–8], and Sun describes the almost-all transition profile of the k -fold divisor function at its critical polylogarithmic window [12, Theorem 1.1 and Corollary 1.2, p. 3]. These results control short-interval averages. They do not by themselves supply a reset-deviation estimate or its unproved weighted-phase substitute, for two reasons.

First, their conclusions hold for almost all x , whereas the proposed application needs every reset row beyond its threshold. An exceptional set of density zero may still contain infinitely many required indices. No density assertion about the reset set, or avoidance of the exceptional set by that orbit, is proved here.

Second, τ is not bounded by 1, and the relevant statistic is the distance to $\mathbb{Z}$ of a *dyadically weighted* sum of τ -values at a precision shrinking exponentially in the window parameter. The first term $(\tau(M + 1) - 1)/2$ has fixed weight $1/2$, independent of the

window length; an unweighted average over that window does not by itself control this term. An application through the phase would need both an exceptional-set argument along the specified indices and control of the weighted sum. It would still have to justify the passage from that sum to the reset deviation, including the skip-set correction.

The required statistic is different. A central limit theorem for prime-factor counts, or an almost-all theorem for unweighted interval averages, is not by itself a bound for the distance to an integer of this dyadically weighted divisor sum. An application must state its precision, its dependence on the growing window length, and its exceptional set. The references cited above do not supply that orbitwise estimate in the form used here. This is a limitation of the proposed application, not a survey claim that every relevant local estimate is unavailable.

The countermodels in the record impose similarly specific limitations. They concern fixed-precision residues ([LEAN SOURCE](#)), a decoder without fresh arithmetic input ([LEAN SOURCE](#)), and fixed-precision signatures ([LEAN SOURCE](#)). They do not partition all possible proof methods. Together with the denominator estimate, they explain why the particular shortcuts examined here are insufficient; they do not prove that a weighted-window method is necessary or impossible.

Possible estimates to prove

A conditional bound along a right-branch run.

Let $r \geq 6$, $L \geq 1$, and suppose that the transitions at rows $r, \dots, r + L - 1$ all take the right branch. Put $w_s = \text{rem}(s) - 2^s$ and let p_s be the lower-prefix pulse, so $0 \leq p_s \leq 2(s - 2)$. Iterating the exact recurrence $w_{s+1} = 4w_s - 4 - p_s$ gives

$$w_{r+L} = 4^L w_r - \sum_{j=0}^{L-1} 4^{L-1-j} (4 + p_{r+j}).$$

The reverse triangle inequality and the finite geometric sum give

$$|w_{r+L}| \geq 4^L \left(|w_r| - \frac{4 + 2(r + L)}{3} \right).$$

This estimate alone does not force a right-branch run to end. For that conclusion one also needs an upper bound on its terminal deviation. More precisely, if $C > 0$ and

$$|w_{r+L}| \leq C 2^{r+L}, \quad |w_r| > \frac{4 + 2(r + L)}{3},$$

then division by 4^L and taking logarithms give

$$L \leq r + \log_2 C - \log_2 \left(|w_r| - \frac{4 + 2(r + L)}{3} \right).$$

Every term and hypothesis is explicit here. A separate window estimate must supply the displayed terminal upper bound; it does not follow from the right-branch recurrence alone. This explains the possible connection between the initial deviation and the run

length without asserting a sharp threshold, a run-length equivalence, or an unproved bound for all later rows. [Math] for this conditional calculation.

An exact finite phase experiment.

For the 200,000 integers $10^6 \leq M < 1.2 \times 10^6$, an integer divisor sieve through 1,200,039 computes $2^{40} \Theta_{40}(M)$ exactly. Reducing that integer modulo 2^{40} gives the following counts and empirical proportions for $\text{dist}(\Theta_{40}(M), \mathbb{Z}) < t$.

t	2^{-6}	2^{-8}	2^{-10}	2^{-12}	2^{-14}
Count	7083	1858	491	125	37
Proportion	0.035415	0.009290	0.002455	0.000625	0.000185
Proportion/ t	2.267	2.378	2.514	2.560	3.031

These counts were reproduced for this revision; the experiment samples consecutive integers, not reset rows. Its median distance is 0.24470815549 ..., compared with 1/4 for a uniform fractional part. A uniform fractional part has probability $2t$ of distance less than t , for $0 < t \leq 1/2$, which explains the comparison in the last row. Agreement of a finite histogram with that model is not a distribution theorem.

Truncation matters at the smallest tested thresholds. With the same starting indices, the counts below 2^{-14} are 22, 36, 37, 37 for $L = 16, 24, 32, 40$, respectively. Thus similar central quantiles do not justify ignoring the truncation error in a rare-event test. More generally, for $L' \geq L \geq 1$ and $B \geq 0$, if $\tau(M + i) - 1 \leq B$ over $L < i \leq L'$, then

$$0 \leq \Theta_{L'}(M) - \Theta_L(M) \leq B(2^{-L} - 2^{-L'}) \leq B2^{-L}.$$

Distance to the nearest integer is 1-Lipschitz. A threshold comparison is therefore guaranteed to remain unchanged when the distance at the shorter cutoff is farther than this error from the threshold. This is a sufficient stability test, not a necessary one. The estimate uses a bound over the specified finite range, not an unproved uniform bound on the divisor function.

A probabilistic application requires further assumptions. For example, if events E_r in a specified probability space satisfied $\mathbb{P}(E_r) \leq C2^{(3-r)/2}$ for every $r \geq R$, the union bound would give

$$\mathbb{P}\left(\bigcup_{r \geq R} E_r\right) \leq \frac{C2^{(3-R)/2}}{1 - 2^{-1/2}}.$$

Summability would also imply only finitely many failures almost surely, not zero failures at every index. Neither the finite histogram nor its fit to a uniform-phase model proves these probability bounds, identifies their events with reset failures, or defines a probability for the actual deterministic orbit. In particular, no membership conclusion is drawn from the fitted small expected counts reported in the earlier record. [Cert] for the finite counts; [Math] for the elementary error and probability inequalities; [Open] for an application to resets.

The reset signs and the exact branch identities. Write p_s for the nonnegative lower-prefix pulse and $w_s = \text{rem}(s) - 2^s$. On a middle transition, where the upper carry does not occur, the update is

$$\text{rem}(s+1) = 4\text{rem}(s) + 2^{s+1} - p_s, \quad w_{s+1} = 4w_s + 2^{s+2} - p_s = 4\text{rem}(s) - p_s.$$

The natural-number subtraction can be read as signed subtraction because $p_s \leq 2(s-2) < 2^{s+1}$ for $s \geq 5$. The middle-branch condition is $4\text{rem}(s) < 2^{s+1} + 4 + p_s$, together with failure of the upper carry. Thus a positive middle-reset deviation is exactly the additional inequality $4\text{rem}(s) > p_s$. The pulse bound and branch inequality alone do not prove it: $s = 5$, $\text{rem} = 1$, $p = 5$ satisfy both numerical bounds but give $4\text{rem} - p = -1$. This is a counterexample to those numerical premises, not a claim that the concrete orbit reaches that state.

On an upper transition, let $o_s > 0$ be the overshoot and $p_s^+ \geq 0$ the upper-prefix pulse. The carry condition pays the subtraction, so

$$w_{s+1} = -(4o_s + p_s^+) < 0.$$

This sign already follows from the exact upper-branch identity; only the middle sign needs further orbit information. On a right transition the corresponding formula is $w_{s+1} = 4w_s - 4 - p_s$. The same middle and upper deviation identities appear in the supplied [release source](#), classified `release_only` by the packet. They separate the sign calculation from the unproved magnitude estimate needed for square-root escape. At the reported middle step $s = 14$, for example, $\text{rem}(14) = 392$, $p_{14} = 3$ and $\text{rem}(15) = 34333$, giving $w_{15} = 1565$.

11.2 Progress to larger endpoint depths

The statement needed

The exact-row transition [LEAN SOURCE](#) ([Lean]) says every exact row $n \geq 6$ either doubles to $2n - 1$ or recycles to $2c - 2$ for some $4 \leq c \leq n$. Cofinality of exact rows; which gives HALF; needs the recycle branch not to return to a bounded endpoint forever. Three sufficient forms:

$$(a) \quad \text{SkippedCoreCriticalQuotientSupply} : \forall D \subseteq [2, c), 4 \leq c, v(D) < \frac{1}{2} \wedge \frac{1}{2} - v(D) < w c$$

$$\Rightarrow 2^{(2c-2)-1} \leq Q(\{c\} \cup D) (2c - 2),$$

where $v(D) := V D$ ([LEAN SOURCE](#));

(b) the crossing rank c produced by the recycle branch does not repeat a bounded value infinitely often;

$$(c) \quad c \geq g(n), \quad 2g(n) - 2 > n \quad \text{at every reachable recycle input } n.$$

Unboundedness of g alone would not suffice in (c). For example, $g(n) = \lfloor n/2 \rfloor + 1$ tends to infinity but allows the fixed transition $n = 6$, $c = 4$, $2c - 2 = 6$. The displayed strict-progress inequality excludes such cycles; the doubling branch already increases every endpoint $n \geq 6$.

What kind of object would furnish it

The real and integer greedy rules must be distinguished here. A real straddling word agrees with the real greedy prefix ([LEAN SOURCE](#)). An exact local row instead satisfies $QDn = 2^{n-1} - 1$. Although floor division is not injective in general, the quotient weights at fixed n have a gap of at least 1 between distinct subset sums. A witness is therefore unique when it exists and is recovered by the integer greedy test (Theorem 6.100 and Propositions 6.101–6.102). There is no exponential family of witnesses to choose from at that depth. What remains to prove is that successful depths are unbounded.

One possible additional ingredient is an integer-valued function $\Phi(n, D) \geq 0$ on reachable exact rows that increases by at least one under each permitted transition and is bounded on rows of bounded endpoint. Along an infinite trajectory, Φ would then tend to infinity, forcing the endpoints to be unbounded. Strict increase only on recycle steps would need a separate check of what doubling does to Φ ; an unspecified real-valued increasing quantity need not tend to infinity. This describes one possible proof mechanism, not a necessary class of all proofs. The corpus already has the sharp capacity test in exact form: by [LEAN SOURCE](#), form (a) is exactly the statement that the $c - 2$ -bit suffix capacity $SD1(2c - 2) < 2^{c-2}$ is met.

The specific obstacle

The formal counterexample concerns exactly the transition schema: that schema *plus* a seed provably does not give cofinality. The theorem [LEAN SOURCE](#) ([Lean]) exhibits the one-point predicate $P(n) := (n = 6)$, which satisfies the seed ([LEAN SOURCE](#), the explicit witness $\{2, 3, 6\}$) and the exact shape of the transition, and is not cofinal. So no argument that uses only the dichotomy and the base case can work, however cleverly the induction is arranged. Any proof must consume arithmetic content of `ExactLocalMersenneHalfRow` that the abstract schema does not see.

There is a second, more specific trap, and it is also formalised. The natural sufficient condition via fractional split mass is *not necessary*: the fixture [LEAN SOURCE](#) ([Lean], at $D = \{2, 3\}$, $c = 5$) exhibits a real crossing core violating it. Thus the fractional-mass bound is not necessary for an individual crossing core. This counterexample does not rule out proving that sufficient bound on a more restricted family.

Possible estimates to prove

The uniquely determined prefix at a crossing. The identification below is proved; the required capacity bound remains open.

Form (a) is universally quantified over cores D , which looks like an exponential obligation. It is not. Its hypotheses are exactly those of [LEAN SOURCE](#) ([Lean]), which forces $D = \text{halfGreedyPrefixSupport}(c - 1)$. So the $\forall D$ collapses to a single orbit and form (a) is a statement to be checked along *one* sequence indexed by c , namely

$\forall c \geq 4$ skipped by the real half-greedy orbit : $S(\text{halfGreedyPrefixSupport}(c - 1))1(2c - 2) < 2^{c-2}$.

The crossing hypothesis has not disappeared: these are the ranks where the real greedy rule skips the next term. Each resulting finite inequality is decidable. Comput-

ing its margin can identify examples to explain a proposed estimate, but a finite table neither selects a proof method nor establishes a cofinal supply.

The progress-function criterion above remains unproved. Endpoint doubling is not a proved doubling law for the largest selected exponent, and an unbounded function of the incoming endpoint alone does not rule out a bounded cycle.

11.3 Prime support and the limits of one formal argument

The base-2 statement

$$\text{Irrational}\left(\sum_{p \text{ prime}} \frac{1}{2^p - 1}\right)$$

is known. Expanding geometrically and collecting divisors gives

$$\sum_{p \text{ prime}} \frac{1}{2^p - 1} = \sum_{n \geq 1} \frac{\omega(n)}{2^n},$$

and Tao–Teräväinen prove the latter series irrational in Theorem 1.3 of *Quantitative correlations and some problems on prime factors of consecutive integers*. This is a cited analytic result, not a theorem formalised in the pinned Lean corpus.

The existing pairwise-coprime Lean theorem does not recover it because that theorem assumes $\sum_{a \in A} 1/a < \infty$, whereas $\sum_p 1/p$ diverges. Its generic block-certificate engine would also demand, in the prime case,

$$2^r \mid \omega(N + r) \quad (1 \leq r \leq K),$$

a simultaneous exact divisibility pattern far stronger than the input used by Tao–Teräväinen. Small computational witnesses for $K = 2, 3$ diagnose that the condition is not vacuous, but they neither prove the known theorem nor define a new open case.

Thus the contribution of this route audit is negative and coordinate-specific: it identifies why the present formal certificate engine fails to reproduce a theorem already available by quantitative correlation estimates. The earlier proposal to prove the prime-support irrationality as new mathematics is withdrawn. Extensions to other bases or prime-power supports are mentioned by the cited authors with proof modifications omitted; this paper promotes only the fully written base-2 theorem.

11.4 Equivalent carry conditions at perfect-square depths

The statement needed

The following five conditions are equivalent to $1/2 \in A$:

- (1) $C_G(N) \leq 2\sqrt{N} + 4$ for every $N \geq 0$.
- (2) The set $\mathbb{N}_{>0} \setminus G$ is infinite.
- (3) For arbitrarily large n there is $D \subseteq \{2, \dots, n\}$ with $Q(D, n) = 2^{n-1} - 1$.
- (4) For arbitrarily large M , $\text{ihc}(G, M) \leq B(M + 1)$.

(5) For arbitrarily large $M \geq 1$ there is $D \subseteq \{2, \dots, M\}$ with $|\text{ihc}(D, M-1)| \leq B(M)$.

Here G is the actual greedy support for $1/2$. The forward directions of (2)–(5) are recorded Lean ([LEAN SOURCE](#), [LEAN SOURCE](#), [LEAN SOURCE](#), [LEAN SOURCE](#), all [Lean]). The reverse implications are ordinary deductions in this paper. The perfect-square argument below supplies the terminal bounds without changing the constant in the original statements.

The carry bound at perfect-square depths

Lemma 11.3 (The terminal bound at square depths). *Let $A \subseteq \mathbb{N}$ with $1 \notin A$ and $X_A(2) = 1/2$. Then for every $k \geq 1$,*

$$(\text{ihc } A (k^2 - 1) : \mathbb{R}) = \text{binaryCoeffTail } (c_A) (k^2) \leq 2k + 4 = B(k^2).$$

Lean: [paper_square_depth_terminal_bound](#), [square_depth_witness](#).

Proof. By [LEAN SOURCE](#) ([Lean]),

$$(\text{ihc } A N : \mathbb{R}) = 2^{N+1} \left(\frac{1}{2} - X_A(2) \right) + \text{binaryCoeffTail } (c_A) (N+1),$$

and the first term vanishes by hypothesis. Take $N = k^2 - 1$, so $N + 1 = k^2$. By [LEAN SOURCE](#) ([Lean], unconditional) the tail at k^2 is at most $2\sqrt{k^2} + 4 = 2k + 4$, and by [LEAN SOURCE](#), $B n = 2\text{Nat.sqrt } n + 4$, which at $n = k^2$ equals $2k + 4$ exactly. Non-negativity of the tail ([LEAN SOURCE](#), [Lean]) gives the lower side. $\square$

[Math]; every input is supplied formal and the composition is four lines.

The former square-root adapter issue is resolved by restricting cofinal witness depths to perfect squares: at $n = k^2$ the two roots coincide exactly and no widening is needed. Applying Lemma 11.3 along $M = k^2 - 1$ gives (4) directly; taking a to be the indicator of an achieving support truncated to $\text{Fin}(M+1)$ at $M = k^2$, and using [LEAN SOURCE](#) together with $1 \notin A$ (forced because $w 1 = 1 > 1/2$), gives the witness required by [LEAN SOURCE](#) and hence (5).

What the equivalences establish

Each biconditional is a theorem about the exact logical content of a proposed hypothesis. It is also a reformulation of the membership question. These are compatible descriptions, not competing evidence classes. The equivalence identifies the missing content but does not show that the reformulation cannot expose a useful invariant. In particular, (3) shows that the existence of finite witnesses at unbounded depths, even without compatibility between those witnesses, is equivalent to membership. The absence of a compatibility requirement may still be useful in constructing witnesses; equivalence alone says nothing about which formulation is easier to prove.

Scope of the equivalences. (1) does *not* extend to every A with $1 \notin A$: without [LEAN SOURCE](#) (which supplies $\delta := 1/2 - S_G \geq 0$ for the greedy support specifically) one obtains only the one-sided equivalence $\text{hbound} \iff (1/2 - S_A) \leq 0$. The mechanism

in all five cases is the same and should be recorded once: $2^{N+1}\delta$ dominates any $\sqrt{N}$ envelope, so a strip hypothesis forces $\delta = 0$; and $\delta = 0$ makes the unconditional tail bound supply the strip. Membership is then read off by [LEAN SOURCE](#) ([Lean]).

11.5 Sparse and dense supports under the certificate criterion

The statement needed

$A \subseteq \mathbb{N}_{>0}, \quad A \text{ infinite} \quad \implies$
 c_A has arbitrarily late super-logarithmic zero windows
 or the specified block certificates exist at every precision.

For a fixed nonempty A , let $a_0 = \min A$. Every interval of a_0 consecutive integers contains a multiple of a_0 , so c_A has no zero window longer than $a_0 - 1$. The first alternative is therefore impossible, independently of rationality. The cited logarithmic result [LEAN SOURCE](#) does not create a sparse subclass here. The separate lcm-gap theorem [LEAN SOURCE](#) remains a valid sufficient criterion, but does not prove this proposed alternative. The proposal would consequently require the following certificates for every infinite support:

$$\exists N, K, L, C, \quad K \leq L, \quad \forall r \in [1, K] : 2^r \mid c_A(N + r),$$

$$\sum_{r=K+1}^L c_A(N + r) 2^{L-r} \leq C, \quad q(C + N + L + 2) < 2^L.$$

What kind of object would furnish it

The identity $c_A = \mathbf{1}_A * \mathbf{1}$ is a Dirichlet convolution, and Möbius inversion gives $\mathbf{1}_A = \mu * c_A$. These identities hold for every support, but do not supply the displayed divisibility or tail estimates. Shared factors restrict which residue prescriptions can be realised; they do not make the Chinese remainder theorem inapplicable. For positive moduli m_i , a finite system $N \equiv a_i \pmod{m_i}$ is solvable precisely when

$$a_i \equiv a_j \pmod{\gcd(m_i, m_j)} \quad \text{for every } i, j.$$

Necessity follows by subtraction. For sufficiency, keep the strongest prime-power modulus for each prime; compatibility makes all weaker prescriptions agree with it, and the ordinary CRT combines the resulting coprime prime powers [23, Theorem 3.1, pp. 2–3]. For example, $N \equiv 1 \pmod{4}$ and $N \equiv 1 \pmod{6}$ give $N \equiv 1 \pmod{12}$, whereas $4 \mid N$ and $2 \nmid N$ are incompatible. A block-certificate construction must satisfy these compatibility conditions and its separate tail budget. Neither Möbius inversion nor CRT alone provides both inputs.

The specific obstacle, proved

Here the honest answer is stronger than an obstacle: the dense branch as stated is *false*, and the falsifier is an explicit, natural, divisor-dense support. This also corrects the reduction recorded earlier in this document, that universal #257 passes “with no loss” to $O1' := \forall A \text{ infinite, Cert}(A)$.

Proposition 11.4 (Squarefree support: exact engine ceiling, not an open value). *Let $A = \{n \geq 2 : n \text{ is squarefree}\}$. Then*

$$c_A(n) = 2^{\omega(n)} - 1,$$

which is odd for every $n \geq 2$. Consequently neither the digitwise nor the carry-aware divisibility-first block-certificate schema has an instance at any even base. The two exact no-go theorems are [LEAN SOURCE](#) and [LEAN SOURCE](#).

Lean: [paper_squarefree_support_engine_ceiling](#).

[Lean] (uniform) *binary digits*.

This does not leave the squarefree value open. Duverney–Tachiya’s Corollary 1.2 and Example 1.1 prove that, for every $h \geq 1$,

$$1, \sum_{n \text{ squarefree}} \frac{1}{2^n - 1}, \dots, \sum_{n \text{ squarefree}} \frac{1}{2^{hn} - 1}$$

are linearly independent over $\mathbb{Q}$. After deleting the rational $n = 1$ term, squarefree support is therefore irrational at every base 2^j . Among the integer bases, the cited theorem leaves the non-powers of two uncovered. This is cited mathematics, not formalised here. Proposition 11.4 is a failure of two proof schemas on a value already known to be irrational.

The coordinate dependence is exact. Adjoin 1 to the support. The series changes by the rational number $(b - 1)^{-1}$, so its irrationality is unchanged, while the incidence becomes $2^{\omega(n)}$: [LEAN SOURCE](#) and [LEAN SOURCE](#). At base 2, CRT then supplies arbitrarily long opening blocks with $\omega(N + r) \geq r$: [LEAN SOURCE](#) and [LEAN SOURCE](#). The middle-window and height conditions remain; no certificate or new irrationality proof follows. The durable conclusion is methodological: test a no-go statement under rational finite changes before attributing it to the underlying value.

11.6 What the experiments and obstructions do not imply

Everything in this subsection is inference from the accumulated failures, not theorem. It is labelled that way throughout, and where the evidence does not determine the answer it says so rather than manufacturing a story.

Finite observations and their limits

The recorded obstructions concern specified coordinate systems. Fixed precision, a bounded carry state and a prescribed suffix can lose information needed by the actual-selector condition. They do not show that the underlying deterministic sequence has no

arithmetic structure. The relevant exact statements remain [LEAN SOURCE](#), [LEAN SOURCE](#) and [LEAN SOURCE](#). Their hypotheses, rather than the fact that an attempted argument failed, determine their scope.

The inherited experiment log reports 200,000 seam rows with counts $R : M : U = 100197 : 49899 : 49898$, a roughly geometric run-length histogram, and a longest equal run of 15 in 40,000 binary digits. It also reports agreement with OEIS A211706 and a near-integer phase constant of about 2.5. These computations and the external digit comparison were not rerun in this revision. They are descriptive reports, not certificates of normality, independence, a random sampling model, or an infinite anti-concentration estimate. A statement about “sampling noise” would require a specified probabilistic model that is not supplied here.

The previously reported exceptional row $r = 7$ and crossing $(s, d) = (10, 7)$ are likewise statements about the tested range only. The exact range and reproduction artefacts must accompany them before they are used as numerical evidence outside this historical record. No claim that this is the only exceptional row of the infinite orbit is made.

These coordinates all descend from the same greedy-carry representation; they are not independent tests of a common random hypothesis. Strict strict tail domination gives an injective coding, but says neither that its digits are random nor that averaging over indices is impossible. The measure-one and nowhere-dense properties of the achievement set do not decide membership of a specified rational. The correct research question is whether a precise arithmetic estimate can be proved for the actual selector, not whether finite data resemble a random model.

Why the support and rational-target problems remain distinct

Proposition 11.4 is a reason to separate the two halves rather than assimilate them. On the universal side the present certificate engine first meets a coordinate-dependent parity obstruction. The squarefree value is nevertheless already known to be irrational at power-of-two bases, and adjoining 1 removes the opening-block obstruction without changing irrationality. Thus the no-go diagnoses this engine; it does not diagnose the value or the universal problem.

Short-window near-integer anti-concentration is one possible missing input for the particular mechanisms analysed here. No equivalence shows that it is necessary for either problem, and no argument identifies it as a common necessary obstruction to both. The squarefree example shows why finite rational changes of support must be checked before an opening-block failure is interpreted as a property of the value.

What the evidence does and does not determine

It does not determine whether a short proof exists. Nothing in a catalogue of failed methods bounds the length of a successful one, and this document should not pretend otherwise.

It records limitations of specified methods. The fixed-precision completion theorem and the balanced-pulse countermodel concern the particular summaries defined above, not all fixed-precision arithmetic or all finite-state methods with additional arithmetic

input. An irrationality-measure bound on E does not classify a specified rational in A . Instantiating the present block-certificate engine at a general support is blocked, on the universal side, by Proposition 11.4. Survival through any tested finite depth does not by itself establish membership; a fatal-gap certificate would establish non-membership. These facts do not decide rational-target membership, and they do not yield an impossibility theorem for all probabilistic, measure-theoretic or averaging methods. Each recorded obstruction is reusable in a way that a failed attempt is not. That is the contribution: not a proof, not a reduction, and not a map of every region in which a proof cannot live.

V. Outstanding obligations and provenance

12 What is open, stated exactly

Erdős #257 asks whether $\sum_{n \in A} 1/(2^n - 1)$ is irrational for *every* infinite $A \subseteq \mathbb{N}_{\geq 1}$. Nothing in this paper decides this, in either the universal form or the single-target half-value form below. Parts I–IV do prove unconditional theorems for specified support classes, as well as uniform finite-denominator and achievement-set results. The unresolved membership conditions must be distinguished from those results: some are equivalences, some are unproved sufficient conditions, and some have only finite computational checks. This section states the exact open targets, in display mathematics, with their Lean sites where formalised and their precise logical relationships to one another. **Read the implication table in §12.8 before attacking any single target below:** several apparent waypoints are *equivalent* to the full problem, not easier approaches to it.

12.1 Notation fixed for this section

The functional `erdosSupportSeries 2 (-)` denotes the base-2 Erdős support series $X_A(2) := \sum_{n \in A} 1/(2^n - 1)$ ([LEAN SOURCE](#)). $c_A(n) := |\{d \mid n : d \in A\}|$ ([LEAN SOURCE](#)) is the divisor-indicator coefficient whose generating tail feeds every certificate below. $A := \mathcal{A} = \{x \mid \exists A, 0 \notin A \wedge x = \text{positiveMersenneSupportValue } A\}$ ([LEAN SOURCE](#)). $G := \text{greedyMersenneSupport}(1/2)$ is the canonical greedy support for target $1/2$; $\text{rem}(s)$ is its integer seam remainder at row s , and a *reset* at row r is a row where the greedy orbit returns to the near- 2^r regime after a run of takes or skips.

12.2 The universal statement

Definition 12.1 (Universal #257).

$$\text{U257} := \forall A \subseteq \mathbb{N}_{\geq 1}, A.\text{Infinite} \implies \text{Irrational}(X_A(2)).$$

This is the literal statement of Erdős #257. [Open]. The corpus’s own sufficient-condition engine is

$$\text{Cert}(A) := \forall q > 0, \exists N, K, L, C, K \leq L, (\forall r \in [1, K], 2^r \mid c_A(N + r)) \wedge$$

$$\sum_{r=K+1}^L c_A(N + r) \cdot 2^{L-r} \leq C \wedge \exists t, 0 < c_A(N + L + 1 + t) \wedge q(C + N + L + 2) < 2^L,$$

with $\text{Cert}(A) \implies \text{Irrational}(X_A(2))$ already proved unconditionally for *every* A ([LEAN SOURCE](#), [Lean], (n/a); the theorem is universally quantified over A ; only its hypothesis is unsupplied). Therefore the stronger statement

$$\text{O1}' := \forall A, A.\text{Infinite} \implies \text{Cert}(A).$$

would imply U257, but the reduction is *not* lossless and O1' is actually false. For $A = \{n \geq 2 : n \text{ squarefree}\}$ the incidence is $2^{\omega(n)} - 1$, hence odd for every $n \geq 2$, and the

digitwise certificate schema has no instance at any even base ([LEAN SOURCE](#), [Lean]). This is a method ceiling, not a counterexample to irrationality: Duverney–Tachiya prove the squarefree value irrational at every power-of-two base. Likewise the pairwise-coprime required input does not cover the primes because $\sum_p 1/p$ diverges, even though Tao–Teräväinen prove the base-2 prime-support value irrational. The honest universal target is U257 itself; O1' is a rejected sufficient strengthening, not an open equivalent normal form.

12.3 Whether one half is represented

Definition 12.2 (Half membership).

$$\text{HALF} := \left(\frac{1}{2} : \mathbb{R} \right) \in A.$$

HALF is a *single instance* of the negation of U257: any witness A with $0 \notin A$, $X_A(2) = 1/2$ is infinite ([LEAN SOURCE](#), [Lean], (uniform) ; finite supports have odd-denominator value, so can never equal a dyadic rational), hence a genuine counterexample to U257. Thus $\text{HALF} \Rightarrow \neg \text{U257}$. The reverse implication is not established: failure of half-membership would not decide U257 if false, since U257 could still fail at a different target $t \in \mathbb{Q}$. HALF is [Open]. It is logically equivalent, by an unconditional Lean iff, to a purely combinatorial statement about one fixed orbit:

$$\text{HALF} \iff (\text{greedyMersenneSkippedSupport}(1/2)).\text{Infinite}$$

([LEAN SOURCE](#), [Lean], (uniform), coordinate *divisor counts and finite sums*, target-generic in the reverse direction; the forward direction uses only that the target is rational, so the same theorem holds verbatim for every rational $t \geq 0$ in place of $1/2$).

12.3.1 Conditions equivalent to half-membership

The following three conditions are equivalent to half-membership. They express the same unresolved assertion in different terms. The proofs of equivalence identify the content still to be established; they do not compare the difficulty of possible proofs.

Proposition 12.3 (Infinitely many positive skips are equivalent to half-membership).
Define

$$\text{CPGS} := \forall N, \exists c \geq N, c \text{ skipped by the rational half-greedy orbit} \wedge 0 < r(1/2)(c - 1)$$

([LEAN SOURCE](#)). The positivity conjunct is unconditionally true for every skipped c ([LEAN SOURCE](#), [Lean], (uniform); an odd-denominator parity fact), hence deletable without changing the Prop's meaning. What survives, $\forall N \exists c \geq N$ skipped, is exactly $(\text{greedyMersenneSkippedSupport}(1/2)).\text{Infinite}$, so

$$\text{CPGS} \iff \text{HALF}.$$

Lean: [paper_cpgs_equiv](#).

[Math]. (cofinal). *divisor counts and finite sums, greedy orbit*. The cited implication [LEAN SOURCE](#) therefore gives an equivalent formulation, not a proof of membership.

Proposition 12.4 (Terminal carry bounds are equivalent to half-membership). *Consider finite sets $D \subseteq \{2, \dots, M\}$ at arbitrarily large depths M with*

$$|\text{ihc}(D, M-1)| \leq 2\lfloor \sqrt{M} \rfloor + 4.$$

This is the exact terminal-strip condition in [LEAN SOURCE](#). It is equivalent to $1/2 \in A$: the forward direction follows from the displayed finite-approximation estimate; for the converse, truncate an achieving support at $M = k^2$ and apply [Lemma 2.6](#). Thus constant 4 already suffices at cofinally many depths. The relaxed constant 6 gives a bound at every depth, using $\sqrt{M} \leq \lfloor \sqrt{M} \rfloor + 1$, but that relaxation is not needed for the cofinal statement.

Lean: [paper_terminal_strip_equiv](#), [paper_relaxed_constant_six_every_depth](#).

[Math]. (cofinal). *support rigidity, half-carry recurrence.* This corrects the claim that the existence condition is logically weaker. An equivalent formulation may still be useful for finding a proof.

Observation 12.5. Propositions [12.3](#) and [12.4](#) do not decide half-membership. They show that the proposed existence conditions are equivalent to it, not that those formulations have no possible use. The upper bound on the centred carry in [LEAN SOURCE](#) has the same equivalence for the greedy support $A = G$. For arbitrary A with $1 \notin A$, an upper bound alone yields only $x_A \geq 1/2$, as the identity $\text{ihc}(A, N) = 2^{N+1}(1/2 - x_A) + T(N+1)$ shows. The greedy inequality $x_G \leq 1/2$ is what turns that conclusion into equality.

12.4 Exact finite sums at unbounded depths

Definition 12.6. For every integer N , require some $n \geq \max\{N, 1\}$ and $D \subseteq \{2, \dots, n\}$ with

$$\sum_{d \in D} \left\lfloor \frac{2^n}{2^d - 1} \right\rfloor = 2^{n-1} - 1.$$

([LEAN SOURCE](#), [Open], (cofinal)). Supports at different depths need not agree. This is the condition in [Definition 6.66](#); no strict weakening of half-membership is asserted.

$\text{CofinalExactLocalMersenneHalfRows} \Rightarrow \text{HALF}$ is proved unconditionally via compactness of A (continuous image of $(\mathbb{N} \rightarrow \text{Fin } 2)$ under the digit-coding map, hence closed) ([LEAN SOURCE](#), [Lean]). The cited theorem already proves that a representing support exists, because that is what membership of A means. It need not choose a support by a specified algorithm from the row witnesses. An ordinary diagonal subsequence argument gives a compatible limiting support: successively fix each coordinate along an infinite subsequence, then use the uniform tail bound to pass to its value. This uses continuity of the digit map, also recorded in [LEAN SOURCE](#). The diagonal argument preserves $2 \in A$, since all sufficiently deep exact rows contain exponent two ([LEAN SOURCE](#)). This describes an ordinary existence proof ([Math]); no additional formal extraction theorem or effective algorithm from purely existential row data is claimed.

[Proposition 2.7\(c\)](#) already identifies $\text{CofinalExactLocalMersenneHalfRows}$ as logically equivalent to $1/2 \in A$. Equivalence of a membership condition, constructive extraction of a coherent support, and a rejected sufficient strengthening are not interchangeable descriptions. Absence of a Lean extraction theorem does not make O3 a

strictly weaker substitute for HALF. The remaining work on O3 is to supply the cofinal exact rows, or any equivalent membership condition.

12.4.1 A capacity condition for exact sums at unbounded depths

The finite construction uses a $(c - 2)$ -bit capacity test. Write

$$S D k M := 2^{M-k} - Q D M - 1$$

([LEAN SOURCE](#)), with natural-number subtraction understood as truncated at zero. The loose bound below requires a critical crossing: $c \geq 4$, $D \subseteq \{2, \dots, c - 1\}$, and $0 < 1/2 - V(D) < w_c$. A merely below-half core is not enough: $D = \emptyset$ gives $S(D, 1, 2c - 2) = 2^{2c-3} - 1$. For a critical-crossing core the two bounds are

$$(\text{loose, proved}) \quad S D 1 (2c - 2) < 2^{c-1} \quad (*)$$

$$(\text{sharp, needed}) \quad S D 1 (2c - 2) < 2^{c-2} \quad (**)$$

The implication $(*)$ from the critical-crossing hypotheses is proved uniformly in c ([LEAN SOURCE](#), [Lean], (uniform)).

$(**)$ is exactly one bit sharper and, via [LEAN SOURCE](#) (under the stated core hypotheses), is equivalent to $Q(D \cup \{c\}, 2c - 2) \geq 2^{2c-3}$. Here c is the added exponent; the quotient is evaluated at depth $2c - 2$, not at depth c . The proof of $(*)$ gives the more informative additive bound $S < 2^{c-2} + |D|$ before using $|D| \leq c - 2$. One possible cofinal supply is therefore the following exclusion of a specified integer interval:

$$\begin{aligned} (\text{O3-supply}) \quad & \forall N \exists c \geq \max(N, 4) \exists D \subseteq \{2, \dots, c - 1\}, \\ & 0 < 1/2 - V(D) < w_c, \quad S(D, 1, 2c - 2) \notin [2^{c-2}, 2^{c-2} + c - 3]. \end{aligned}$$

The excluded band contains $c - 2$ integers. The sharp capacity bound itself allows 2^{c-2} nonnegative integer values. The numerical widths differ from those in Section 12.5, but the conditions concern different quantities. Width alone does not prove a logical implication between them.

The sharp-fill implication [LEAN SOURCE](#) ([Lean]) assumes $(**)$ for a below-half core $D \subseteq \{2, \dots, c - 1\}$, but does not require a critical crossing or identify that core with a real greedy prefix. This gives another formulation in which to seek the required finite rows. It does not establish a strict logical weakening of half-membership, nor does it supply cores satisfying the bound at unbounded ranks. Moreover, the integer gap separation fixes the resulting exact row whenever it exists; the formulation does not create multiple exact rows at a fixed depth.

The corpus's complete inductive skeleton from the endpoint-six seed to O3 is already assembled ([LEAN SOURCE](#), [Lean]); it consumes the sharp capacity input, in its reduced canonical form

$$c \geq 4, c \notin G \implies Q((G \cap \{2, \dots, c - 1\}) \cup \{c\}, 2c - 2) \geq 2^{2c-3}.$$

([LEAN SOURCE](#)), at essentially every step ([Open], (cofinal)). This condition is required along the specified greedy prefixes. That restriction on the witnesses does not by itself prove a strict logical ordering with the cofinal condition above.

12.5 Reset bounds and runs of right transitions

Definition 12.7 (Reset $\sqrt{\text{-escape}}$ from row 10).

$$\text{SQRTESC} := \forall r \geq 10 \text{ that are resets, } |\text{rem}(r+1) - 2^{r+1}| > 2^{(r+5)/2} \quad (\text{equivalently } \Delta_{r+1}^2 > 2^{r+5}).$$

The conditional implication uses the integer bound $B(r) = 2^{\lfloor (r+4)/2 \rfloor} + 2r+3$ described in Theorem A: $B(r)^2 \leq 2^{r+5}$ for $r \geq 10$, while a right branch at the critical crossing would force $|\Delta_{r+1}| \leq B(r)$. Thus the actual implication chain is

$$\text{SQRTESC} \Rightarrow \text{LargestSkipLateStepSocket} \Rightarrow \text{HALF}.$$

The last implication is [LEAN SOURCE](#); the supplied release proof of the first is identified in Theorem A. These are sufficient implications, not an equivalence with half-membership.

For comparison one may ask for

$$\text{RUNBOUND} : \quad \forall \text{ reset } r \geq 31, \quad L_r < (r-3)/2,$$

where L_r counts the consecutive right branches after reset r . Along L such steps the exact relation is

$$w_{r+1+L} = 4^L w_{r+1} - \sum_{j=0}^{L-1} 4^{L-1-j} (4 + p_{r+1+j}).$$

The pulse sum and the threshold ranges cannot be dropped to claim $\text{RUNBOUND} \Leftrightarrow \text{SQRTESC}$. The pointwise run/deviation equivalence stated in earlier source notes is not established by the bound used in the conditional proof. The observation at $r = 7$ (deviation 9, square-root threshold 64) is outside both proposed late ranges; that fact alone does not justify replacing the initial threshold 10 by 31.

Certified evidence, not a proof. Exact integer arithmetic to row 200,000: TARGET holds at every reset row except $r = 7$; the maximum observed R-run is 19, at row 158,096, against a required threshold there of $\approx 79,046$; a margin of $\sim 4000\times$ ([Cert], (fixed)). This is a finite verified list, *not* a cofinal supply, and must never be read as one: it is the exact instance of the difference between a finite and an unbounded parameter range this programme's own doctrine warns against.

Observation 12.8 (The sign of the reset deviation). At all 1209 observed resets, the sign of the deviation is perfectly predicted by branch type ($M \Rightarrow \text{dev} > 0$, $U \Rightarrow \text{dev} < 0$, zero exceptions), and the margin over the threshold grows from 1.11 bits at $r = 14$ to ~ 1246 bits by row 2500 ([Cert]). Proving the sign law reduces to the middle inequality $4\text{rem}(r) > p_r$. For an upper reset the sign already follows from $w_{r+1} = -(4o_r + p_r^+) < 0$, with positive overshoot and nonnegative pulse. The middle inequality remains [Open]. Even after both signs are known, a separate lower bound for their magnitudes is required for SQRTESC.

Observation 12.9 (Dangerous-reset rigidity; argument B). If reset r is *dangerous* ($|\text{rem}(r+1) - 2^{r+1}| \leq 2^{(r+5)/2}$) and the preceding reset r_0 has a pure R-run of length $L' = r - r_0 - 1 > (r+5)/4$, then w_{r_0+1} is confined to at most two adjacent integers, determined by the

divisor-pulse stream alone. A chain of n consecutive dangerous resets with long runs forces $n - 1$ nested exact congruence towers. This shows the failure mode is a tower of exact integer coincidences, not drift; it is [Math], and it does *not* by itself exclude the two pinned values; an isolated dangerous reset after a short run is untouched by this theorem.

12.6 The remaining $-2, -1$ middle-cell exclusion

Definition 12.10. `SeamTwoSidedDyadicCellEscape` excludes three integer values of the middle coordinate, $4 \cdot \text{rem}(s) - \text{belowPulse}(s) - 4 \notin \{-3, -2, -1\}$, plus one right-pulse-leak bound $4 \cdot \text{overshoot}(s) + \text{abovePulse}(s) \leq 2^{s+2}$ under $\text{overshoot}(s) \leq 2^s$ ([LEAN SOURCE](#)).

Granted this, induction from a decide-checked row-5 base propagates $\forall s \geq 5, \text{rem}(s) \leq 2^s \vee \text{overshoot}(s) \leq 2^s$ ([LEAN SOURCE](#), [Lean] given the hypothesis, (uniform in the consequence)). One of the three cells, -3 , is already closed unconditionally in the all-right-tail configuration ([LEAN SOURCE](#), [Lean]). The two remaining cells, -2 and -1 , are [Open]: what would close O5 is their exclusion for the actual quotient-greedy orbit (not a hypothetical one), together with the right-pulse-leak inequality. This excludes specified integer values in one coordinate. The square-root reset condition instead bounds the magnitude of a different quantity by a rank-dependent threshold. Their numerical widths do not order the two hypotheses, and no implication between them is established here.

A distinct sufficient condition concerns upper resets only. For every $d \geq 13$ at which the upper carry occurs, set $E_d = 4\text{overshoot}(d) + \text{abovePulse}(d)$ and require, for every $0 \leq j \leq d$,

$$2^{d-j+1} < E_d \quad \text{or} \quad E_d + 2(d+j) \leq 2^{d-j+1}.$$

This is the condition named `SeamUpperResetDyadicBandEscape` ([LEAN SOURCE](#), [Open], proved [Cert] for rows $13 \leq d \leq 30$ by the cited finite computation ([LEAN SOURCE](#))). This band has linear width $O(d)$, whereas the other bound uses a width of order $2^{d/2}$. The quantities being bounded are different, so this comparison of widths does not prove a logical implication or strict weakening. The linear-width condition suffices for half-membership by [LEAN SOURCE](#). It is proved logically equivalent to a single-critical-index check per row ([LEAN SOURCE](#), [Lean]), which is a genuine quantifier-complexity reduction, not a strength reduction.

12.7 Whether one twenty-first is represented

For the second rational target, the recorded equivalence is

$$\frac{1}{21} \in \mathcal{A} \iff \neg \text{TwentyOneFatalAlignedBranch}.$$

[LEAN SOURCE](#) [Lean]. The fatal branch means a finite fatal greedy gap followed by cofinite exact selection, eventual quotient/rational alignment, and eventual visits to every doubling block. If it occurs, the canonical quotient remainder is eventually strictly

larger than 2^R and follows one explicit affine recurrence: [LEAN SOURCE](#) and [LEAN SOURCE](#). Any unbounded sequence of closed returns instead proves membership.

No finite support on ranks at least 2 sums to $1/21$ ([LEAN SOURCE](#)). Thus membership would already be an infinite-support rational counterexample to U257. The missing step is to exclude the fatal aligned affine-supercapacity branch (or force unbounded closed returns); neither is proved. [Open]

12.8 Implication table

Read left to right: $\Rightarrow$ records a proved implication with the hypotheses stated in the relevant section, and $\Leftrightarrow$ records an equivalence. Source proofs and ordinary arguments retain the evidence distinctions explained above; the table is not a fresh kernel replay. “No implication stated” means that no comparison is established here. Quantitative bounds are not ranked solely by their widths.

Target	Relation to its endpoint	Status
Universal irrationality	Would exclude both rational targets; the stronger universal certificate condition is false	[Open]
Half-membership	Would give a counterexample to universal irrationality	[Open]
1/21 membership	Equivalent to excluding the specified failure branch; would give a counterexample	[Open]
Infinitely many greedy skips (Proposition 12.3)	$\Leftrightarrow$ O2	[Open], equivalent condition
Terminal-only strip (Prop. 12.4)	$\Leftrightarrow$ O2	[Open], equivalent condition
The centred square-root bound for G	$\Leftrightarrow$ O2	[Open], equivalent condition
Exact finite sums at unbounded depths	$\Leftrightarrow$ O2 (Prop. 2.7(c)); Lean lands $\Rightarrow$	[Open], equivalent condition
(O3-supply) linear-width capacity band	$\Rightarrow$ O3 (Lean, given lemma)	[Open]
The quotient condition on skipped greedy prefixes	$\Rightarrow$ O3 (Lean)	[Open]; no strict comparison asserted
Reset square-root bound	$\Rightarrow$ O2 (release-source proof); no run-length equivalence asserted	[Open]
The reset sign condition	no implication to half-membership stated	[Open]
Rigidity after long runs	partial rigidity, no exclusion	[Open]
The remaining middle-cell exclusions and right-branch bound	$\Rightarrow$ two-sided bound (Lean, given hyp.)	[Open]

The upper-reset band $\Rightarrow$ O2 (Lean), linear-width condition

[Open], cited check for $13 \leq r \leq 30$

12.9 Possible inputs and their scope

The following comparisons identify missing hypotheses in the recorded arguments. They neither predict the length of a proof nor classify all possible methods.

O1 (universal). The question is U257 itself. The stronger schema $O1' = \forall A \text{ Cert}(A)$ is a rejected sufficient strengthening, not future work: Proposition 11.4 already exhibits a squarefree support on which neither the digitwise nor the carry-aware block-certificate required input exists, while Duverney–Tachiya still prove that squarefree value irrational. Supplying $\text{Cert}(A)$ for arbitrary infinite supports is therefore not an open equivalent normal form of U257. Remaining work on O1 is the universal irrationality statement, or a different method that covers supports the certificate schemas cannot.

O2/O3/O5 (Mersenne-specific). The sharp-capacity gap $(*) \rightarrow (**)$ and the $-2, -1$ middle-cell exclusion are both, at bottom, requests for anti-concentration of an explicit divisor-count quantity (S , or $4 \cdot \text{rem} - \text{belowPulse} - 4$) away from specified short integer windows. A congruence argument, a quantitative distribution estimate, or another method would have to supply the stated inequalities. The reformulations do not decide which kind of argument can do so.

O4 (reset separation). The specified quotient-row argument needs $|\text{rem}(r+1) - 2^{r+1}| > 2^{(r+5)/2}$ at its reset rows. A weighted divisor window is one proposed source of information. After division by 2^{r+1} , the displayed condition becomes $|\lambda_{r+1}| > 2^{(3-r)/2}$. A bound for $\text{dist}(\Theta_L(M), \mathbb{Z})$ concerns a different quantity; Section 11.1 displays the additional skip-set term that a transfer argument would need to control. No such implication or uniform estimate is supplied. Part IV compares this with a totient-window hypothesis for a particular argument about Problem 249. The analogy does not identify the two statements, transfer a theorem between them, or show that either hypothesis is necessary for its original problem.

12.10 Summary

The results here do not decide the universal statement ($O1 = \text{U257}$) or the two rational membership targets $1/2$ (O2) and $1/21$ (O2b). The conditions surveyed here have different logical roles. Those equivalent to O2, including O3 by Proposition 2.7(c), do not establish membership merely by reformulating it. The sufficient conditions O4, O5, and O3-supply have not been supplied at the required cofinal scale, and no converse is proved for them. The proposed universal certificate strengthening $O1'$ is a separate, rejected statement: it is already false on squarefree support. Finite computational checks, however large their margins, are distinct from all these quantified assertions. The sufficient support theorems in the first part are proved for their stated classes. The identities, finite checks and equivalences in this section organise the remaining questions but do not resolve either rational target or the universal problem.

13 Logarithmic cost under arithmetic sampling

Positive fractional covers retain their geometric-kernel estimate, but their optimised cost cannot be bounded uniformly by a positive logarithmic divisor cost. The obstruction occurs on complete arithmetic periods, rather than only in incomplete-period errors. We give the construction and the surviving initial-interval estimate here; the short note records their consequences beside its cube-family calculation.

The proofs in this section are ordinary arguments from the 17 September research note [19, Theorem 1, Corollary 3 and Proposition 4]. They have not had independent review or a fresh Lean verification. For the underlying Lambert subseries, Kovač–Tao treat the same fixed-base object in [18, Section 2.1.2 and Remark 4.1]; their strict-tail argument gives distinct support coding and Cantor topology. Van Doorn–Kovač treat a freely chosen denominator sequence. Their convention is $n_{i+1}/n_i \geq \lambda$ at every index. For each $1 < \lambda < 2$, they construct such a sequence whose finite reciprocal sums include every rational in $[0, 2]$; no sequence with $n_{i+1}/n_i \geq 2$ can do this on any nonempty open interval [20, Theorem 1, Lemma 4 and Corollary 5]. This is a statement about the rational points of an interval, not an interval consisting entirely of finite sums; there are only countably many finite subsets. Those results provide context, not an arithmetic logarithmic estimate or a solution of fixed-base rational membership. The proof dependencies below are positive divisor majorisation, elementary prime selection, the Chinese remainder theorem, and the original kernel estimate (12). The short note’s logarithmic lower bound and cube-family calculation concern the same fractional-cover cost; their structured upper bound is not a uniform logarithmic converse.

13.1 The two finite costs

Let F be a finite nonempty subset of $\mathbb{N}_{>0}$, and write

$$Q = \text{lcm}(F), \quad f_F(n) = \#\{a \in F : a \mid n\}, \quad U_F(N) = \sum_{r \geq 1} 2^{-r} f_F(N + r) = \sum_{a \in F} \frac{2^{N \bmod a}}{2^a - 1}.$$

For $0 < t \leq 1$ define, with natural logarithms,

$$\kappa_1(F; t) = \min \left\{ \sum_{d \mid Q} \frac{c_d}{d} : c_d \geq 0, \quad \log(1 + f_F(s)/t) \leq \sum_{d \mid s} c_d \quad (s \mid Q) \right\}. \quad (14)$$

The constraints also hold at every positive integer n , by replacing n with $\gcd(n, Q)$. The programme is finite and feasible; a bounded cost sublevel bounds each coefficient, so its minimum is attained. For $L \geq 1$, let $\mathbb{P}_L$ be uniform sampling of $N = Lm$ over a period of $U_F(Lm)$, for which $Q/\gcd(Q, L)$ is a permissible length. This notation does not assert that the displayed period is least.

We use $K_*(F)$ for the infimum of the existing fractional-cover cost

$$K = \sum_j \frac{C_j \eta_j^{-\alpha_j}}{2^{\alpha_j} - 1}, \quad C_j = \sum_{d \geq 1} \frac{c_{j,d}}{d}, \quad (15)$$

over finite or countable families of finite sets F_j of positive integers covering F , exponents $0 < \alpha_j \leq 1$, weights $\eta_j > 0$ with $\sum_j \eta_j = 1$, and coefficients $c_{j,d} \geq 0$ satisfying $f_{F_j}(n)^{\alpha_j} \leq \sum_{d|n} c_{j,d}$ for every $n \geq 1$. The covering sets need not be disjoint or contained in F . Covers of infinite cost do not lower the infimum, and a finite-cost cover exists, for example the single set F with exponent one. This is the manuscript's fractional cost, not a redefinition of the logarithmic cost.

Theorem 13.1 (arithmetic logarithmic counterexample). *For every integer $H \geq 2$ and every real $A_0 \geq 0$, there are a squarefree positive integer L and a finite nonempty set F of distinct squarefree positive integers such that*

$$\min F > \max\{L, A_0\}, \quad \kappa_1(F; 1) \leq \frac{30 \log 2}{H}, \quad \mathbb{P}_L(U_F > 1) \geq 1 - e^{-1}.$$

For these F, L , some $R \geq 0$ satisfies $\mathcal{D}_{L,R,L} \mathbf{1}_{\{U_F > 1\}} > 1/2$. Consequently, no absolute constant C satisfies

$$\mathcal{D}_{L,R,M} \mathbf{1}_{\{U_F > t\}} \leq C(1 + L/M) \kappa_1(F; t) \quad (\text{E})$$

uniformly in finite F , positive integers L, M , integers $R \geq 0$, and $0 < t \leq 1$.

Lean: [arithmetic_logarithmic_counterexample, no_absolute_dyadic_kappaOne_constant](#).

Proof. The modulus makes many small divisor conditions automatic. Fresh tag primes then make the remaining incidences independent, and taking all least common multiples turns their number into an exponential incidence. A logarithm pays only for the individual tags, while the shifted potential can exceed one at many different offsets.

Choose a finite set of primes P such that, with

$$L = \prod_{p \in P} p, \quad D = \tau(L), \quad B = \prod_{p \in P} (1 + 1/p),$$

we have $B \geq H^2$. Such a set exists because the sum of reciprocal primes diverges and $\log(1 + 1/p) \geq 1/(2p)$. Here τ counts positive divisors and σ denotes their sum. Complementary divisors satisfy

$$\frac{\sigma(r)}{r} \frac{\sigma(L/r)}{L/r} = B \quad (r \mid L).$$

Thus the set $G = \{r \mid L : \sigma(r)/r \geq H\}$ has at least $D/2$ elements. Also $D \geq B \geq H^2$.

Choose a finite set $\mathcal{B}$ of parent primes, all exceeding $\max\{L, H, D, A_0\}$, with

$$\frac{4}{D} \leq \sum_{q \in \mathcal{B}} \frac{1}{q} \leq \frac{5}{D}.$$

Divergence of the prime reciprocal sum after any finite cutoff permits stopping at the first crossing of $4/D$; the overshoot is less than $1/D$. For each $q \in \mathcal{B}$ and $d \mid L$, choose a finite set $T_{q,d}$ of tag primes, all greater than H , disjoint from P , from $\mathcal{B}$, and from every other tag set, such that

$$\frac{4d}{H} \leq S_{q,d} := \sum_{p \in T_{q,d}} \frac{1}{p} \leq \frac{5d}{H}.$$

The same stopping argument works successively: each new reciprocal is less than $1/H \leq d/H$, and only finitely many primes have been excluded. No bound on the size of these prime sets is needed.

Let $\mathcal{I}_q = \{(d, p) : d \mid L, p \in T_{q,d}\}$, and define

$$F_q = \{q \operatorname{lcm}\{dp : (d, p) \in I\} : I \subseteq \mathcal{I}_q\}, \quad F = \bigcup_{q \in \mathcal{B}} F_q,$$

where the least common multiple of the empty set is one. Every tag prime occurs in just one labelled pair. Its presence in the least common multiple therefore recovers whether that pair was selected. Hence different subsets give distinct exponents. All exponents are squarefree; different parent primes give disjoint sets, since no parent prime occurs in any other set. Their minimum is greater than $\max\{L, A_0\}$. The nonempty tag sets with $d = L$ also show $L \mid Q$.

For an integer $n \geq 1$ put $Z_q(n) = \sum_{(d,p) \in \mathcal{I}_q} \mathbf{1}_{\{dp \mid n\}}$. The least-common-multiple encoding gives the exact identity

$$f_{F_q}(n) = \mathbf{1}_{\{q \mid n\}} 2^{Z_q(n)}.$$

Indeed, when $q \mid n$, a selected least common multiple divides n exactly when every selected pair does. Consequently

$$\log(1 + f_{F_q}(n)) \leq (\log 2) \left(\mathbf{1}_{\{q \mid n\}} + \sum_{(d,p) \in \mathcal{I}_q} \mathbf{1}_{\{qdp \mid n\}} \right).$$

All moduli in this majorant divide $\operatorname{lcm}(F_q)$. Its cost is at most

$$\frac{\log 2}{q} \left(1 + \sum_{d \mid L} \frac{S_{q,d}}{d} \right) \leq \frac{\log 2}{q} (1 + 5D/H). \quad (16)$$

Now sample N uniformly modulo Q , subject to $L \mid N$. For $r \in G$, the event $q \mid N + r$ has probability $1/q$. For a fixed parent these events are disjoint as $1 \leq r \leq L < q$. Conditional on one such event, the tag-prime events remain independent by the Chinese remainder theorem. Since $d \mid N + r$ is equivalent to $d \mid r$, the mean of $Z_q(N + r)$ is

$$\mu_r = \sum_{d \mid r} S_{q,d} \geq \frac{4}{H} \sigma(r) \geq 4r,$$

and its variance is at most μ_r . Chebyshev's inequality gives

$$\mathbb{P}_L(Z_q(N + r) < r \mid q \mid N + r) \leq \frac{\mu_r}{(\mu_r - r)^2} \leq \frac{4}{9r} \leq \frac{4}{9}.$$

Thus success has conditional probability at least $1/2$. On success the r th term of $U_{F_q}(N)$ is at least $2^{-r} 2^r = 1$. Later positive terms, for example those supplied by the

exponent $q \in F_q$, make $U_{F_q}(N) > 1$ strictly. The event E_q of success for some $r \in G$ consequently satisfies

$$\mathbb{P}_L(E_q) \geq \frac{|G|}{2q} \geq \frac{D}{4q}.$$

Different parent groups depend on disjoint prime coordinates, so the E_q are independent under this conditional uniform measure. It follows that

$$\mathbb{P}_L(U_F > 1) \geq 1 - \prod_{q \in \mathcal{B}} (1 - \mathbb{P}_L(E_q)) \geq 1 - \exp\left(-\frac{D}{4} \sum_{q \in \mathcal{B}} \frac{1}{q}\right) \geq 1 - e^{-1}.$$

Since $f_F = \sum_q f_{F_q}$ and $\log(1 + \sum_q x_q) \leq \sum_q \log(1 + x_q)$ for $x_q \geq 0$, we may sum the positive majorants in (16):

$$\kappa_1(F; 1) \leq \frac{5 \log 2}{D} (1 + 5D/H) \leq \frac{30 \log 2}{H}.$$

Since $L \mid Q$, write $P_0 = Q/L$ and $p = \mathbb{P}_L(U_F > 1)$. The bounded periodic sequence $\mathbf{1}_{\{U_F(Lm) > 1\}}$ satisfies, for every integer $T \geq 1$,

$$\left| \frac{1}{T} \sum_{m=1}^T \mathbf{1}_{\{U_F(Lm) > 1\}} - p \right| \leq \frac{P_0}{T}.$$

Summing over the dyadic window gives

$$\left| \mathcal{D}_{L,R,M} \mathbf{1}_{\{U_F > 1\}} - p \right| \leq \frac{2P_0}{M2^R}.$$

Take $M = L$ and $2^R \geq 16P_0$. The error is at most $1/8$, whereas $p \geq 1 - e^{-1} > 5/8$, so the mean exceeds $1/2$. A putative bound (E) would instead give at most $60C \log 2/H$. Choosing an integer $H > 120C \log 2$ contradicts it. $\square$

13.2 The arithmetic lower bound for every fractional cover

Corollary 13.2 (finite-functional separation). *For every finite nonempty F , with $Q = \text{lcm}(F)$,*

$$K_*(F) \geq \max_{\ell \mid Q} \mathbb{P}(U_F(N) > 1 \mid \ell \mid N), \quad (17)$$

where N is uniform modulo Q . The supports of Theorem 13.1 satisfy $K_*(F) \geq 1 - e^{-1}$ and $\kappa_1(F; 1) \leq 30 \log 2/H$. In particular, no absolute C gives $K_*(F) \leq C\kappa_1(F; 1)$ for all finite nonempty F , even after all covering sets, exponents, coefficients and positive weights have been optimised.

Lean: `sup_condExceedProb_le_paperCoverCost, exists_support_paperCoverCost_ge, no_absolute_paperCoverCost_kappaOne_constant.`

Proof. Fix one finite-cost cover in (15), and write $U_j = U_{F_j}$ and $B_j = 2^{\alpha_j}$. Coverage gives $U_F \leq \sum_j U_j$. If $U_F > 1$, then $U_j > \eta_j$ for at least one index. Subadditivity and the divisor majorants therefore give the pointwise estimate

$$\mathbf{1}_{\{U_F > 1\}} \leq \sum_j \eta_j^{-\alpha_j} U_j^{\alpha_j} \leq \sum_{j,d} \eta_j^{-\alpha_j} c_{j,d} w_{B_j,d}.$$

Nonnegative interchange and the original geometric-kernel estimate (12), with its factor $1 + 4L/M$, imply

$$\mathcal{D}_{L,R,M} \mathbf{1}_{\{U_F > 1\}} \leq (1 + 4L/M)K.$$

For fixed L, M , let R tend to infinity. Periodicity of the left-hand indicator gives $\mathbb{P}_L(U_F > 1) \leq (1 + 4L/M)K$. Then let M tend to infinity, with L fixed, to obtain $\mathbb{P}_L(U_F > 1) \leq K$. For $L = \ell \mid Q$ this is the conditional probability in the statement. It holds for every admissible cover, so take their infimum and the maximum over the finitely many divisors ℓ . Finally use the constructed $L \mid Q$ and let H tend to infinity. $\square$

13.3 What survives without arithmetic restriction

Proposition 13.3 (ordinary initial intervals). *For every finite nonempty F , integer $X \geq 1$, and $0 < t \leq 1$,*

$$\frac{1}{X} \#\{1 \leq N \leq X : U_F(N) > t\} \leq \frac{2}{\log(4/3)} \kappa_1(F; t).$$

Lean: [logarithmic_initial_interval](#).

Proof. Choose a positive logarithmic divisor majorant $g(n) = \sum_{d \mid n} c_d$ of cost $K = \sum_{d \mid Q} c_d/d$, putting $c_d = 0$ when $d \nmid Q$, and let $a = \log(4/3)$. For positive integers n, Y ,

$$g(n) \leq Kn, \quad \sum_{n=1}^Y g(n) = \sum_{d \mid Q} c_d \lfloor Y/d \rfloor \leq KY.$$

If $U_F(N) > t$, there is an $r \geq 1$ with $\log(1 + f_F(N+r)/t) > ar$. Otherwise every $f_F(N+r)/t \leq (4/3)^r - 1$, and

$$\frac{U_F(N)}{t} \leq \sum_{r \geq 1} 2^{-r} ((4/3)^r - 1) = 1,$$

a contradiction. If $K \geq a/2$ the required counting bound is trivial. For $K < a/2$, a witnessing integer $n = N+r$ satisfies $a(n-N) < g(n) \leq Kn$, whence $n < aN/(a-K) \leq 2X$ when $N \leq X$. A fixed n can witness at most $g(n)/a$ integer values of N , since $1 \leq n-N < g(n)/a$. The union bound thus gives

$$\#\{1 \leq N \leq X : U_F(N) > t\} \leq \frac{1}{a} \sum_{n=1}^{2X} g(n) \leq \frac{2KX}{a}.$$

Take the infimum over the admissible majorants. $\square$

The proposition also bounds every $L = 1$ dyadic observation window, since such a window averages ordinary initial-interval means. It does not bound conditional arithmetic sampling uniformly in L . Freezing a finite initial support in the irrationality argument requires exactly that arithmetic restriction. The proposed logarithmic-tail irrationality corollary therefore loses this proof; its conclusion is not disproved. The corollary above separates finite functionals, not infinite-support irrationality classes. The weighted, fractional-cover, mixed and finite-period noncollapse results are unchanged. No conclusion here decides universal #257 or membership of $1/2$ or $1/21$.

Two proof details and source comparators

The elementary atom comparison can be sharpened to $\Delta_{b,A}(N) \leq \Delta_{2,A}(N)$ for all real $b \geq 2$. After cancelling $b - 1$, differentiate the ratio $\sum_{i < r} b^i / \sum_{i < d} b^i$: its numerator is $\sum_{i < r \leq j < d} (i - j)b^{i+j-1} \leq 0$. Only the rational-lattice step needs an integer base.

For the scalar positive-cover cost, put $z = \log_2 t$. Direct differentiation of $y^z / (y - 1)$ on $1 < y \leq 2$ gives

$$\Psi(t) = t \quad (1 \leq t \leq 4), \quad \Psi(t) = \frac{z^z}{(z - 1)^{z-1}} \quad (t > 4).$$

These are ordinary proof details, not additional replayed declarations.

The original periodic and linear-independence papers of Luca and Tachiya are now distinguished from their later overview [13, 14, 3]. Their theorem formulations used in this revision were read in the supplied author-written overview; the two original publisher PDFs were not successfully retrieved in this revision. Van Assche's nonvanishing and small-linear-form argument for the full Lambert series [15, Lemma 1 and (34)–(35)] is a comparator for the missing upper-height control, not a result for arbitrary thinnings.

For achievement-set background, Hornich's original bibliographic record [16] is separated from the exposition actually used here [4]. Głab and Prus-Wiśniowski's survey [17] provides a modern map of geometry and cardinality questions, not a rational-membership theorem for the present set.

Acknowledgements

I thank Wouter van Doorn for advice on mathematical exposition, especially on making restrictive hypotheses intelligible, avoiding unnecessary notation, and writing for a first-time reader. His comments concerned a note on Problem 243; this acknowledgement does not imply review of the mathematics of the present paper.

Statements and declarations

Artefact and data availability. The linked Lean declarations, fixed toolchain, and library manifest are published in the companion repository. This manuscript provides navigation and exposition rather than proof authority.

Authority boundary. Kernel checking establishes that a proposition was proved; it does not authorise the exposition, literature judgements, or any claim that Problem 257 is solved.

Funding and competing interests. This work received no external funding. The author declares no competing interests.

References

- [1] P. Erdős, *On arithmetical properties of Lambert series*, J. Indian Math. Soc. 12 (1948), 63–66.

- [2] P. Erdős, *On the irrationality of certain series*, Math. Student 36 (1968), 222–226 (issued 1969); [five-page scan](#). The theorem on p. 222 treats pairwise-coprime support with convergent reciprocal sum at every integer base $b \geq 2$; the claimed removal of pairwise coprimality is stated without proof. The same page says the reciprocal-sum condition could be replaced by a weaker but more complicated condition, and p. 226 suggests $\sum_{n_i < x} 1/n_i = o(\log \log x)$ for pairwise coprime supports.
- [3] F. Luca and Y. Tachiya, *Linear independence results for the values of divisor functions series*, RIMS Kôkyûroku No. 2014 (2017), 138–150. Theorem A (p. 139) restates Theorem 1.1 of their paper *Irrationality of Lambert series associated with a periodic sequence*, Int. J. Number Theory 10 (2014), no. 3, 623–636, DOI: for a purely periodic integer sequence $a(n)$, not identically zero, $\sum_{n \geq 1} a(n)/(q^n - 1)$ is irrational for every integer q with $|q| > 1$. Example 2 (p. 140) treats the odd support.
- [4] Z. Nitecki, *Subsum sets: intervals, Cantor sets, and Cantorvals*, arXiv:1106.3779v2 (2013). Theorem 4 (p. 9) shows that when every term exceeds its tail the subsum set is a Cantor set of Lebesgue measure $\lim_n 2^n X_n$, where X_n is the n th tail, and credits this to H. Hornich (1941).
- [5] H. Kaneko, Y. Suzuki, and Y. Tachiya, *Refinements of Erdős’s irrationality criterion for certain sparse infinite series*, arXiv:2601.20743v1 (2026). The averaged tail $R_c(q, x, z)$, a sum over offsets $j \geq z$, is (1.7) and Theorem 1 is on p. 3; its rational-integer form, Theorem 3, is on p. 5. Lemma 1 (pp. 6–7) derives irrationality from arbitrarily small nonzero scaled tails, and Lemma 2 (pp. 7–8) shows that most scaled tails in a range are small under a counting condition on the support. Their criteria assume sparse coefficient supports.
- [6] D. Duverney and Y. Tachiya, *Refinement of the Chowla–Erdős method and linear independence of certain Lambert series*, Forum Math. 31 (2019), no. 6, 1557–1566, DOI. Page numbers refer to the linked author preprint. Theorem 1.1 (pp. 2–3) is the refined Chowla–Erdős criterion; its proof in Section 2 (pp. 5–7) selects an index by minimising a local coefficient mass along an arithmetic progression, (2.3)–(2.9). Corollary 1.2 gives the general $F_s(E)$ theorem and its monomial images under $|q| \operatorname{lcm}(1, \dots, \ell) \leq s$, and at every integer base when $s = \infty$; Example 1.1 gives the joint squarefree family at all bases 2^j , $j \geq 1$. Both are on p. 4, which also relates Corollary 1.2 to the conjecture of Erdős and Graham; the proof of Corollary 1.2 is on pp. 10–11.
- [7] T. Tao and J. Teräväinen, *Quantitative correlations and some problems on prime factors of consecutive integers*, arXiv:2512.01739v2 (submitted December 2025, revised April 2026). Theorem 1.3 (p. 4) proves $\sum_{n \geq 1} \omega(n)/2^n = \sum_p (2^p - 1)^{-1}$ irrational, settling the prime-support case of #257 at base 2; the paragraph after it states that the prime-support result extends to every integer base $b \geq 2$, and separately that full prime-power support is irrational at base 2. The modifications for those two assertions are left to the reader; arbitrary subsets of the primes or prime powers are not treated there. Theorem 3.1 is on p. 24, and the proof of Theorem 1.3 is Section 5, pp. 44–56.
- [8] W. Zudilin, *Heine’s basic transform and a permutation group for q -harmonic series*, Acta Arith. 111 (2004), no. 2, 153–164, DOI. Theorem 1 (p. 154) gives $\mu(h_p(1)) \leq 2.46497868 \dots$ for $h_p(1) = \sum_{n \geq 1} 1/(p^n - 1)$, uniformly in the integer p with $|p| \geq 2$; the remark at the end of Section 3 (p. 159) corrects the computation of the author’s earlier paper.
- [9] J. M. Campbell, *On the binary digits of the Erdős–Borwein constant*, arXiv:2605.24160v1 (2026). Theorem 1 (p. 12) proves that the block 11 occurs infinitely often in the binary expansion of $\sum_{n \geq 1} 1/(2^n - 1)$, answering a question of Crandall recalled in the introduction.

- [10] K. Matomäki and M. Radziwiłł, *Multiplicative functions in short intervals*, Ann. of Math. (2) 183 (2016), no. 3, 1015–1056, DOI. Page numbers refer to arXiv:1501.04585v4.
- [11] A. P. Mangerel, *Divisor-bounded multiplicative functions in short intervals*, Res. Math. Sci. 10 (2023), no. 1, Paper No. 12, DOI. Page numbers refer to arXiv:2108.11401v2.
- [12] Y.-C. Sun, *The critical-window profile for the k -fold divisor function in almost all short intervals*, arXiv:2401.08432v3 (2026).
- [13] F. Luca and Y. Tachiya, *Irrationality of Lambert series associated with a periodic sequence*, International Journal of Number Theory 10 (2014), no. 3, 623–636. doi:10.1142/S1793042113501121.
- [14] F. Luca and Y. Tachiya, *Linear independence of certain Lambert series*, Proceedings of the American Mathematical Society 142 (2014), no. 10, 3411–3419. doi:10.1090/S0002-9939-2014-12102-2.
- [15] W. Van Assche, *Little q -Legendre polynomials and irrationality of certain Lambert series*, The Ramanujan Journal 5 (2001), 295–310. doi:10.1023/A:1012930828917; arXiv:math/0101187v1.
- [16] H. Hornich, *Über beliebige Teilsummen absolut konvergenter Reihen*, Monatshefte für Mathematik und Physik 49 (1941), 316–320. doi:10.1007/BF01707309.
- [17] S. Głąb and F. Prus-Wiśniowski, *Achievement sets – current results and open problems*, arXiv:2512.17285v1 (2025), preprint. <https://arxiv.org/abs/2512.17285v1>.
- [18] V. Kovač and T. Tao, *On several irrationality problems for Ahmes series*, Acta Math. Hungar. 175 (2025), no. 2, 572–608, DOI. Page numbers refer to arXiv:2406.17593v4. Remark 4.1 (p. 13) records the strict tail inequality and the Cantor structure. Theorem 2.3 (p. 5; proof pp. 13–14) constructs rational merged sums from several bases under its mass hypothesis; the construction merges several bases and leaves fixed-base #257 untouched.
- [19] *The logarithmic endpoint fails under arithmetic sampling*, AI-assisted ordinary proof note, 17 September 2026. Theorem 1, Corollary 3 and Proposition 4. Unpublished working note; independent review and fresh Lean verification are outstanding.
- [20] W. van Doorn and V. Kovač, *Lacunary sequences whose reciprocal sums represent all rational numbers in an interval*, Acta Arith. 223 (2026), 275–295, DOI. The theorem and page references used here refer to arXiv:2509.24971v3 (3 December 2025), <https://arxiv.org/abs/2509.24971v3>. Theorem 1 is on p. 2; Lemma 4 and Corollary 5 are on pp. 5–7.
- [21] P. Erdős, *Some problems and results on the irrationality of the sum of infinite series*, J. Math. Sci. 10 (1975), 1–7.
- [22] K. Barreto, J. Kang, S.-H. Kim, V. Kovač, and S. Zhang, *Irrationality of rapidly converging series: a problem of Erdős and Graham*, arXiv:2601.21442v3 (2026), to appear in Bull. London Math. Soc.
- [23] K. Conrad, *The Chinese remainder theorem*, expository notes, Theorem 3.1 (pp. 2–3). The compatibility criterion for noncoprime moduli used above is obtained by the displayed prime-power reduction.
- [24] P. Erdős, *Beweis eines Satzes von Tschebyschef*, 1932, pp. 194–198, especially §§4–5.